

Designing a Data Warehouse on the Microsoft SQL Server Platform

UNDERSTANDING THE FUNDAMENTAL CONCEPTS OF
A DATA WAREHOUSE



Ana Voicu

@ana_voicu



Overview



Goals and purpose of a data warehouse

Introducing dimensional modeling

- What are facts and fact tables?
- What are dimensions and dimension tables?
- Star schemas

Putting it all together

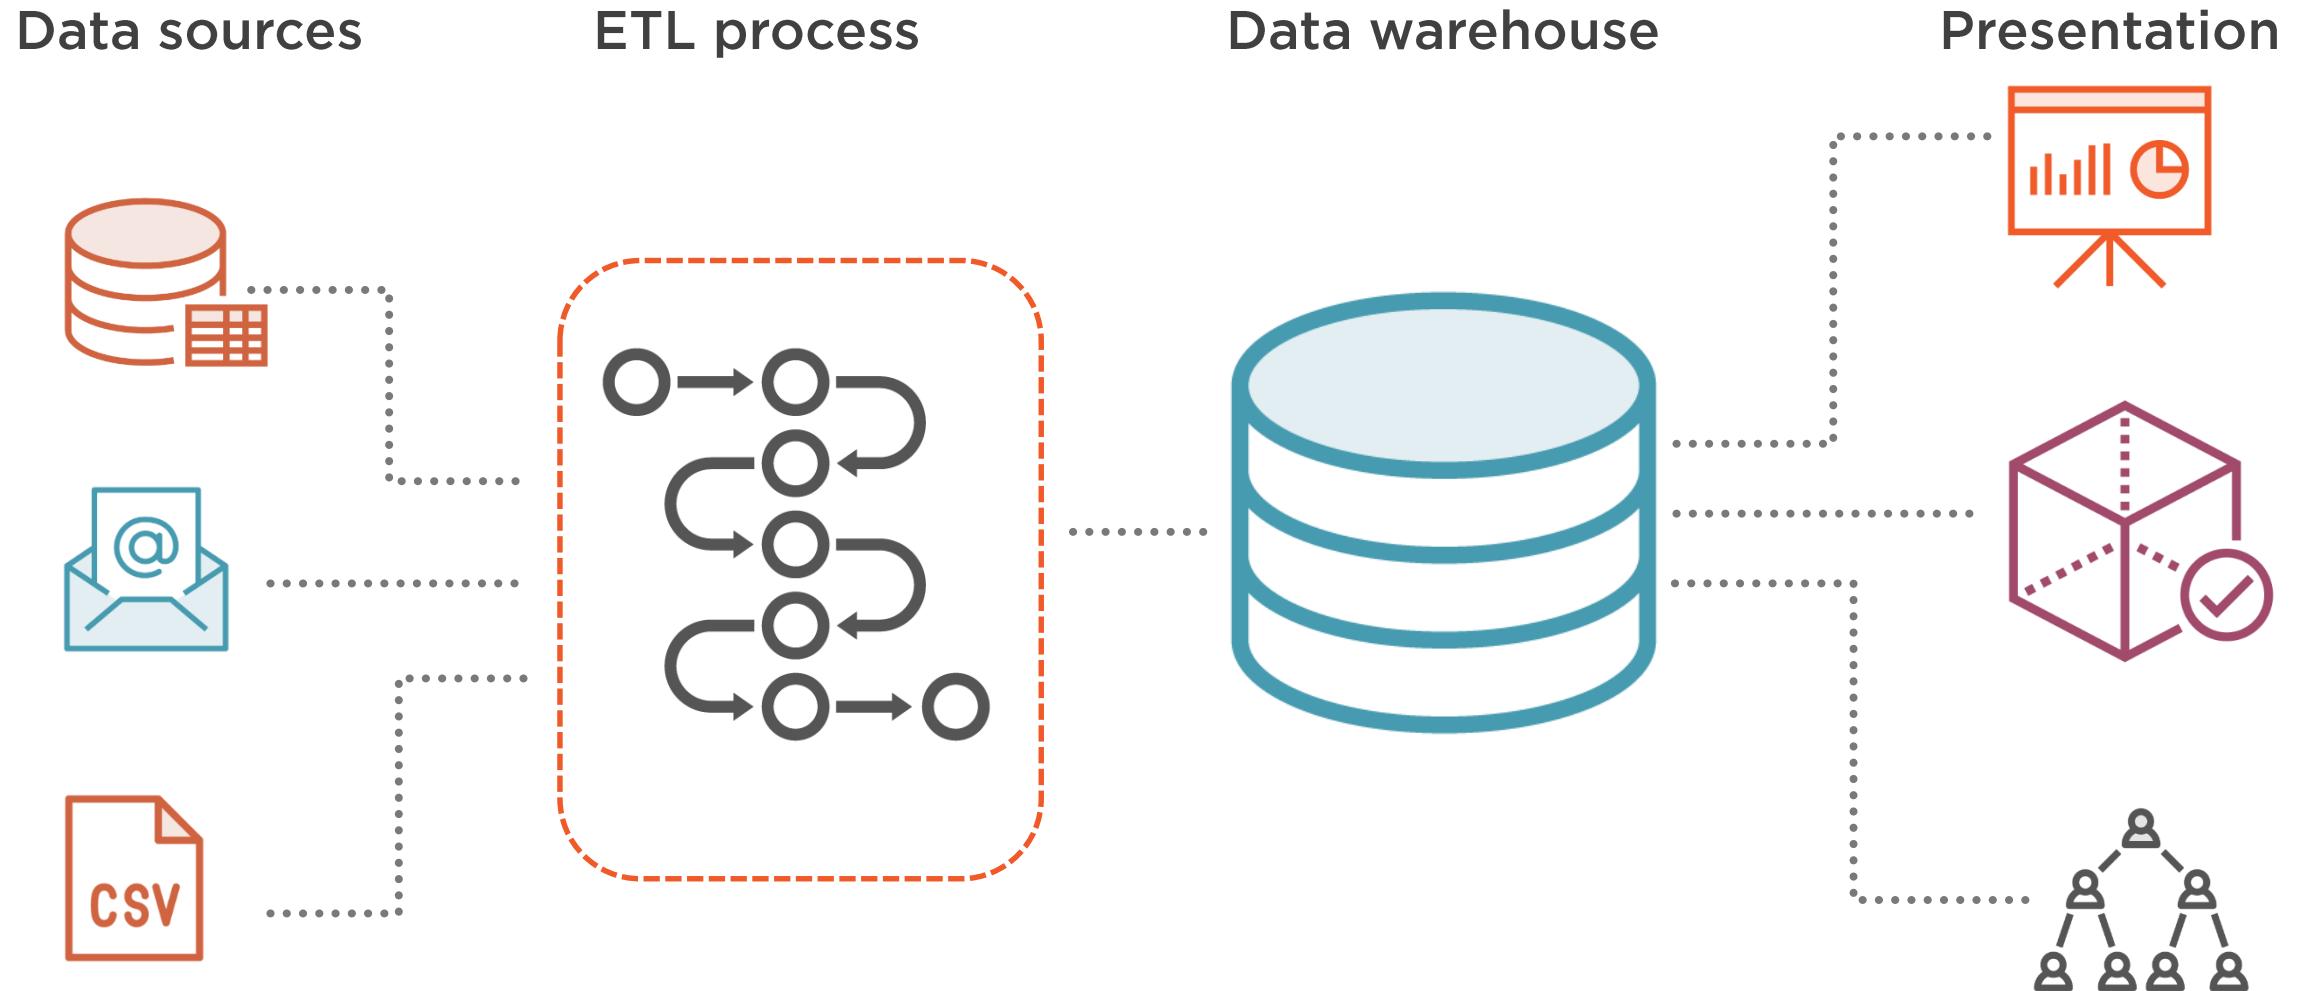
- The 4-step dimensional design process



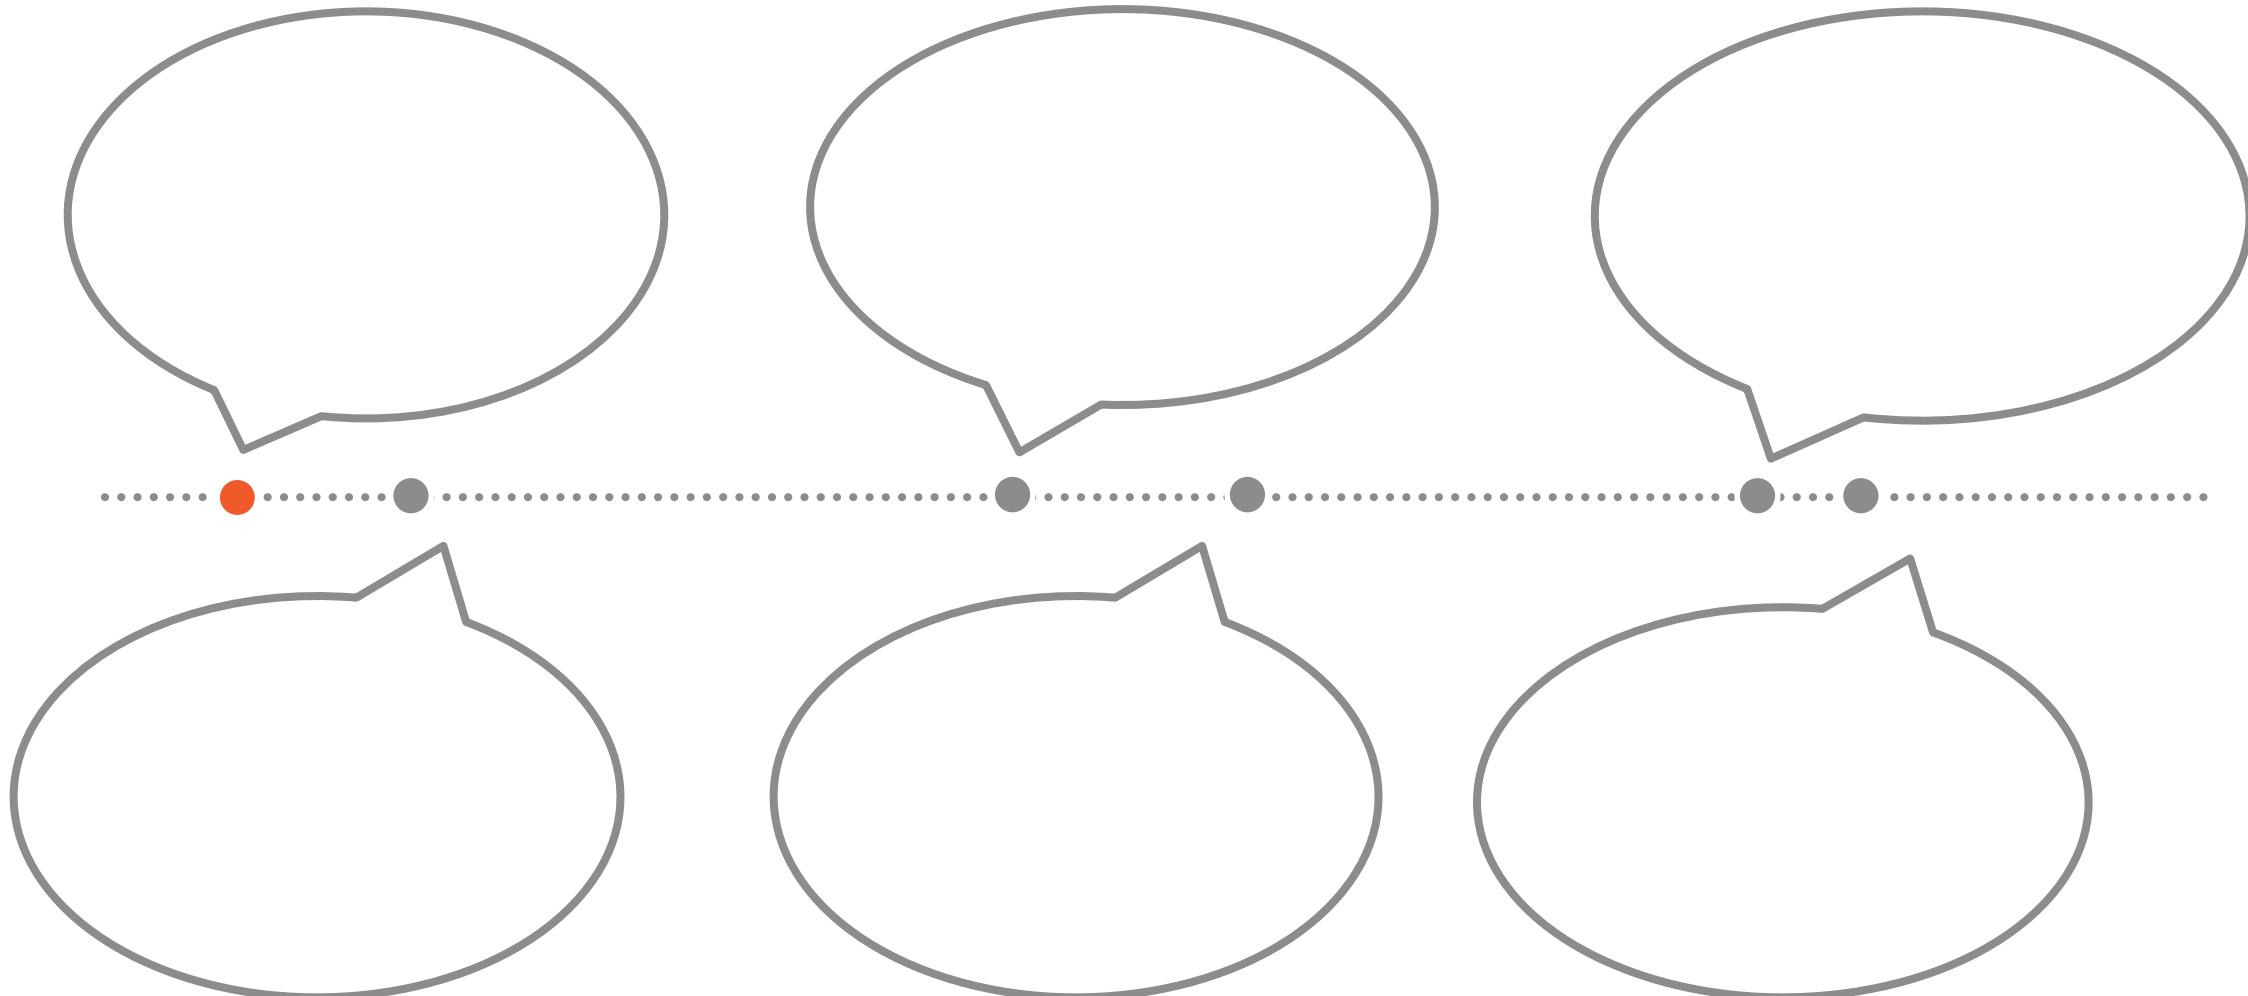
Goals of a Data Warehouse Solution



What Is a Data Warehouse?



Problems a Data Warehouse Can Solve



Requirements of a Data Warehouse Solution

Easily accessible

Fast

Consistent

Flexible

Secure

Foundation for
decision-making



Responsibilities of a Data Warehouse Designer

Understand the
business users

Deliver high-quality,
relevant, and
accessible
information

Sustain the DW
environment



Introduction to Dimensional Modeling



Dimensional Modeling



Database design method optimized for data warehouse solutions

Popular technique because it addresses two important requirements:

- Deliver data in an understandable format
- Deliver fast query performance

Key word is “simplicity”



Elements of a Dimensional Model



Facts (the measurements/metrics or facts from your business process)



Dimensions (for providing the context of a business process event)



Attributes (the various characteristics of a dimension)



Star schema (and/or OLAP cubes)



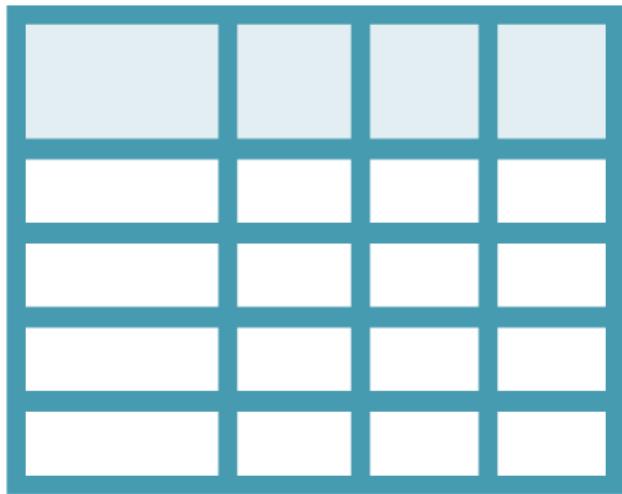
“We sell ice-cream and other products in various locations and measure our achievements over time.”

Jane Poppins

Happy Scoopers CEO



Fact Tables and Facts



Fact table = table that stores the performance measurements resulting from an organization's business process events



Fact = a business measure

- Sales
- Profit
- Volume
- Number of transactions



Fact Tables and Facts



Facts answer questions like:

- “What are we doing?” (sell, buy, count)
- “What do we want to achieve?” (more sales, bigger profit)

1 row in the fact table is 1 measurement in real life

Fact columns in a fact table should be additive

Facts make sense in combination with dimensions

- Linked with foreign keys
- Date/Time dimension is present in most data warehouses



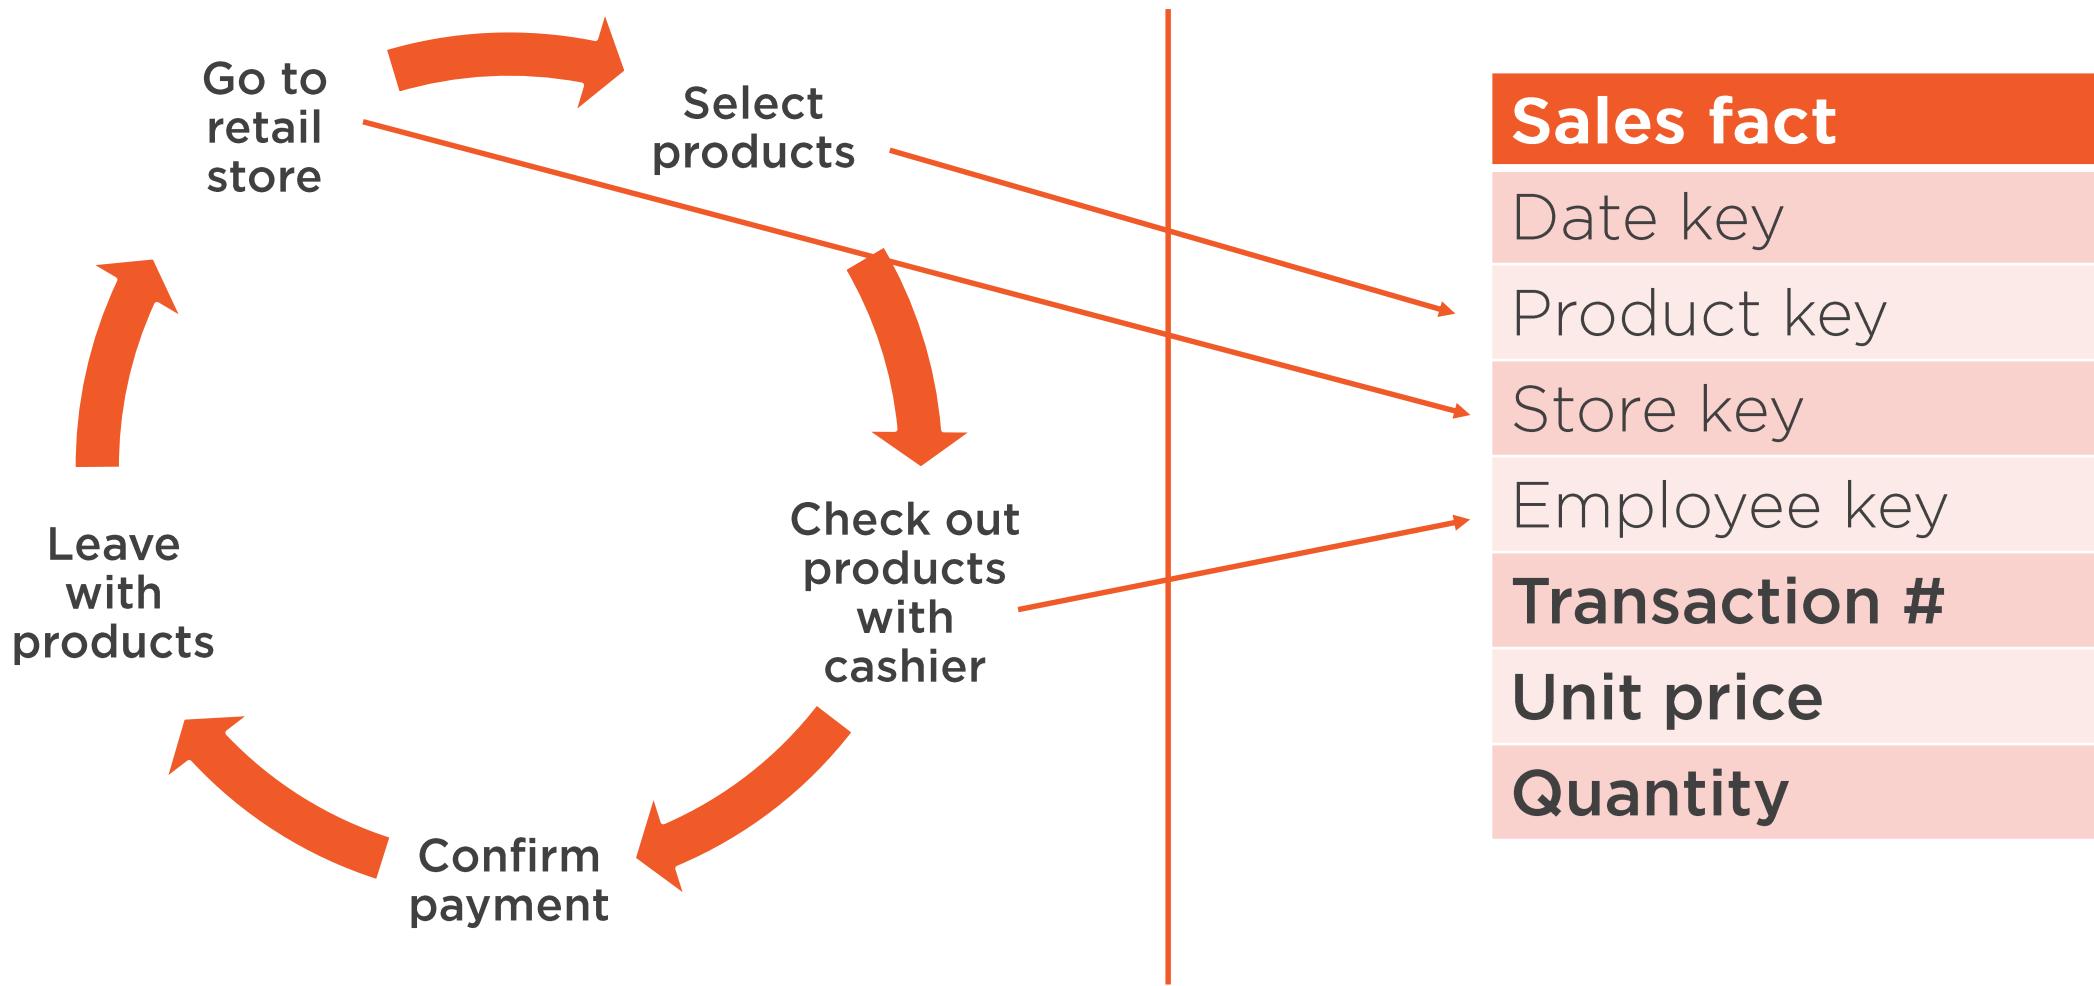
Example of a Fact Table



Sales fact
Date key
Product key
Store key
Employee key
Transaction #
Unit price
Quantity



Example of a Fact Table



Example of a Fact Table

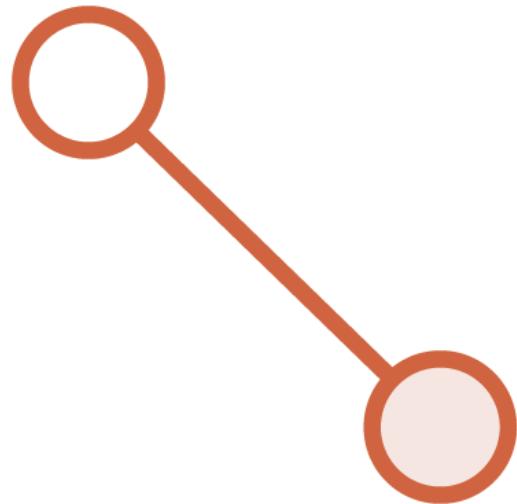


Sales Fact
Date key
Product key
Store key
Employee key
Transaction #
Unit price
Quantity

C
O
M
P
O
S
I
T
E
K
E
Y



Characteristics of Fact Tables



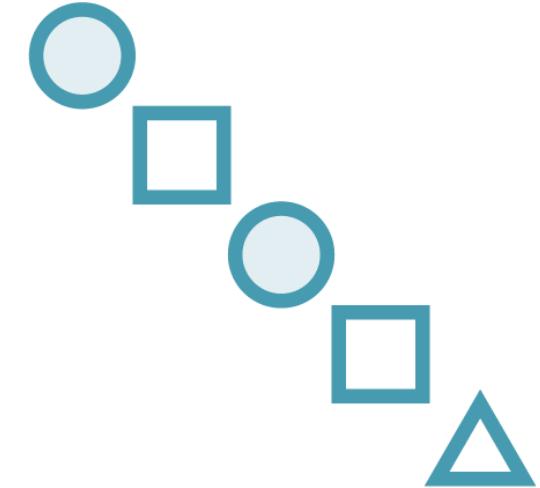
1:1 relationship
between fact table
row and real-world
event



Most facts
should be
additive



Foreign keys to
dimension
tables



Composite key
as PK for fact
table



What Are Dimensions?



Companions to a fact table

Textual context associated with a business process measurement event



Questions Answered by Dimension Tables



Who



What



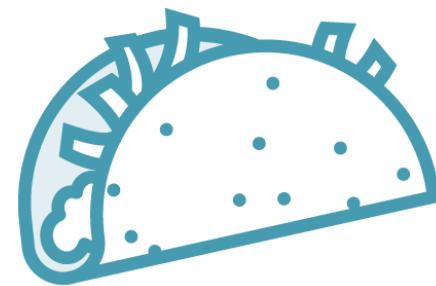
Where



When



How



Why



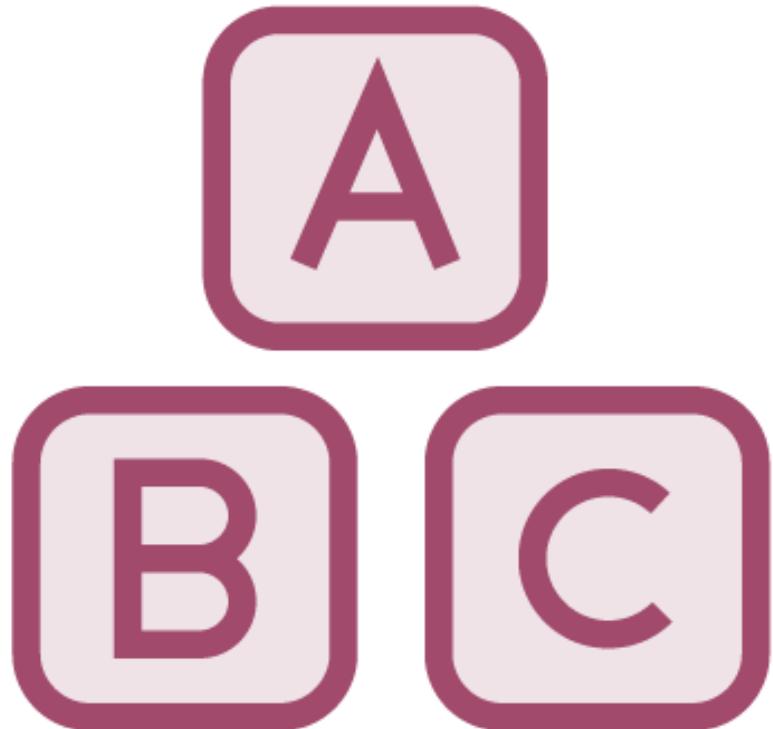
Example of a Dimension Table



Product Dimension
Product key
Product name
Brand name
Category name
Subcategory name
Package type
Package size
Weight
Weight unit of measure



Characteristics of Dimension Tables



No limit for the number of attributes in a dimension table

- Common to have tables with 50 to 100 attributes
- Some dimension tables have only a handful of attributes

Have fewer rows than fact tables

- But can be much wider

Defined by a single primary key

- Basis for the referential integrity with the fact table

Denormalized

- Flattened many-to-one relationships within a single dimension table



Dimension Attributes

The primary source of

- Query constraints
- Groupings
- Report labels

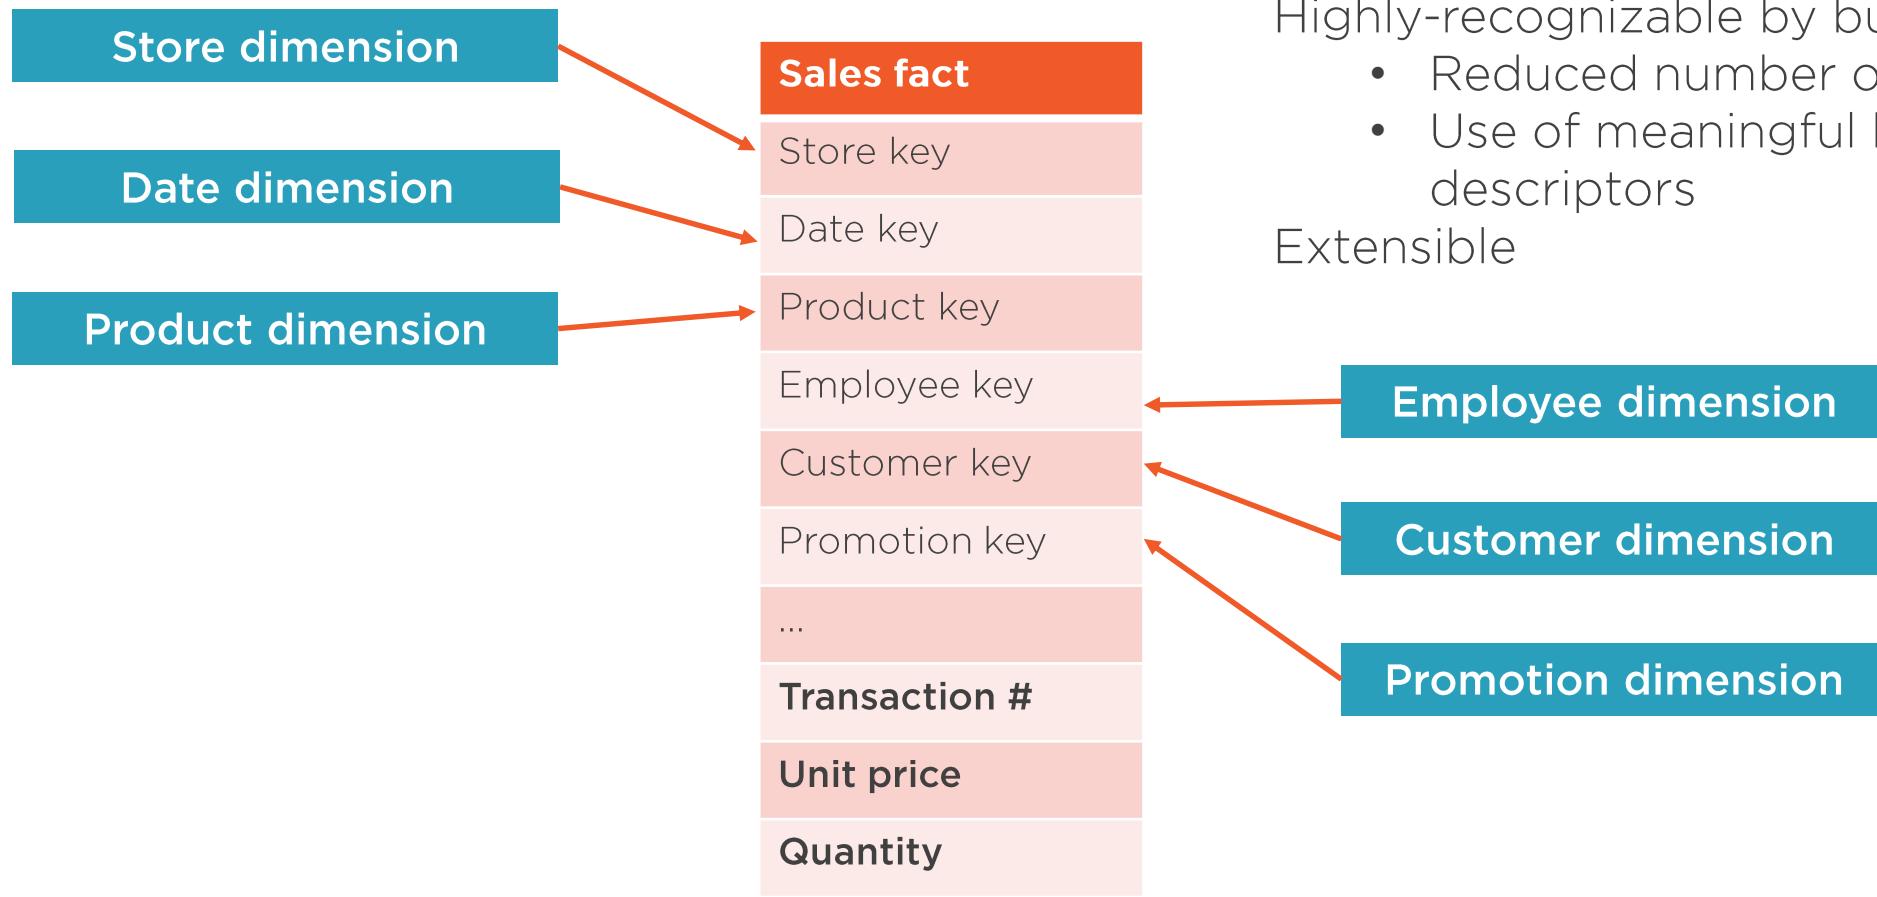
Quality of attributes \propto quality of the system

- Use real words vs. cryptic abbreviations
- Minimize codes in the dimension tables

Product Key	Name	Category name	Subcategory name
1	Beery cotton candy	Candy	French candy
2	Cotton candy	Candy	French candy
3	Peppermint candy (seasonal)	Candy	Fudge
4	Green tea ice cream	Ice-cream	Reduced fat
5	Chocolate chip cookie dough ice cream	Ice-cream	Fat-free frozen dairy
6	Neapolitan ice cream	Ice-cream	Lactose-free
7	Cantuccini	Cookie	Biscotti
8	Chocolate mint cookie	Cookie	Retro snacks
9	Lemon cookie	Cookie	Fruity cookies



Dimensional Model



Characteristics:

Simple

Symmetric

Highly-recognizable by business users

- Reduced number of tables
- Use of meaningful business descriptors

Extensible

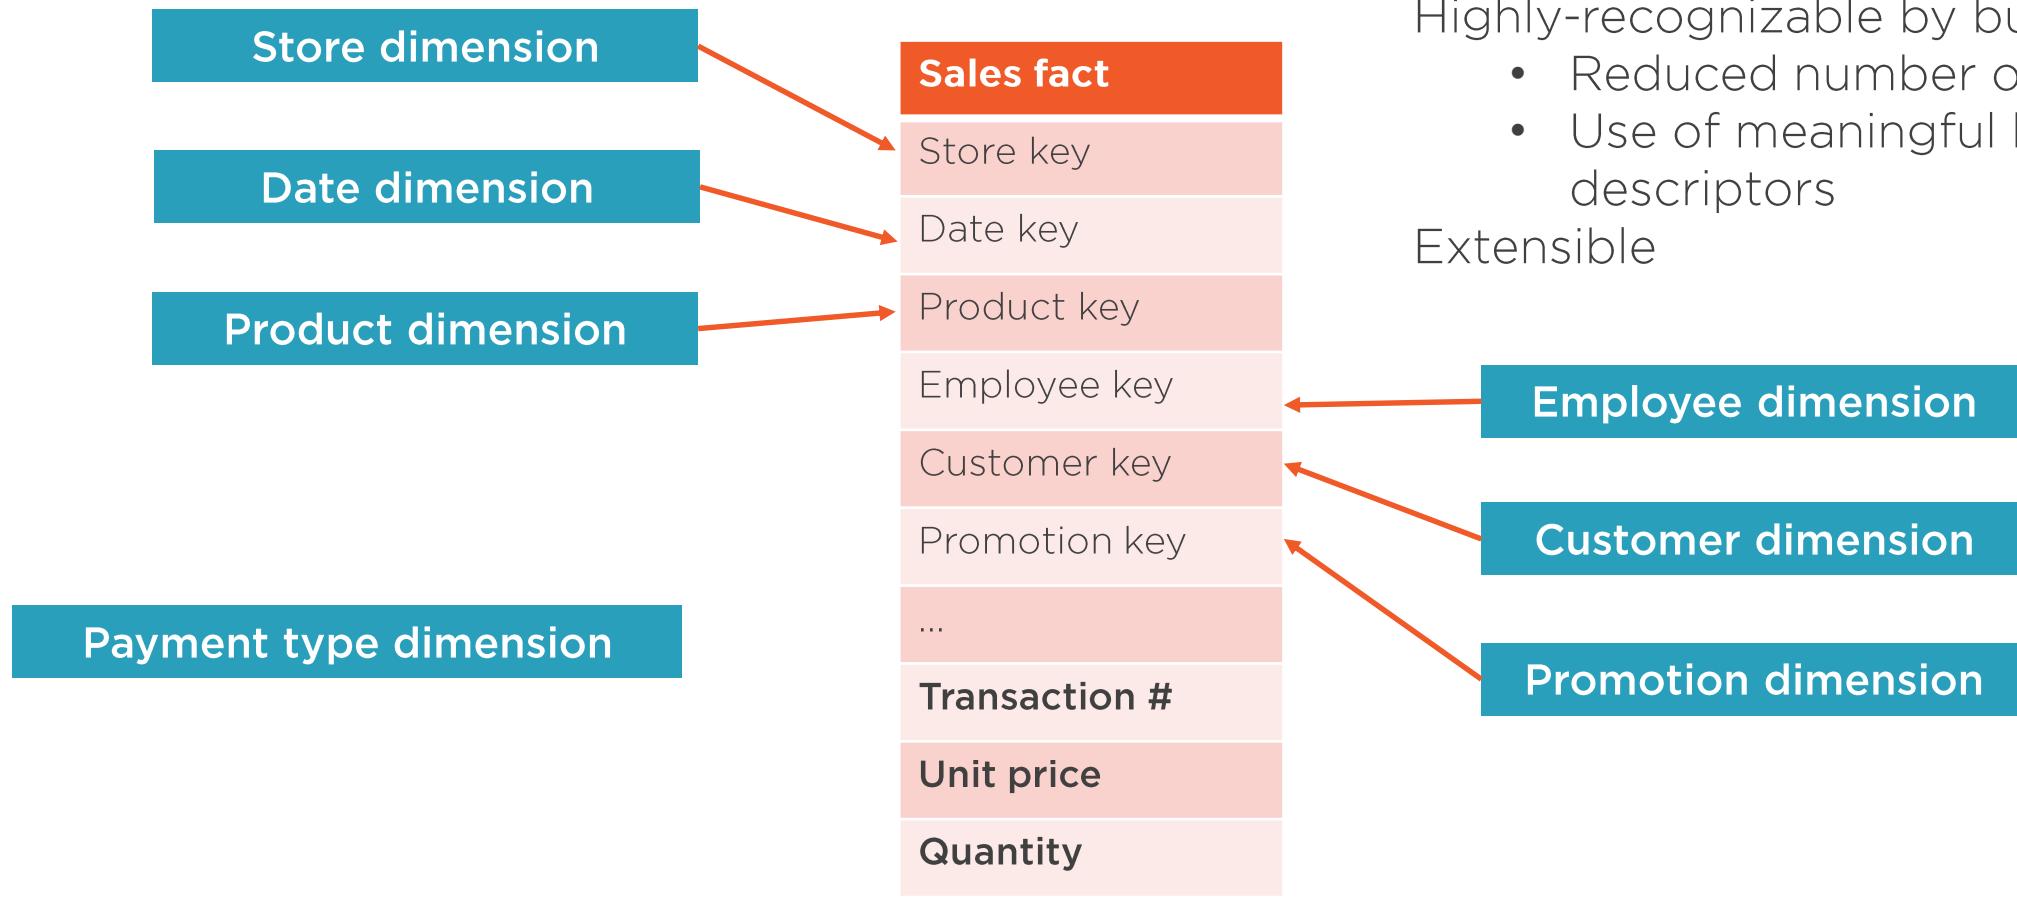
Employee dimension

Customer dimension

Promotion dimension



Dimensional Model



Characteristics:

Simple

Symmetric

Highly-recognizable by business users

- Reduced number of tables
- Use of meaningful business descriptors

Extensible

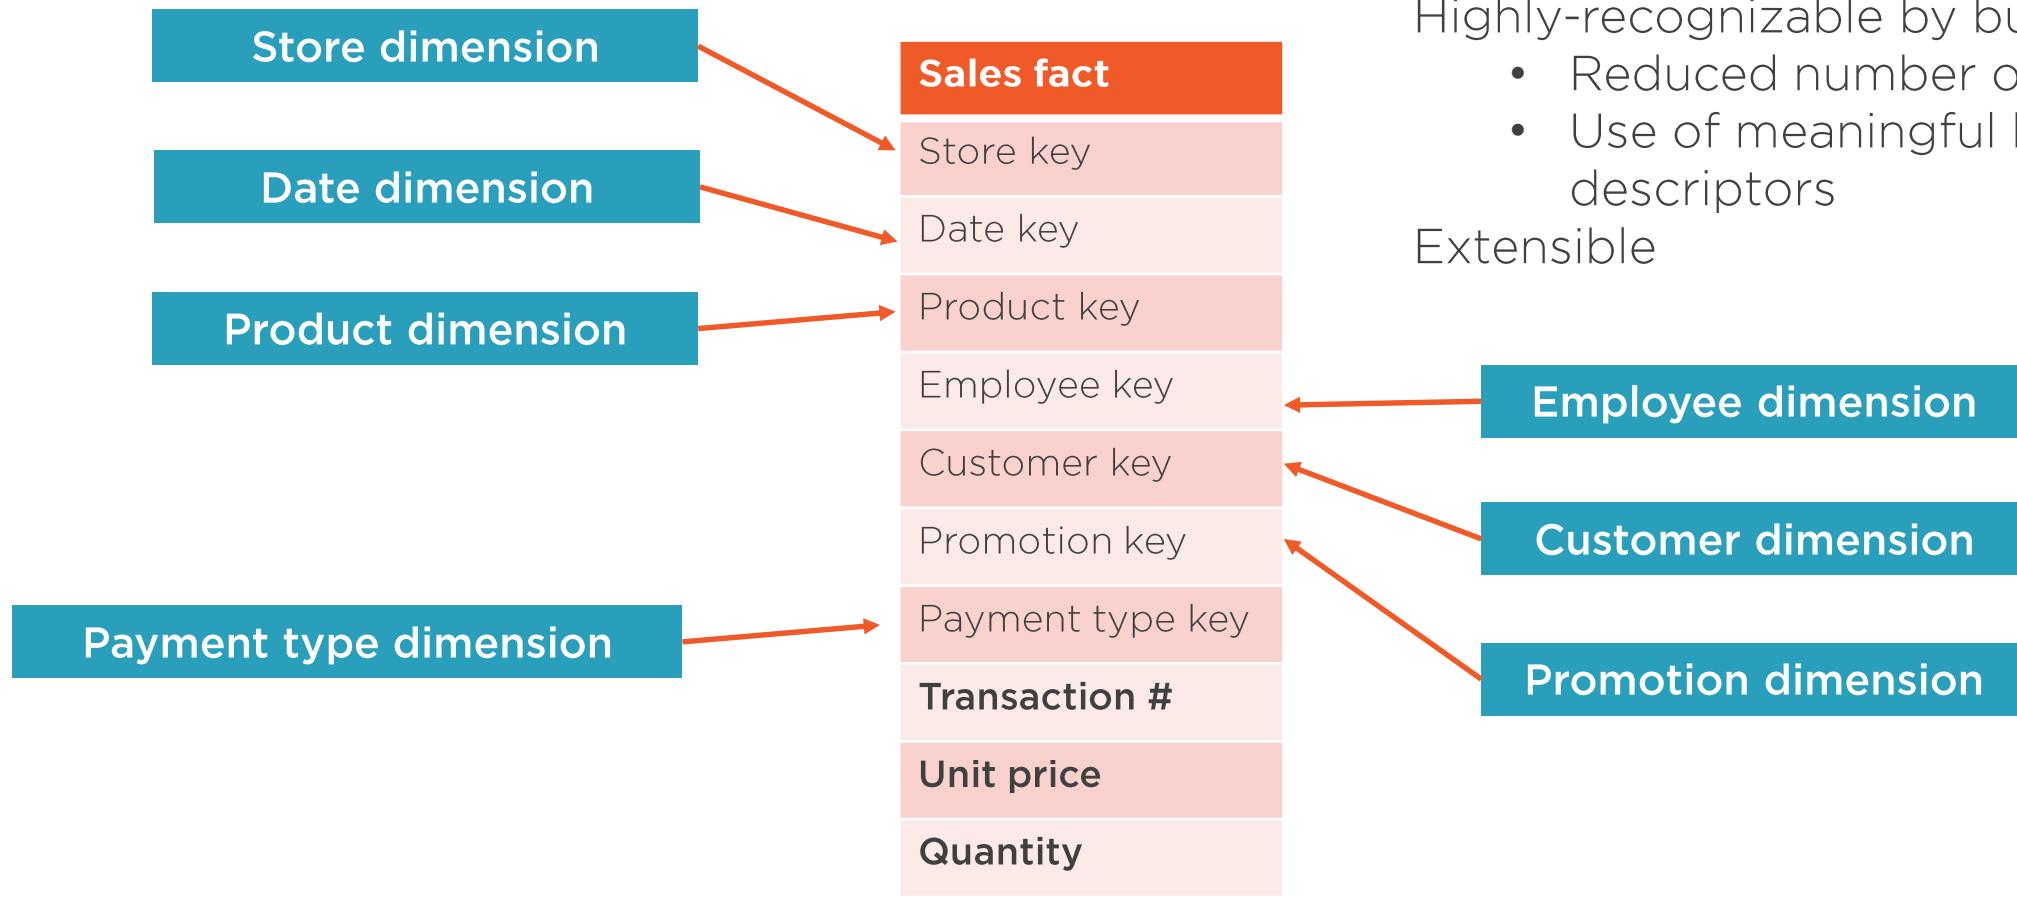
Employee dimension

Customer dimension

Promotion dimension



Dimensional Model



Characteristics:

Simple

Symmetric

Highly-recognizable by business users

- Reduced number of tables
- Use of meaningful business descriptors

Extensible

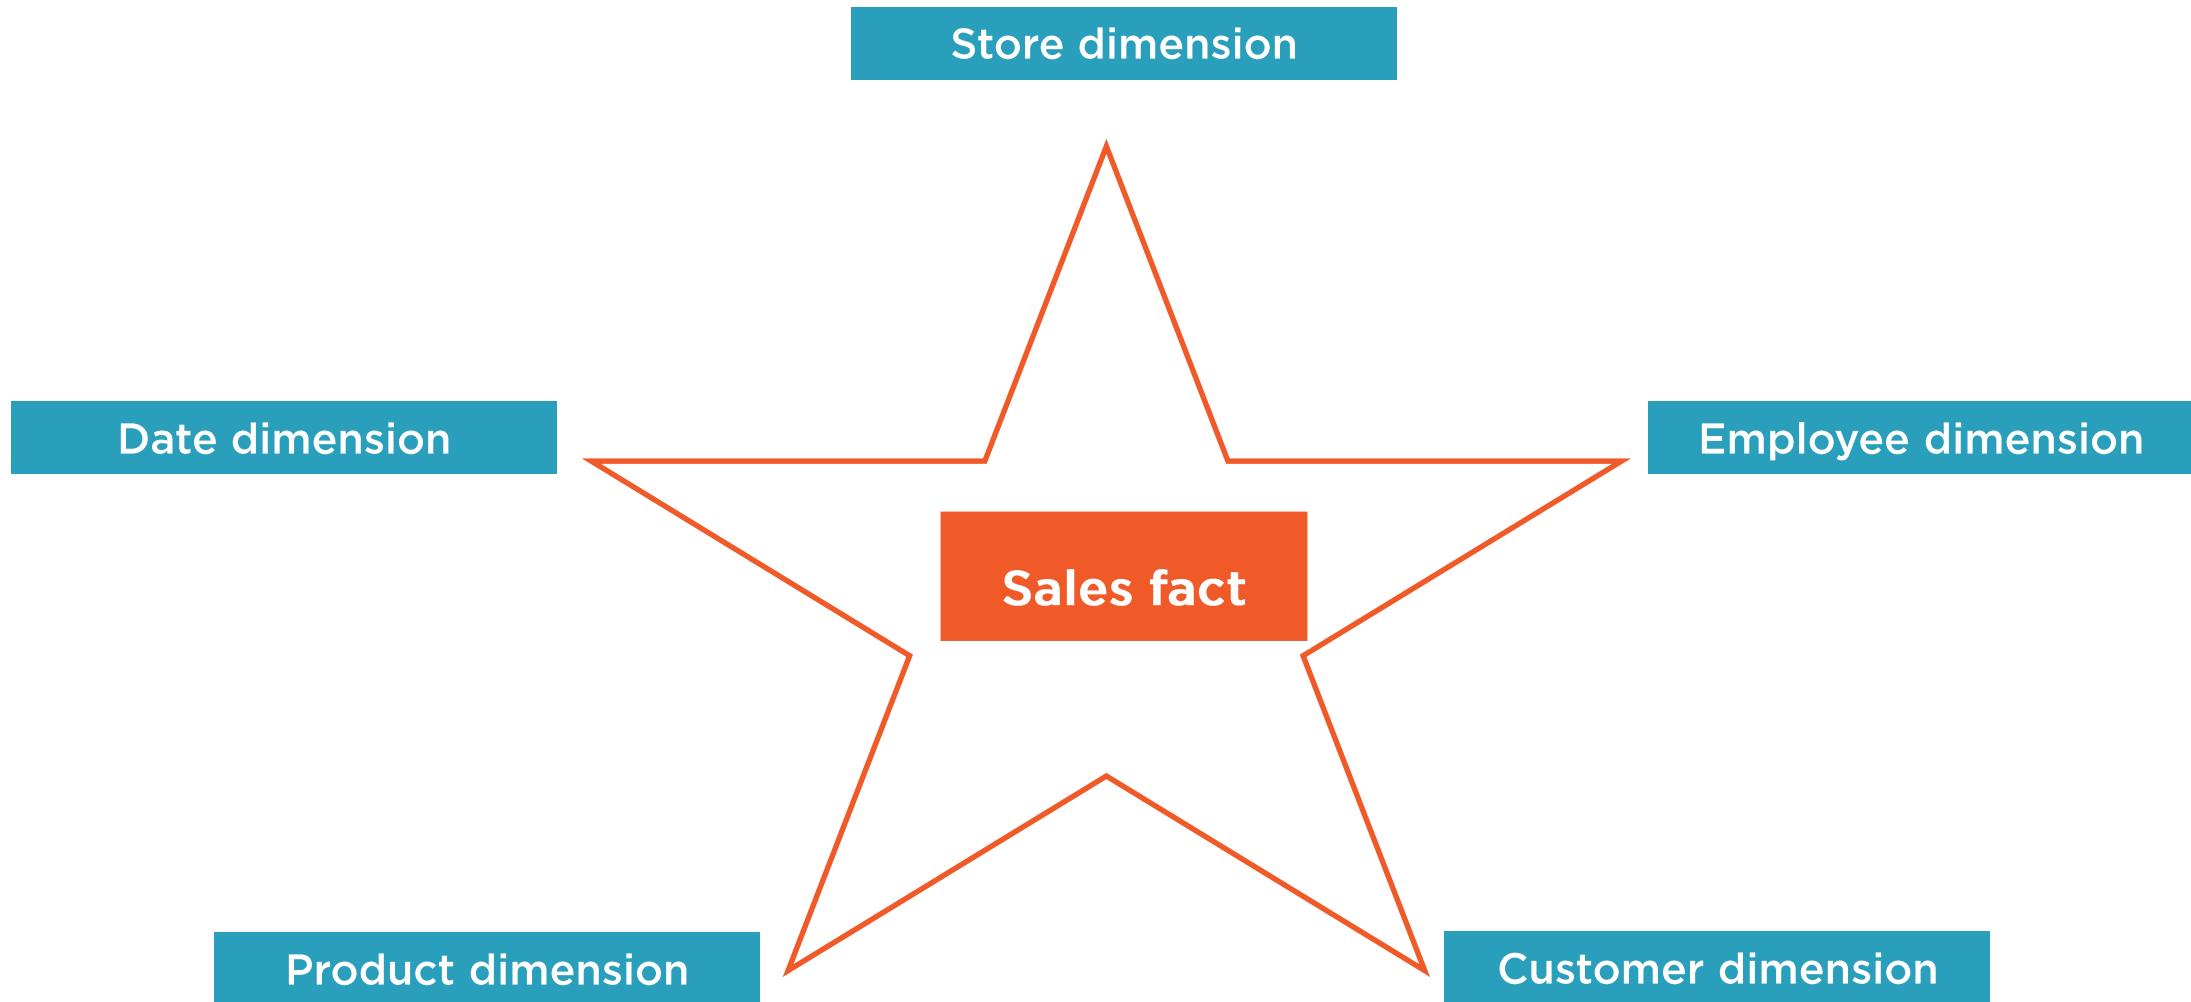
Employee dimension

Customer dimension

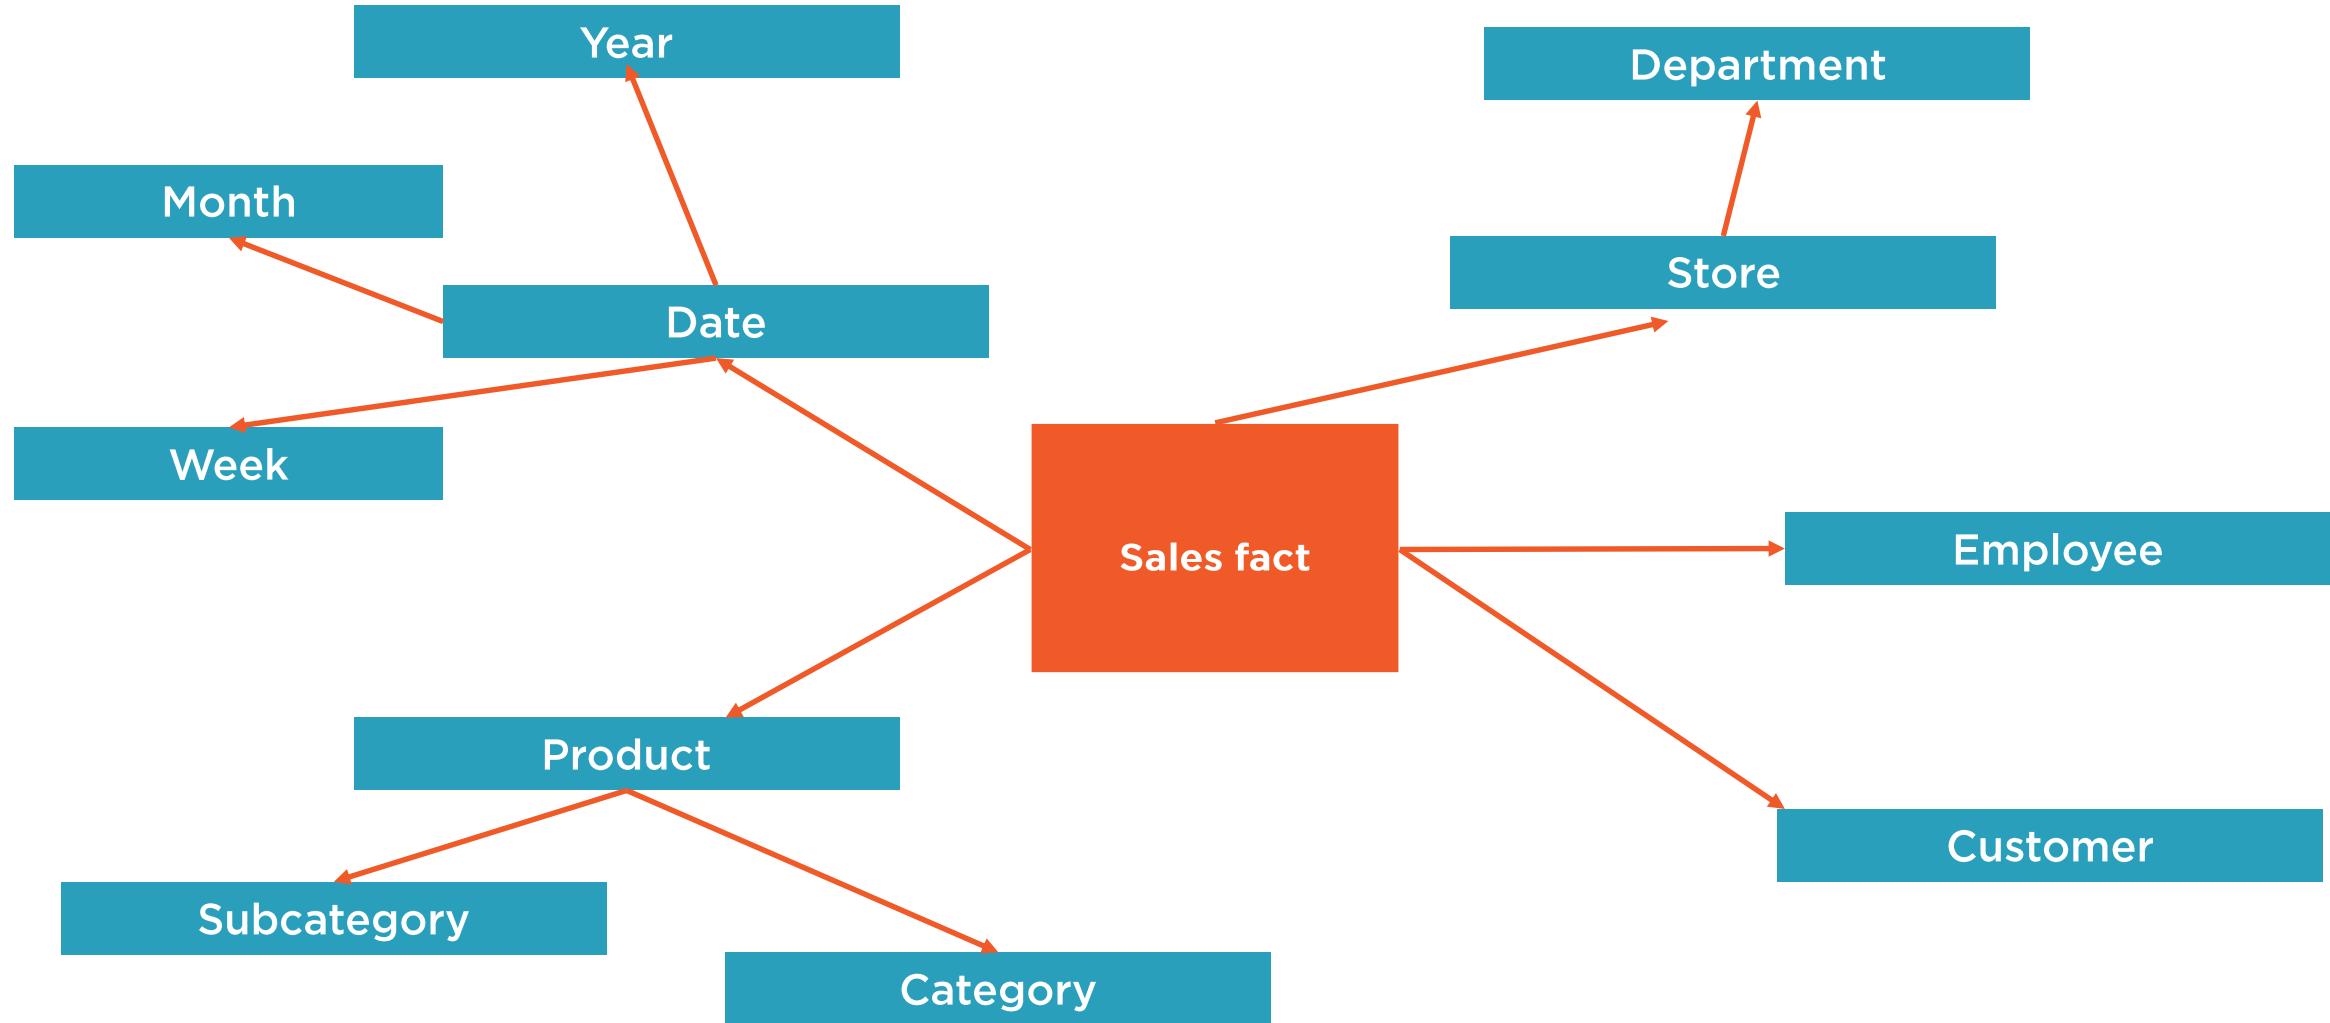
Promotion dimension



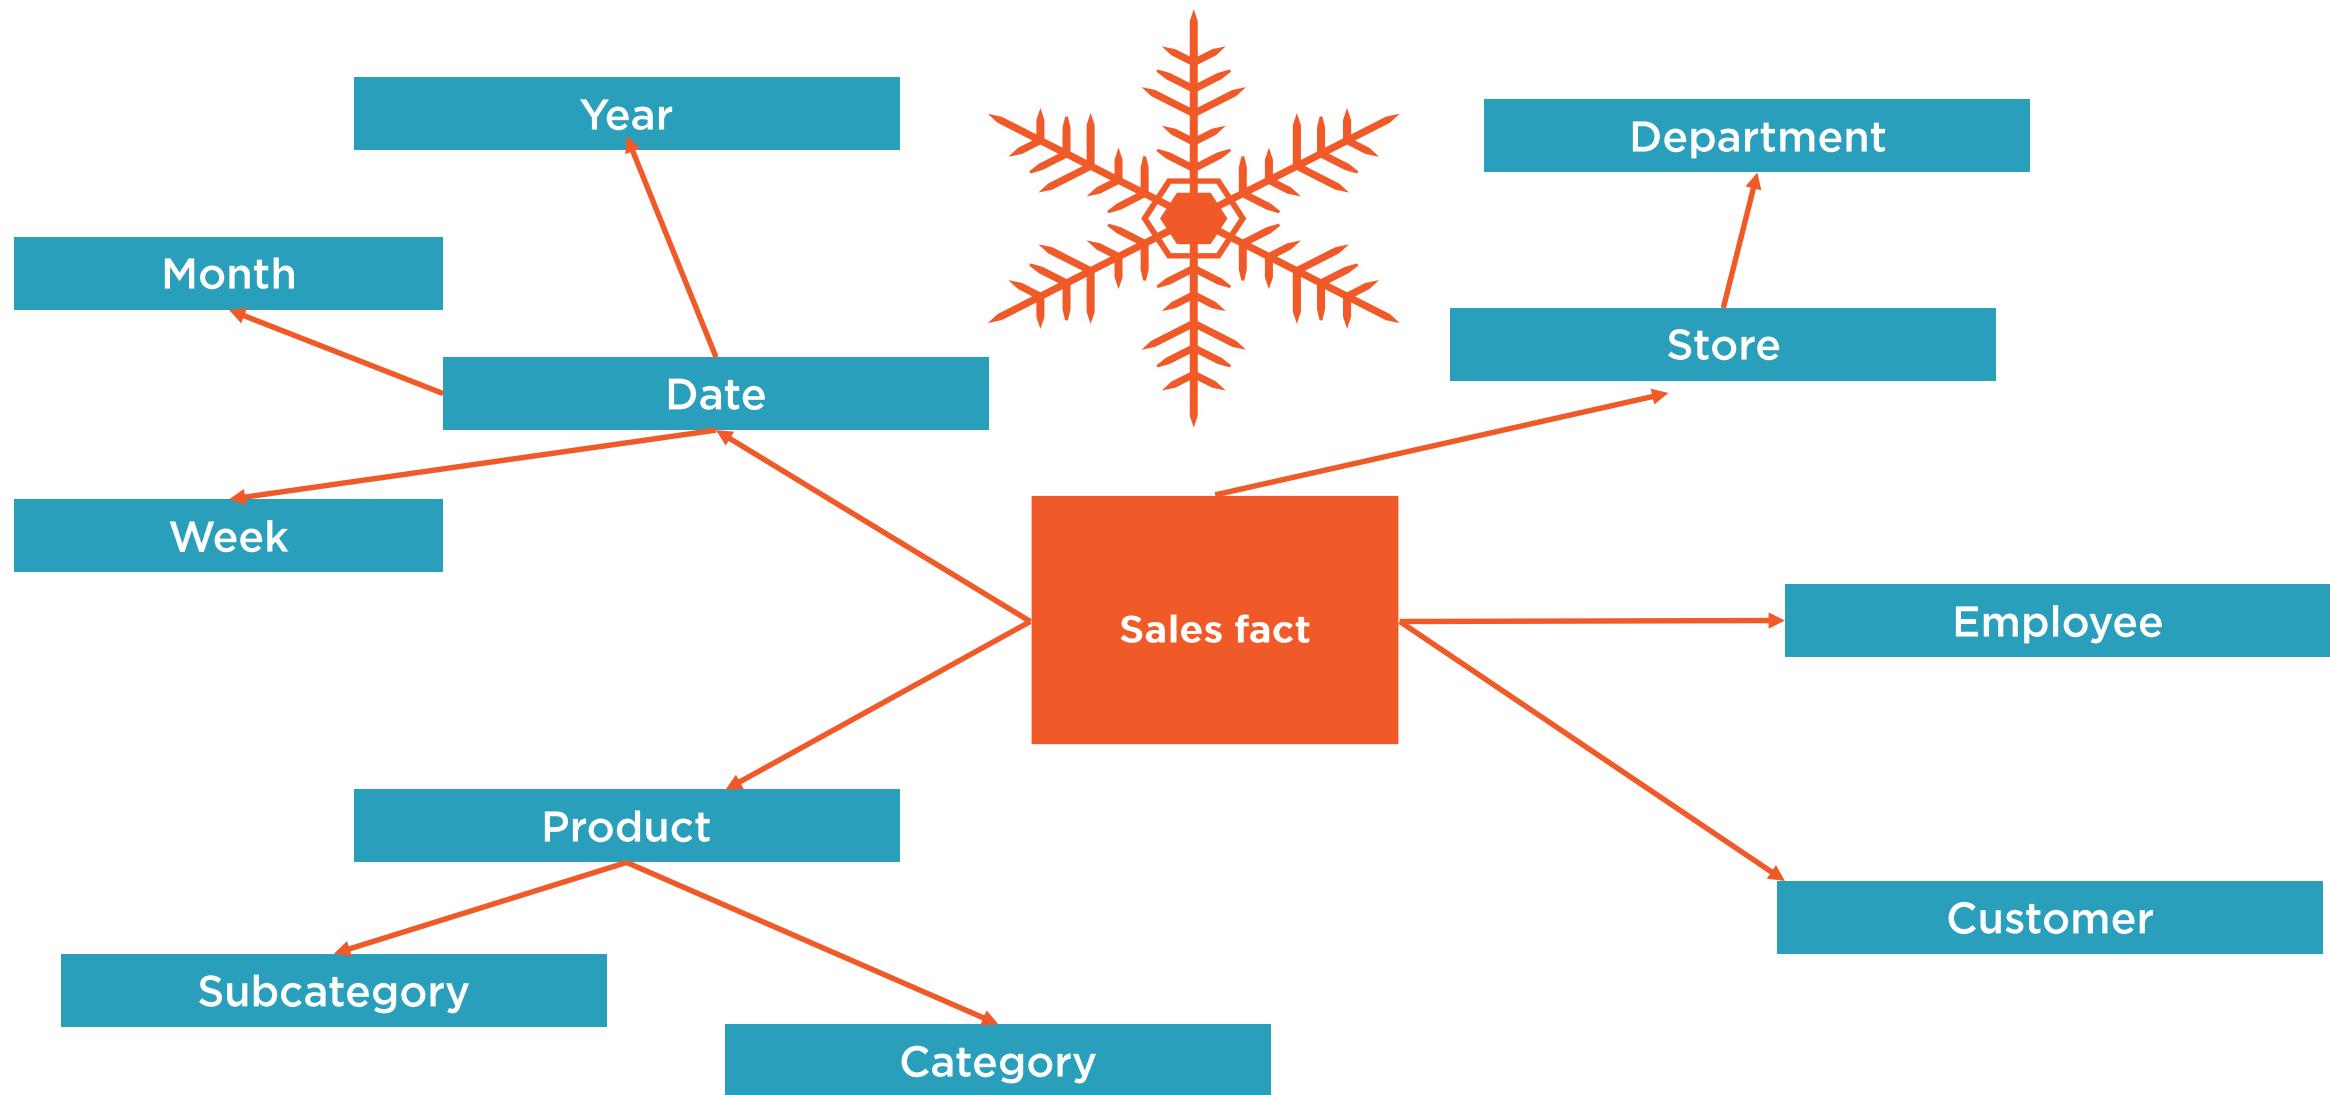
Dimensional Model as a Star Schema



Dimensional Model as a Snowflake



Dimensional Model as a Snowflake



The Four-step Dimensional Design Process



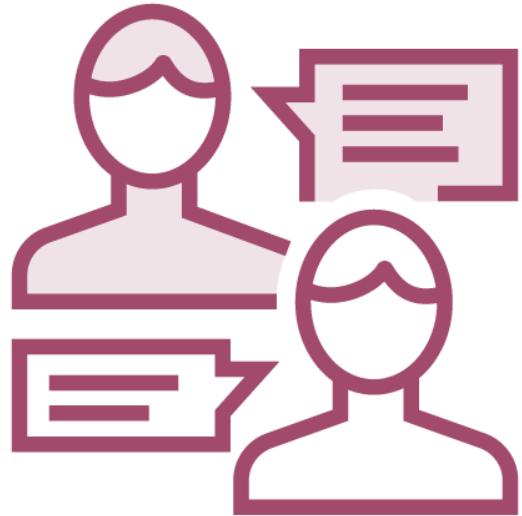
Dimensional Design



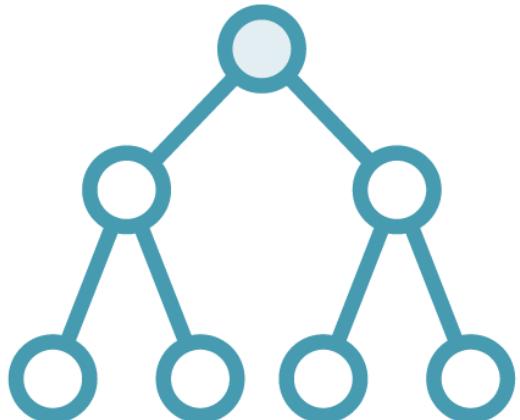
- Is done “with pen and paper”**
- Focuses on understanding the deliverables of the project**
- Consists of four steps**



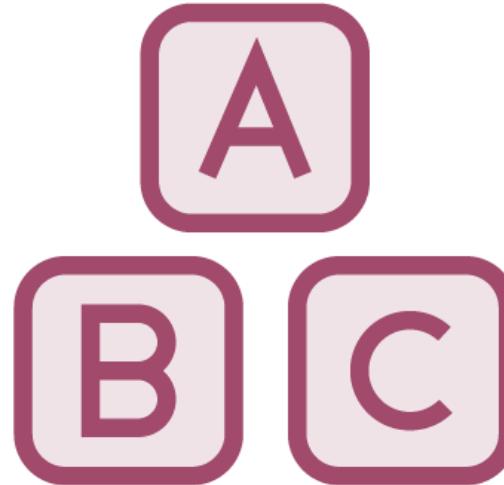
Steps of Dimensional Design



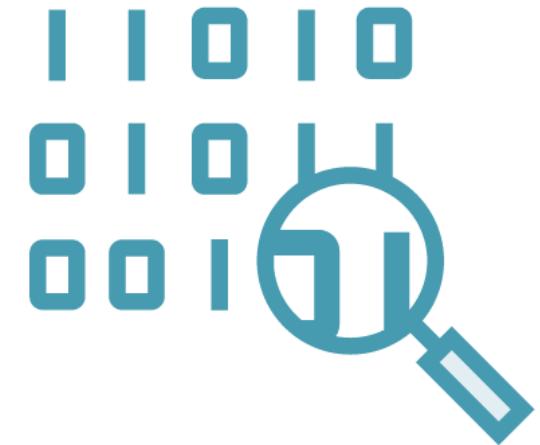
Select the
business
process



Declare the
grain



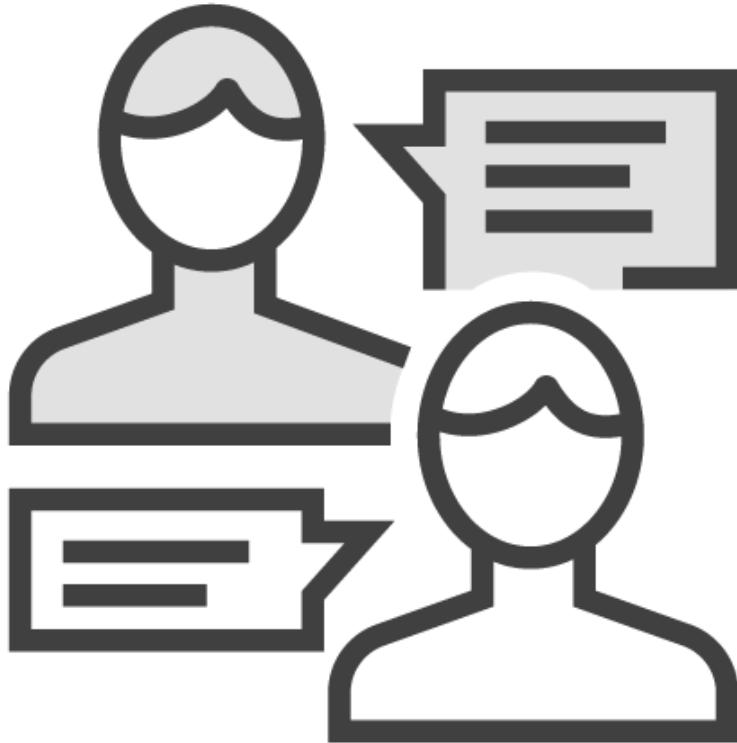
Identify the
dimensions



Identify the
facts



Step 1: Select the Business Process



Low-level activities performed by an organization

Are identified by listening carefully to the business users

Characteristics:

- Expressed as verbs
- Are supported by an operational system
- Generate KPIs



Step 2: Identify the Grain



Specify the detail level of a business process we want to measure

Example	Grain	Questions to ask
Sales/day	One row at the end of the day	How much was payed by each customer? Is this important information?
Sales/day/order	One row for every finished order	What was the most sold product? Is this important information?
Sales/day/order/product	One row for every product sold within an order	Is this enough information?



Step 2: Identify the Grain



Specify the detail level of a business process we want to measure

“How do you describe a single row in the fact table?”

Grain declarations are expressed in business terms



Steps 3 and 4

Identify the dimensions

“How do business users describe the data resulting from the process?”

“who, what, where, when, why, how”

Examples:

- Date
- Product
- Customer
- Employee

Identify the facts

“What is the process measuring?”

All candidate facts must be true to the grain for that fact table

Facts with different grains are split in separate tables

Examples:

- Sales price
- Sales quantity (or Units sold)



Both business requirements and
the realities of the source data
should be considered when
designing the dimensional model



Dimension Table Design Techniques



Ana Voicu
@ana_voicu



Dimension Table Techniques



What is specific to a dimension table?

Technical guidelines for designing a table

- Working with keys
 - Surrogate key
 - Business key
- Populating the dimension's attributes with data

Examples of common dimensions

- Date dimension
- Product dimension



What Is Specific to a Dimension Table?



Example: Sources for the Product Dimension

Administration system

Inventory changes

- Normalized model
- Products, departments, categories, subcategories

Software product for storing recipes

Used by the cook for daily recipes

POS used by staff

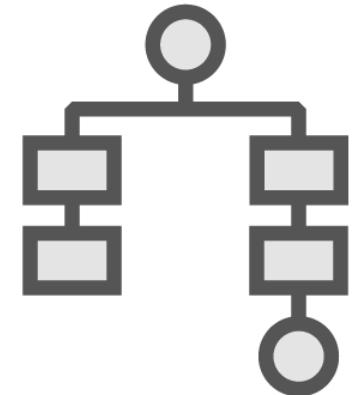
Take and fulfill orders

Flat file with information from an old accounting system

The system is not in use anymore

- But kept for historic information

Product dimension



Requirements of a Dimension Table



How will you build a table meeting these requirements?

- Dimension table can be linked to fact tables
- Keep track of the source system for data
- Data is consistent and readable
- Keep history of data changes from the operational system



What Is Specific to a Dimension Table?



A dimension table is different than a relational table

- The normalization rules don't apply
- Not optimized for data modification aspects
- Integrates in a consistent way multiple source data
- Data is easy to read by the user



Surrogate Keys



Relationship between Fact and Dimension

One-to-many relationship

Dim Product	Fact Sales
PK	PK
Attribute1	FK product
Attribute2	FK date
Attribute3	...
Attribute4	Fact1
Attribute5	Fact2
...	...

Relationship components

- 1:M example
 - 1 product (chocolate chip cookie)
 - 2 sales (transactions)
- Keys involved
 - Primary key of dimension table
 - Foreign key from fact table
- It's important to create it correctly
 - Plays an important role in the BI solution
 - Should be unique across all data sources for the dimension



Creating the Primary Key

Bad practice

Use the PK from the source systems

Can lead to overlapping rows

Source system 1

Id	Name	...
3	Chocolate chip cookie	...

Source system 2

Id	Name	...
3	Potato	...

Good practice

Generate an artificial key

- Simple integer column
- Automatically incremented
- Doesn't require careful maintenance

In DW terminology, this is the surrogate key



Synonyms for Surrogate Key



Meaningless key

Integer key

Non-natural key

Artificial key

Synthetic key



Advantages of Using a Surrogate Key



Advantages of Using a Surrogate Key



Integrate multiple source systems

Advantages of Using a Surrogate Key



Integrate multiple source systems

- Example: gathering data for the Product dimension in a restaurant or store



Integrate Data from Multiple Sources



Integrate Data from Multiple Sources

Source system 1

ID	Name	Description	UM	Color
123	Milk	Non-fat	L	White
776	Sugar	Brown sugar	KG	Brown



Integrate Data from Multiple Sources

Source system 1

ID	Name	Description	UM	Color
123	Milk	Non-fat	L	White
776	Sugar	Brown sugar	KG	Brown

Source system 2

ID	Name	Description	UM	Color
776	Butter	Non-fat	G	Yellow
778	Sugar	Normal sugar	KG	White



Integrate Data from Multiple Sources

Source system 1

ID	Name	Description	UM	Color
123	Milk	Non-fat	L	White
776	Sugar	Brown sugar	KG	Brown

Source system 2

ID	Name	Description	UM	Color
776	Butter	Non-fat	G	Yellow
778	Sugar	Normal sugar	KG	White



Integrate Data from Multiple Sources

Source system 1

ID	Name	Description	UM	Color
123	Milk	Non-fat	L	White
776	Sugar	Brown sugar	KG	Brown

Source system 2

ID	Name	Description	UM	Color
776	Butter	Non-fat	G	Yellow
778	Sugar	Normal sugar	KG	White

Data warehouse – Product dimension

Dim Key	Original Key	Name	Description	UM	Color
1	123	Milk	Non-fat	L	White
2	776	Sugar	Brown sugar	KG	Brown
3	776	Butter	Non-fat	G	Yellow
4	778	Sugar	Normal sugar	KG	White



Advantages of Using a Surrogate Key



Integrate multiple source systems

Keep track of attribute changes over time



Advantages of Using a Surrogate Key



Integrate multiple source systems
Keep track of attribute changes over time

- Use surrogate keys to handle multiple versions of a row



Keep Track of Attribute Changes over Time



Keep Track of Attribute Changes over Time

Before 09-16-2018

Key	Name	Description	UM	Color
387	Cherry toffee	Great cherry taste, sticky, chewy consistency	Piece	Pink



Keep Track of Attribute Changes over Time

Before 09-16-2018

Key	Name	Description	UM	Color
387	Cherry toffee	Great cherry taste, sticky, chewy consistency	Piece	Pink

After 09-16-2018

Key	Name	Description	UM	Color
387	Super cherry toffee	Great cherry taste, sticky, chewy consistency	Piece	Pink



Keep Track of Attribute Changes over Time

Before 09-16-2018

Key	Name	Description	UM	Color
387	Cherry toffee	Great cherry taste, sticky, chewy consistency	Piece	Pink



After 09-16-2018

Key	Name	Description	UM	Color
387	Super cherry toffee	Great cherry taste, sticky, chewy consistency	Piece	Pink



Keep Track of Attribute Changes over Time

Before 09-16-2018

Key	Name	Description	UM	Color
387	Cherry toffee	Great cherry taste, sticky, chewy consistency	Piece	Pink

After 09-16-2018

Key	Name	Description	UM	Color
387	Super cherry toffee	Great cherry taste, sticky, chewy consistency	Piece	Pink

Data warehouse - Product dimension

Dim Key	Orig. Key	Name	Description	UM	Color
112	387	Cherry toffee	Great cherry taste, sticky, chewy consistency	Piece	Pink
244	387	Super cherry toffee	Great cherry taste, sticky, chewy consistency	Piece	Pink



Keep Track of Attribute Changes over Time

Before 09-16-2018

Key	Name	Description	UM	Color
387	Cherry toffee	Great cherry taste, sticky, chewy consistency	Piece	Pink

After 09-16-2018

Key	Name	Description	UM	Color
387	Super cherry toffee	Great cherry taste, sticky, chewy consistency	Piece	Pink

Data warehouse - Product dimension

Dim Key	Orig. Key	Name	Description	UM	Color	Valid From	Valid To
112	387	Cherry toffee	Great cherry taste, sticky, chewy consistency	Piece	Pink	01-01-2017	09-16-2018
244	387	Super cherry toffee	Great cherry taste, sticky, chewy consistency	Piece	Pink	09-17-2018	12-31-9999



Advantages of Using a Surrogate Key



- Integrate multiple source systems
- Keep track of attribute changes over time
- Protect the data warehouse from operational changes



Protect the Data Warehouse from Operational Changes



Characteristics of an operational system:

- Logs and processes daily activities/current actions
- Some data history is preserved
- Obsolete data is deleted after a period of time
- Codes/IDs can be reassigned to new data



Protect the Data Warehouse from Operational Changes



Protect the Data Warehouse from Operational Changes

Before 01-01-2018

Key	Name	Description	UM	Color
432	Chocolate chip cookie	Grandma's cookie recipe	Piece	Brown
776	Sugar	Brown sugar	KG	Brown



Protect the Data Warehouse from Operational Changes

Before 01-01-2018

Key	Name	Description	UM	Color
432	Chocolate chip cookie	Grandma's cookie recipe	Piece	Brown
776	Sugar	Brown sugar	KG	Brown

After 01-01-2018

Key	Name	Description	UM	Color
776	Sugar	Brown sugar	KG	Brown



Protect the Data Warehouse from Operational Changes

Before 01-01-2018

Key	Name	Description	UM	Color
432	Chocolate chip cookie	Grandma's cookie recipe	Piece	Brown
776	Sugar	Brown sugar	KG	Brown



After 01-01-2018

Key	Name	Description	UM	Color
776	Sugar	Brown sugar	KG	Brown



Protect the Data Warehouse from Operational Changes

Before 01-01-2018

Key	Name	Description	UM	Color
432	Chocolate chip cookie	Grandma's cookie recipe	Piece	Brown
776	Sugar	Brown sugar	KG	Brown

After 01-01-2018

Key	Name	Description	UM	Color
776	Sugar	Brown sugar	KG	Brown

Data warehouse – Product dimension

Dim Key	Orig. Key	Name	Description	UM	Color	Valid From	Valid To
83	432	Chocolate chip cookie	Grandma's cookie recipe	Piece	Brown	01-01-1753	01-01-2018
2	776	Sugar	Normal sugar	KG	Brown	01-01-1753	12-31-9999



Advantages of Using a Surrogate Key



- Integrate multiple source systems**
- Keep track of attribute changes over time**
- Protect the data warehouse from operational changes**
- Handle null or unknown conditions**



Advantages of Using a Surrogate Key



Integrate multiple source systems

Keep track of attribute changes over time

Protect the data warehouse from operational changes

Handle null or unknown conditions

- Empty row technique



Advantages of Using a Surrogate Key



Integrate multiple source systems

Keep track of attribute changes over time

Protect the data warehouse from operational changes

Handle null or unknown conditions

- Empty row technique
- Handle situations when a link does not exist between fact and dimensions



Advantages of Using a Surrogate Key



Integrate multiple source systems

Keep track of attribute changes over time

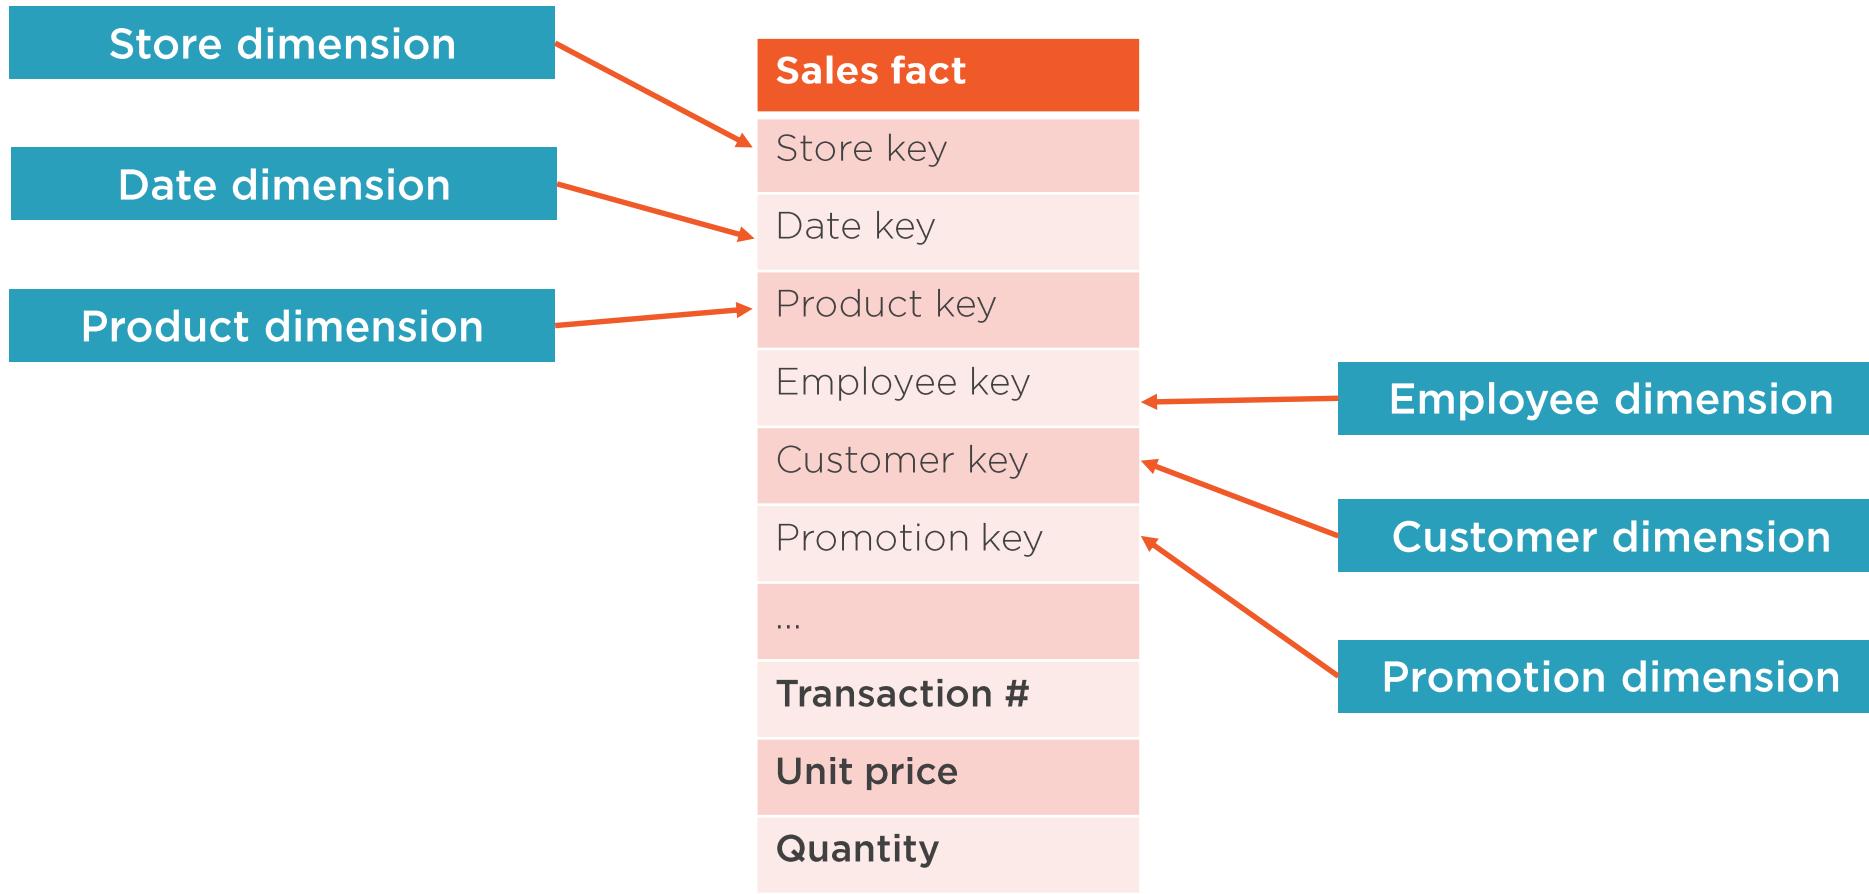
Protect the data warehouse from operational changes

Handle null or unknown conditions

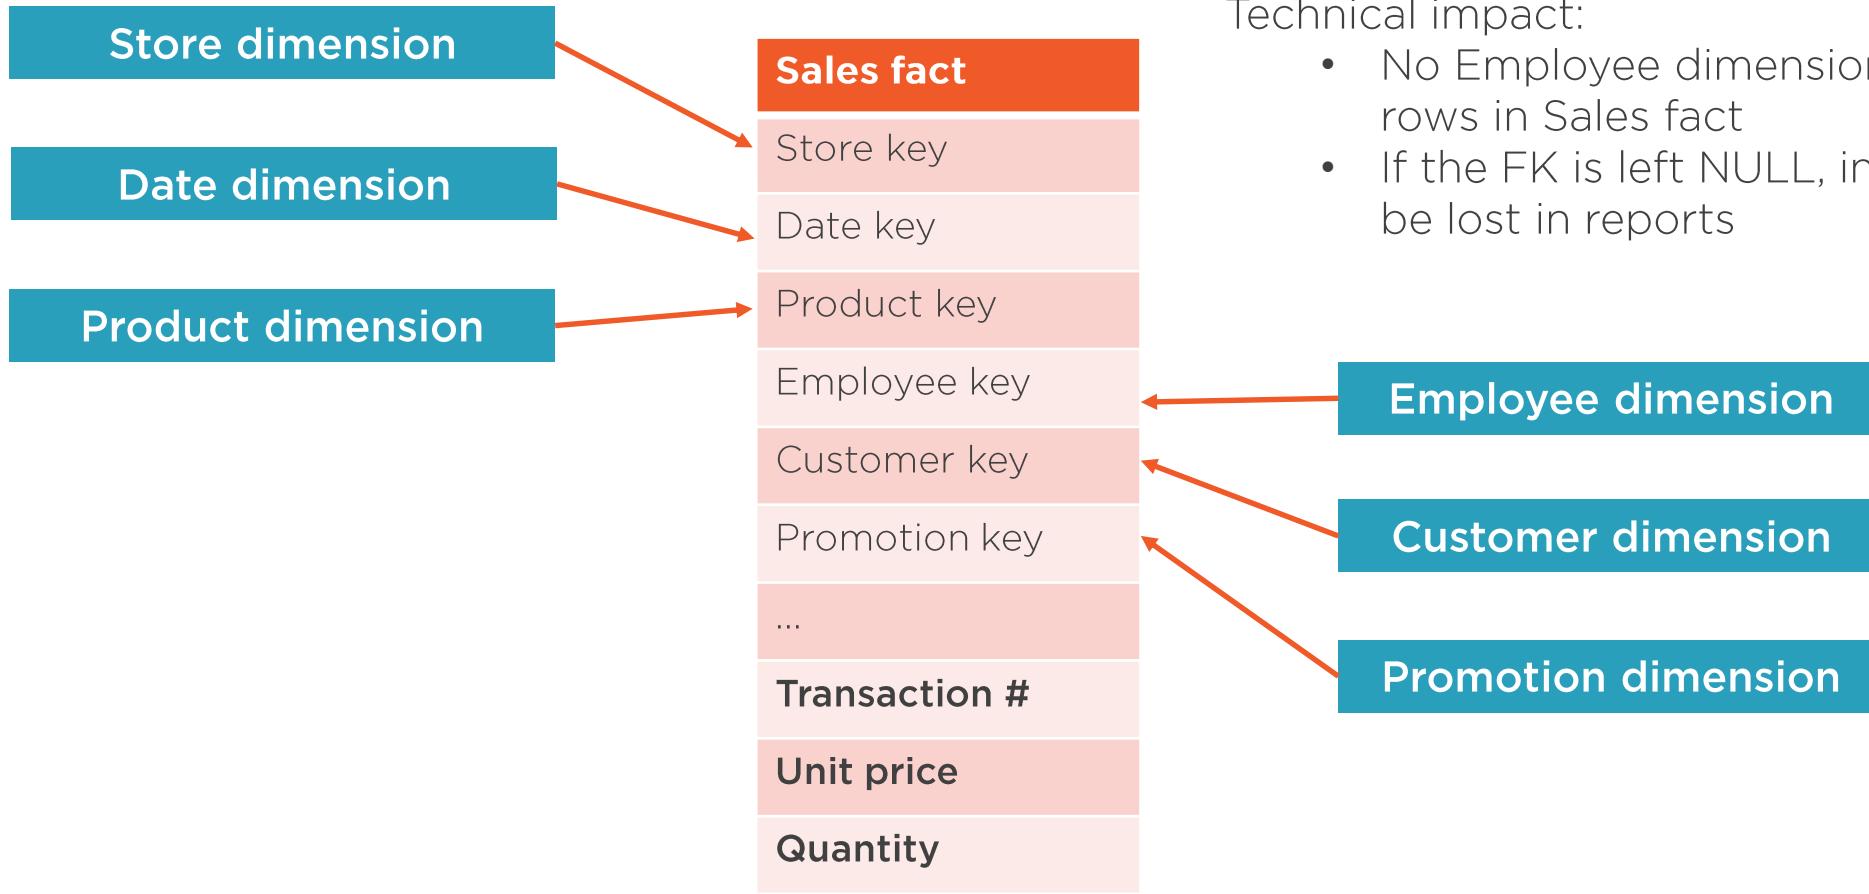
- Empty row technique
- Handle situations when a link does not exist between fact and dimensions



Scenario for Null or Unknown Conditions



Scenario for Null or Unknown Conditions



Self-checkout kiosks were introduced
Customers pay by themselves
Employees are not involved for this transaction
Technical impact:

- No Employee dimension key for these rows in Sales fact
- If the FK is left NULL, information may be lost in reports



Handle Null or Unknown Conditions

Adding the “empty row” to each dimension table

Key	Name	Birthday	Gender	PhoneNr.
1	Unknown	01-01-1753	Unknown	Unknown
32	Alice Cohen	03-05-1989	Female	Unknown



Handle Null or Unknown Conditions

Adding the “empty row” to each dimension table

Key	Name	Birthday	Gender	PhoneNr.
1	Unknown	01-01-1753	Unknown	Unknown
32	Alice Cohen	03-05-1989	Female	Unknown

Data warehouse – Sales fact table

Fact Key	Product Key	Customer Key	Employee Key	Transaction Nr	Amt.
...	432	10	32	#1050	\$5
...	776	12	1	#2367	\$3



Advantages of Using a Surrogate Key



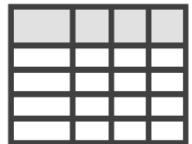
- Integrate multiple source systems**
- Keep track of attribute changes over time**
- Protect the data warehouse from operational changes**
- Handle null or unknown conditions**
- Improve performance**



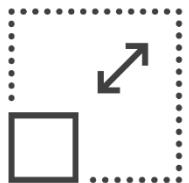
Improve Performance



The surrogate key should be a small number, preferably an integer



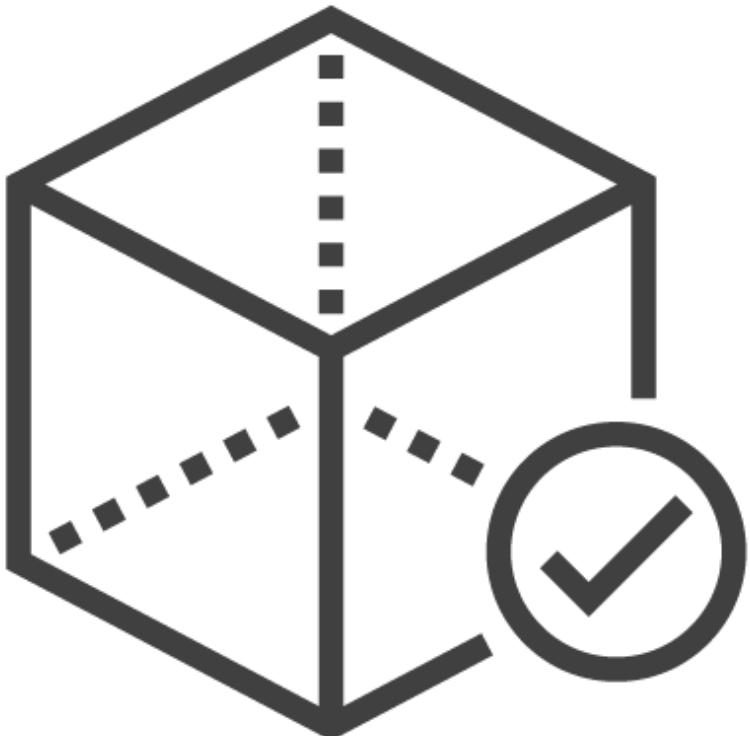
You need to anticipate the growth of your dimension table (a 4-byte int can handle ~2 billion positive values and ~4 billion values in total)



The smaller the surrogate key, the smaller the fact table



Dimensions of a Fact Table



Fact tables

- Have FKs to a lot of dimension tables
- Grow very fast
- Every byte counts!

FK length	Fact table growth	Fact table storage
X bytes	1 billion rows	Y GB
X + 1 bytes	1 billion rows	Y + 1 GB



Advantages of Using a Surrogate Key



- Integrate multiple source systems**
- Keep track of attribute changes over time**
- Protect the data warehouse from operational changes**
- Handle null or unknown conditions**
- Improve performance**



Business Keys



Business Keys



The primary key of the source system

Also known as

- Natural keys
- Production keys
- Operational keys

Business Keys Properties

**Attribute in the dimension
table**



Business Keys Properties

**Attribute in the dimension
table**

**Can be prefixed with a source
system's code
(POS|432 or CSV|635)**



Business Keys Properties

Attribute in the dimension table

**Can be prefixed with a source system's code
(POS|432 or CSV|635)**

Can cause duplicated items in the dimensions



Business Keys Properties

Attribute in the dimension table

**Can be prefixed with a source system's code
(POS|432 or CSV|635)**

Can cause duplicated items in the dimensions

If the BK is composed of meaningful codes, it should be split and each code stored in separate columns



Keys in a Dimension Table



Surrogate key

- Primary key, no business value

Business key

- Original primary from the source system



Other Dimensional Design Considerations



Characteristics of a Dimension Table



Companion of a fact table

Used to describe the business in clear terms

Design implications

- Store descriptive words instead of codes
- Replace flags or indicators with descriptions
- Replace null values with meaningful words



Use Descriptions Instead of Codes



Each code stored should be accompanied by a descriptive decode

Each attribute should be easy to interpret by people



Example of Using Descriptions Instead of Codes

Country of import for the most sold products

# Sold products	Name	Imported from
5000	Zywiec Beer	PL
4856	Pan Beer	HR
3991	Ursus Beer	RO
2674	Mort Subite Beer	BE

Codes are not easy to interpret by everyone

Id	Name	Imported from
5000	Zywiec Beer	Poland
4856	Pan Beer	Croatia
3991	Ursus Beer	Romania
2674	Mort Subite Beer	Belgium

Decodes can be retrieved from the operational system as well



Use Textual Attributes Instead of Flags and Indicators



Flags and operational indicators should be supplemented with meaningful text words

Examples from the Date dimension:

Column	Meaningful name	Indicator
Weekday indicator	Weekday/Weekend	1/0
Holiday indicator	Holiday/Non-holiday	true/false



Replace Null Values in Dimensions



Replace null values with meaningful words, like

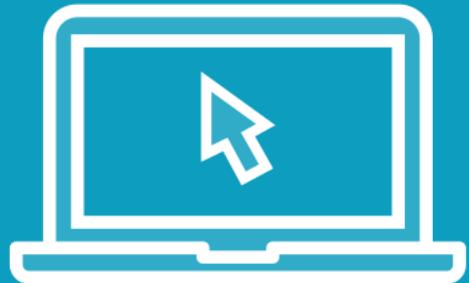
- Unknown
- Not applicable

Null values in data warehouse tables have unexpected behaviors

Nulls are not present in drop-down menus, so data may be missing from reports



Demo



Creating the Date dimension

- What is a Date dimension?
- Why is it important?



What Is a Date Dimension?



A table that stores dates

Is populated with a large number of dates
(10 years, 100 years)

As a dimension table, it is not large

It is loaded once, not periodically



Why Is It Important?



Is present in almost all data warehouses

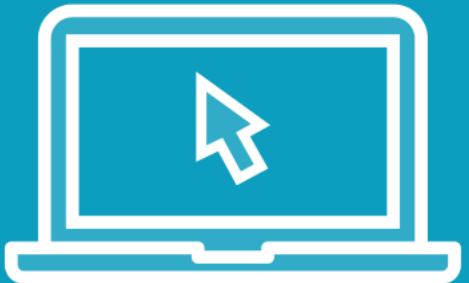
Helps performing time analysis

- Products sold this year vs. last year
- Products sold per month compared to last year
- Per week, per weekend, per holiday..

Contains pre-calculated information about dates



Demo

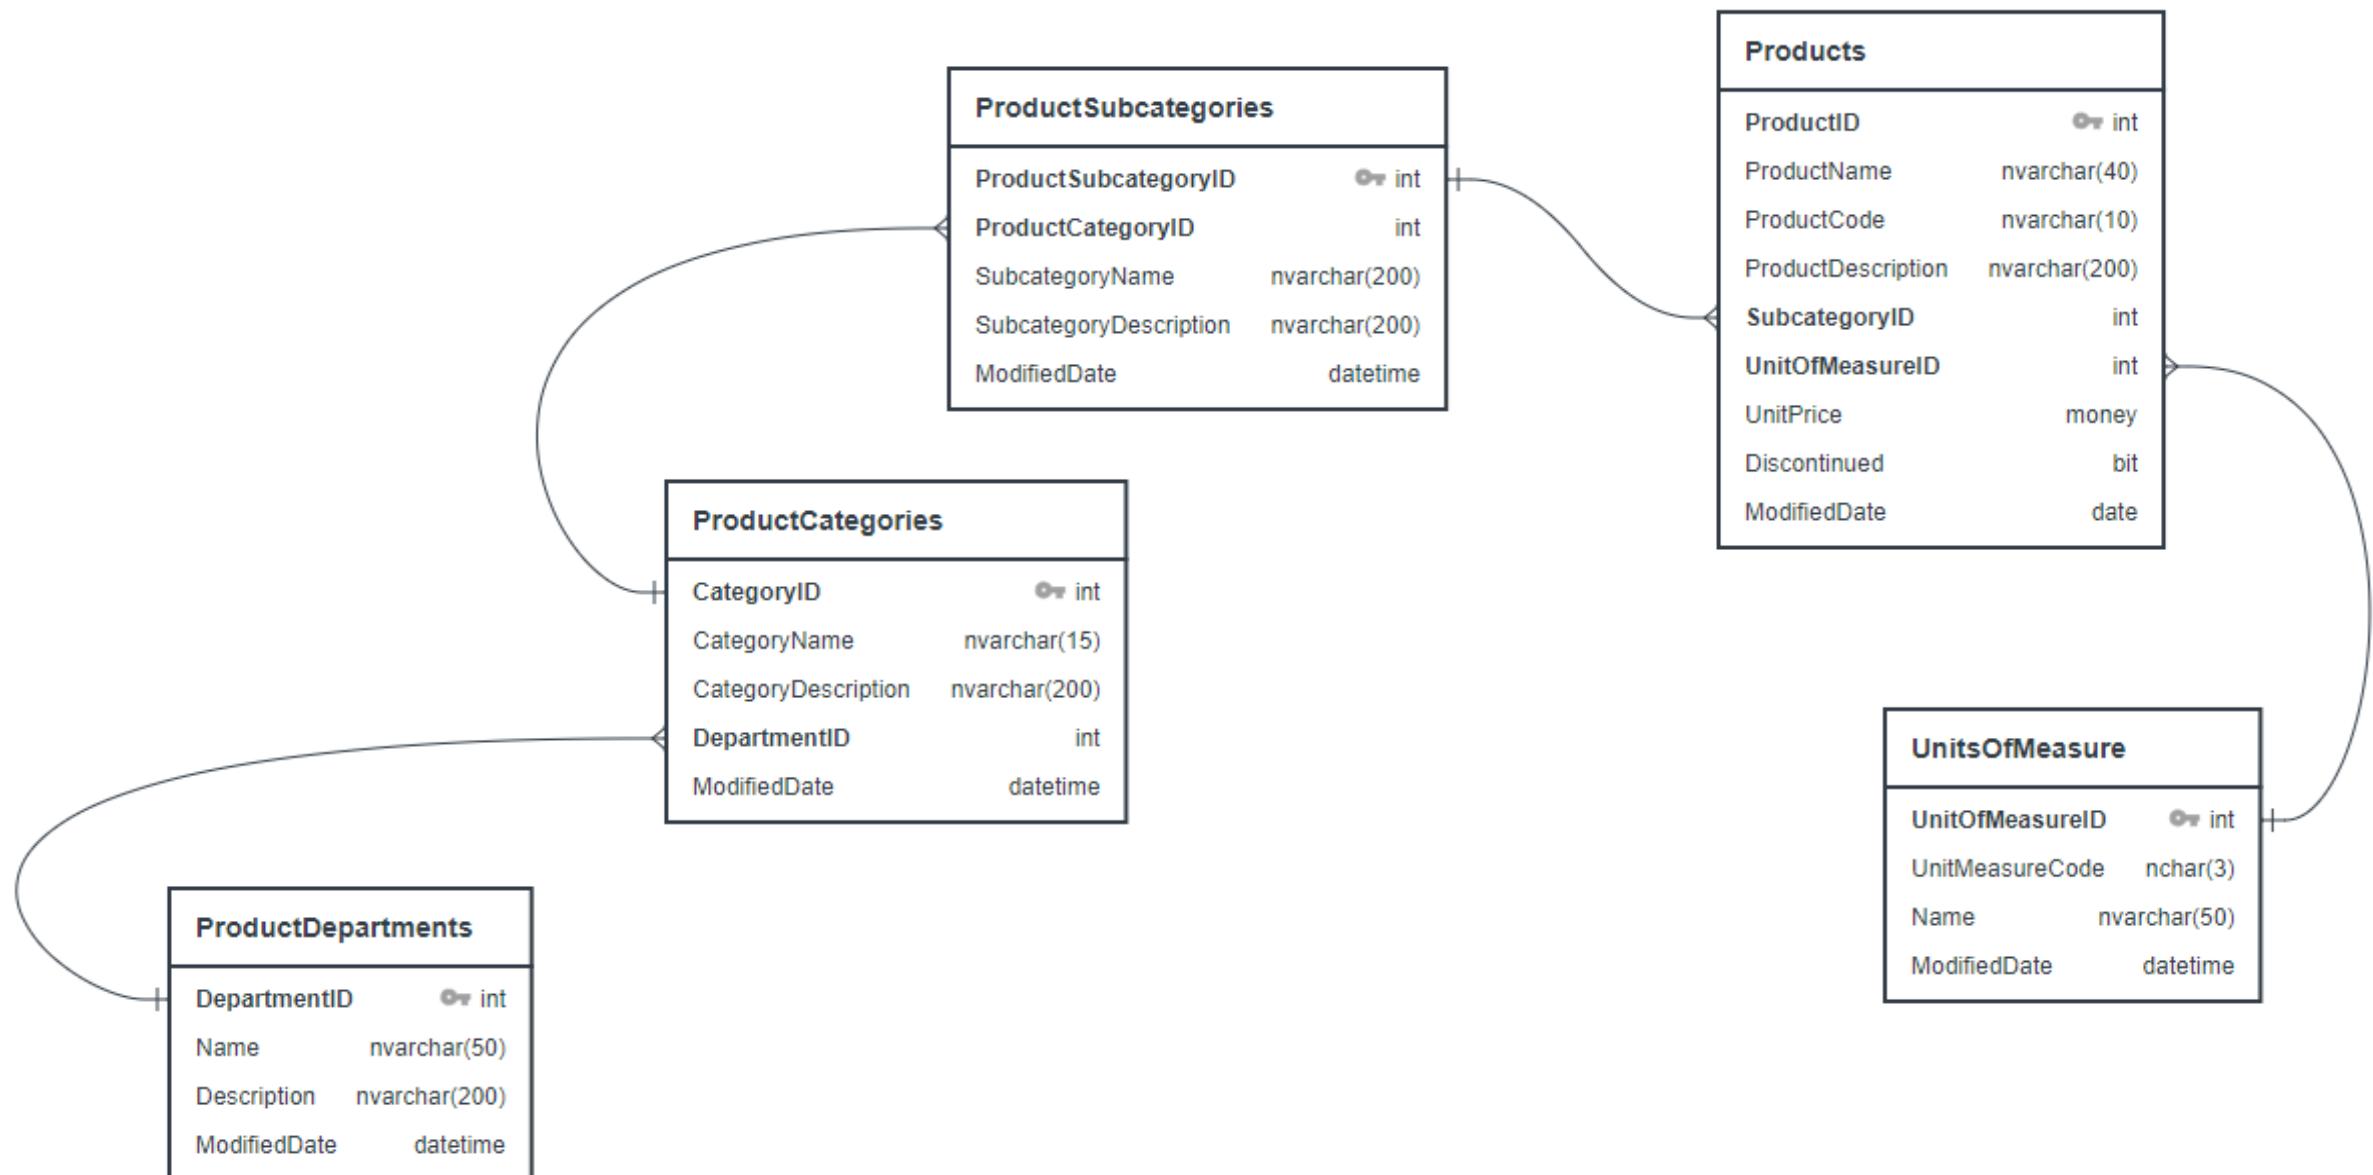


Creating the Product dimension

- Product dimension is present in many data warehouse projects
- It doesn't have a standard set of attributes



Product-related Tables



Summary



Characteristics of a dimension table

- Is different than a relational table
- Integrates data from multiple sources
- Data is easy to read by the user

Keys in a dimension table

- Surrogate key
- Business key

Guidelines for implementing a dimension table

- Store descriptive words instead of codes
- Replace flags or indicators with descriptions
- Replace null values with meaningful words



Creating and Working with Hierarchies



Ana Voicu
@ana_voicu



Overview



In short, a hierarchy is:

- A data structure
- Created with attributes from a dimension
- Used for data aggregation

Topics elaborated in this chapter:

- What is a hierarchy?
- Why is it useful?
- What does drilling-down mean?
- Implementing a hierarchy in a data warehouse



What Is a Hierarchy?

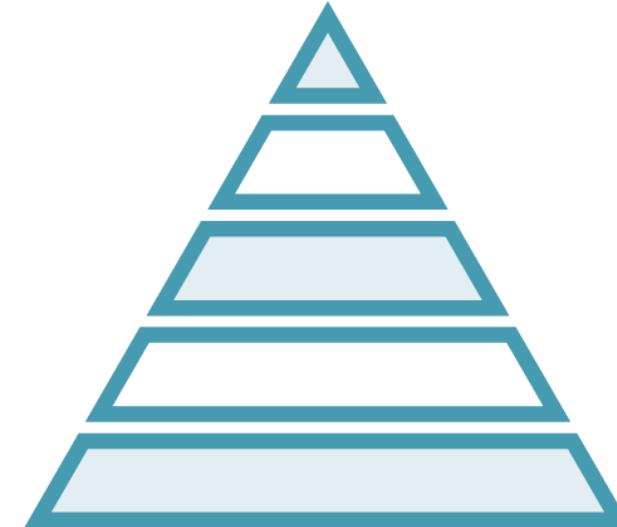


What Is a Hierarchy?



Data structure with multiple levels

- The levels are dimension attributes
- The elements on each level are called nodes

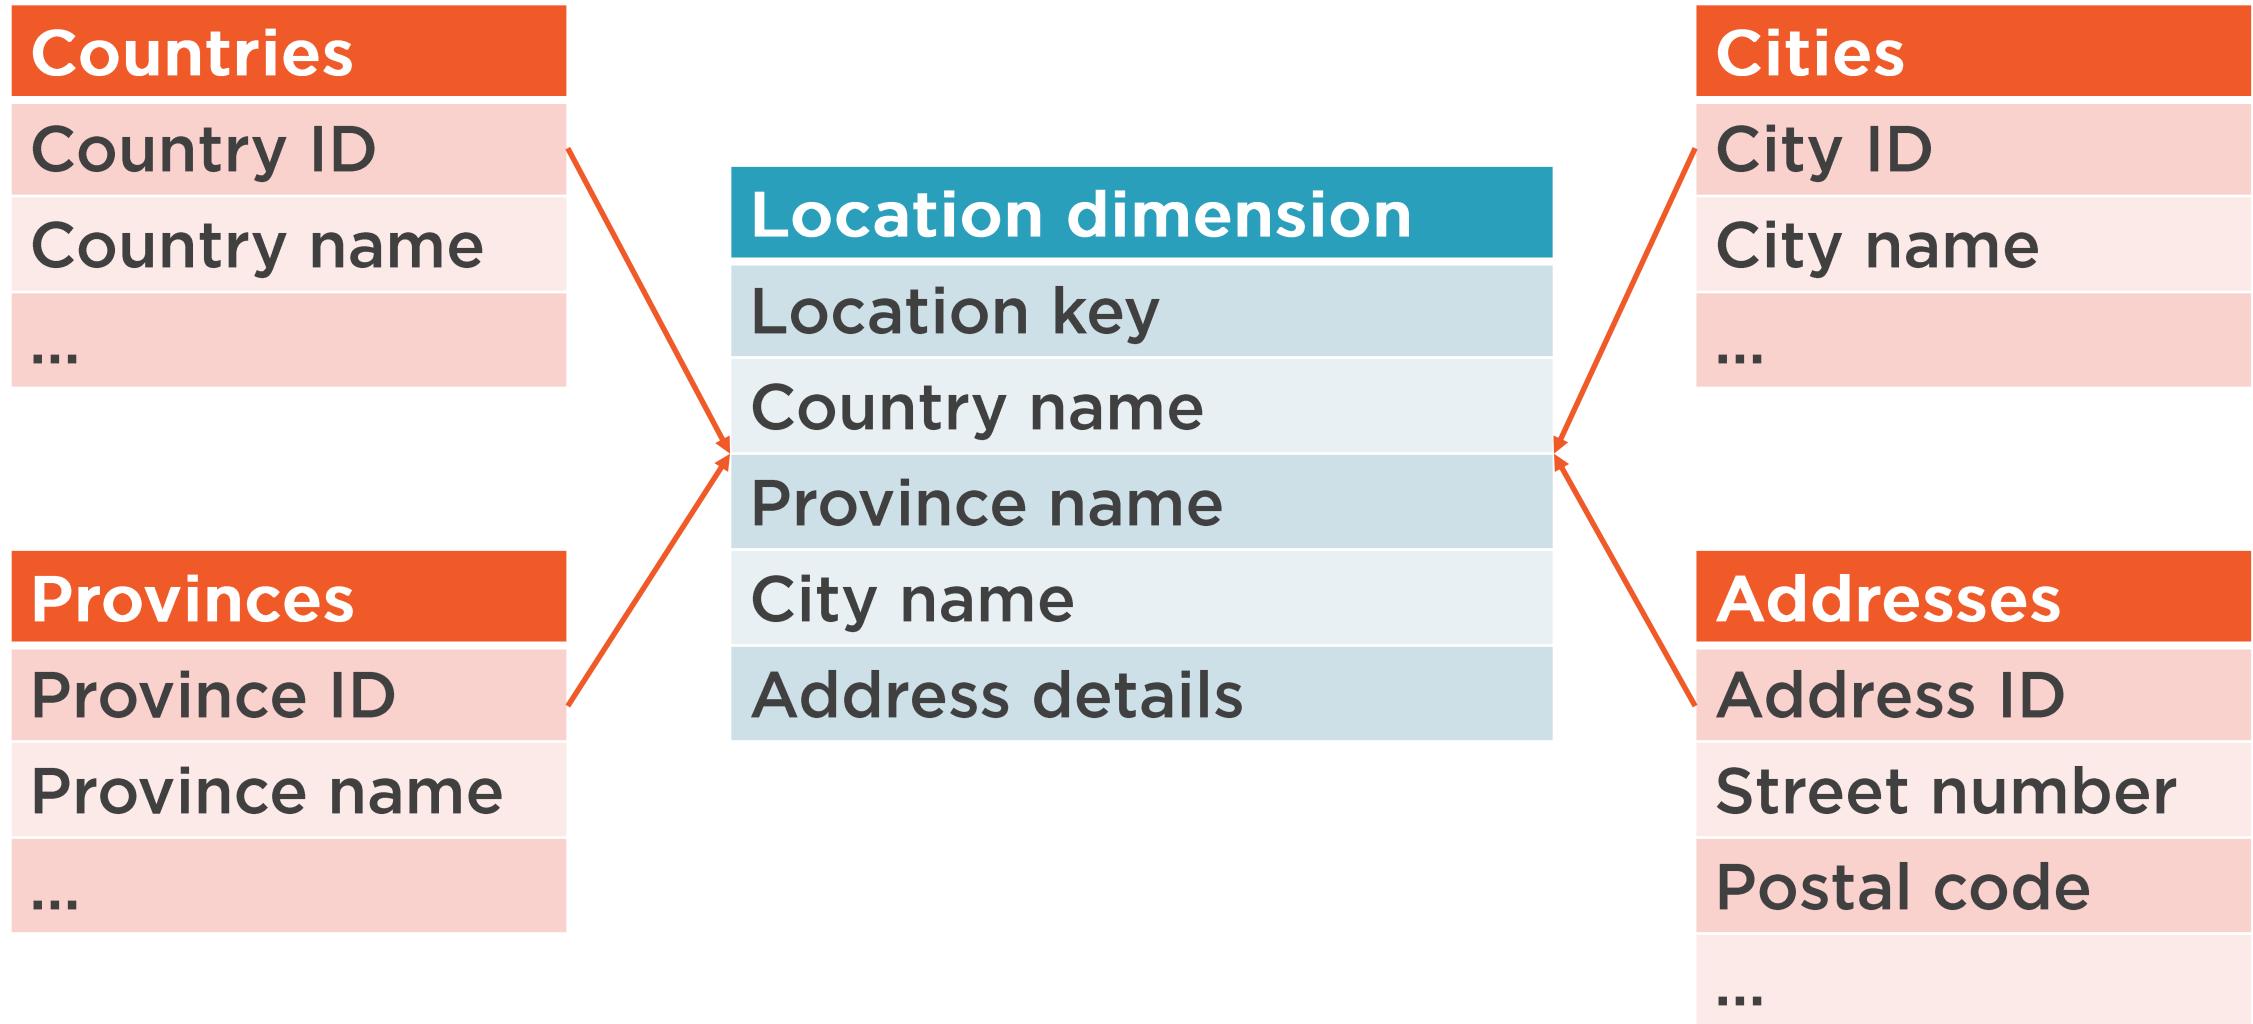


Similar to a pyramid

- Bottom level is the weakest
- Highest level is at the top



A Hierarchy in the Location Dimension



A Hierarchy in the Location Dimension

Bucharest

Vienna

New York

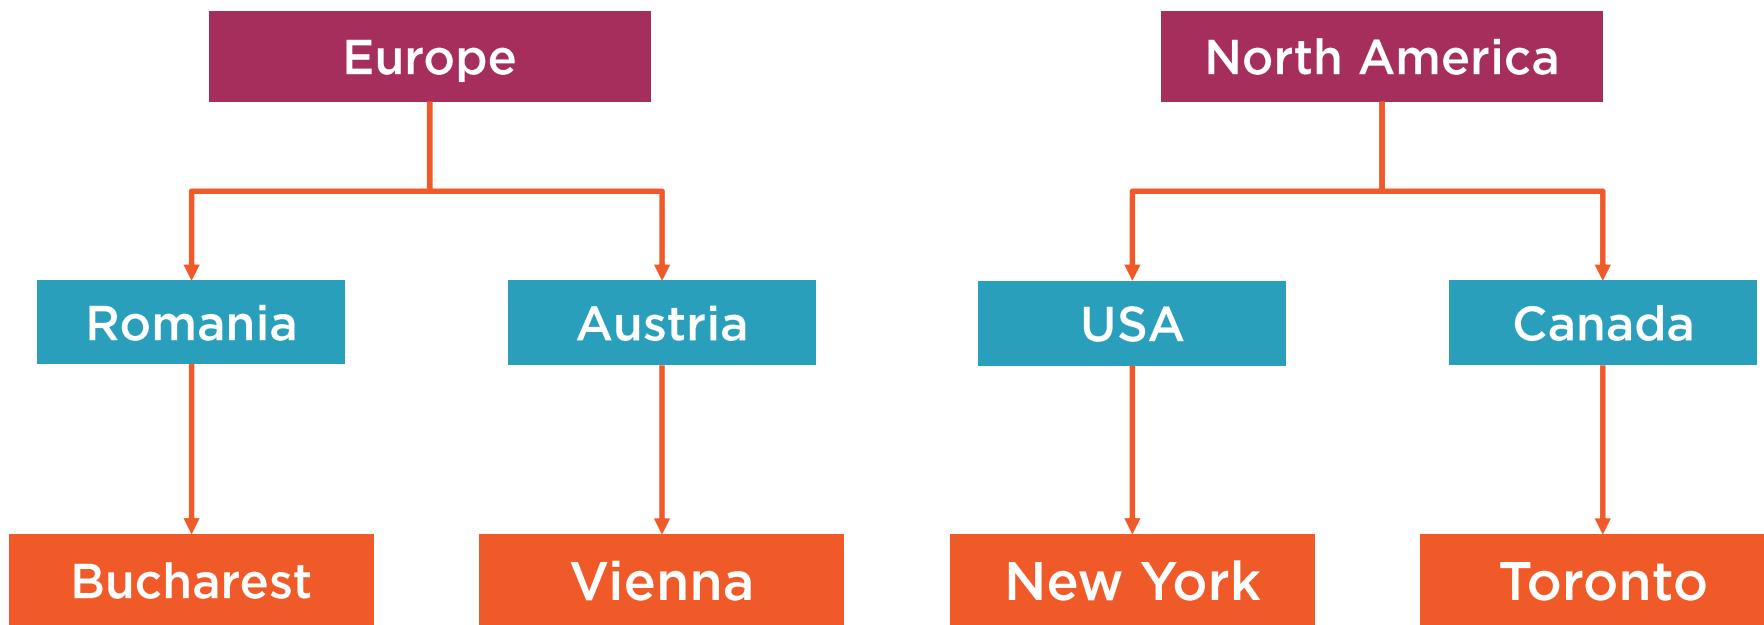
Toronto



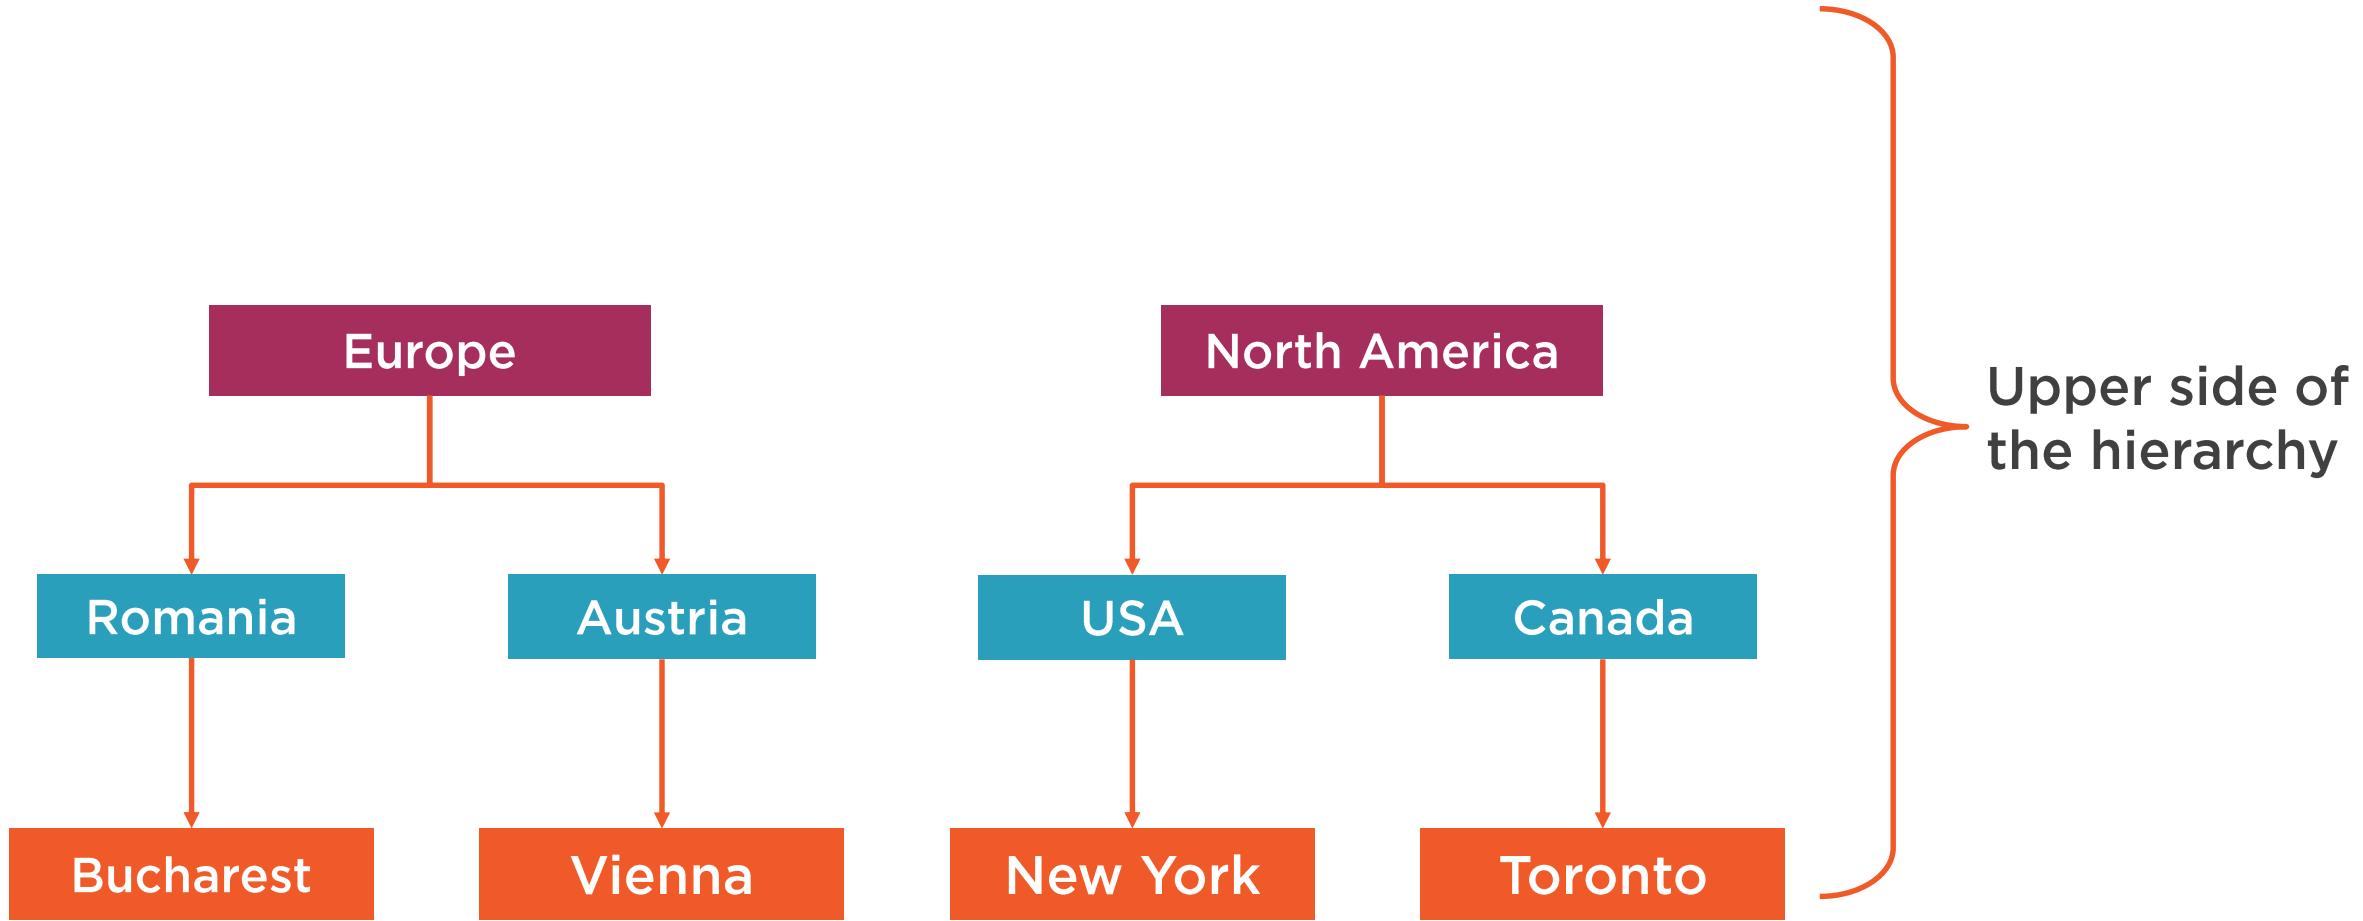
A Hierarchy in the Location Dimension



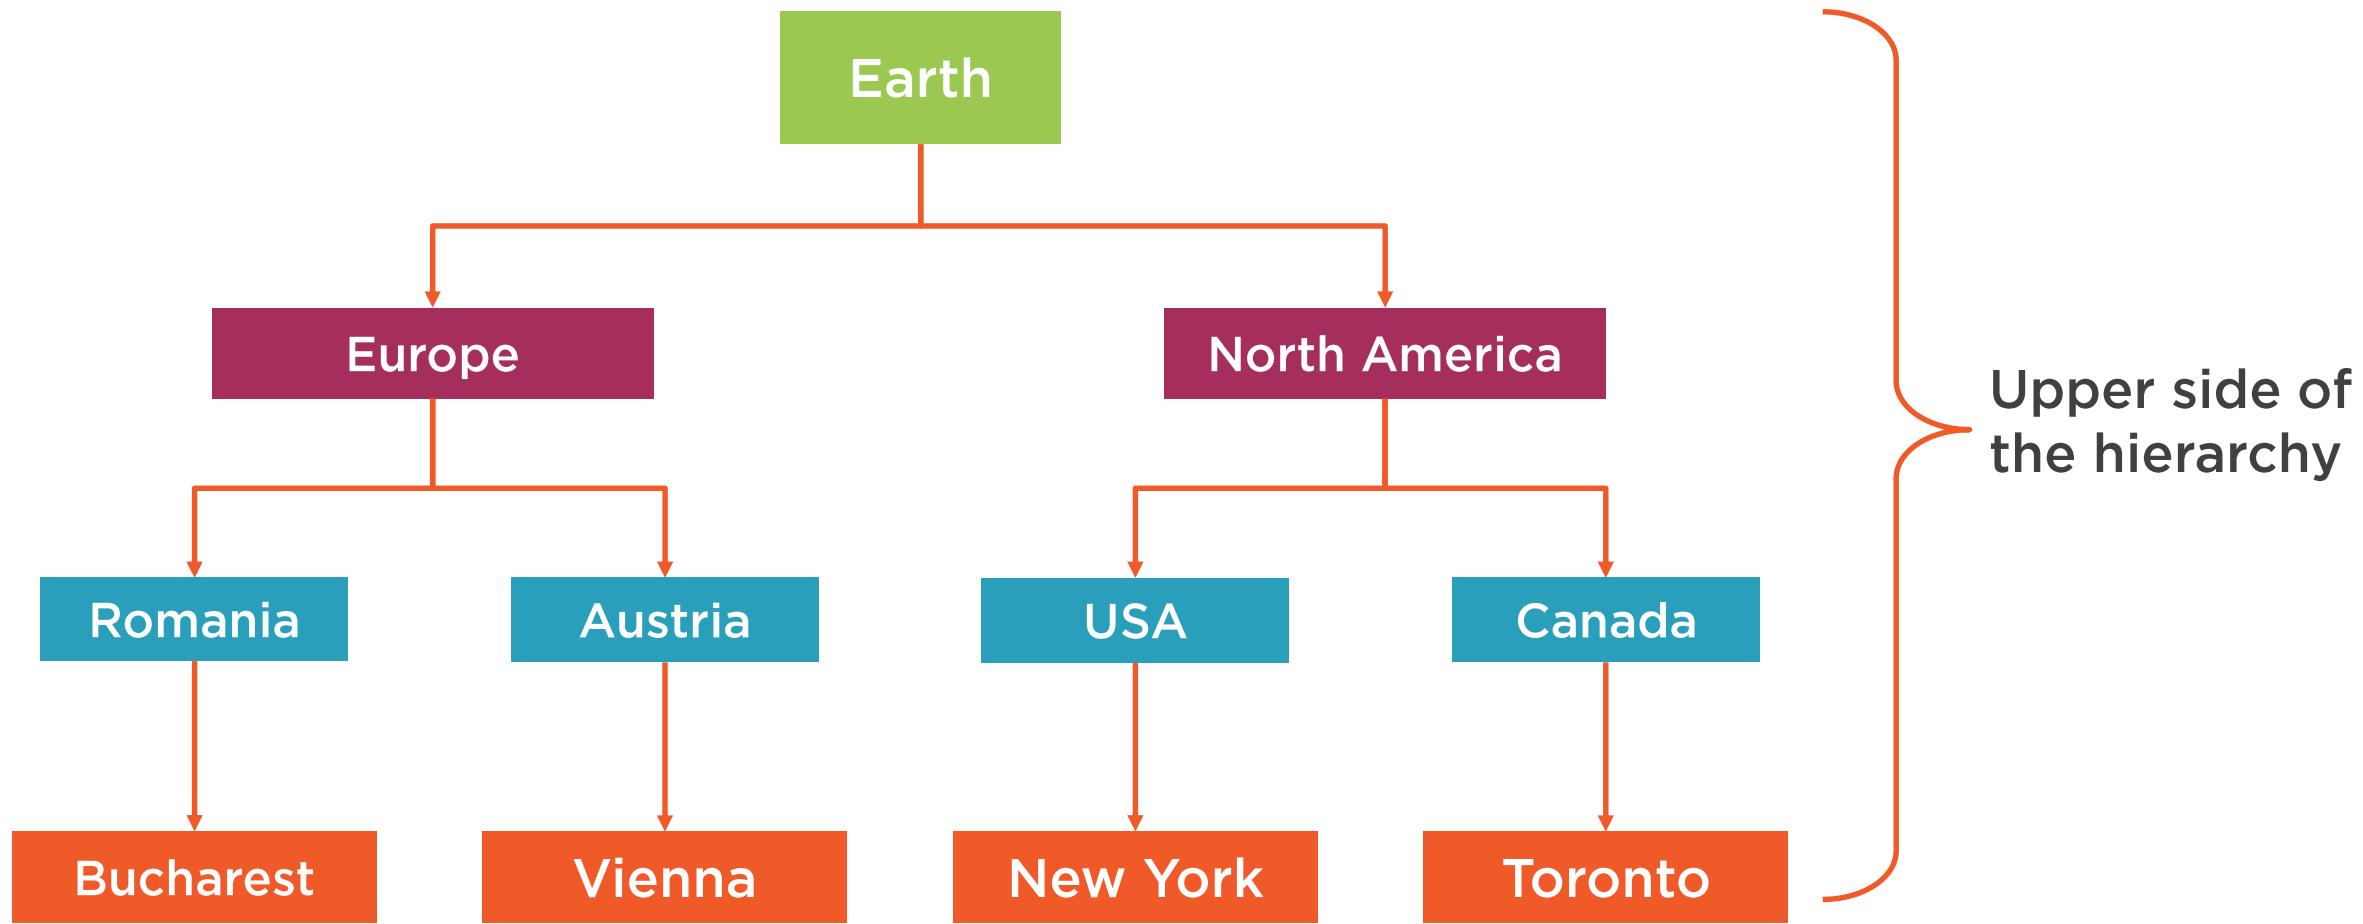
A Hierarchy in the Location Dimension



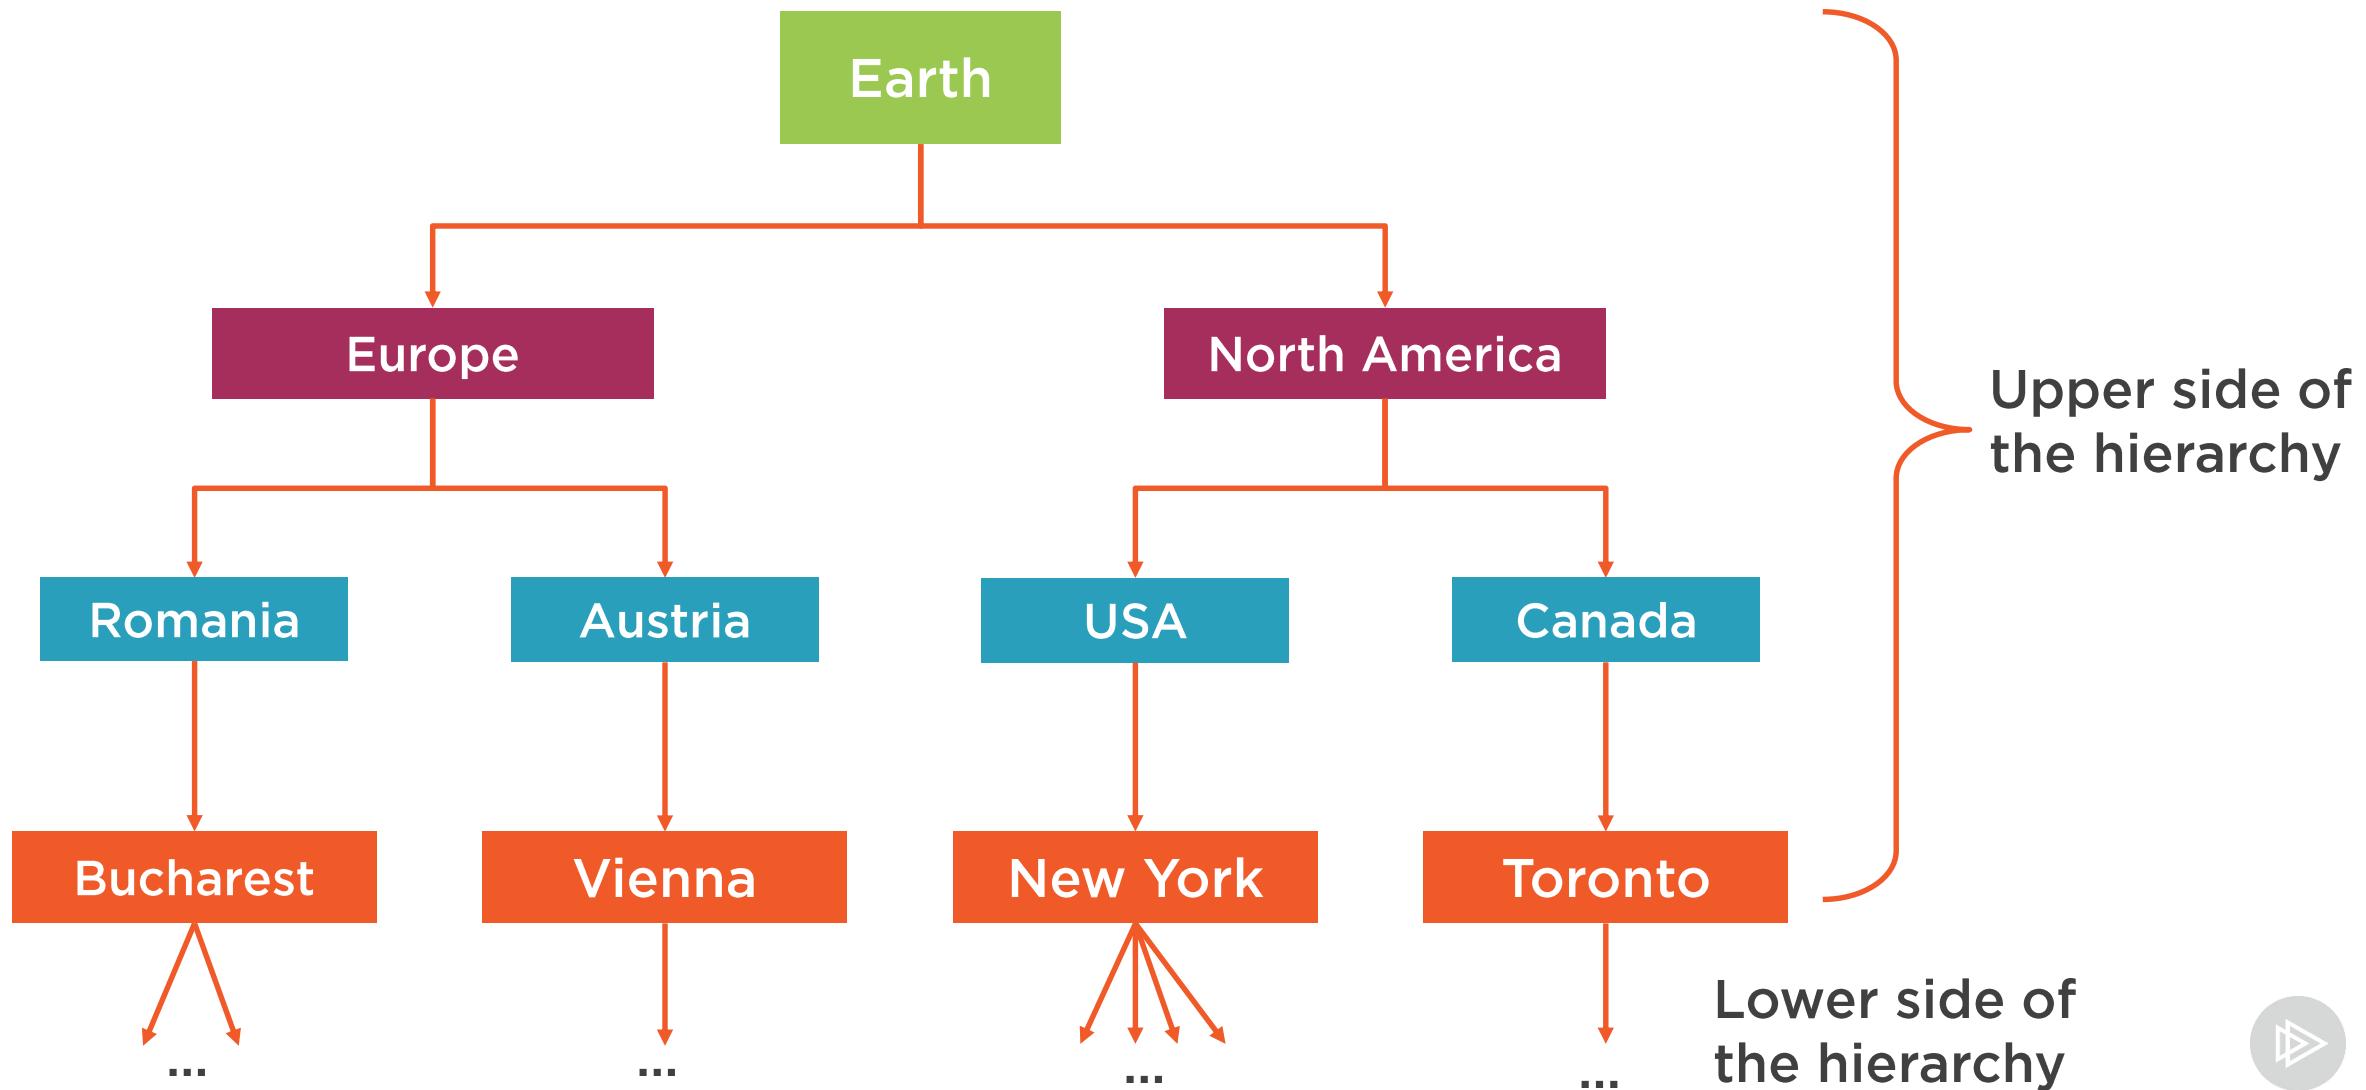
A Hierarchy in the Location Dimension



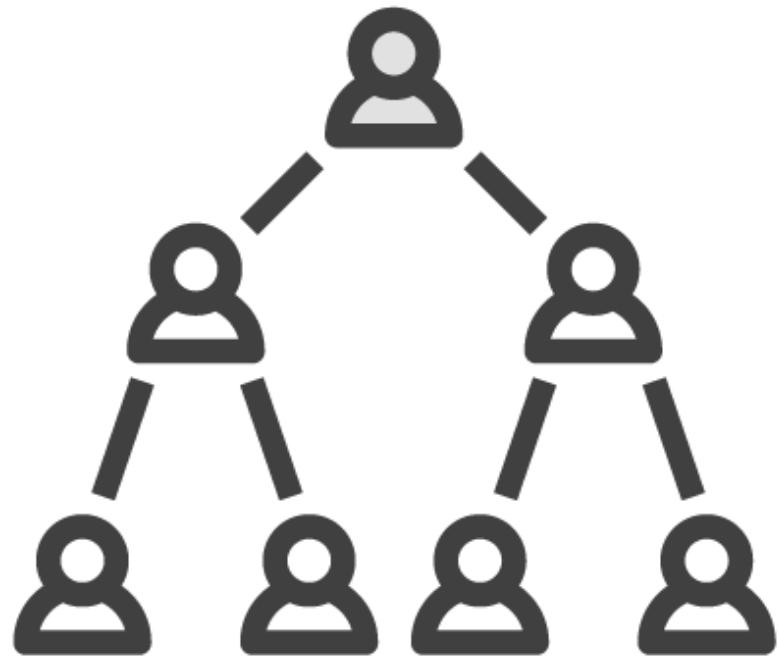
A Hierarchy in the Location Dimension



A Hierarchy in the Location Dimension



What Is a Hierarchy?



Data structure

- Attributes of a dimension are organized together
- Going down the hierarchy nodes -> more details
- Going up the hierarchy -> summarized data



Advantages of Using Hierarchies

Structure

- Important dimensions have plenty attributes
- Analyzing important data can become overwhelming
- Hierarchies provide order in the data

Multiple perspectives

Merchandise hierarchy



Advantages of Using Hierarchies

Structure

- Important dimensions have plenty attributes
- Analyzing important data can become overwhelming
- Hierarchies provide order in the data

Multiple perspectives

Merchandise hierarchy

- Department
 - Category
 - Subcategory
 - Product name



Advantages of Using Hierarchies

Structure

- Important dimensions have plenty attributes
- Analyzing important data can become overwhelming
- Hierarchies provide order in the data

Multiple perspectives

Merchandise hierarchy

- Department
 - Category
 - Subcategory
 - Product name

Packaging hierarchy



Advantages of Using Hierarchies

Structure

- Important dimensions have plenty attributes
- Analyzing important data can become overwhelming
- Hierarchies provide order in the data

Multiple perspectives

Merchandise hierarchy

- Department
- Category
- Subcategory
- Product name

Packaging hierarchy

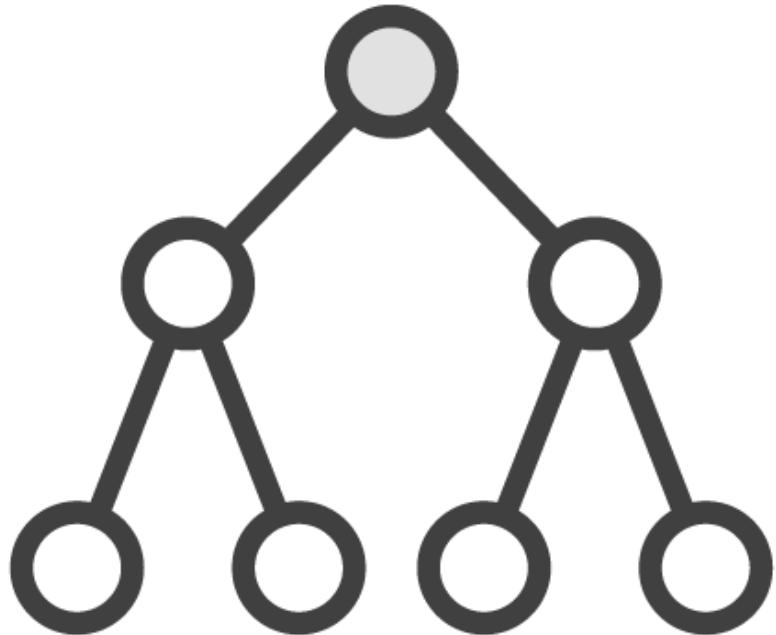
- Package type
- Package size
- Product name

Zoom in/zoom out

- Visualize summarized data and detailed data
- Go as deep as the business requires it



Types of Hierarchies



Fixed-depth (balanced)

- Fixed number of levels
- Easy to create and work with

Variable-depth (unbalanced or ragged)

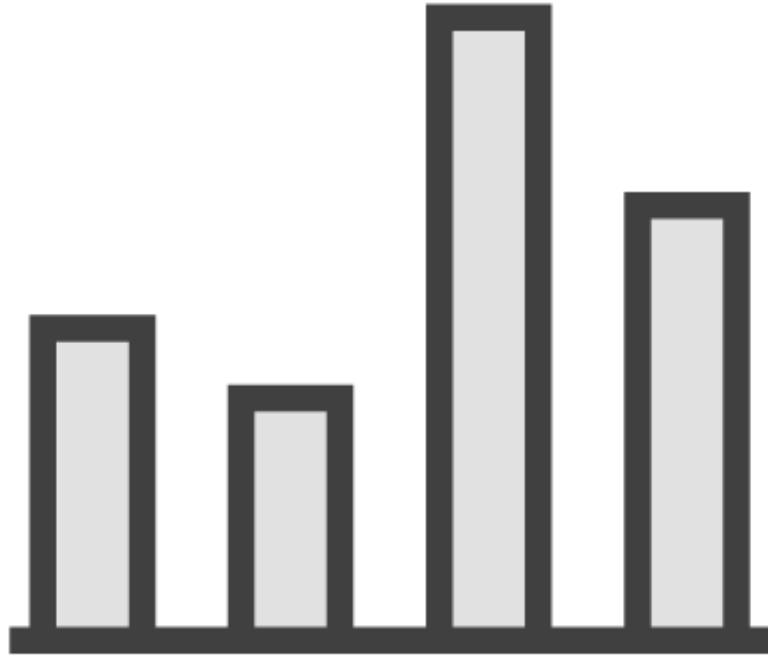
- Uneven number of levels
- Creating them is more complex task



Drilling Down a Hierarchy



Data Warehouse Analysis



Minimum requirements for doing data warehouse analysis

- One fact table
- One dimension table

Example: sales per product report

- Sales fact
- Product dimension



Example of Drilling Down

Sales fact

Store key

Date key

Product key

Employee key

Customer key

...

Transaction

Unit price

Quantity

Amount



Example of Drilling Down

Sales fact	Product dimension
Store key	Product key
Date key	Product name
Product key	Department
Employee key	Category
Customer key	Subcategory
...	Package size
Transaction #	Package type
Unit price	Description
Quantity	Unit of measure
Amount	...



Example of Drilling Down

Sales fact
Store key
Date key
Product key
Employee key
Customer key
...
Transaction #
Unit price
Quantity
Amount

Product dimension
Product key
Product name
Department
Category
Subcategory
Package size
Package type
Description
Unit of measure
...



Merchandise hierarchy

- Department
 - Category
 - Subcategory
 - Product name



Example of Drilling Down

Sales

80.000

Department

Sales

Bakery	Sweets	Beverages
25.000	40.000	15.000



Example of Drilling Down

Department
Category
Sales

Bakery			Sweets			Beverages	
Bread	Pie	Croissant	Cookie	Cake	Candy	Juice	Tea
7.000	13.000	5.000	8.700	6.300	5.000	7.000	6.600



Example of Drilling Down

Department

Category

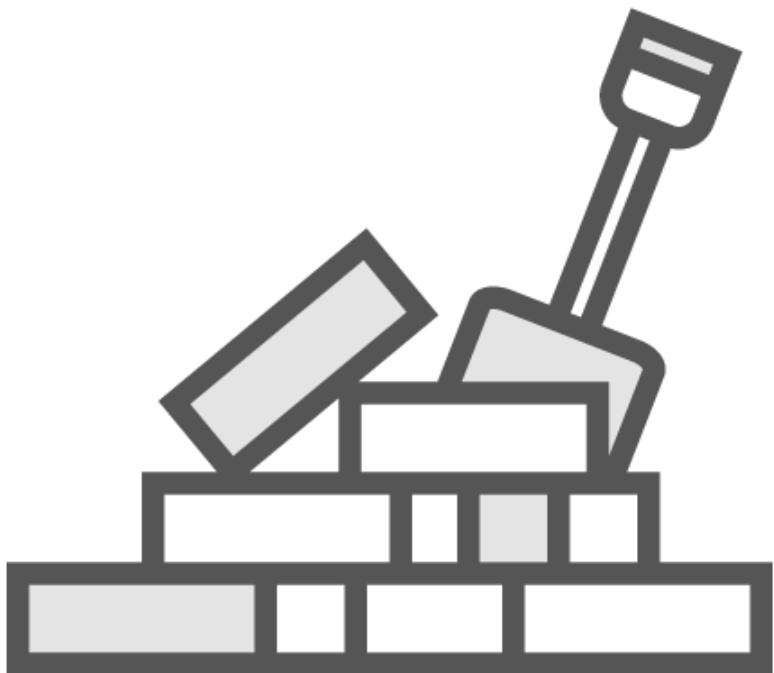
Subcategory

Sales

Bakery							
Bread			Pie			Croissant	
Baguette	Pita	Banana br.	Cream	Fruit	Custard	Sweet	Salty
2.000	3.000	2.000	2.500	2.200	8.300	3.000	2.000



Drilling Down - Summary



Drilling down

- Adding another member of the dimension to the report
- The member doesn't need to be part of a hierarchy

Drilling up/rolling up

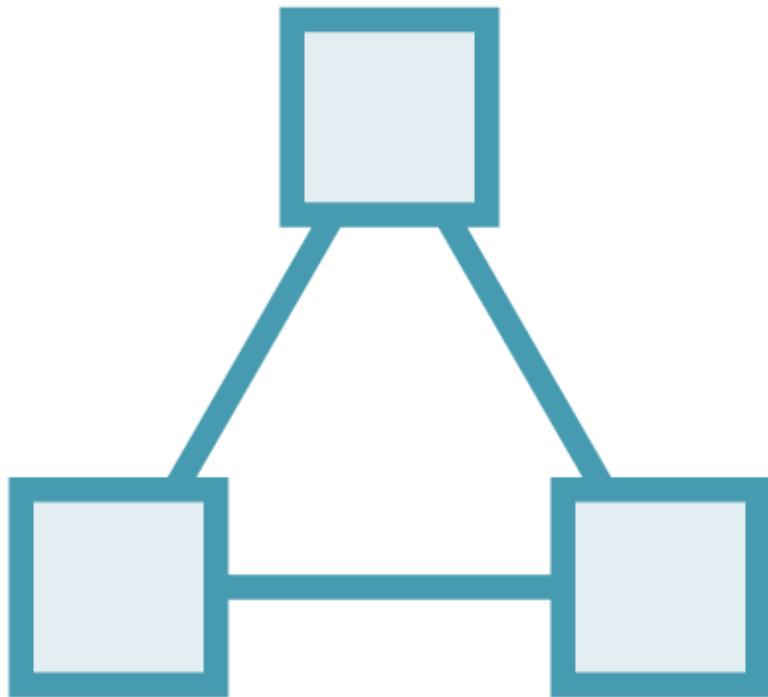
- Taking out an attribute of a dimension from a report



Fixed-depth Positional Hierarchies



Fixed-depth Positional Hierarchies



The number of levels is known upfront

The levels are attributes in the dimension table

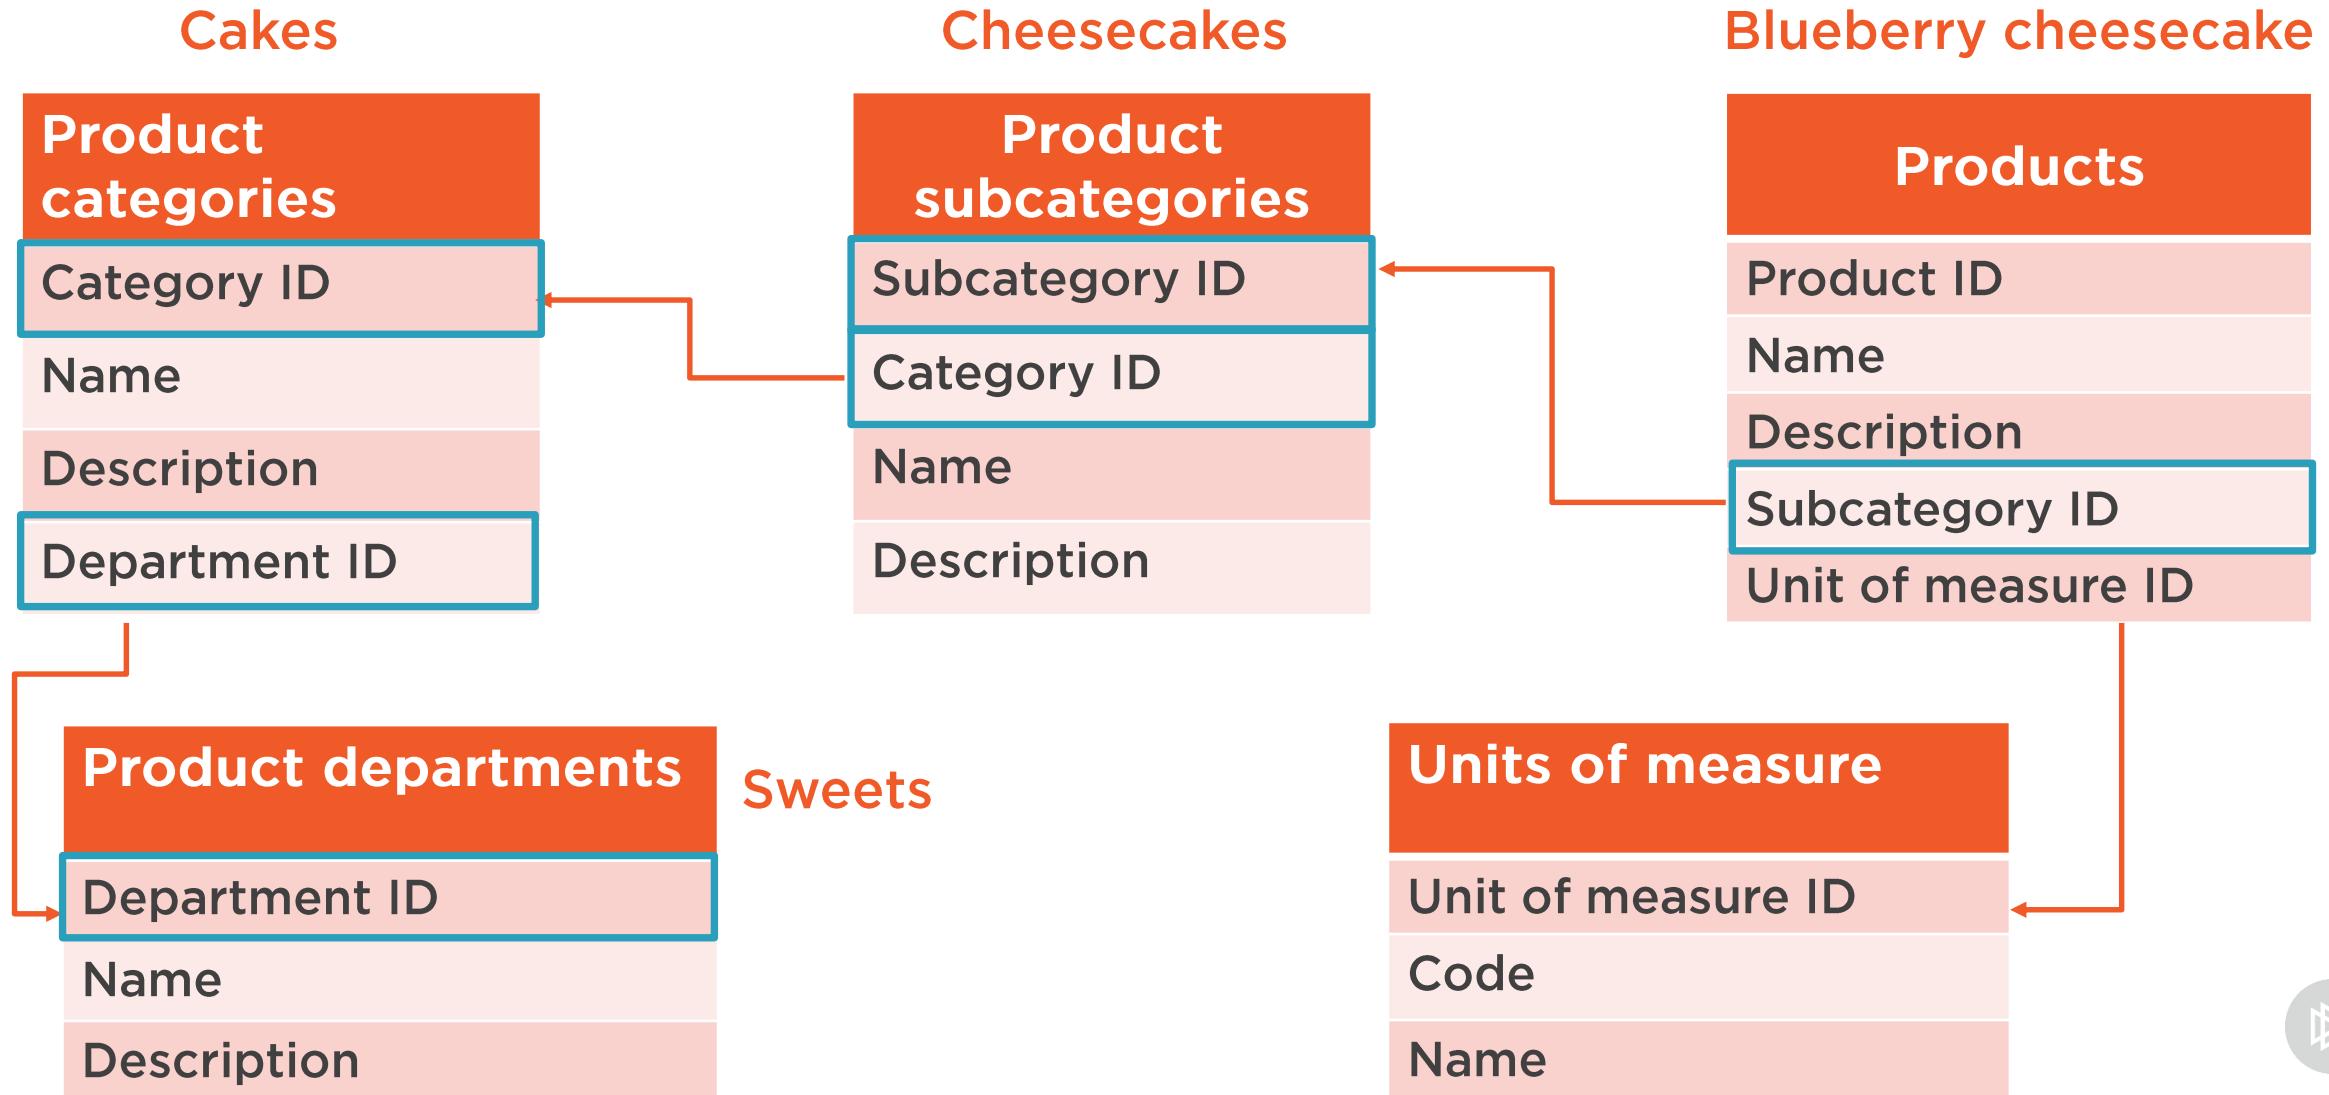
It is a series of many-to-one relationships

Advantages

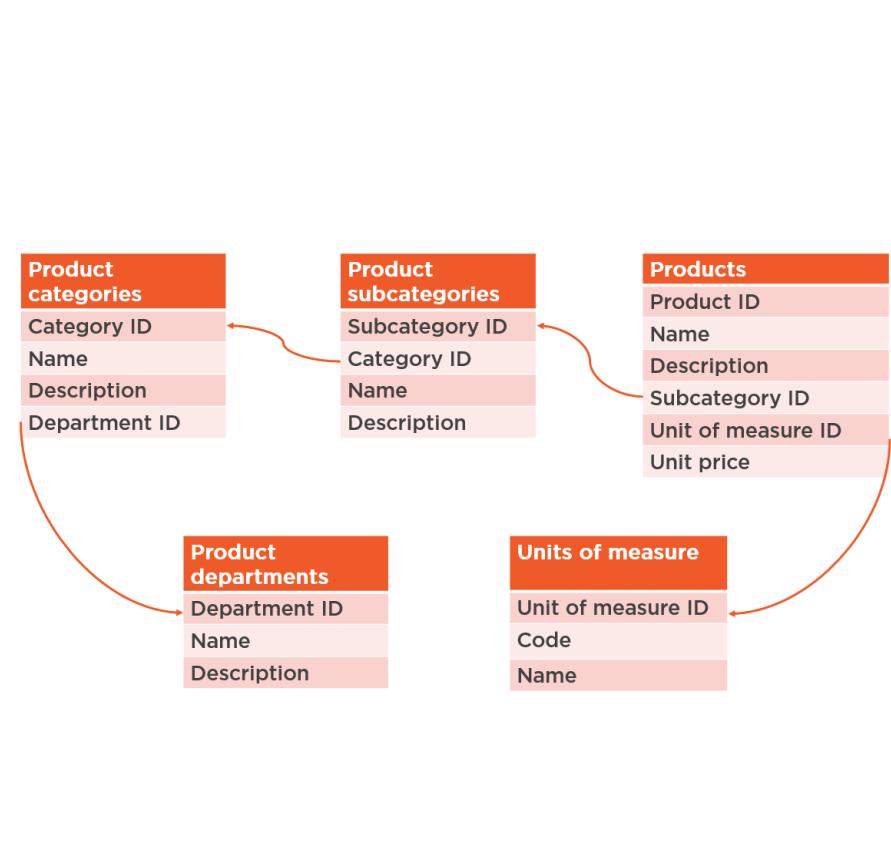
- Easy to navigate
- Offers predictable results
- No impact on performance



Product Information from Different Tables



Product Information from the Data Warehouse



Keep in mind:

- Don't create many snowflake designs
- Most relationships should be from fact tables to the dimensions



Creating the Merchandise Hierarchy

Department

Bakery

Category

Bread

Pie

Subcategory

Baguette

Banana
bread

Custard

Cream

Product

Le Petit
Francais

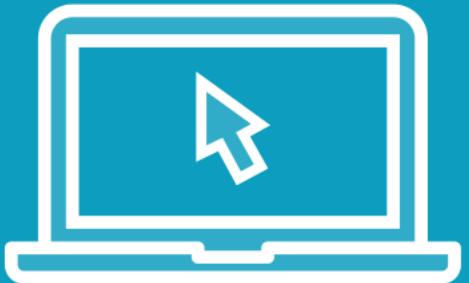
The Big
Banana
Bread

Eagle
Brand
Coconut
Pie

Eagle
Brand
Banana
Cream Pie



Demo



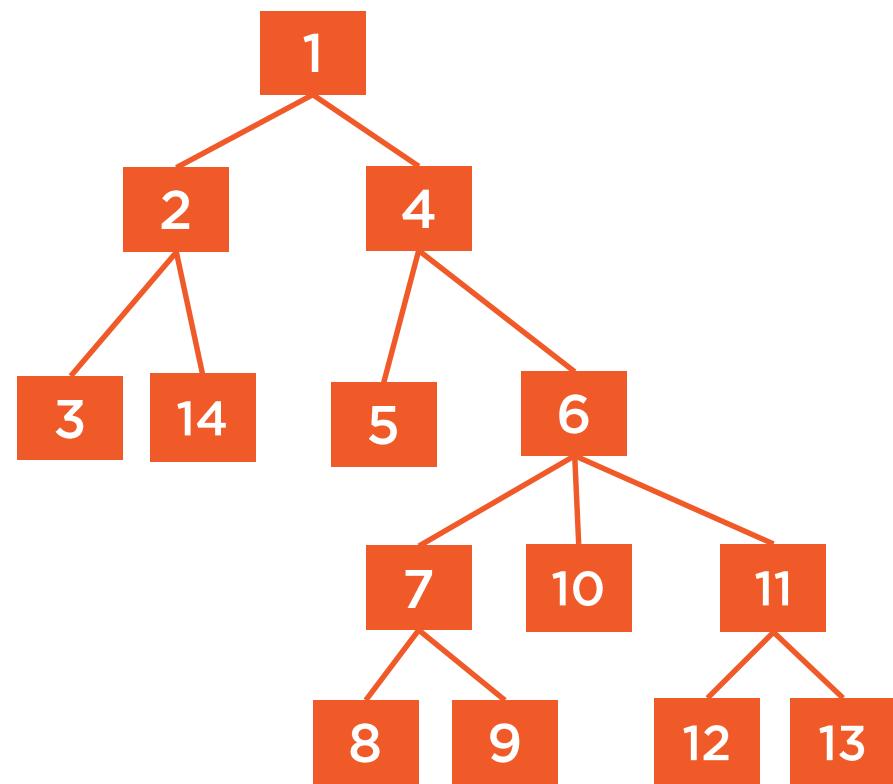
Creating a fixed-depth positional hierarchy



Variable-depth Positional Hierarchies



Characteristics of Variable-depth Hierarchies



The number of levels is not known at design time

- Example: an organizational chart

Are more complex structures (compared to the fixed-depth hierarchies)

Should be used with moderation



Classification of Variable-depth Hierarchies



Slightly ragged

Ragged, created with a hierarchy bridge

Ragged, created with pathstring attributes



Slightly Ragged Hierarchies



- The number of levels is not known beforehand
- The range in depth is small
- Geographic hierarchies are slightly ragged



Example of a Slightly Ragged Hierarchy

Location hierarchy

- Country
 - Province (or state)
 - City
 - Neighborhood
 - Address

Examples of data with missing levels

Singapore (country and city)

Vatican (independent city-state)

Small cities, that don't have neighborhoods



Fitting Data into a Slightly Ragged Hierarchy



Step 1: Create all possible levels of the hierarchy

Step 2: Fill in the missing values per each level with:

- An expression, similar to “not applicable”
- The value of the next parent member

Step 3: Handle the ragged hierarchy as a fixed-depth one



Populating a Slightly Ragged Hierarchy

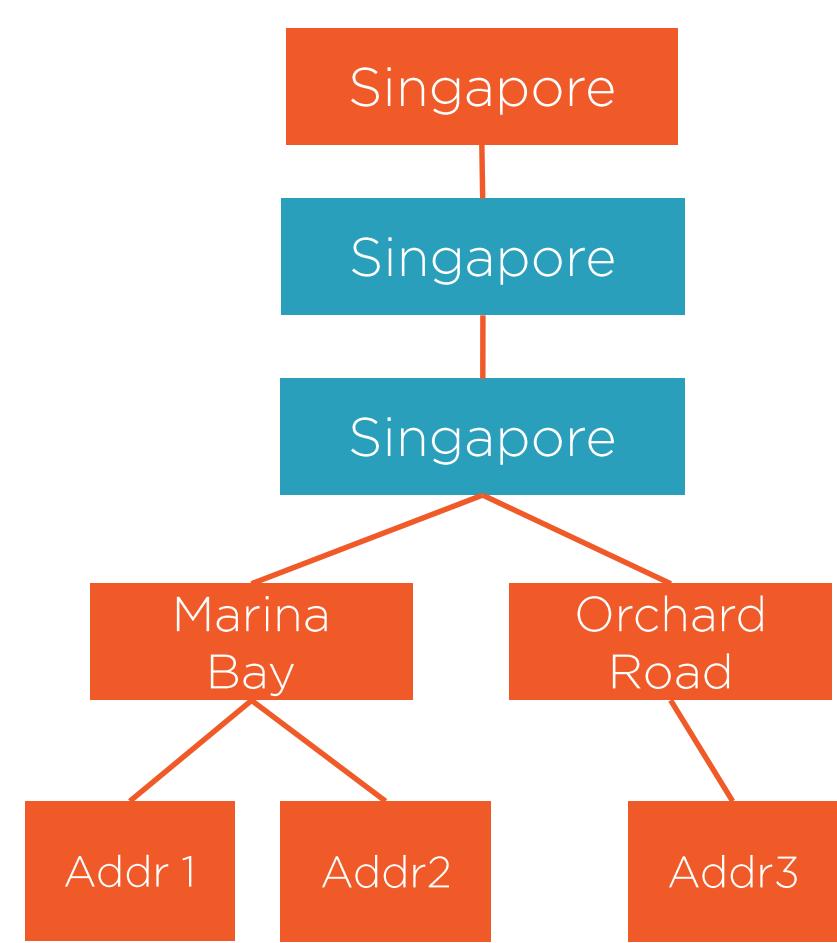
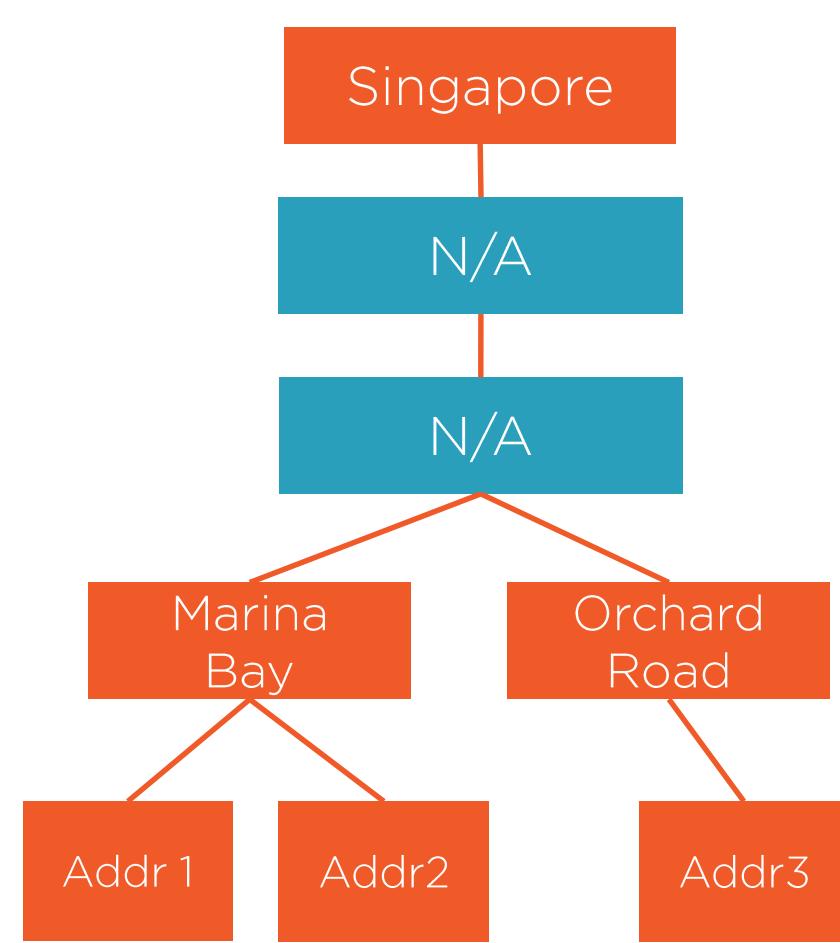
Country

↳ **Province**

↳ **City**

↳ **Neighbor**

↳ **Address**



Demo



Creating and working with a slightly ragged hierarchy

- Based on the Location dimension

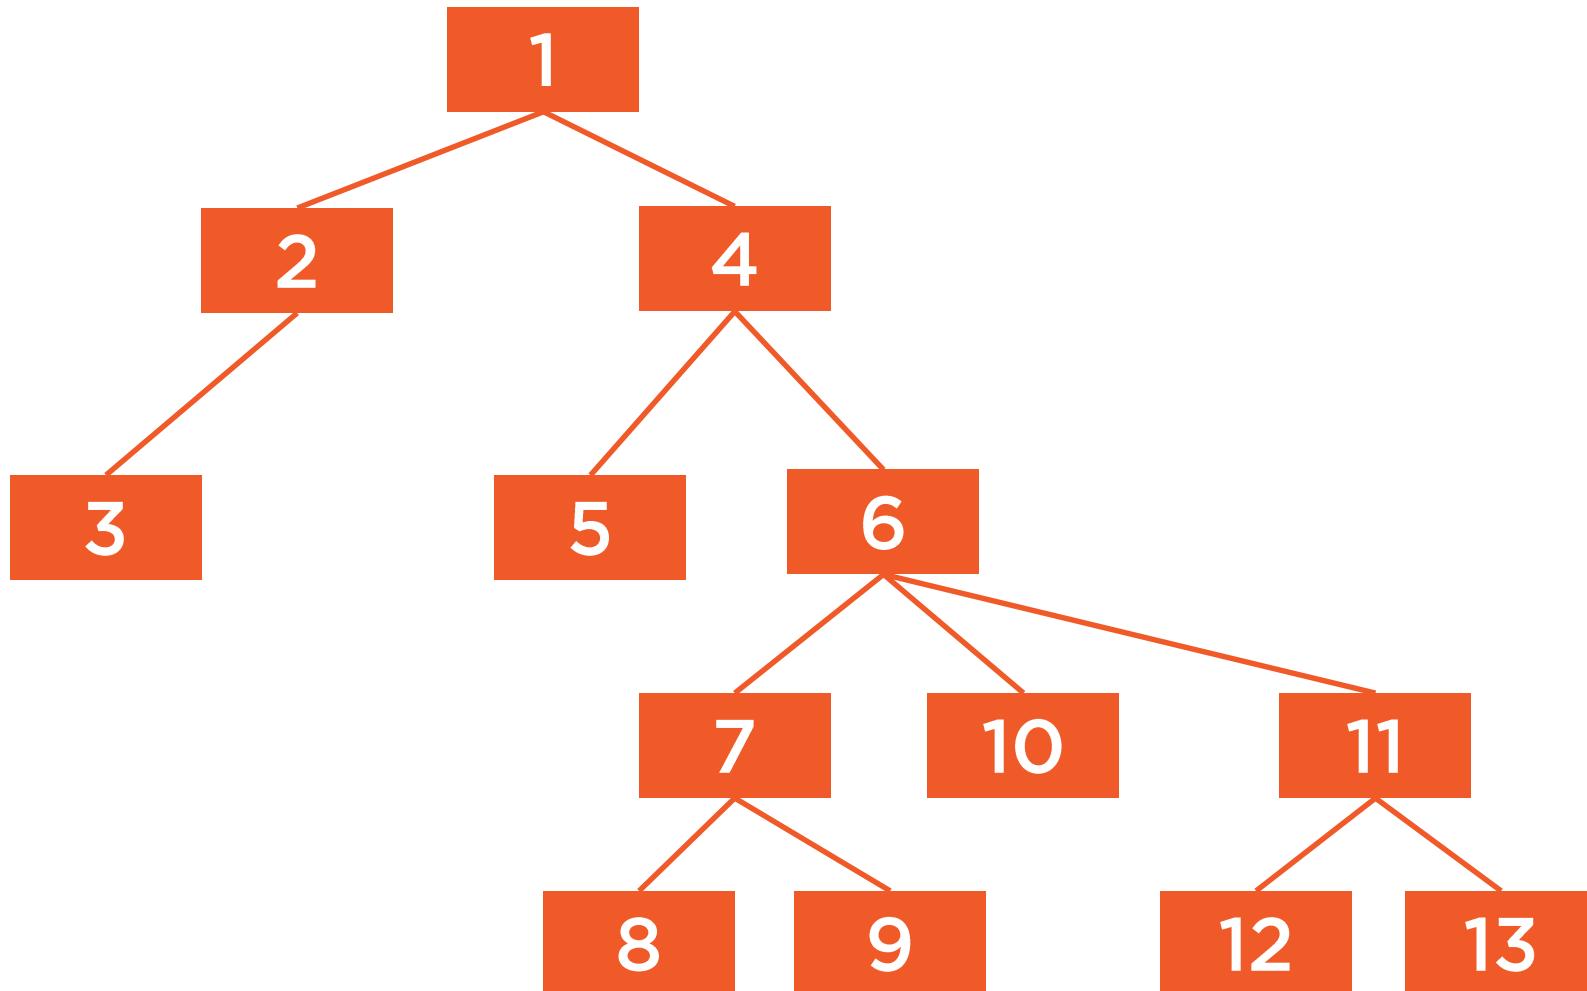


Ragged Hierarchies



Example of a Ragged Hierarchy

The organizational chart



Implementing a Ragged Hierarchy Using a Bridge



A row for each possible path in the hierarchy is stored in a table

Columns in the bridge table

- ID of the parent
- ID of the child
- Number of levels between them
- Whether the node it a top node or bottom node
- Other information relevant for analysis

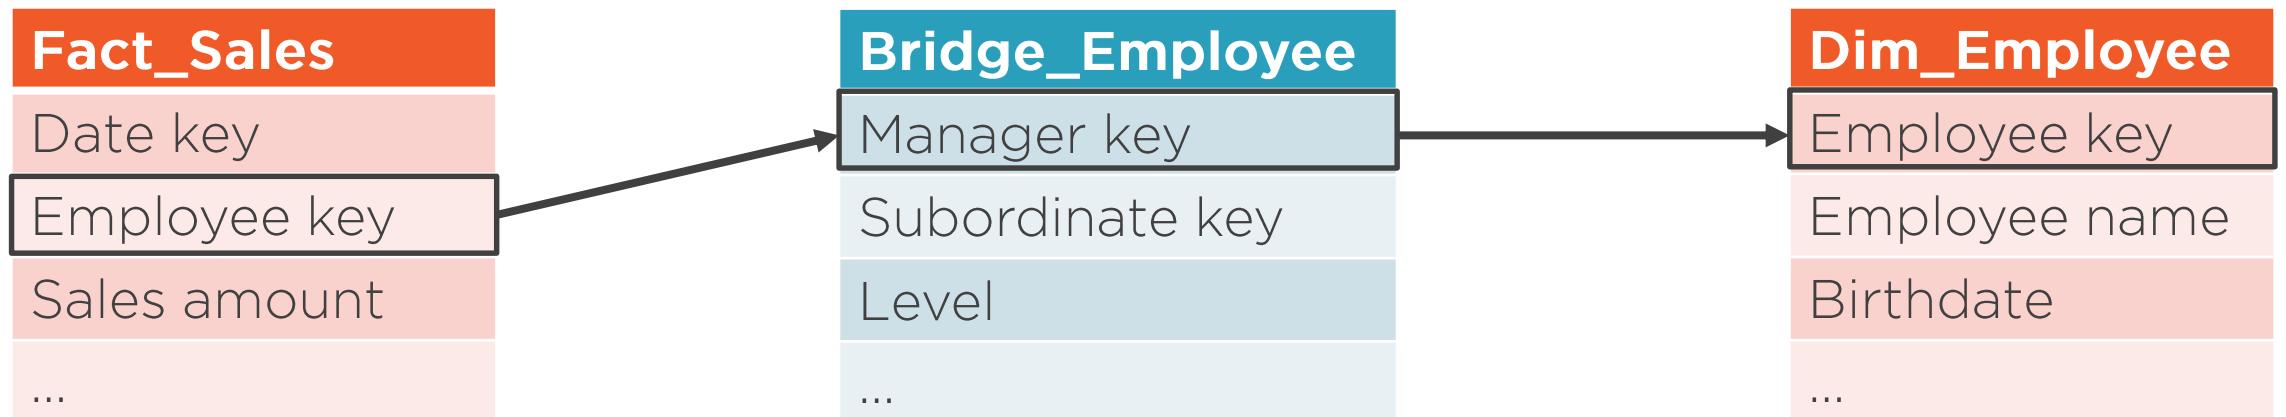
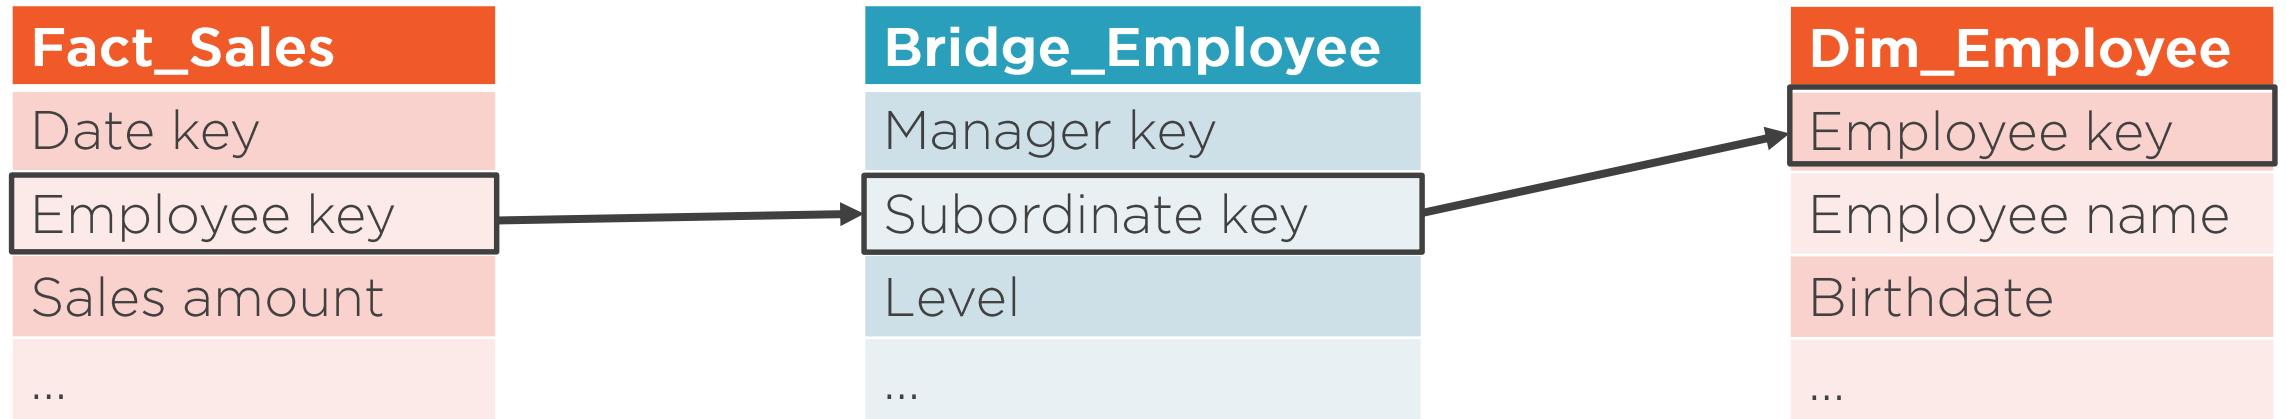


Example of a Hierarchy Bridge Table

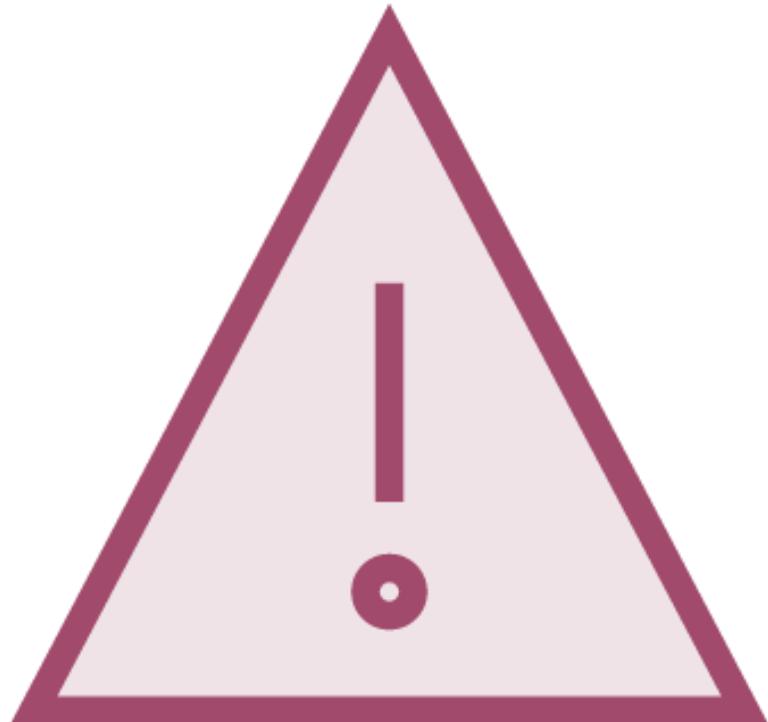
Parent key	Child Key	Level	Top flag	Bottom flag
Julia	Julia	0	Y	N
Julia	Marc	1	N	N
Julia	Theodora	2	N	Y
Julia	Greg	2	N	Y
Marc	Marc	0	Y	N
Marc	Theodora	1	N	Y
Marc	Greg	1	N	Y
Theodora	Theodora	0	Y	Y
Greg	Greg	0	Y	Y



Linking the Fact and the Dimension Table



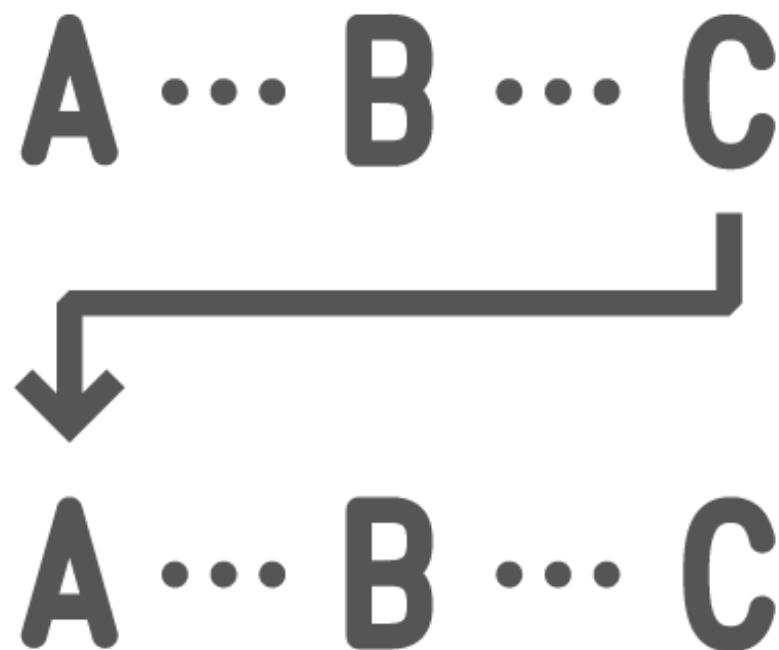
Limitations of the Hierarchy Bridge



**Can grow a lot in size
Performance impact
Difficult to work with**



Ragged Hierarchies Created with Pathstring Attributes



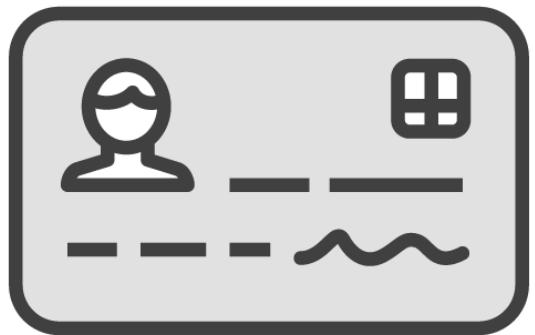
Alternative to the bridge table

The pathstring attribute:

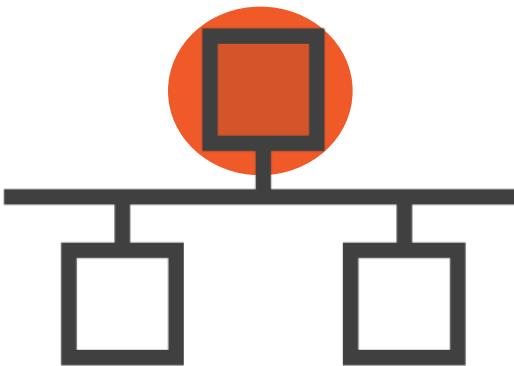
- A special attribute created in the dimension
- A string of characters
- Consists of all the parents of a member from the top of the hierarchy



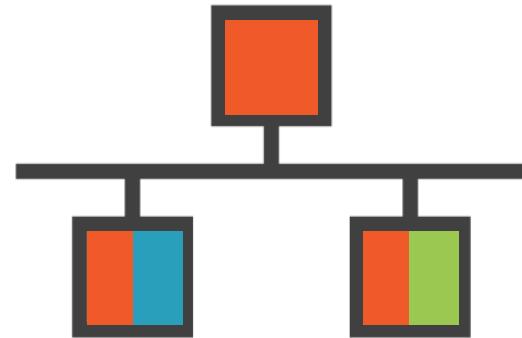
Creating a Hierarchy with Based on Pathstring Attribute



Each node is labeled with a unique value



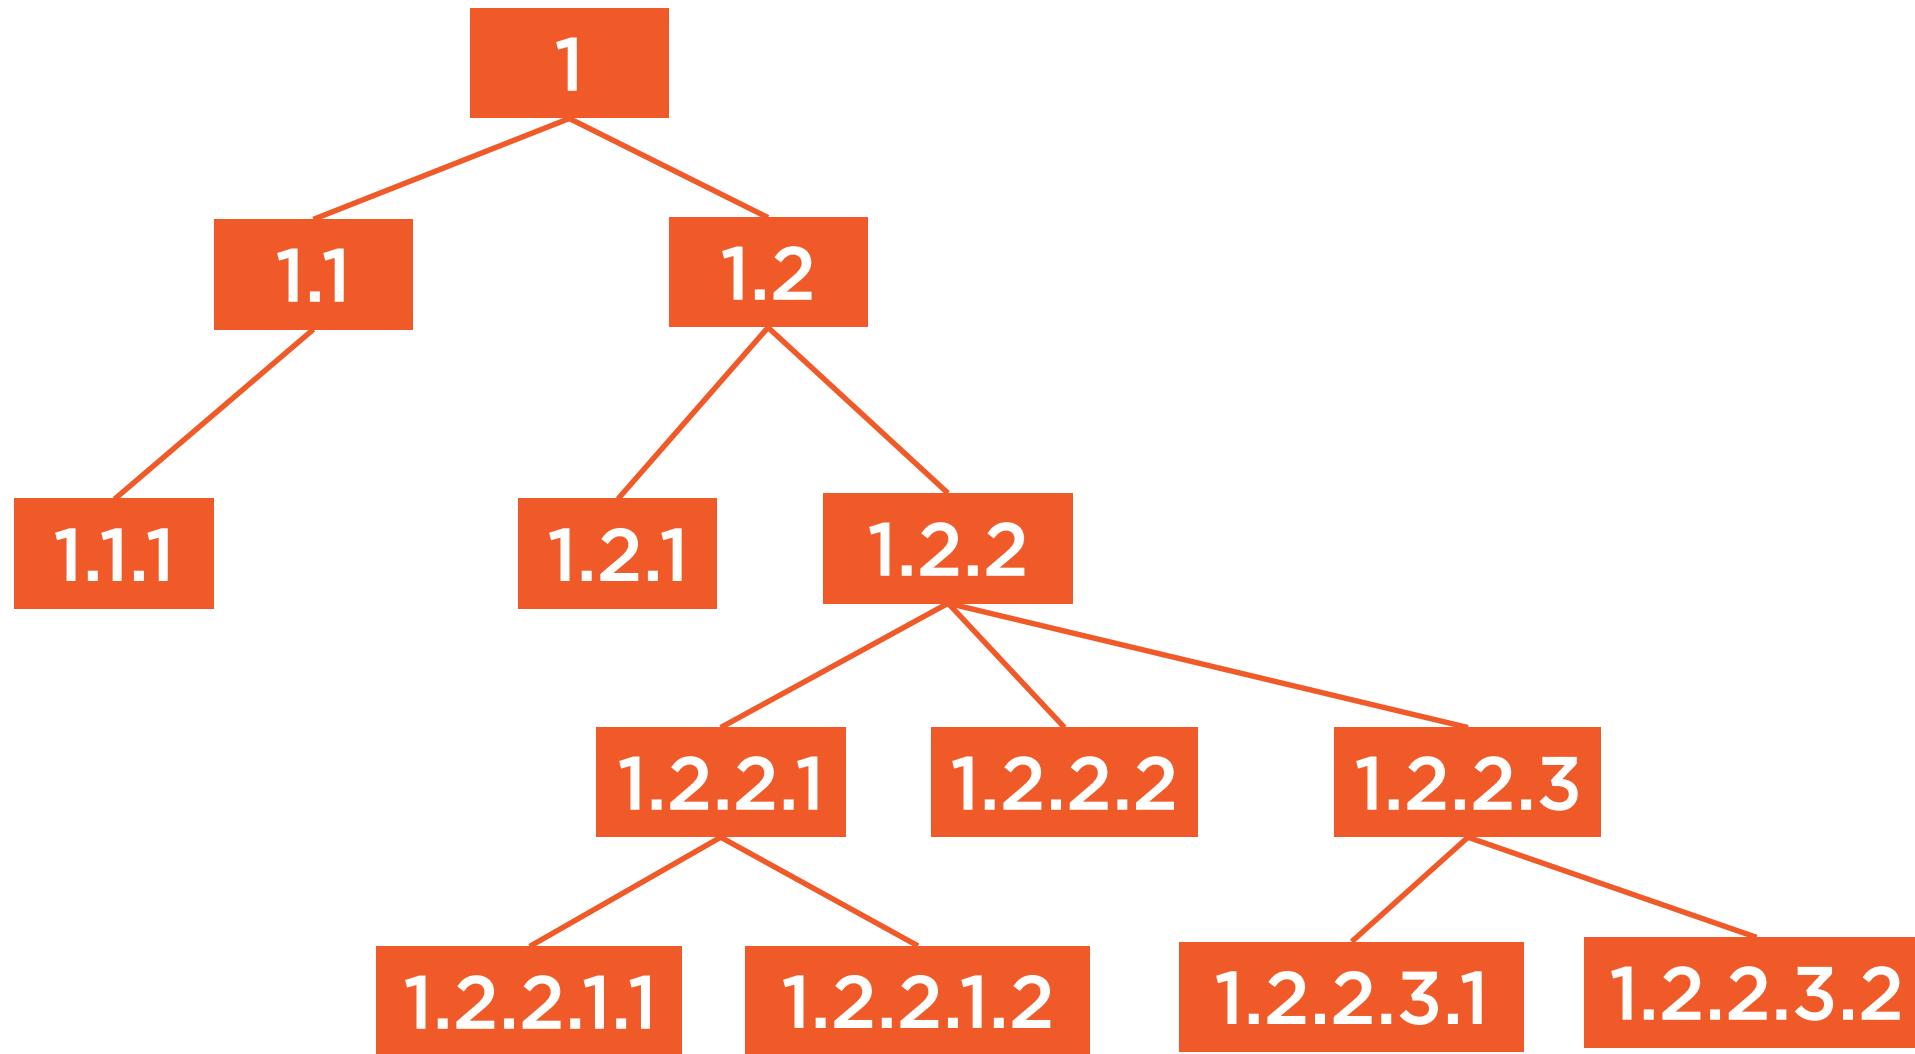
The pathstring of the root node is its unique label



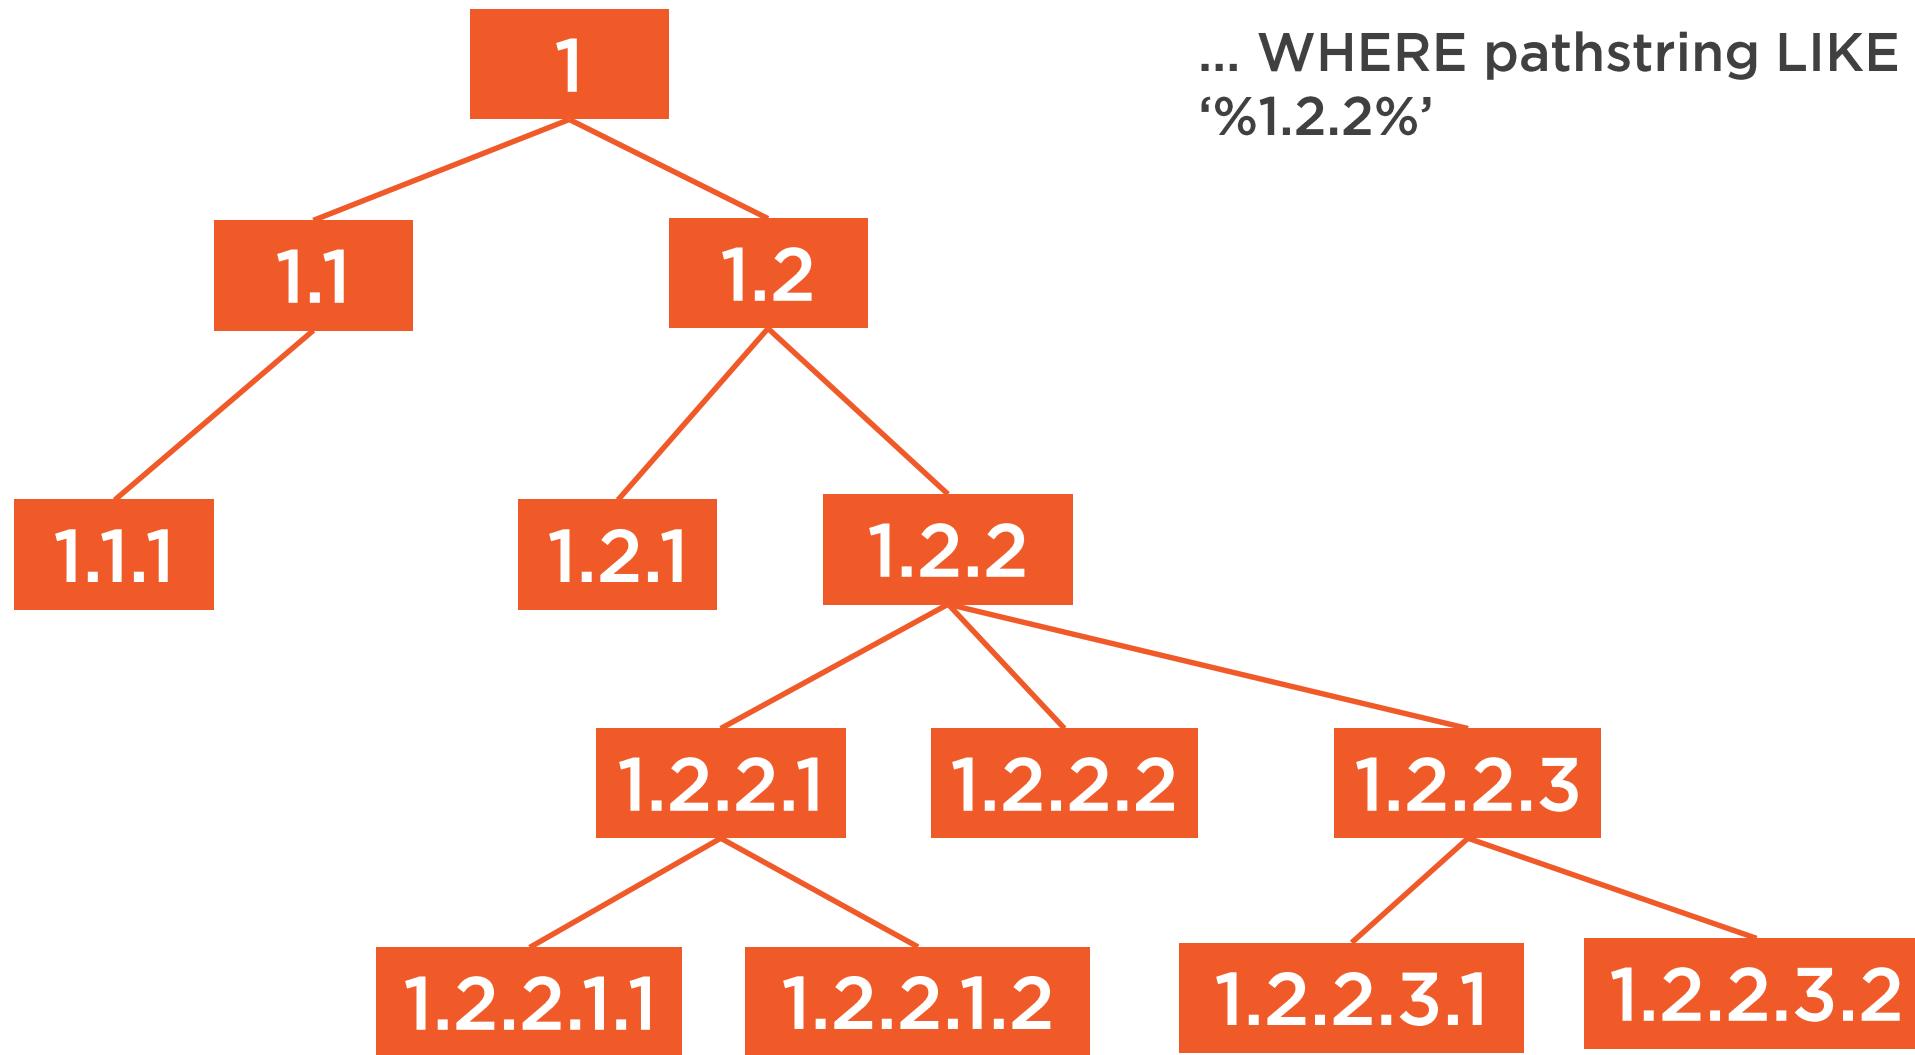
The pathstring on one level includes the pathstring of the parent



Hierarchy Based on Pathstring Attribute



Hierarchy Based on Pathstring Attribute



Limitations of the Pathstring Attribute



Vulnerable to structure changes

- If a member is moved within the organization (or added or deleted)
- Entire hierarchy must be relabeled

Complicated to use by business users

- Accessing the database directly
- Using the 'LIKE' operator in SQL queries



Summary



There is no “best solution” to this problem

A good enough solution is generated by having a clear understanding of

- The data available
- The business requirements



Summary



Hierarchy definition

- Data structures with multiple levels
- Levels are formed with attributes of a dimension

Advantages of hierarchies

- Structure the data
- Multiple perspectives of the same data
- Easily visualize summarized or detailed data

Types of hierarchies

- Fixed-depth
- Variable-depth
 - Slightly ragged
 - Ragged



Creating and Working with Hierarchies



Ana Voicu
@ana_voicu



Overview



In short, a hierarchy is:

- A data structure
- Created with attributes from a dimension
- Used for data aggregation

Topics elaborated in this chapter:

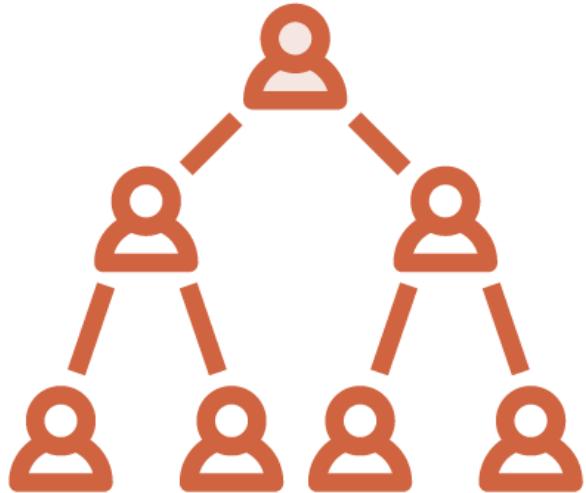
- What is a hierarchy?
- Why is it useful?
- What does drilling-down mean?
- Implementing a hierarchy in a data warehouse



What Is a Hierarchy?

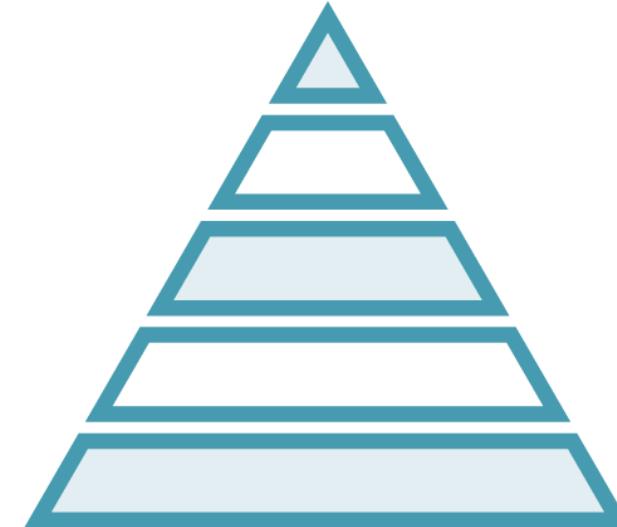


What Is a Hierarchy?



Data structure with multiple levels

- The levels are dimension attributes
- The elements on each level are called nodes

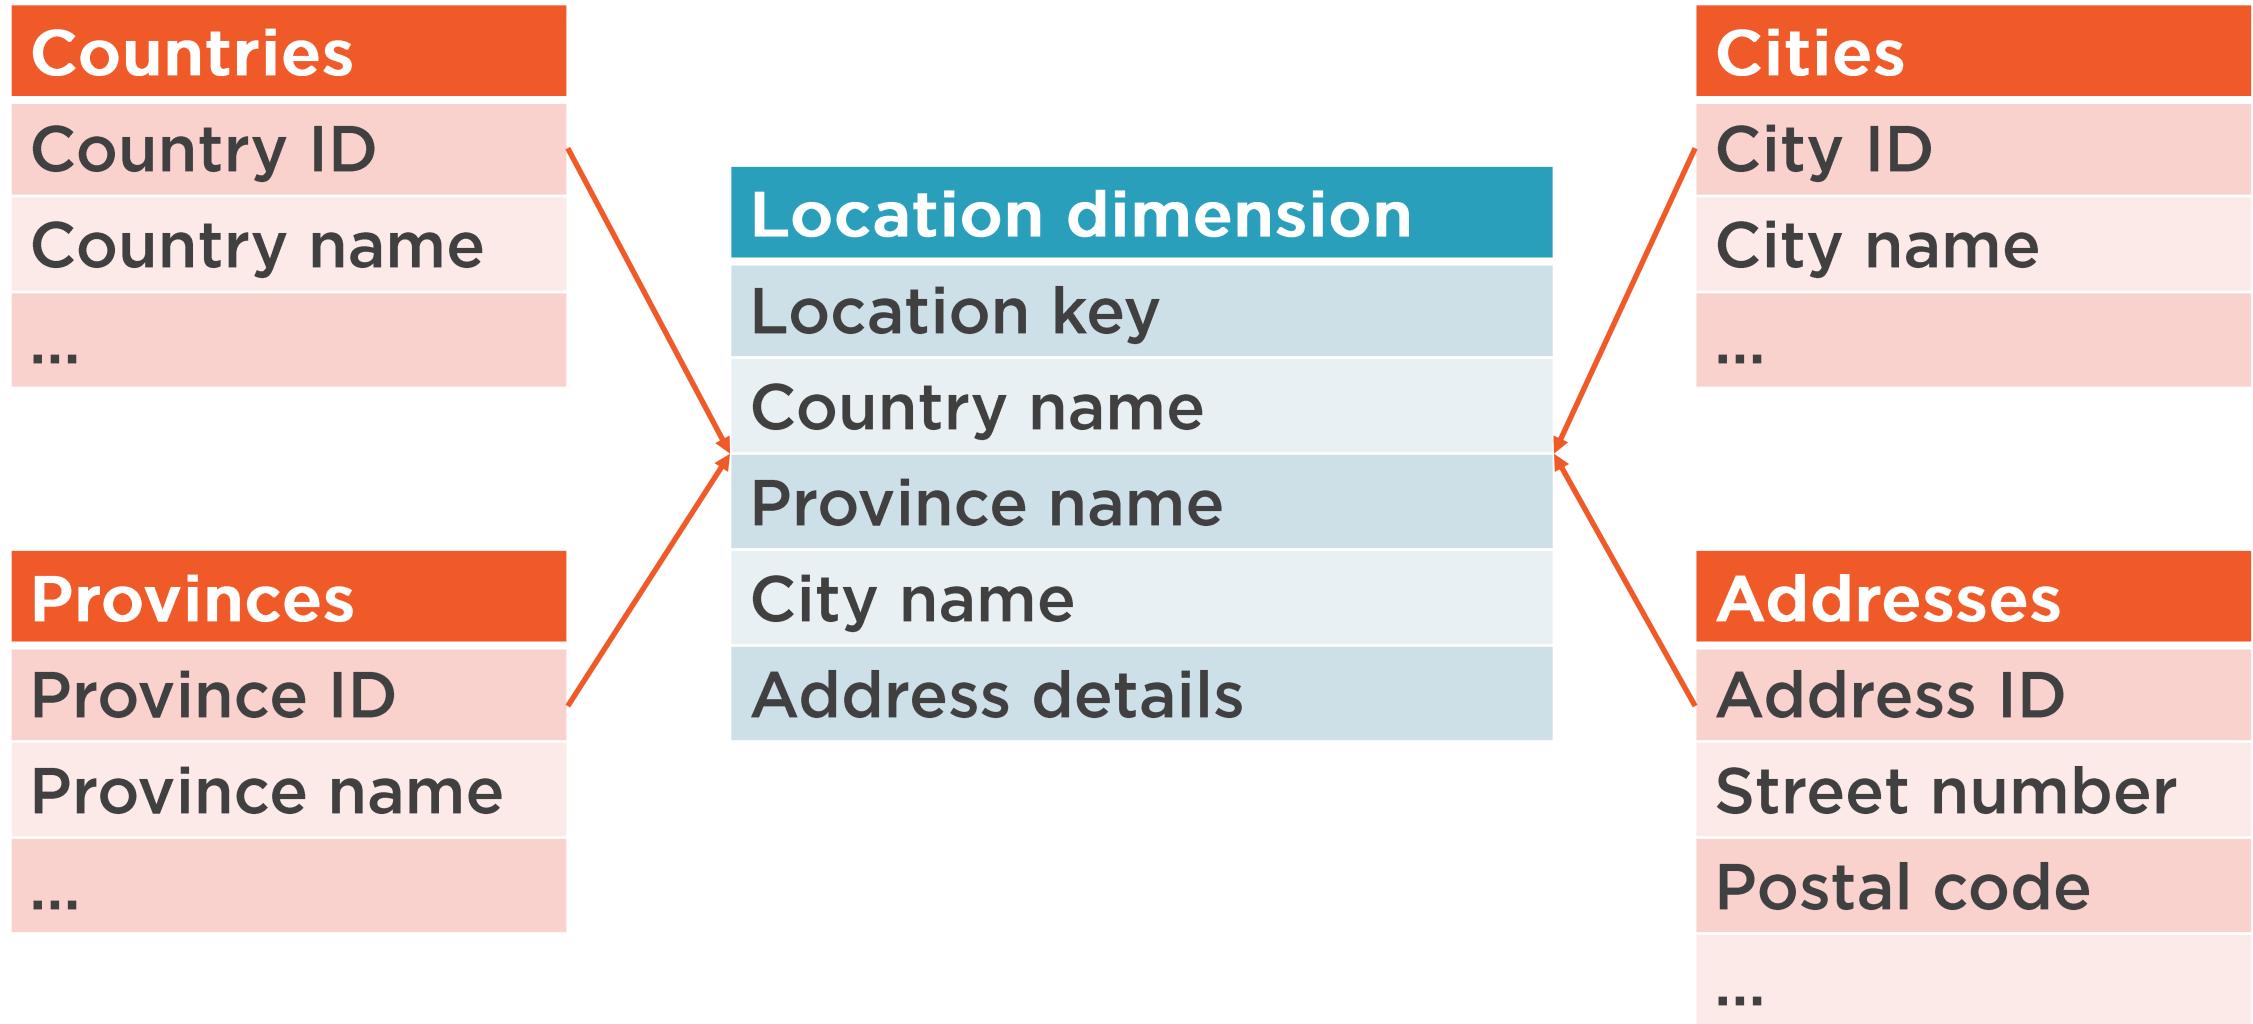


Similar to a pyramid

- Bottom level is the weakest
- Highest level is at the top



A Hierarchy in the Location Dimension



A Hierarchy in the Location Dimension

Bucharest

Vienna

New York

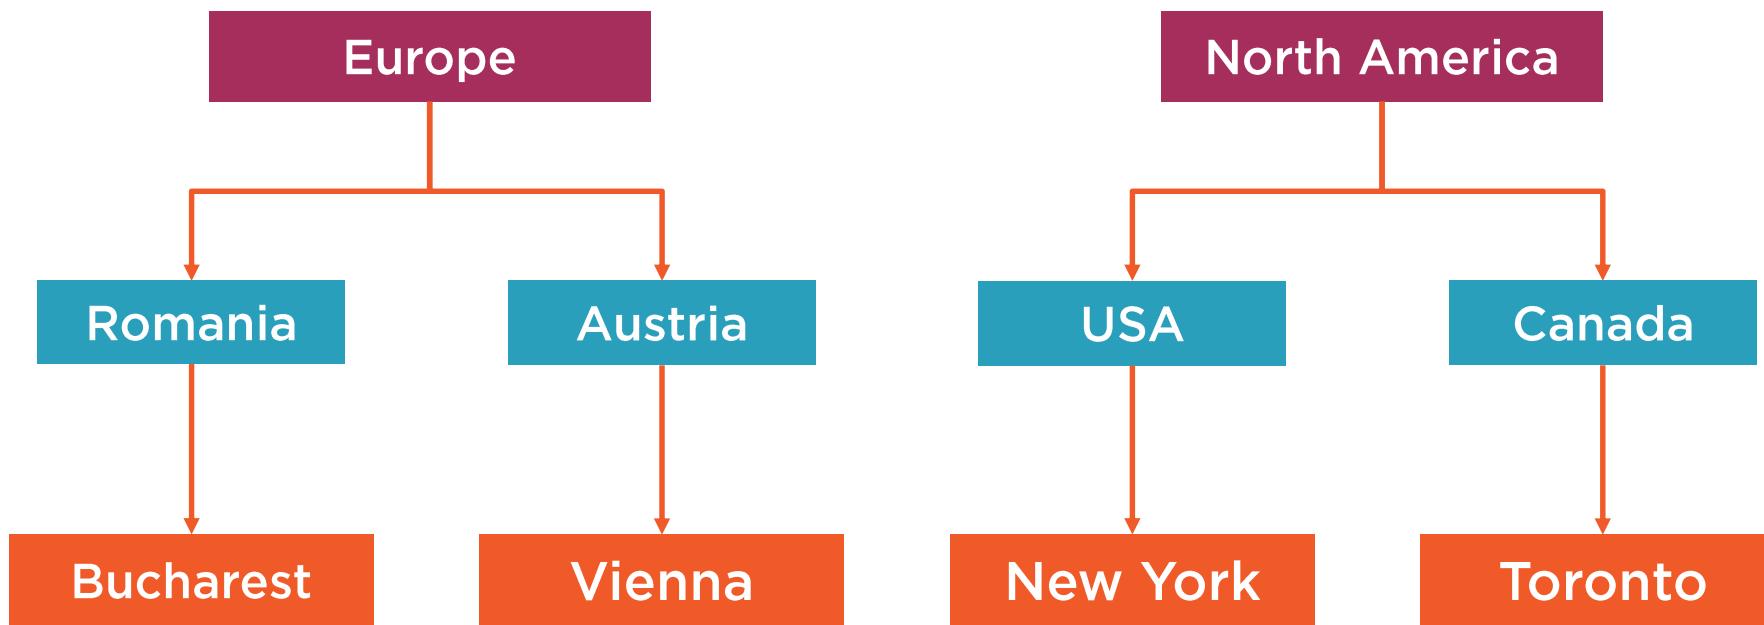
Toronto



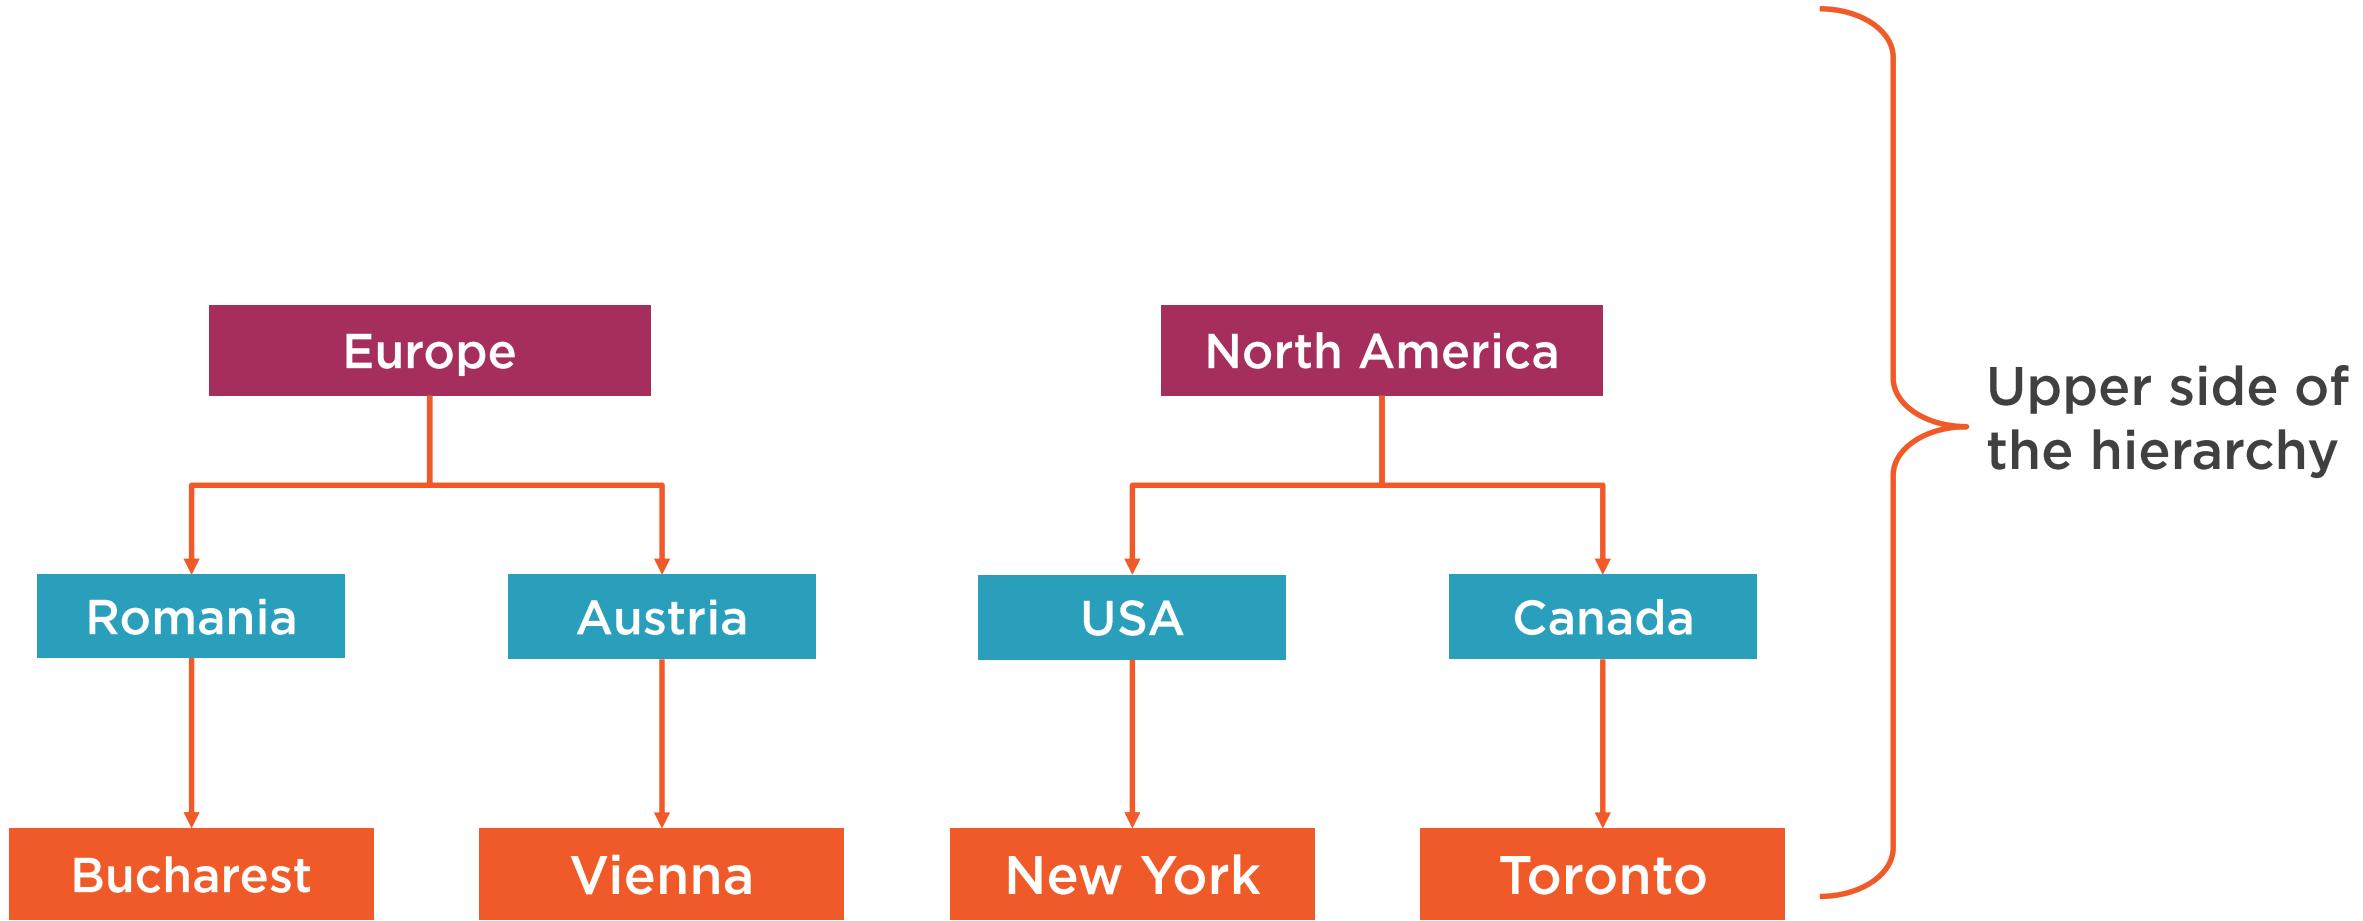
A Hierarchy in the Location Dimension



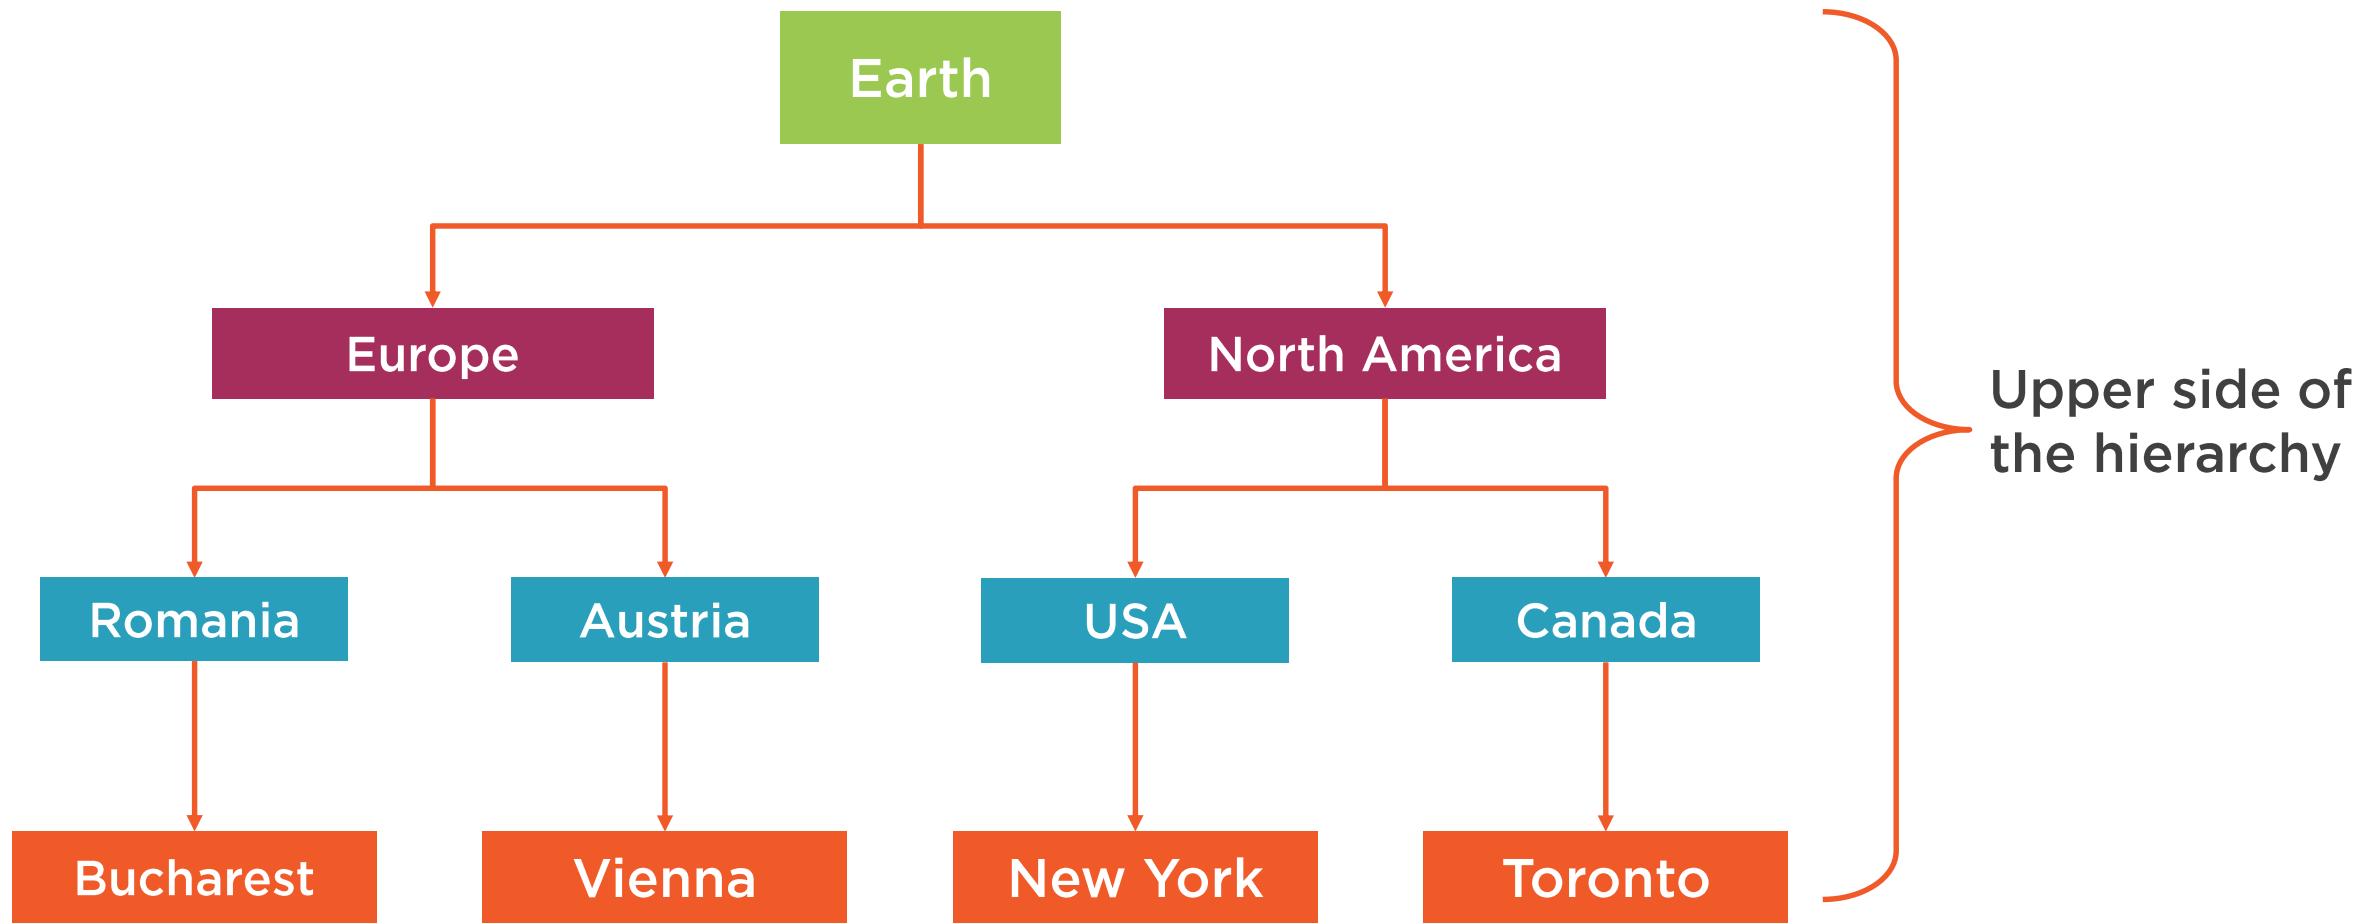
A Hierarchy in the Location Dimension



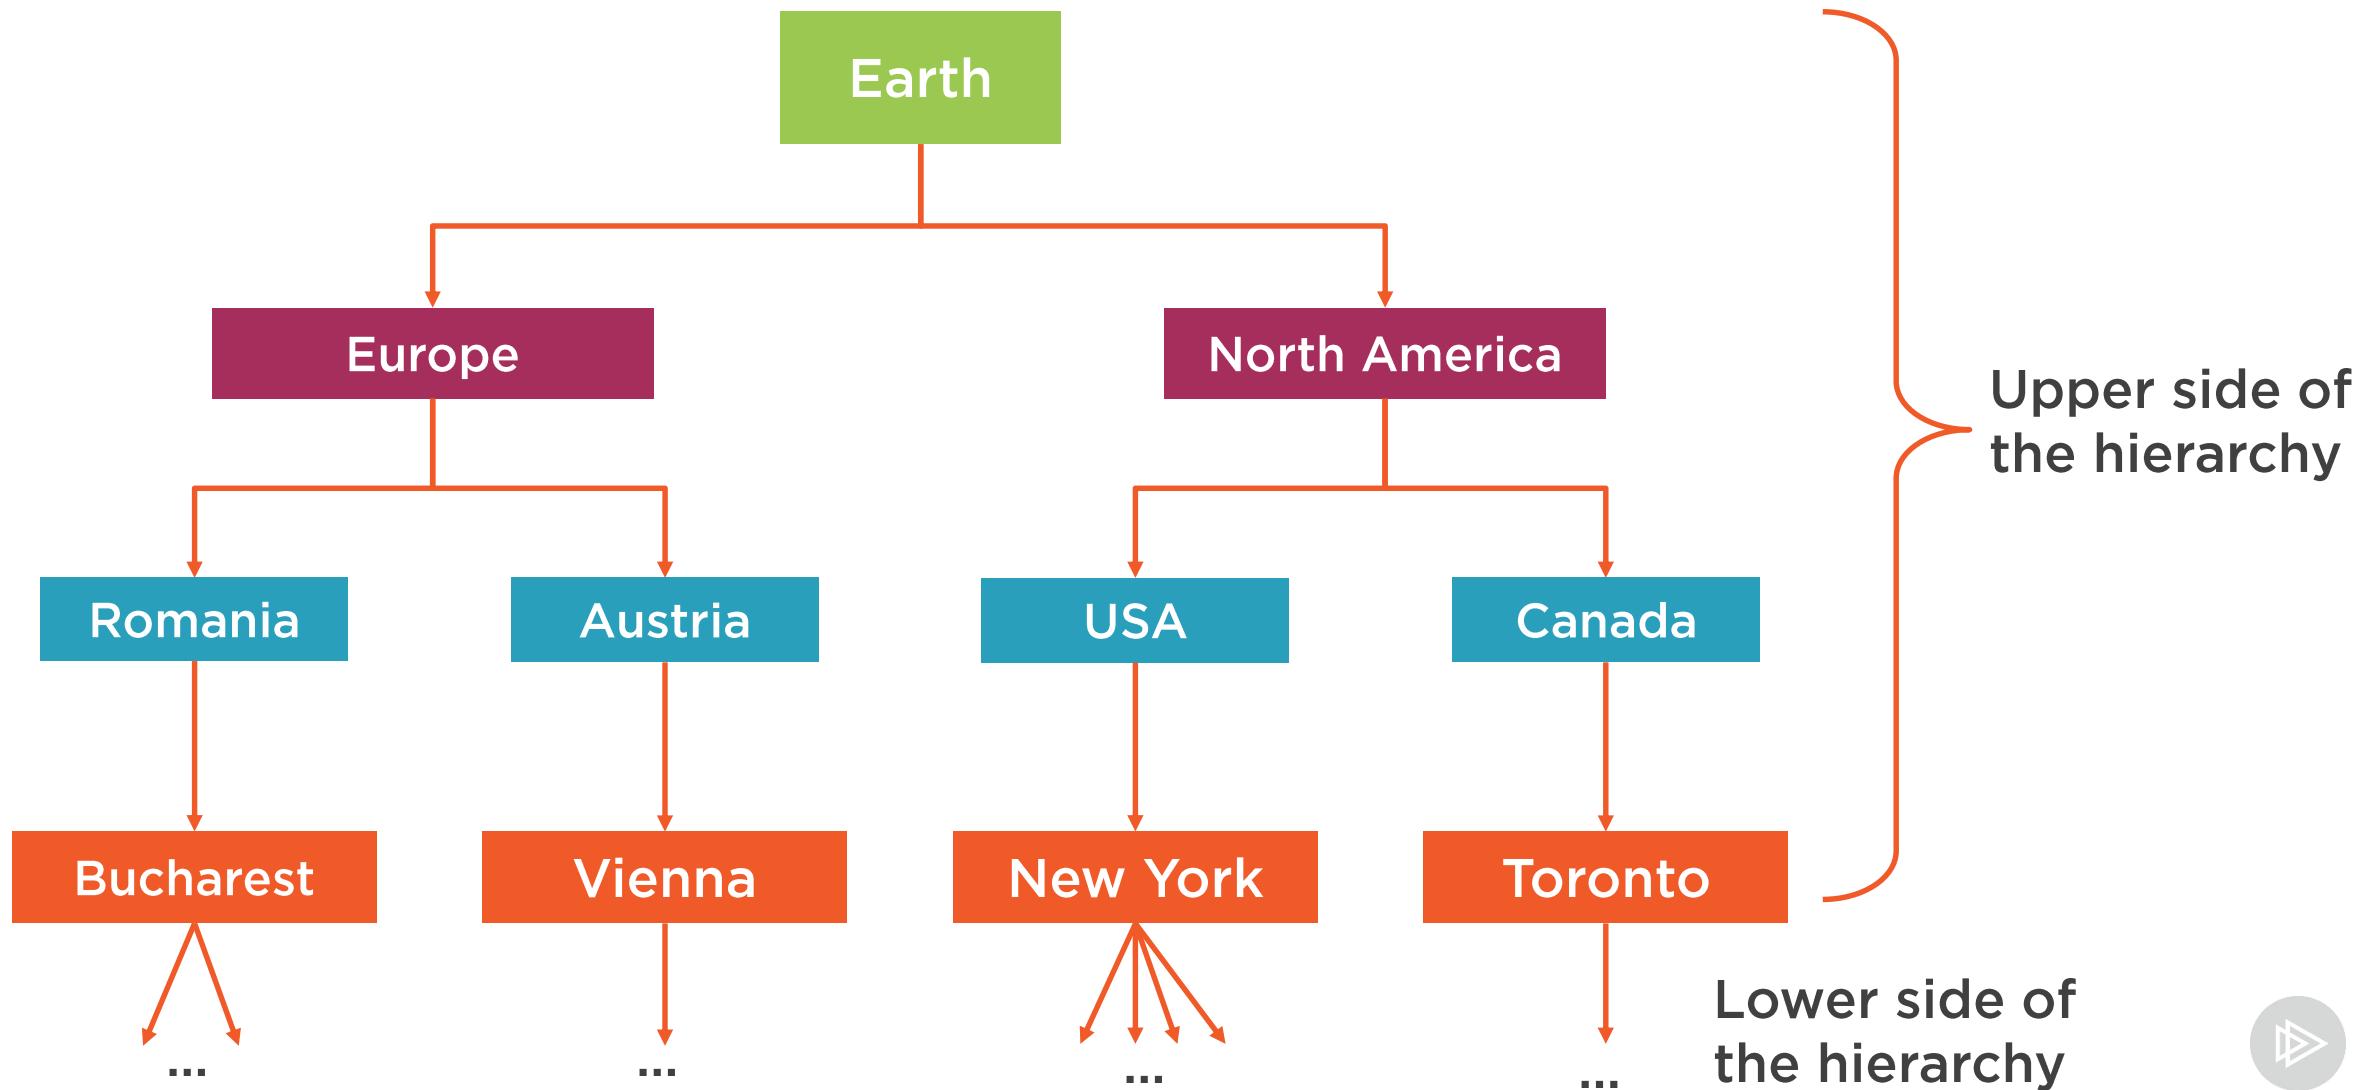
A Hierarchy in the Location Dimension



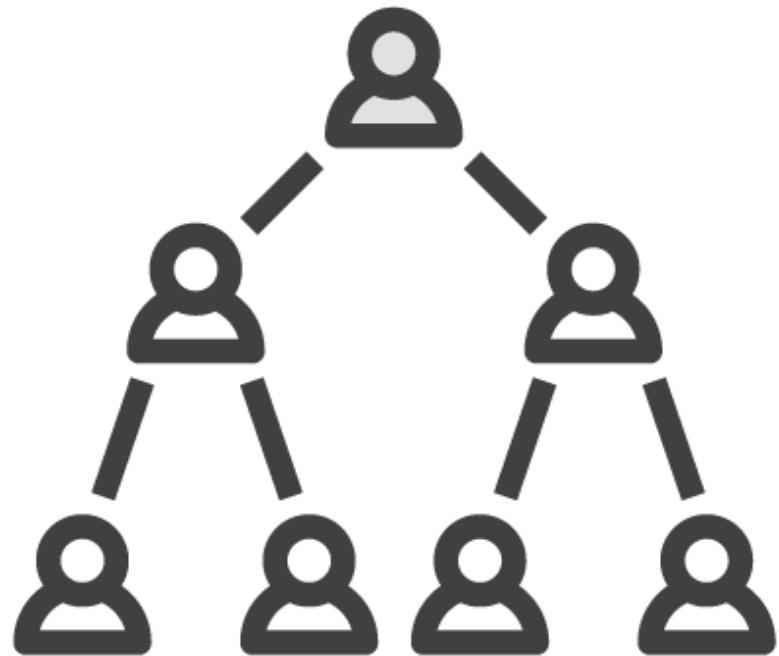
A Hierarchy in the Location Dimension



A Hierarchy in the Location Dimension



What Is a Hierarchy?



Data structure

- Attributes of a dimension are organized together
- Going down the hierarchy nodes -> more details
- Going up the hierarchy -> summarized data



Advantages of Using Hierarchies

Structure

- Important dimensions have plenty attributes
- Analyzing important data can become overwhelming
- Hierarchies provide order in the data

Multiple perspectives

Merchandise hierarchy



Advantages of Using Hierarchies

Structure

- Important dimensions have plenty attributes
- Analyzing important data can become overwhelming
- Hierarchies provide order in the data

Multiple perspectives

Merchandise hierarchy

- Department
 - Category
 - Subcategory
 - Product name



Advantages of Using Hierarchies

Structure

- Important dimensions have plenty attributes
- Analyzing important data can become overwhelming
- Hierarchies provide order in the data

Multiple perspectives

Merchandise hierarchy

- Department
 - Category
 - Subcategory
 - Product name

Packaging hierarchy



Advantages of Using Hierarchies

Structure

- Important dimensions have plenty attributes
- Analyzing important data can become overwhelming
- Hierarchies provide order in the data

Multiple perspectives

Merchandise hierarchy

- Department
- Category
- Subcategory
- Product name

Packaging hierarchy

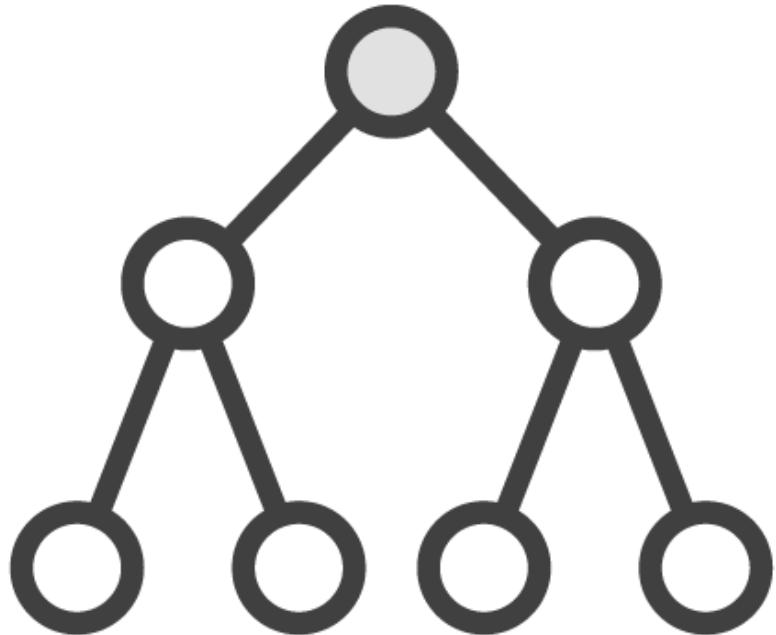
- Package type
- Package size
- Product name

Zoom in/zoom out

- Visualize summarized data and detailed data
- Go as deep as the business requires it



Types of Hierarchies



Fixed-depth (balanced)

- Fixed number of levels
- Easy to create and work with

Variable-depth (unbalanced or ragged)

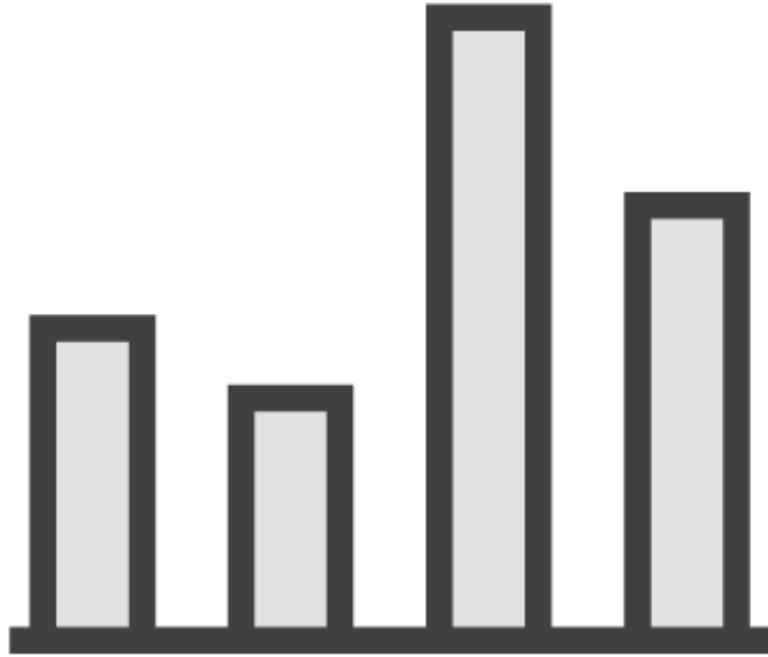
- Uneven number of levels
- Creating them is more complex task



Drilling Down a Hierarchy



Data Warehouse Analysis



Minimum requirements for doing data warehouse analysis

- One fact table
- One dimension table

Example: sales per product report

- Sales fact
- Product dimension



Example of Drilling Down

Sales fact

Store key

Date key

Product key

Employee key

Customer key

...

Transaction

Unit price

Quantity

Amount



Example of Drilling Down

Sales fact	Product dimension
Store key	Product key
Date key	Product name
Product key	Department
Employee key	Category
Customer key	Subcategory
...	Package size
Transaction #	Package type
Unit price	Description
Quantity	Unit of measure
Amount	...



Example of Drilling Down

Sales fact
Store key
Date key
Product key
Employee key
Customer key
...
Transaction #
Unit price
Quantity
Amount

Product dimension
Product key
Product name
Department
Category
Subcategory
Package size
Package type
Description
Unit of measure
...



Merchandise hierarchy

- Department
 - Category
 - Subcategory
 - Product name



Example of Drilling Down

Sales

80.000

Department

Sales

Bakery	Sweets	Beverages
25.000	40.000	15.000



Example of Drilling Down

Department
Category
Sales

Bakery			Sweets			Beverages	
Bread	Pie	Croissant	Cookie	Cake	Candy	Juice	Tea
7.000	13.000	5.000	8.700	6.300	5.000	7.000	6.600



Example of Drilling Down

Department

Category

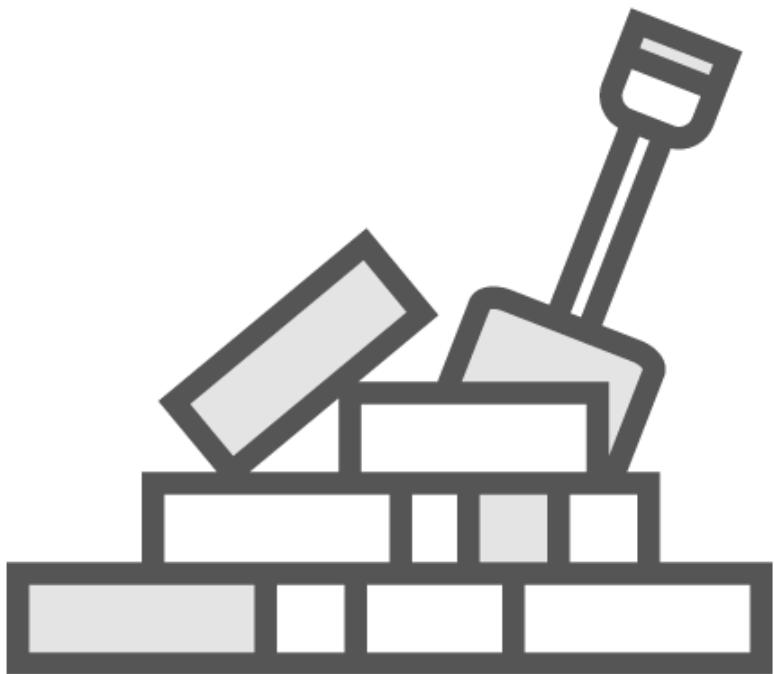
Subcategory

Sales

Bakery							
Bread			Pie			Croissant	
Baguette	Pita	Banana br.	Cream	Fruit	Custard	Sweet	Salty
2.000	3.000	2.000	2.500	2.200	8.300	3.000	2.000



Drilling Down - Summary



Drilling down

- Adding another member of the dimension to the report
- The member doesn't need to be part of a hierarchy

Drilling up/rolling up

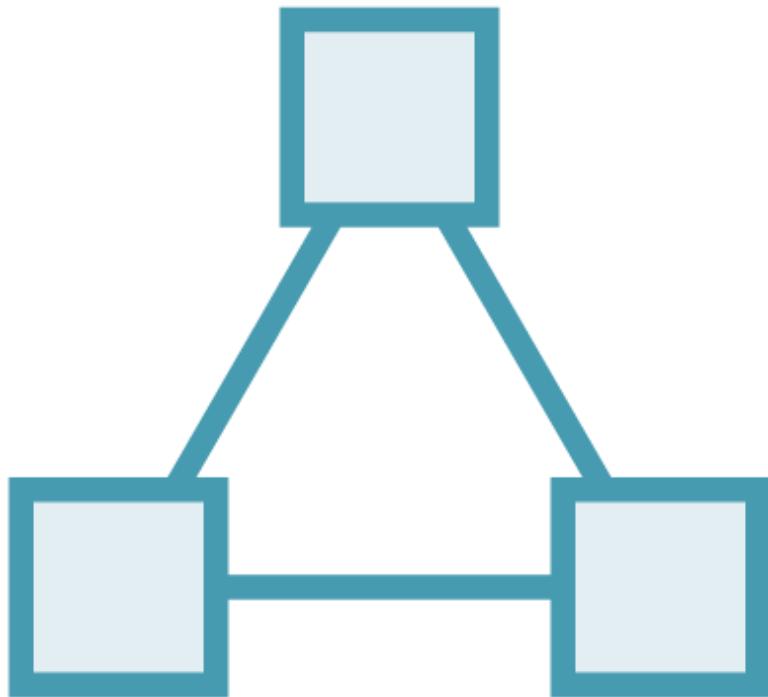
- Taking out an attribute of a dimension from a report



Fixed-depth Positional Hierarchies



Fixed-depth Positional Hierarchies



The number of levels is known upfront

The levels are attributes in the dimension table

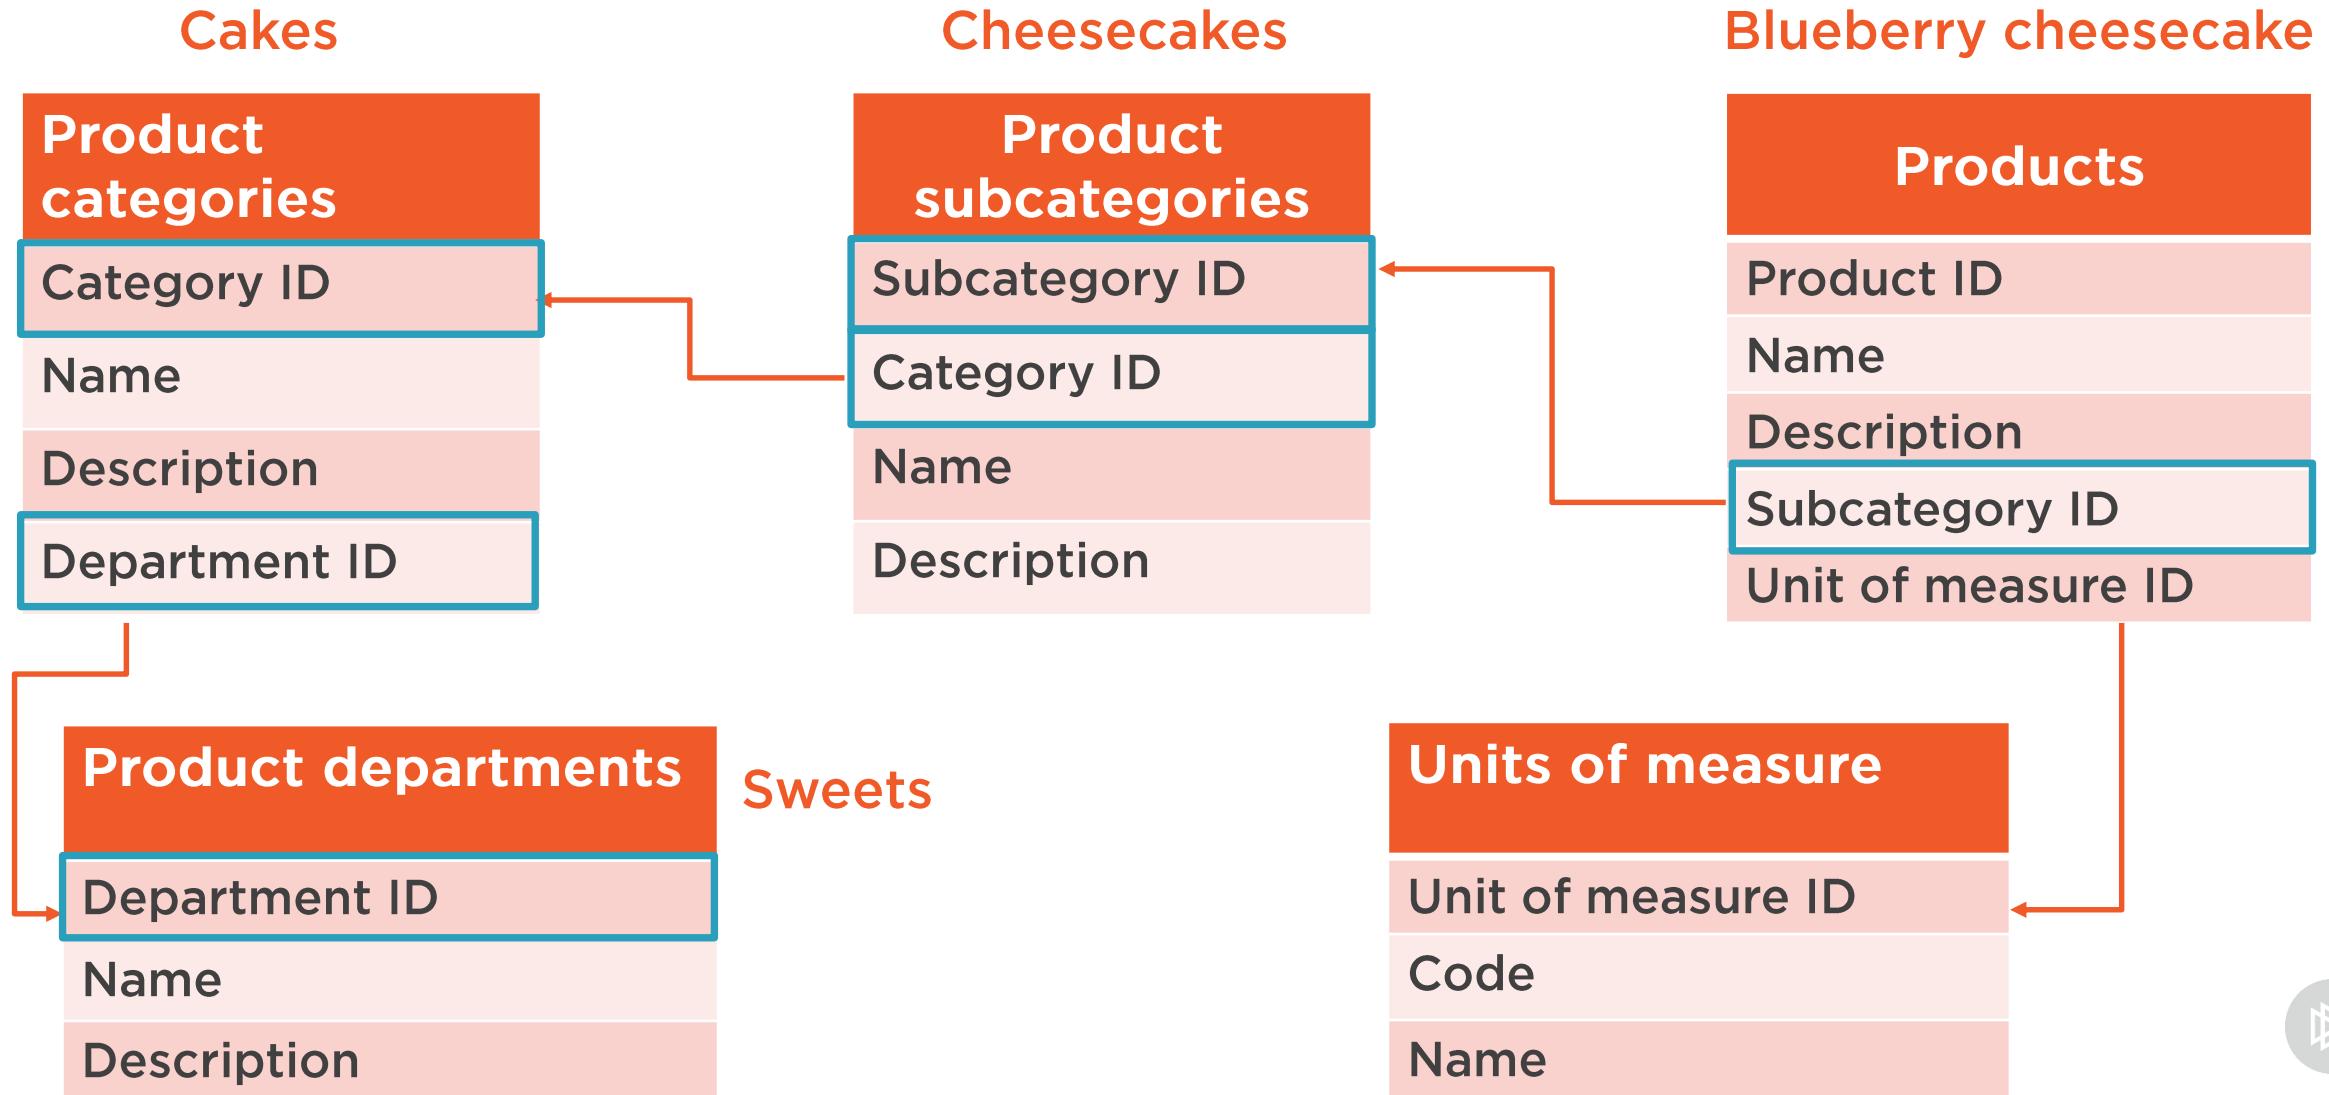
It is a series of many-to-one relationships

Advantages

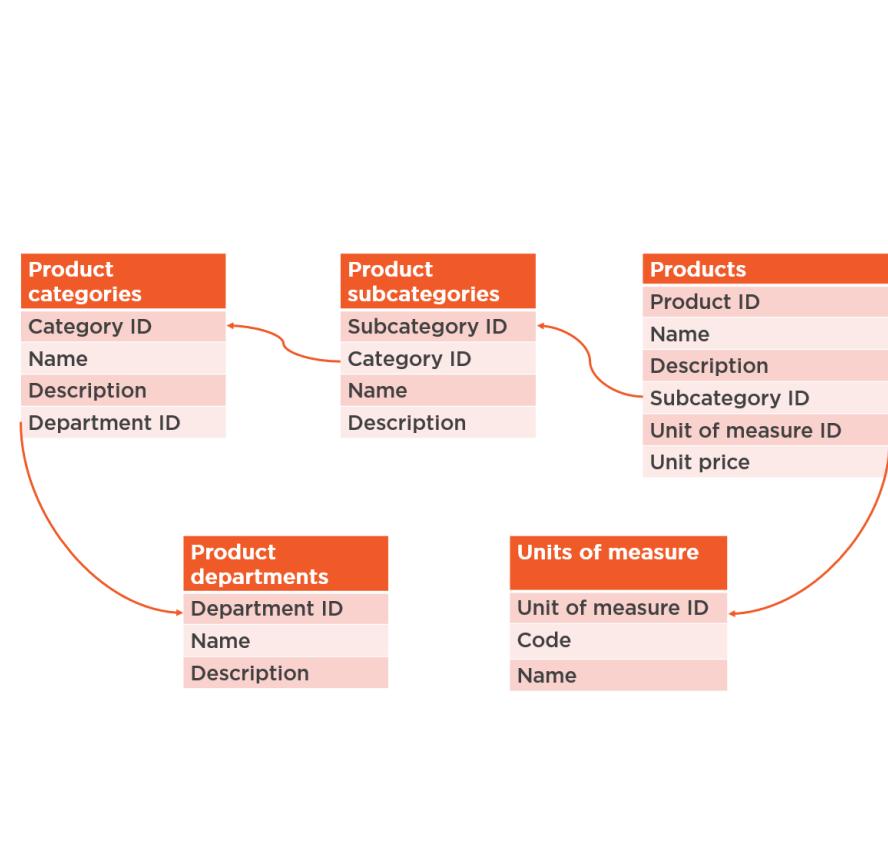
- Easy to navigate
- Offers predictable results
- No impact on performance



Product Information from Different Tables



Product Information from the Data Warehouse



Keep in mind:

- Don't create many snowflake designs
- Most relationships should be from fact tables to the dimensions



Creating the Merchandise Hierarchy

Department

Bakery

Category

Bread

Pie

Subcategory

Baguette

Banana
bread

Custard

Cream

Product

Le Petit
Francais

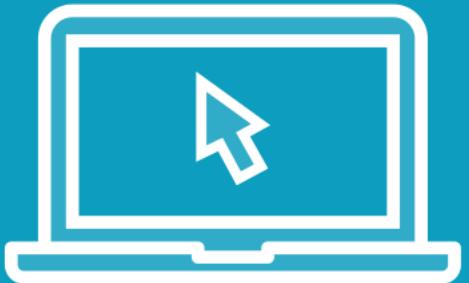
The Big
Banana
Bread

Eagle
Brand
Coconut
Pie

Eagle
Brand
Banana
Cream Pie



Demo



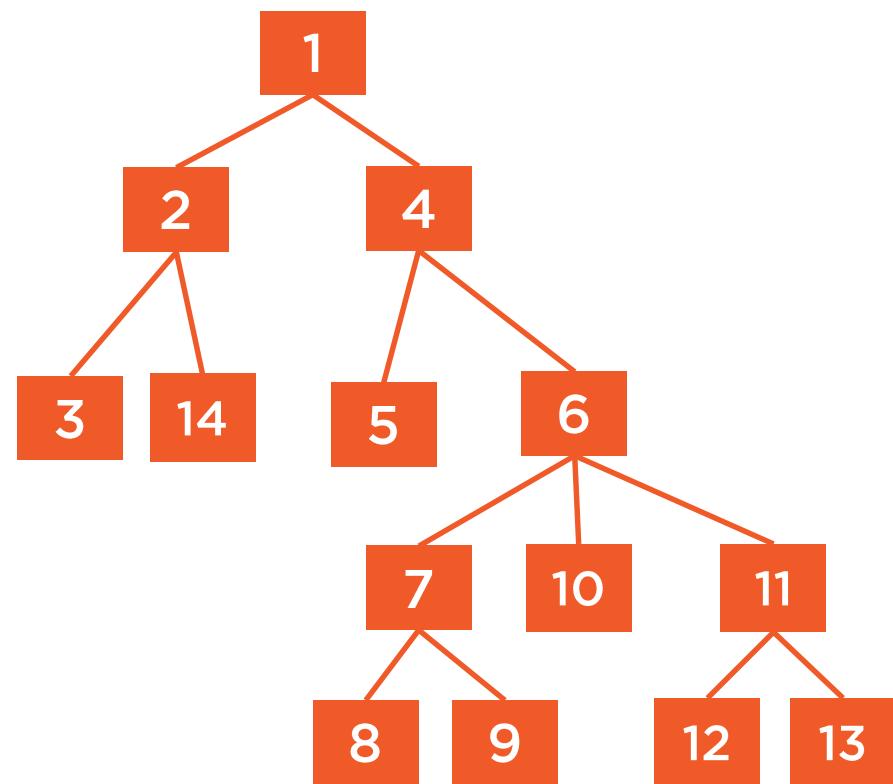
Creating a fixed-depth positional hierarchy



Variable-depth Positional Hierarchies



Characteristics of Variable-depth Hierarchies



The number of levels is not known at design time

- Example: an organizational chart

Are more complex structures (compared to the fixed-depth hierarchies)

Should be used with moderation



Classification of Variable-depth Hierarchies



Slightly ragged

Ragged, created with a hierarchy bridge

Ragged, created with pathstring attributes



Slightly Ragged Hierarchies



- The number of levels is not known beforehand
- The range in depth is small
- Geographic hierarchies are slightly ragged

Example of a Slightly Ragged Hierarchy

Location hierarchy

- Country
 - Province (or state)
 - City
 - Neighborhood
 - Address

Examples of data with missing levels

Singapore (country and city)

Vatican (independent city-state)

Small cities, that don't have neighborhoods



Fitting Data into a Slightly Ragged Hierarchy



Step 1: Create all possible levels of the hierarchy

Step 2: Fill in the missing values per each level with:

- An expression, similar to “not applicable”
- The value of the next parent member

Step 3: Handle the ragged hierarchy as a fixed-depth one



Populating a Slightly Ragged Hierarchy

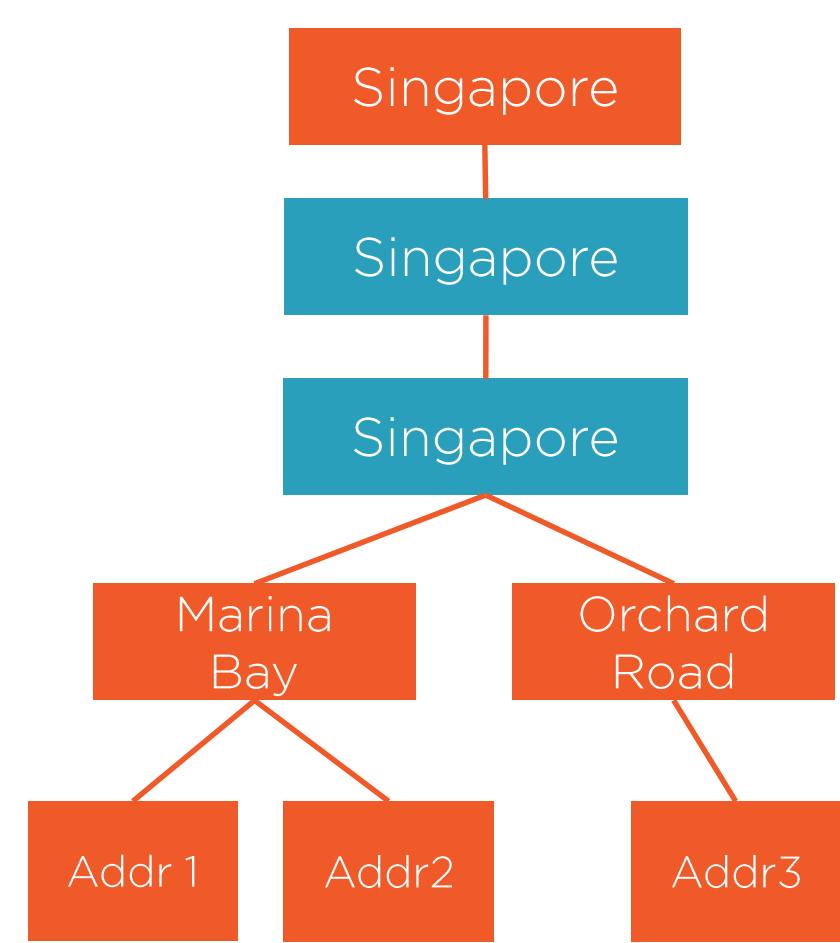
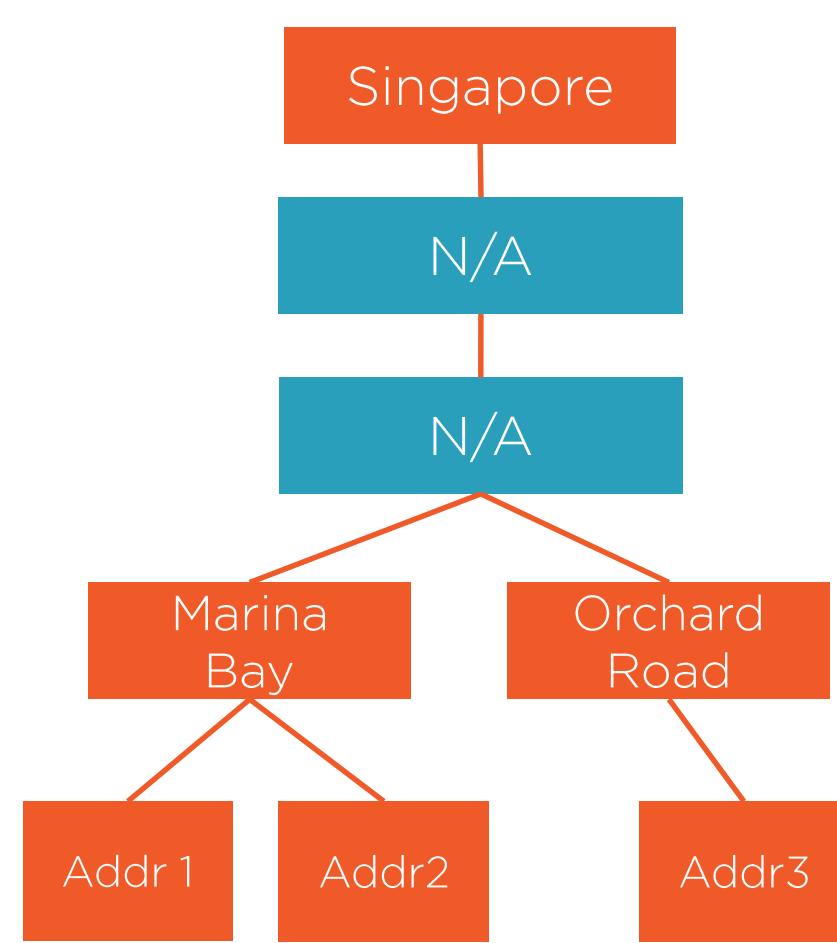
Country

↳ **Province**

↳ **City**

↳ **Neighbor**

↳ **Address**



Demo



Creating and working with a slightly ragged hierarchy

- Based on the Location dimension

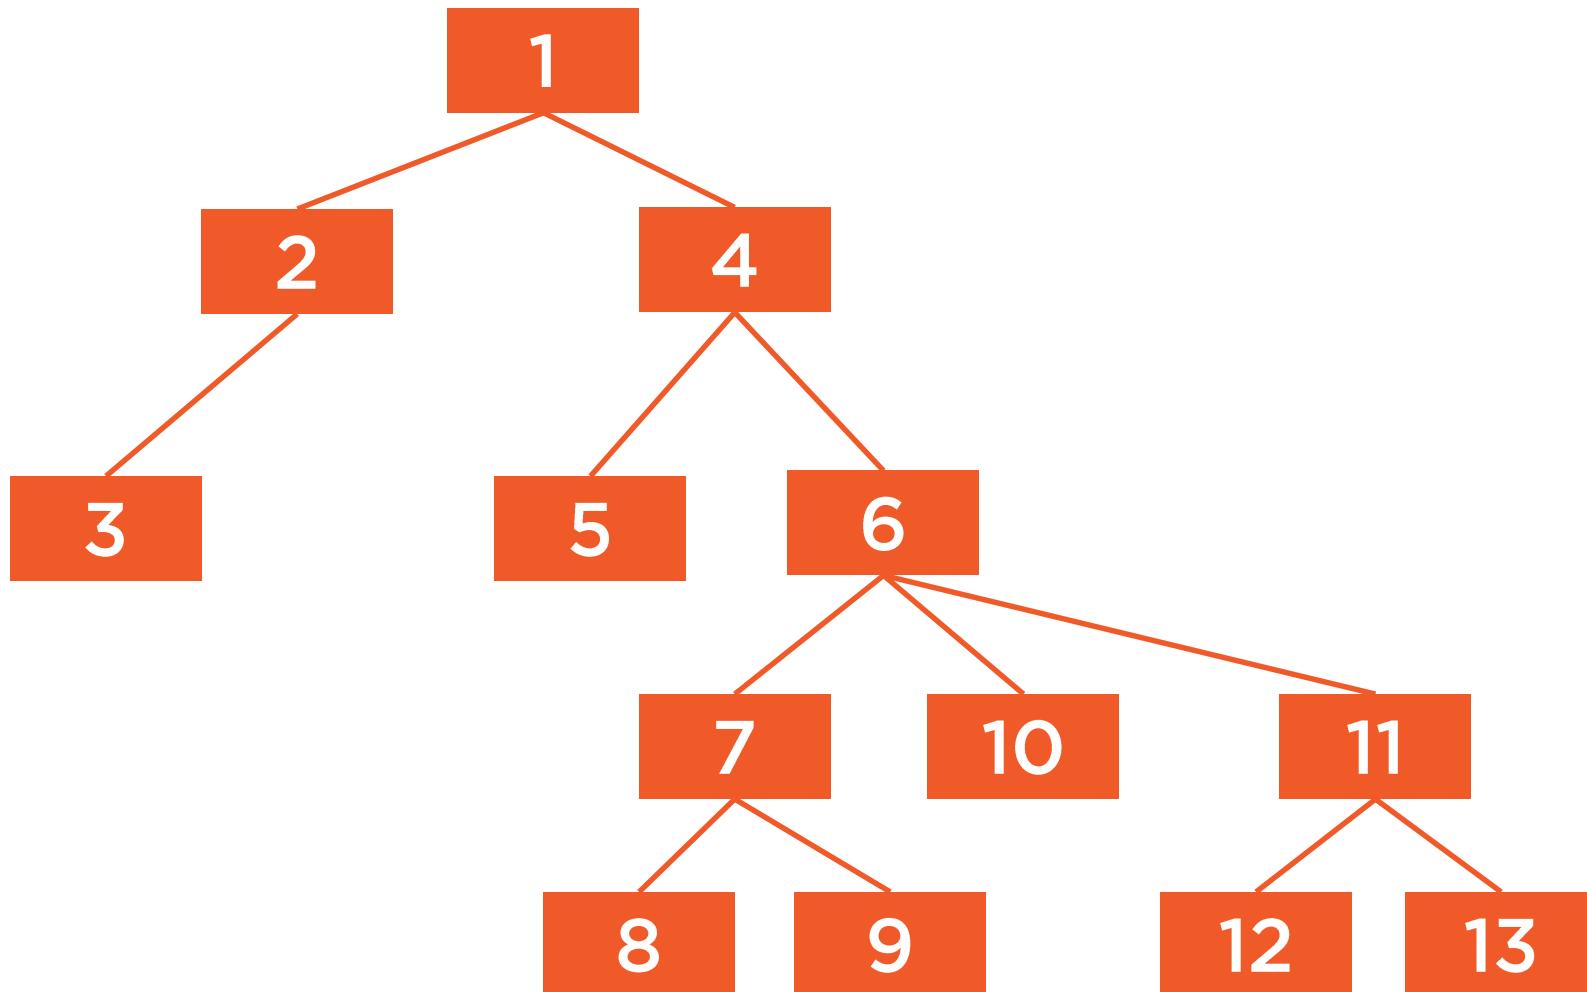


Ragged Hierarchies



Example of a Ragged Hierarchy

The organizational chart



Implementing a Ragged Hierarchy Using a Bridge



A row for each possible path in the hierarchy is stored in a table

Columns in the bridge table

- ID of the parent
- ID of the child
- Number of levels between them
- Whether the node it a top node or bottom node
- Other information relevant for analysis

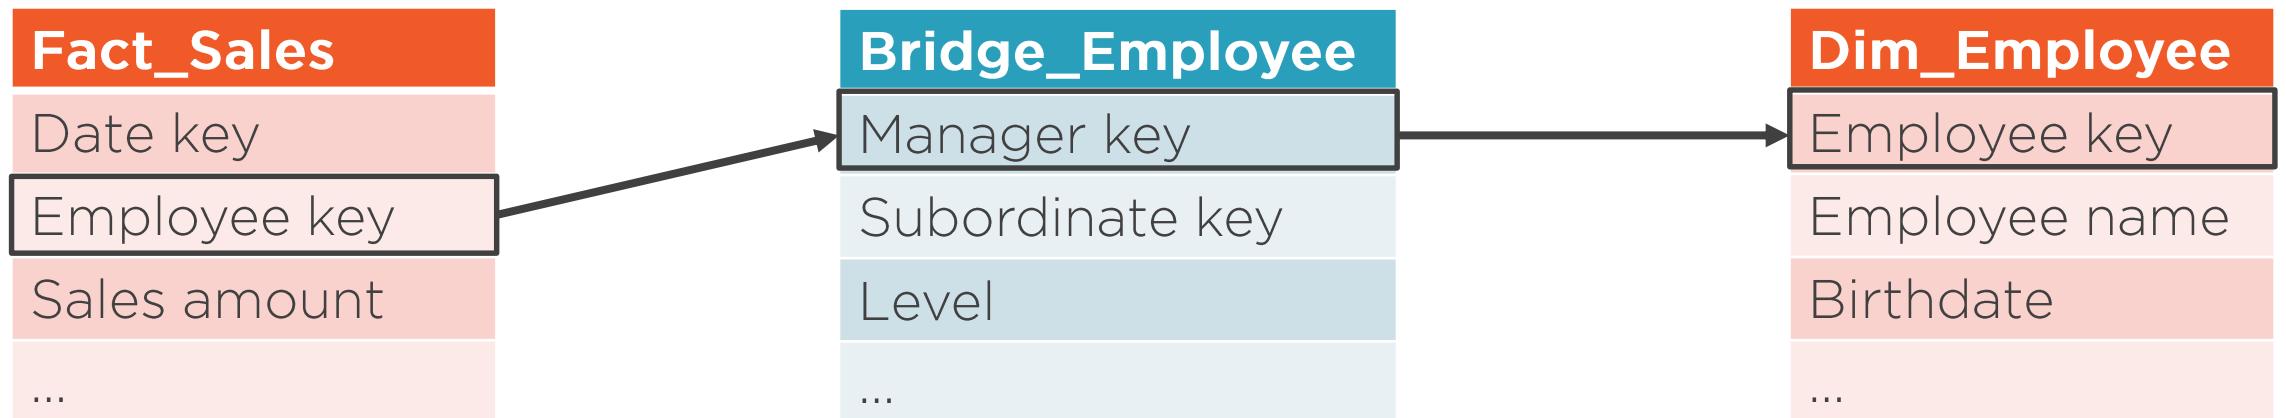
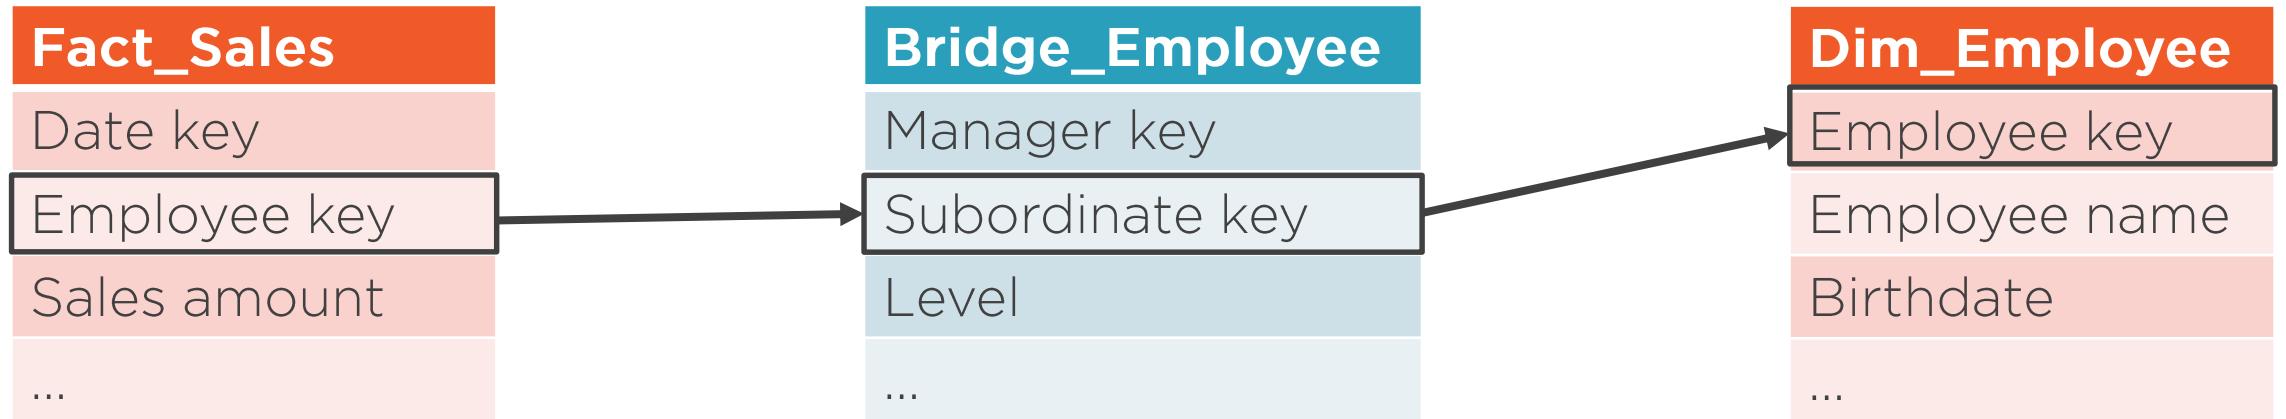


Example of a Hierarchy Bridge Table

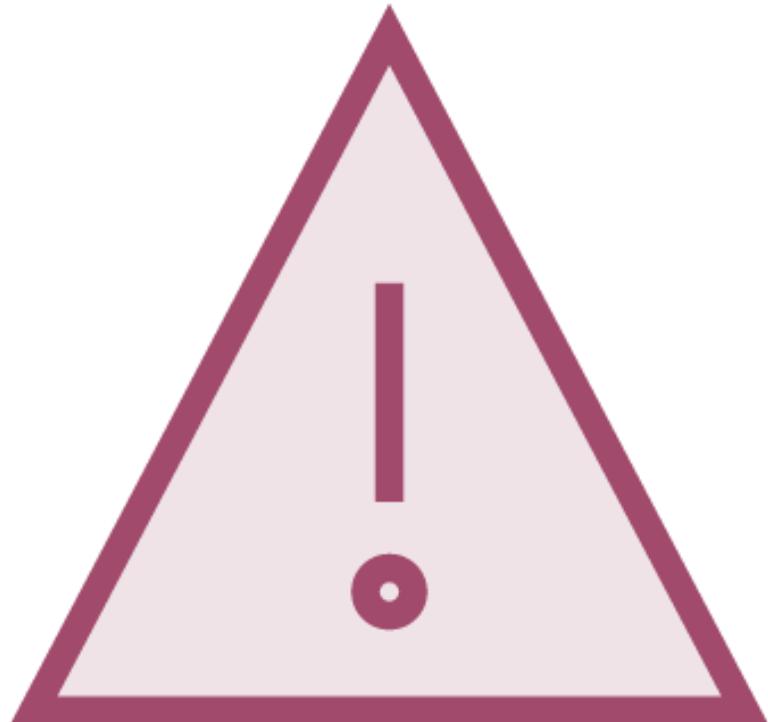
Parent key	Child Key	Level	Top flag	Bottom flag
Julia	Julia	0	Y	N
Julia	Marc	1	N	N
Julia	Theodora	2	N	Y
Julia	Greg	2	N	Y
Marc	Marc	0	Y	N
Marc	Theodora	1	N	Y
Marc	Greg	1	N	Y
Theodora	Theodora	0	Y	Y
Greg	Greg	0	Y	Y



Linking the Fact and the Dimension Table



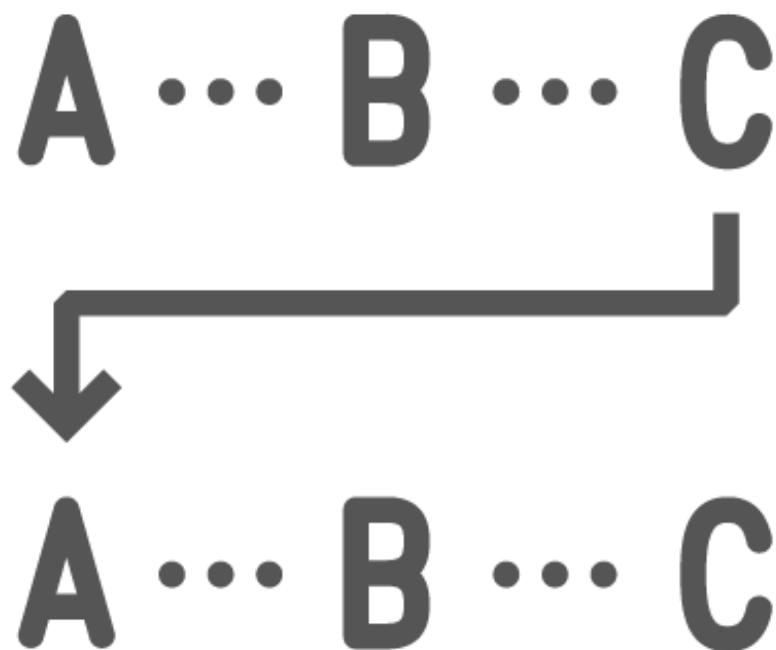
Limitations of the Hierarchy Bridge



**Can grow a lot in size
Performance impact
Difficult to work with**



Ragged Hierarchies Created with Pathstring Attributes



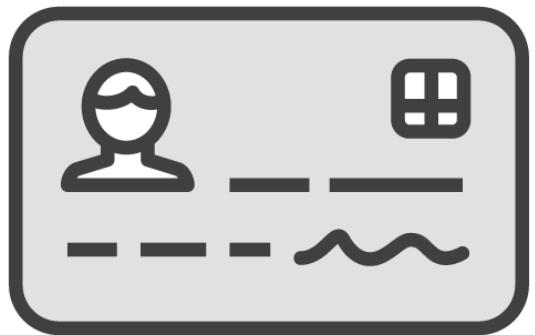
Alternative to the bridge table

The pathstring attribute:

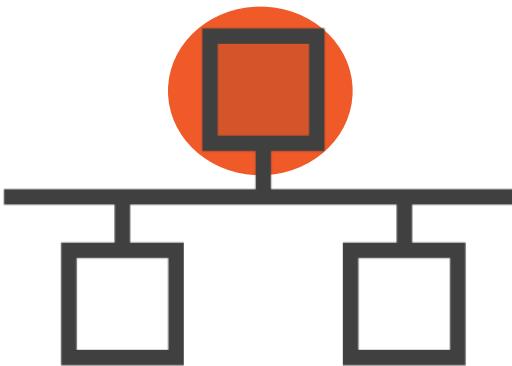
- A special attribute created in the dimension
- A string of characters
- Consists of all the parents of a member from the top of the hierarchy



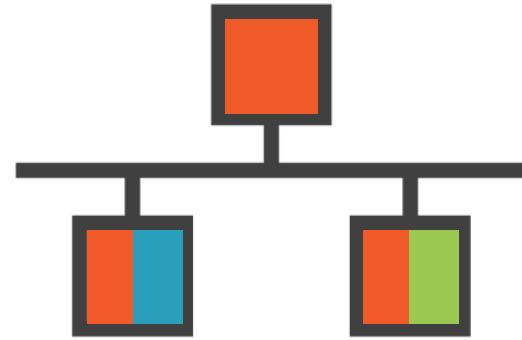
Creating a Hierarchy with Based on Pathstring Attribute



Each node is labeled with a unique value



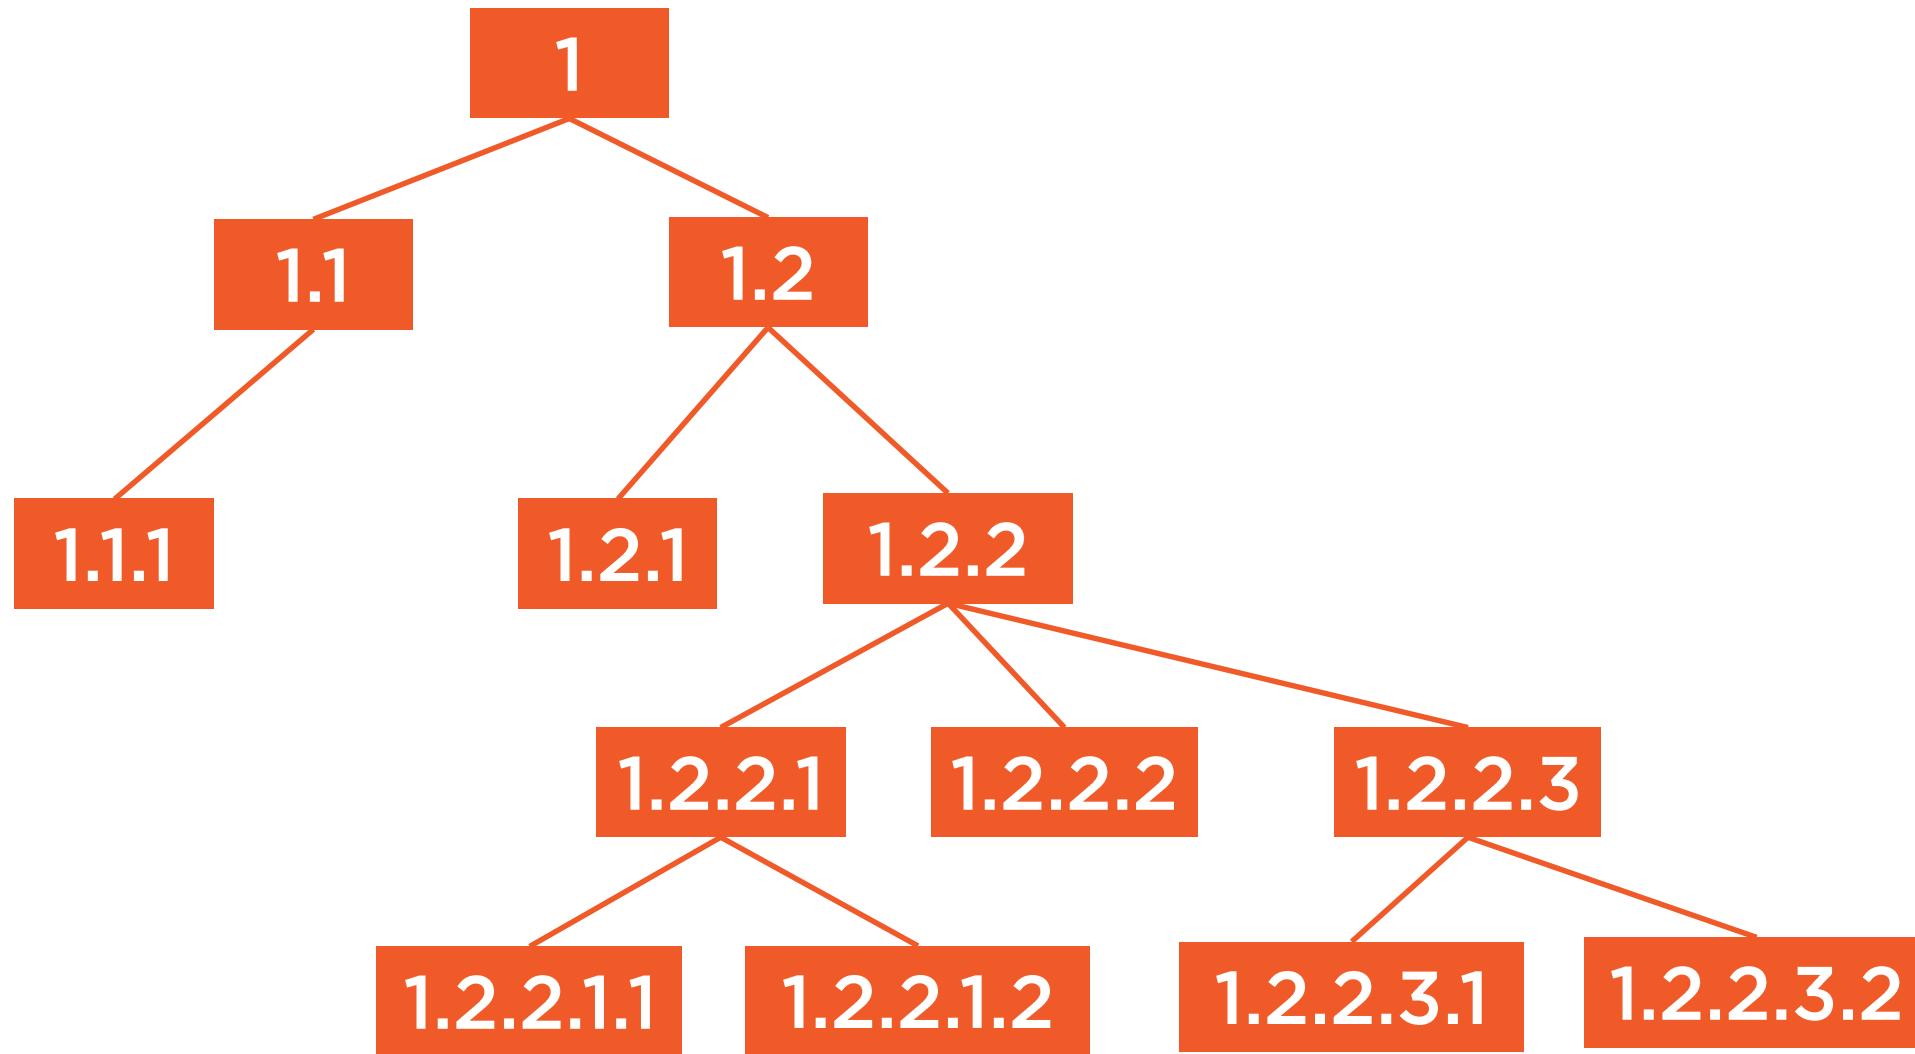
The pathstring of the root node is its unique label



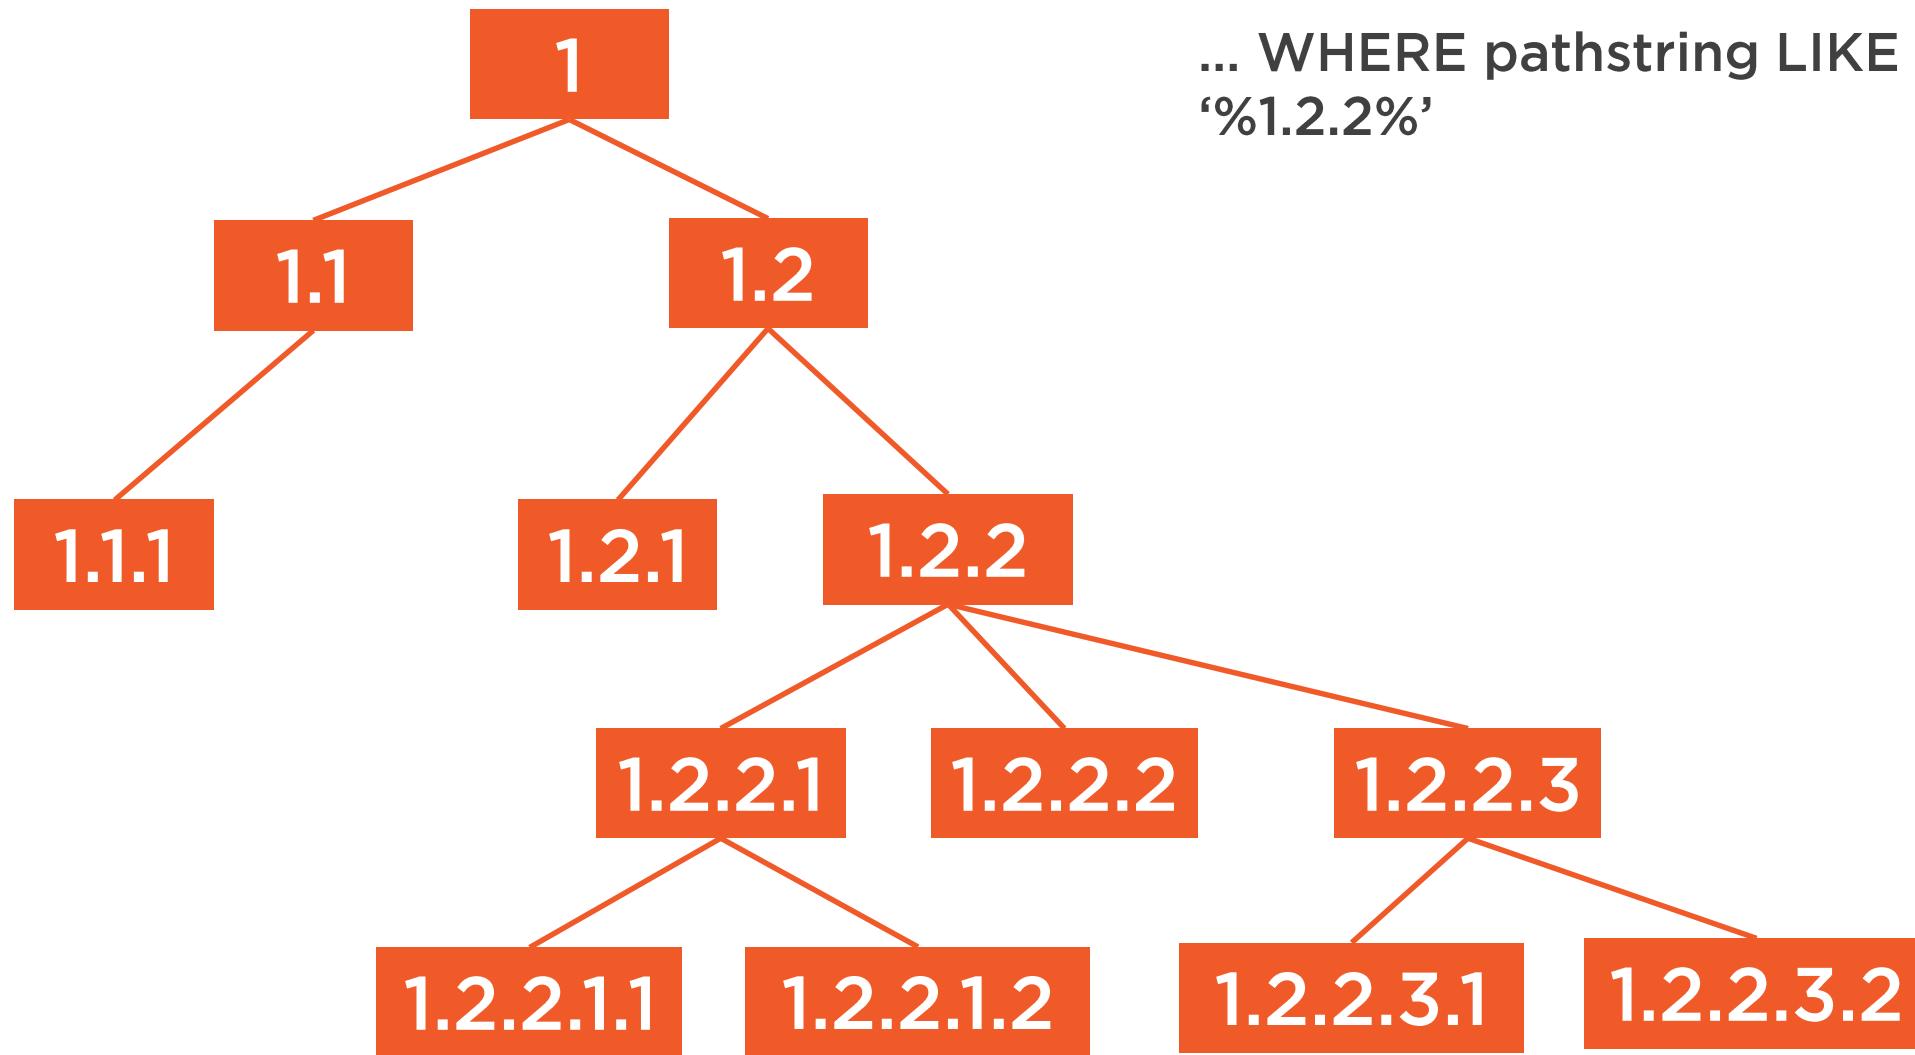
The pathstring on one level includes the pathstring of the parent



Hierarchy Based on Pathstring Attribute



Hierarchy Based on Pathstring Attribute



Limitations of the Pathstring Attribute



Vulnerable to structure changes

- If a member is moved within the organization (or added or deleted)
- Entire hierarchy must be relabeled

Complicated to use by business users

- Accessing the database directly
- Using the 'LIKE' operator in SQL queries



Summary



There is no “best solution” to this problem

A good enough solution is generated by having a clear understanding of

- The data available
- The business requirements



Summary



Hierarchy definition

- Data structures with multiple levels
- Levels are formed with attributes of a dimension

Advantages of hierarchies

- Structure the data
- Multiple perspectives of the same data
- Easily visualize summarized or detailed data

Types of hierarchies

- Fixed-depth
- Variable-depth
 - Slightly ragged
 - Ragged



Implementing Fact Tables



Ana Voicu
@ana_voicu



Overview



Structure of a fact table

The most common types of fact tables

- Transaction
- Periodic snapshot
- Accumulating snapshot

Handling nulls in fact tables

Demo: Creating and working with fact tables



Structure of a Fact Table



Examining a Fact Table's Purpose



Examining a Fact Table's Purpose

Company



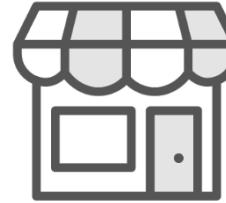
Examining a Fact Table's Purpose

Business processes



Sales

Company



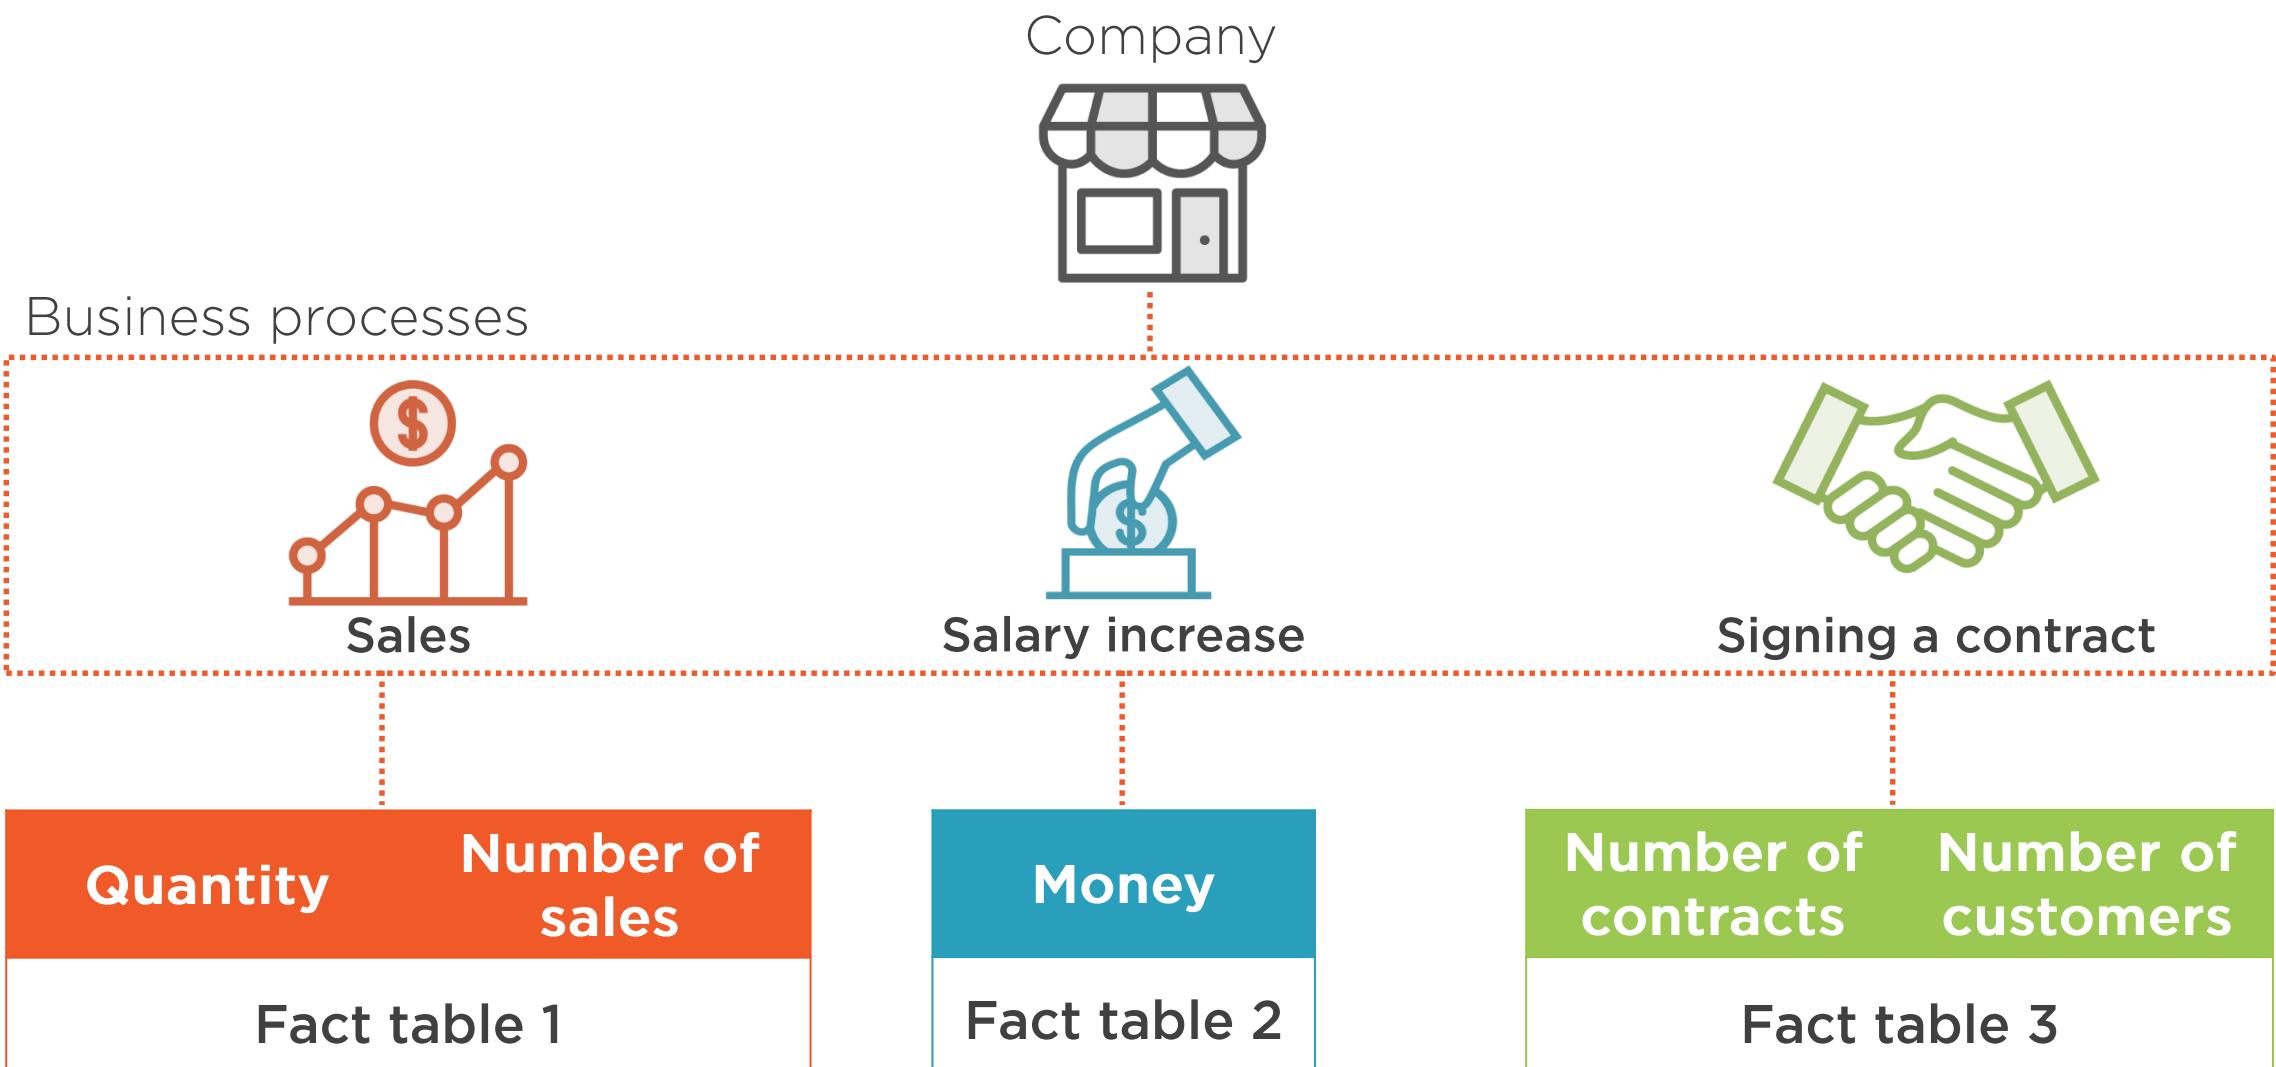
Salary increase



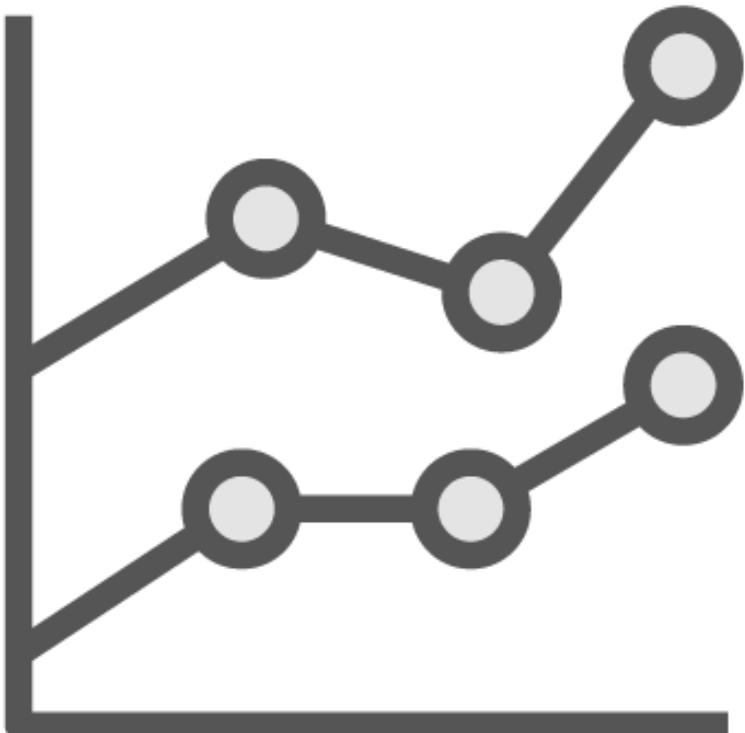
Signing a contract



Examining a Fact Table's Purpose



Fact Table Characteristics



Used with dimensions for:

- Reporting purposes
- Measuring performance
- Analyzing trends in time

Contain a lot of measurements

They tend to grow very fast

Performance is an important design factor



Walmart Case Study

As of April 2019

Number of stores 11,300

Annual revenue in previous fiscal year 514 billion dollars

Days in a year 365

Avg. revenue per day 1.4 billion dollars

Avg. revenue per day 124,000 dollars per store

Avg. price of a banana 19 cents



Walmart Case Study

As of April 2019

Number of stores 11,300

Annual revenue in previous fiscal year 514 billion dollars

Days in a year 365

Avg. revenue per day 1.4 billion dollars

Avg. revenue per day 124,000 dollars per store

Avg. price of a banana 19 cents



Walmart Case Study

As of April 2019

Number of stores 11.300

Annual revenue in previous fiscal year 514 billion dollars

Days in a year 365

Avg. revenue per day 1.4 billion dollars

Avg. revenue per day per store 124.000 dollars

Avg. price of a banana 19 cents

How many transactions are recorded per day?

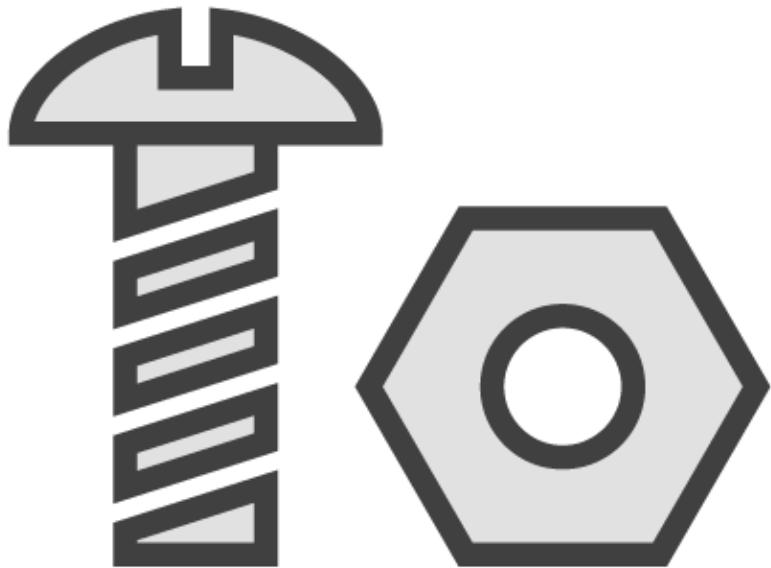
What are the best sold products from each store?

Data warehouse fact tables

- Sales fact table
- Profit and loss
- Periodic revenue per period



Main Elements of a Fact Table



Facts

- Columns storing the measurements of the process

Primary key

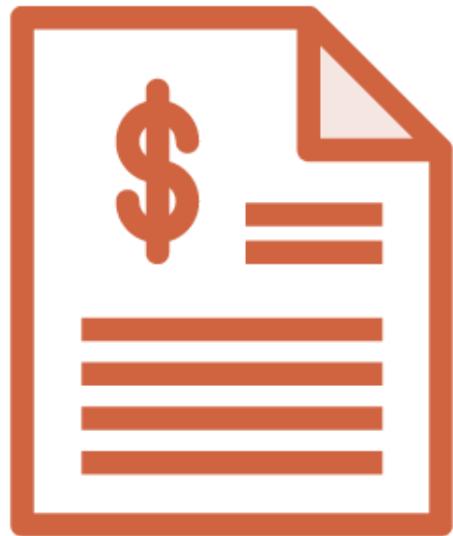
- Uniquely identifies rows

Foreign keys

- Linking the fact and the dimension tables



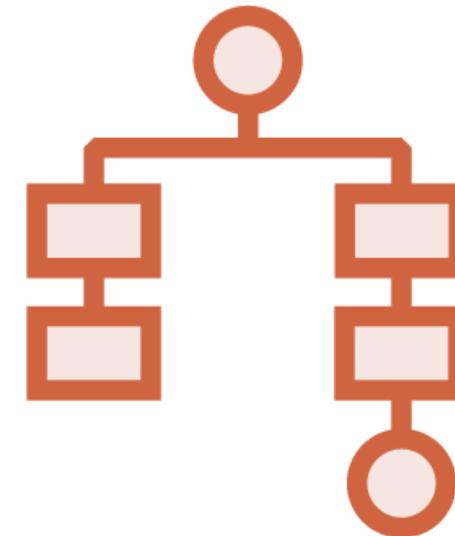
Steps for Identifying Facts



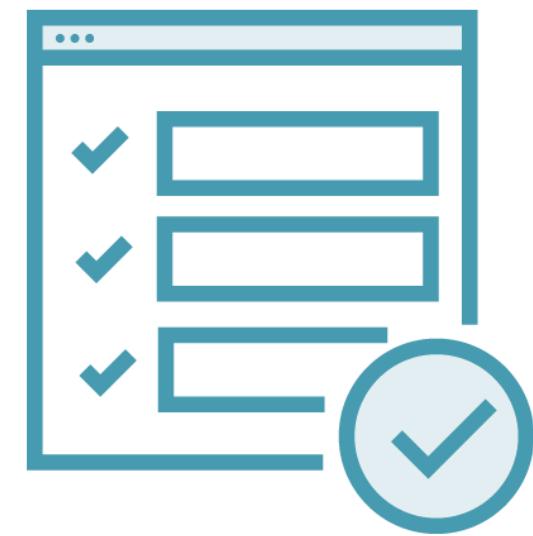
Analyze what is important in a business process



Identify multiple facts



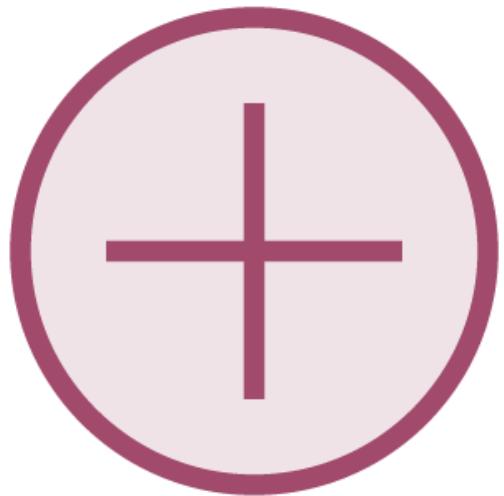
Group facts by granularity



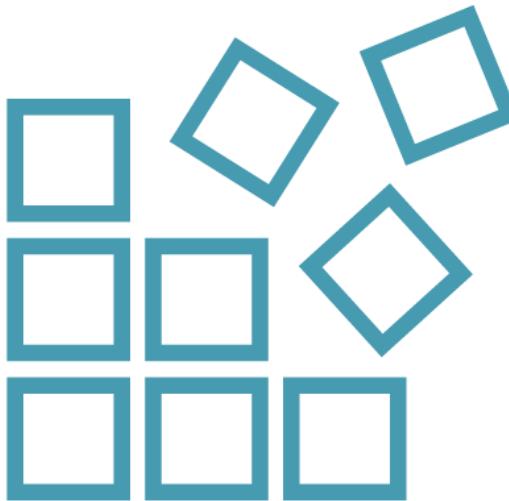
Double-check the facts



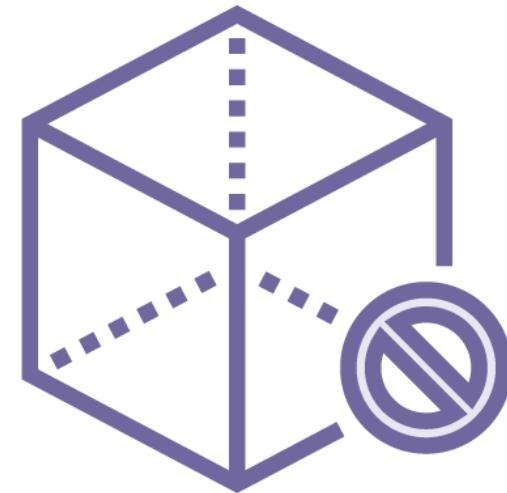
Types of Facts



Additive

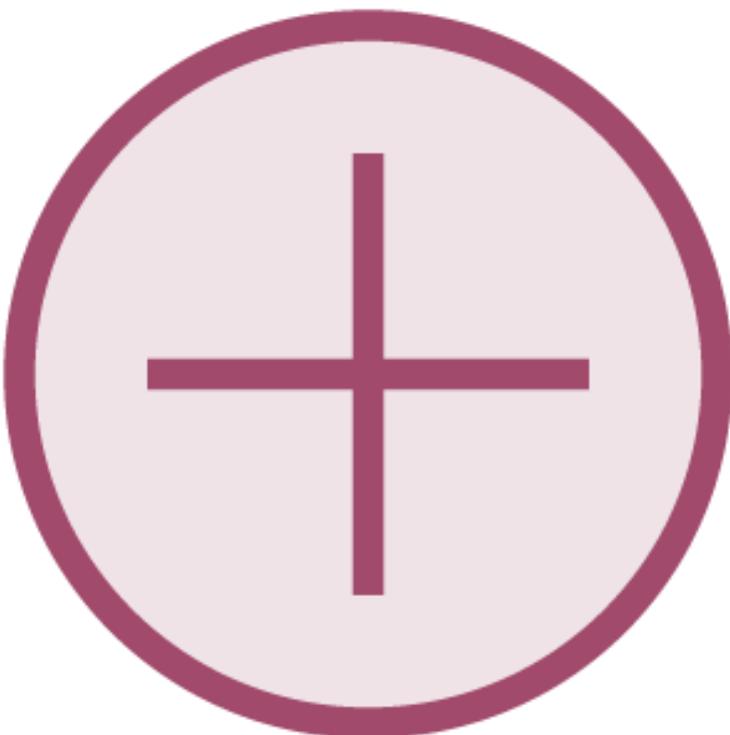


Semi-additive



Non-additive





Additive Facts

Can be summarized across all dimensions from a fact

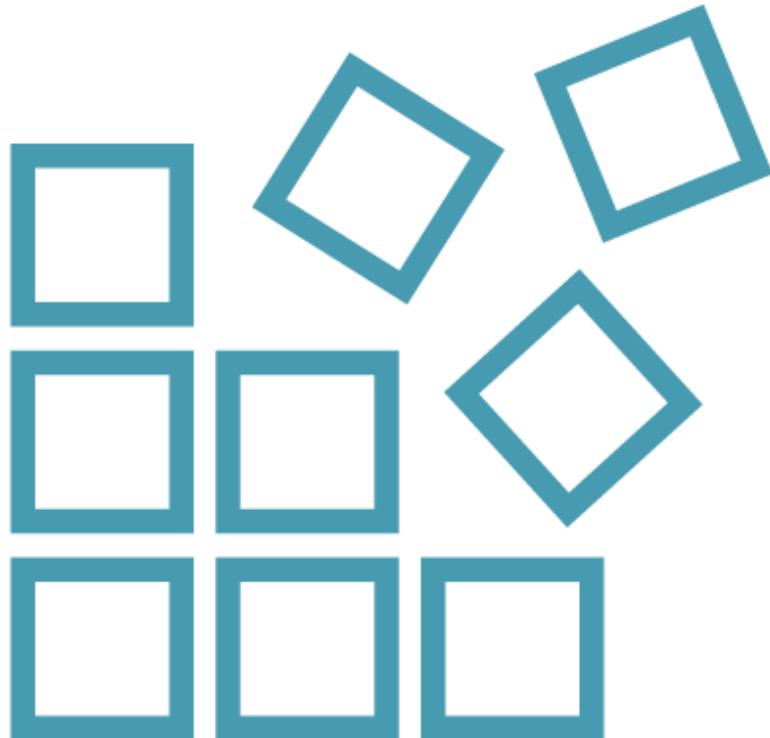
The most common type of fact

Examples:

- Sales amount
- Sales quantity
- Purchase amount
- Number of transactions



Semi-additive Measures



Can be summarized across *some* dimensions

Cannot be summarized across time



Example of Semi-additive Measures



Example of Semi-additive Measures

Date key	Product key	Store key	Quantity	Inventory value at price (\$)
04-03-2019	777	1	200	400
04-04-2019	777	1	100	200



Example of Semi-additive Measures

Date key	Product key	Store key	Quantity	Inventory value at price (\$)
04-03-2019	777	1	200	400
04-04-2019	777	1	100	200
04-03-2019	80	1	50	50
04-04-2019	80	1	25	25



Example of Semi-additive Measures

Date key	Product key	Store key	Quantity	Inventory value at price (\$)
04-03-2019	777	1	200	400
04-04-2019	777	1	100	200
04-03-2019	80	1	50	50
04-04-2019	80	1	25	25

How many units were in the store on the 3rd of April?

What was their value?



Example of Semi-additive Measures

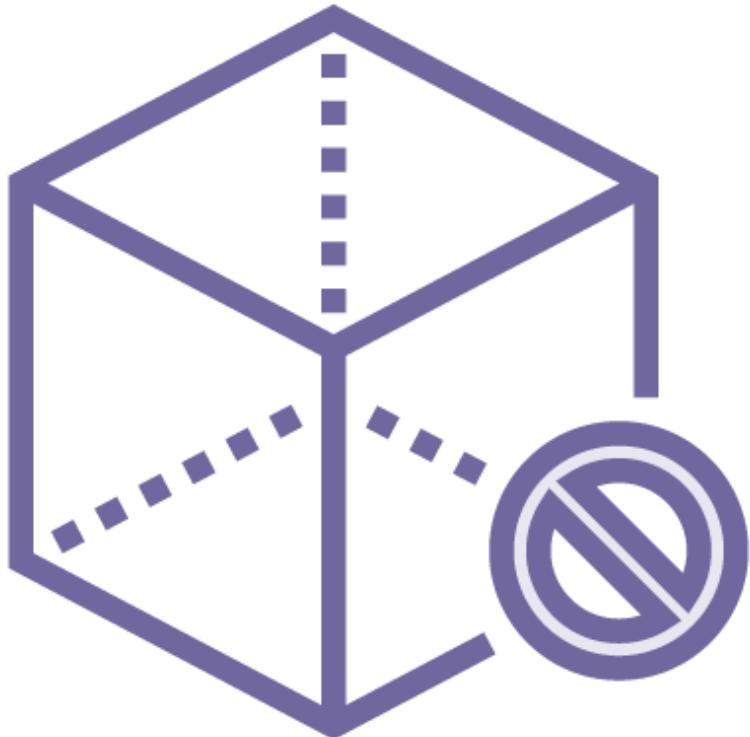
Date key	Product key	Store key	Quantity	Inventory value at price (\$)
04-03-2019	777	1	200	400
04-04-2019	777	1	100	200
04-03-2019	80	1	50	50
04-04-2019	80	1	25	25

Examples of semi-additive measures:

- Inventory information
- Financial account balances
- Water levels on rivers
- Temperature



Non-additive Measures



Ratios or percentages
Cannot be summarized across any dimension



Example of Non-additive Measures



Example of Non-additive Measures

Date key	Product key	Store key	Initial quantity	Removed quantity
04-03-2019	777	1	200	100
04-04-2019	777	1	100	40
04-03-2019	80	1	50	25
04-04-2019	80	1	25	9



Example of Non-additive Measures

Date key	Product key	Store key	Initial quantity	Removed quantity	Removed quantity (%)
04-03-2019	777	1	200	100	50
04-04-2019	777	1	100	40	40
04-03-2019	80	1	50	25	50
04-04-2019	80	1	25	9	36

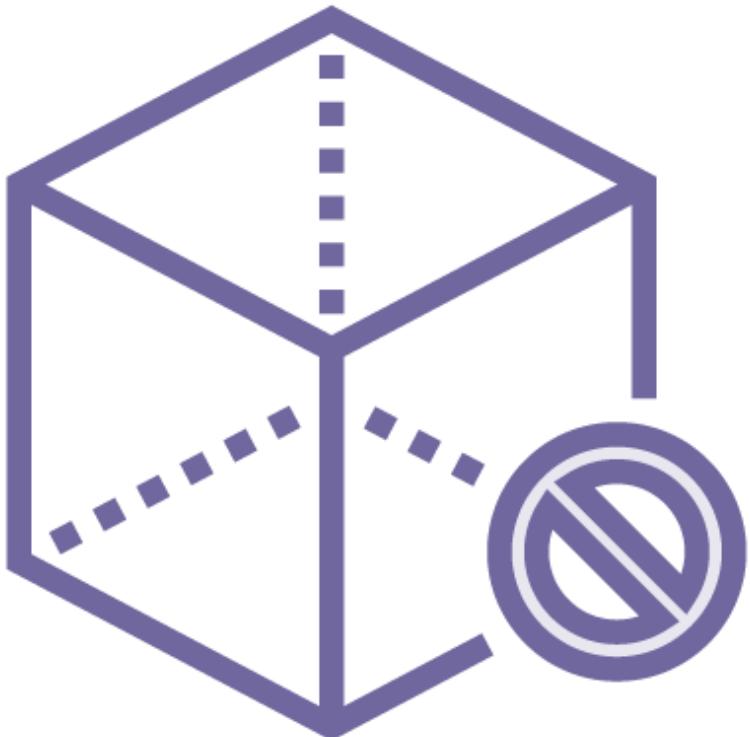


Example of Non-additive Measures

Date key	Product key	Store key	Initial quantity	Removed quantity	Removed quantity (%)	Unit price
04-03-2019	777	1	200	100	50	2
04-04-2019	777	1	100	40	40	2
04-03-2019	80	1	50	25	50	1
04-04-2019	80	1	25	9	36	1



Non-additive Measures



Ratios or percentages

Cannot be summarized across any dimension

Why are non-additive measures considered facts?

- Design preferences
 - Storing them in fact tables vs. in dimension tables, as attributes
- They are calculated based on columns from the fact table

They should be used with care by users



Keys in a Fact Table



Identifying the Keys



Primary key

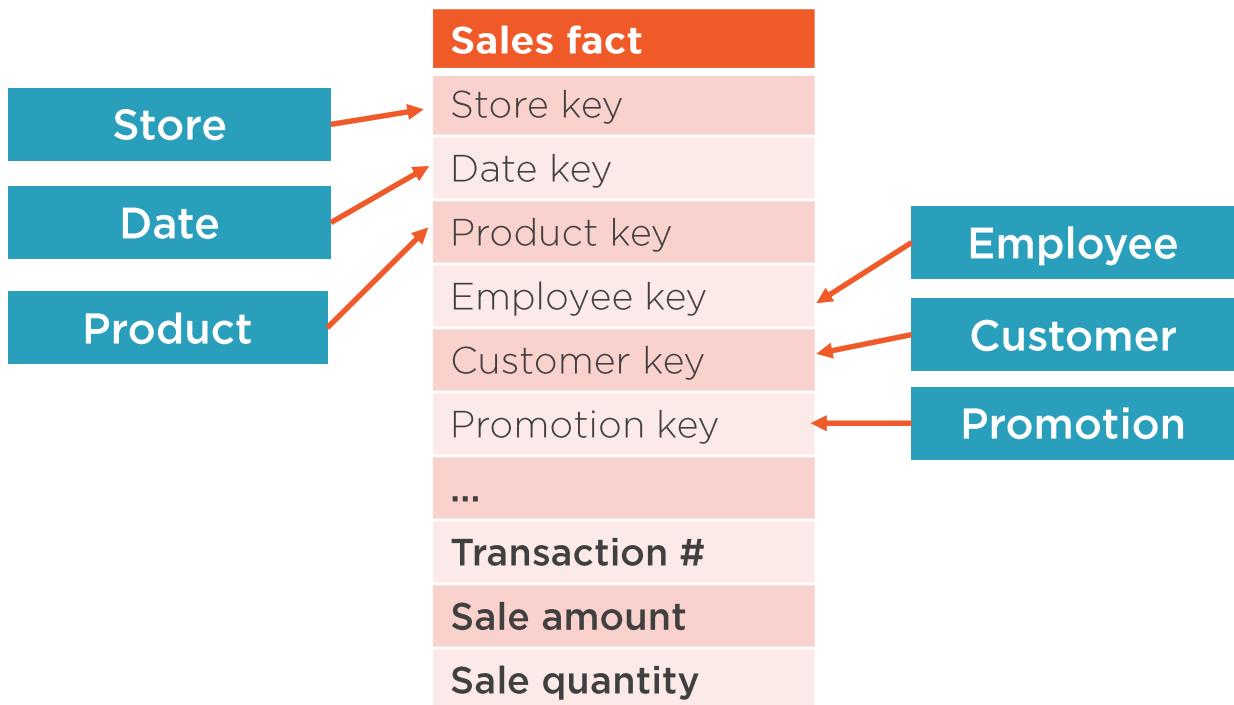
- Composite key
- Surrogate key

Foreign keys for dimensions

Identifying the Primary Key



Identifying the Primary Key



Relationships between fact and dimension are many-to-many

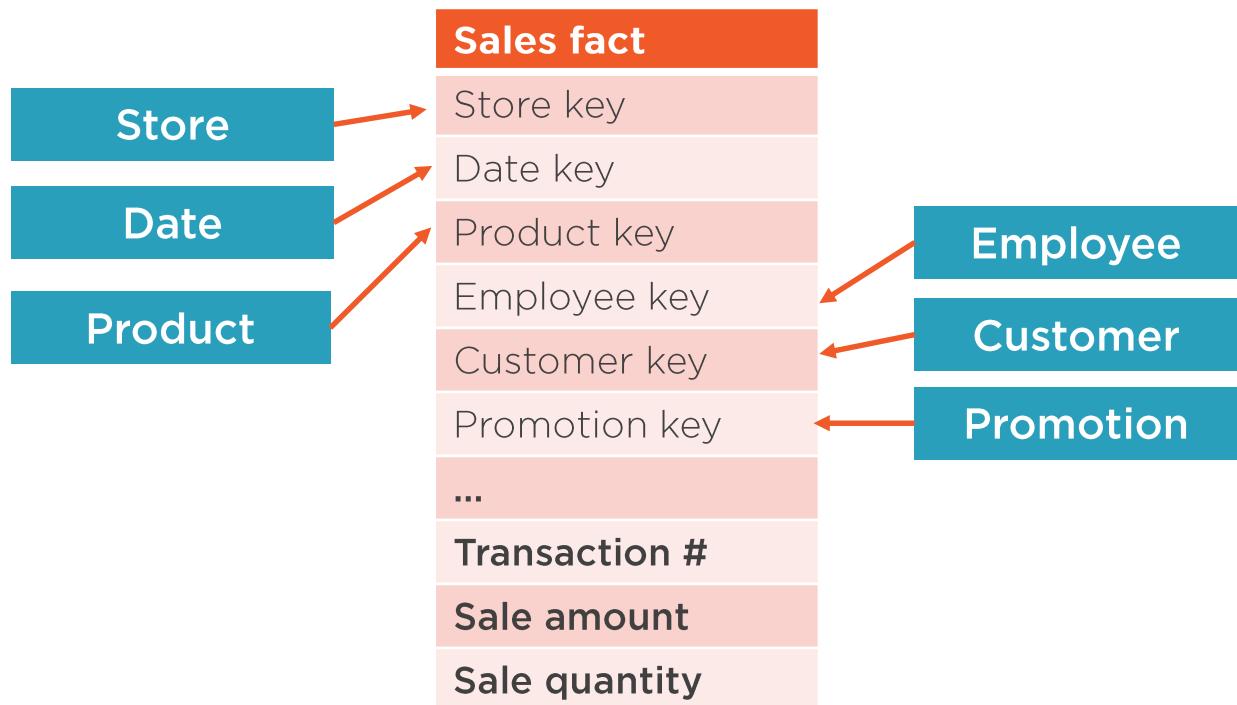
Examples

- 1.a. One sale can have multiple products
- 1.b. One product can be part of multiple sales

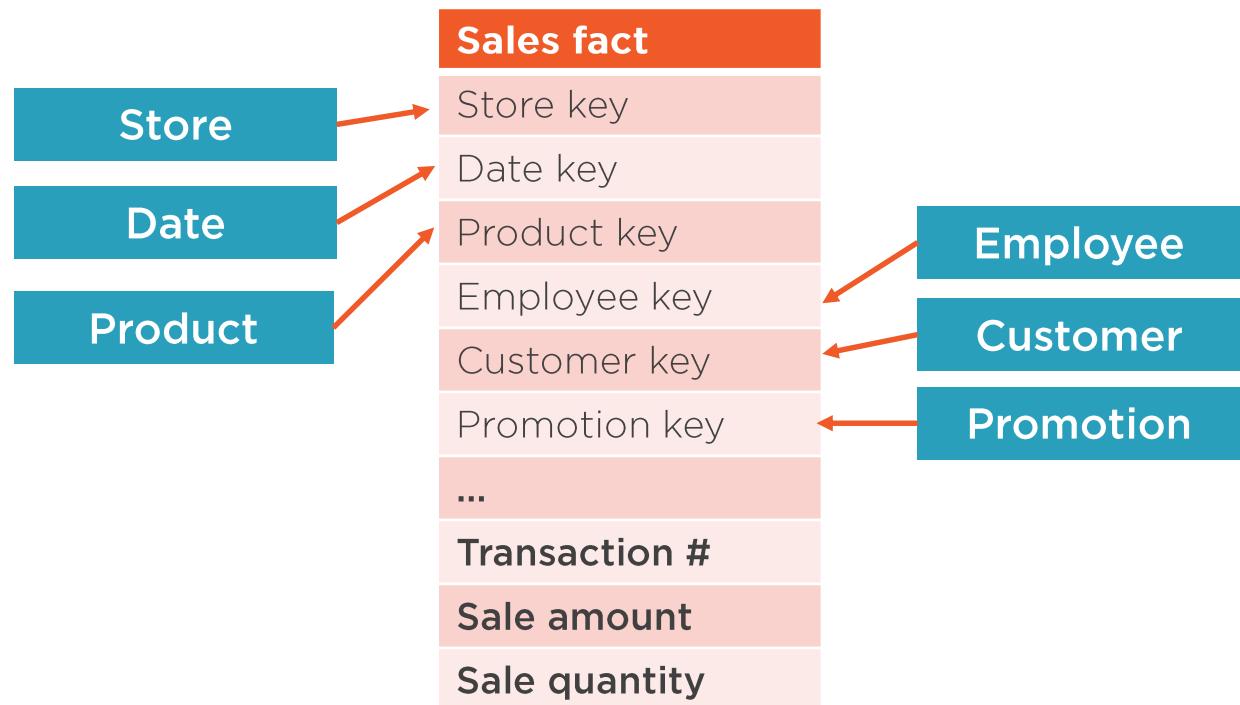
- 2.a. One sale can have multiple promotions applied
- 2.b. One promotion can be applied to many sales



Identifying the Primary Key



Identifying the Primary Key



1. Combine a subset of the foreign keys

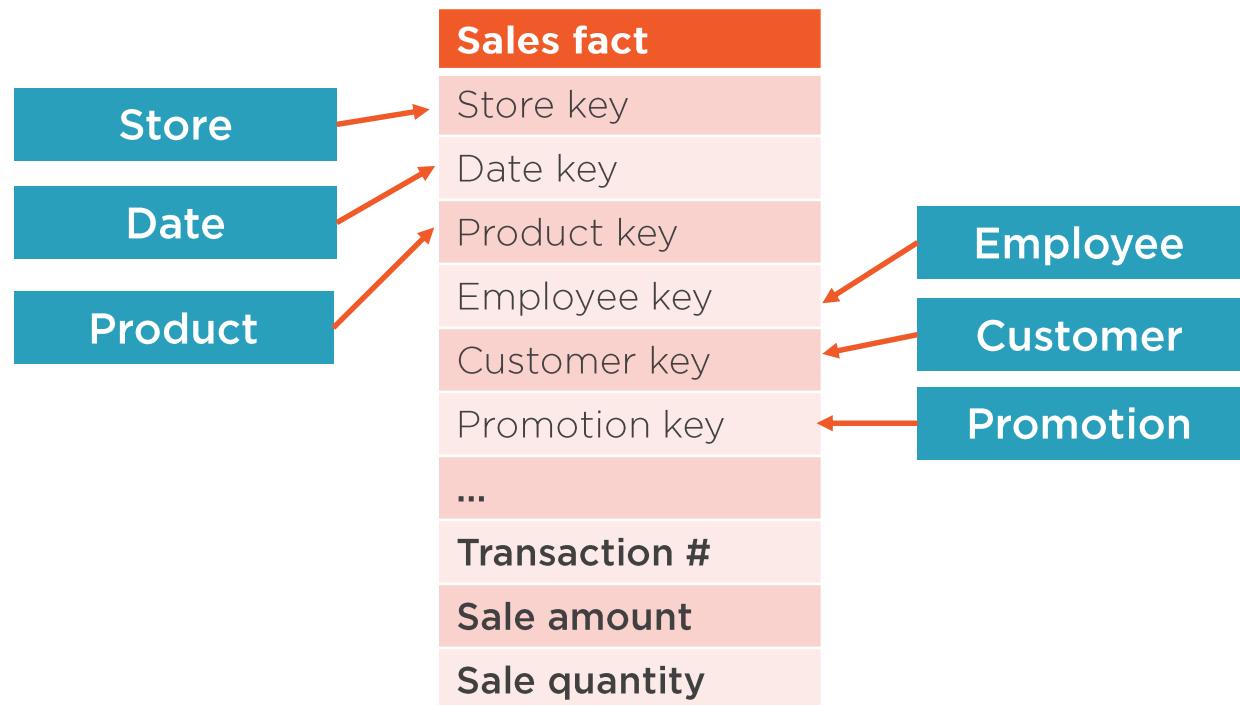
Also called

- Composite key
- Compound key
- Concatenated key

The combination must ensure row uniqueness



Identifying the Primary Key



1. Combine a subset of the foreign keys

Also called

- Composite key
- Compound key
- Concatenated key

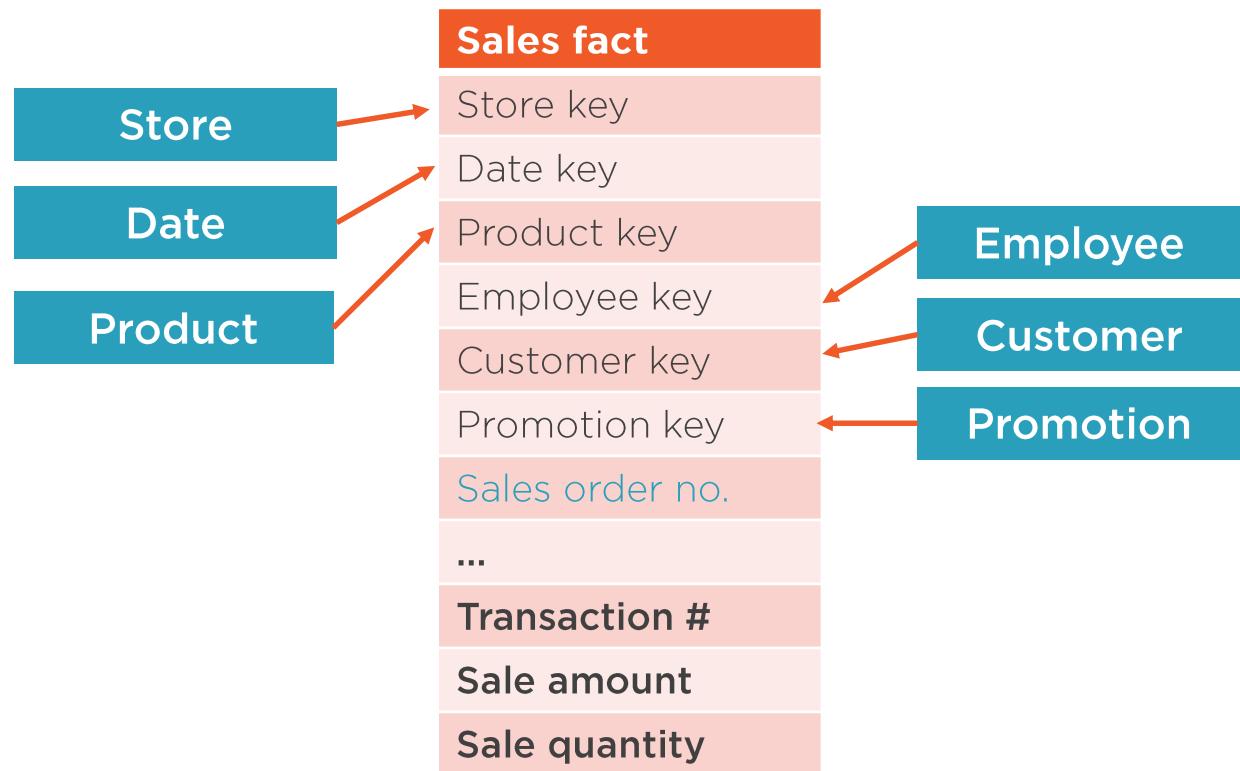
The combination must ensure row uniqueness

2. Use another column from the source system: the primary key for the transactions

Merge the new column to the subset of FKS



Identifying the Primary Key



1. Combine a subset of the foreign keys

Also called

- Composite key
- Compound key
- Concatenated key

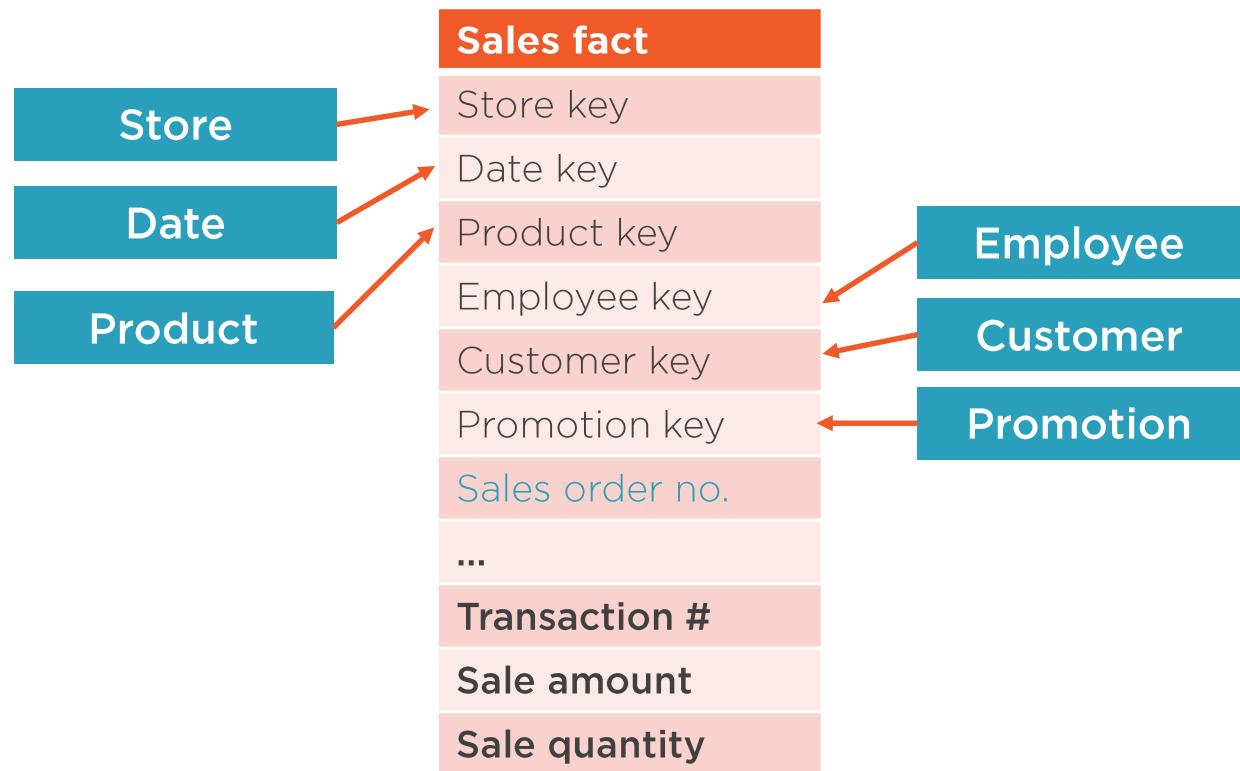
The combination must ensure row uniqueness

2. Use another column from the source system: the primary key for the transactions

Merge the new column to the subset of FKS



Identifying the Primary Key



1. Combine a subset of the foreign keys

Also called

- Composite key
- Compound key
- Concatenated key

The combination must ensure row uniqueness

2. Use another column from the source system: the primary key for the transactions

Merge the new column to the subset of Fks

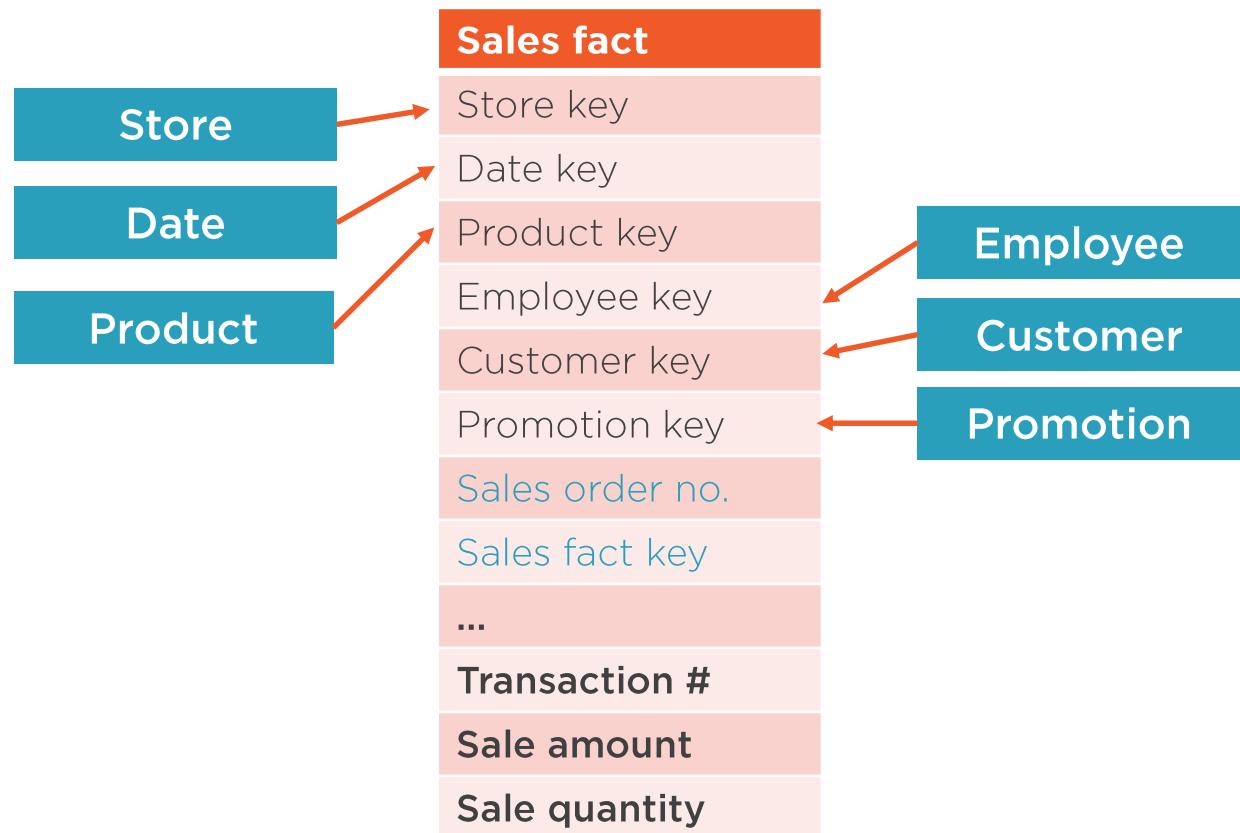
3. Create a surrogate key

- Small integer
- Auto-incremented
- No business value

The Fks don't play a special role anymore



Identifying the Primary Key



1. Combine a subset of the foreign keys

Also called

- Composite key
- Compound key
- Concatenated key

The combination must ensure row uniqueness

2. Use another column from the source system: the primary key for the transactions

Merge the new column to the subset of FKS

3. Create a surrogate key

- Small integer
- Auto-incremented
- No business value

The FKS don't play a special role anymore



Advantages of Using a Surrogate Key



Easier to identify a row in the fact table

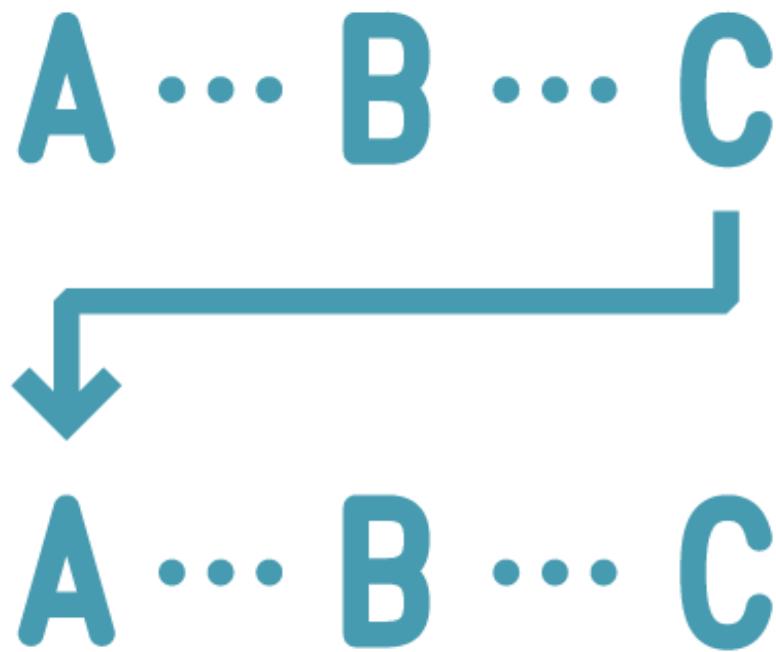
Debugging a load that stopped midway

**Transforming an update into a combination
of**

- Inserts
- Deletes



Creating the Foreign Keys



Each dimension has

- A surrogate key (PK)
- A business key (PK of the source system)

A fact table has

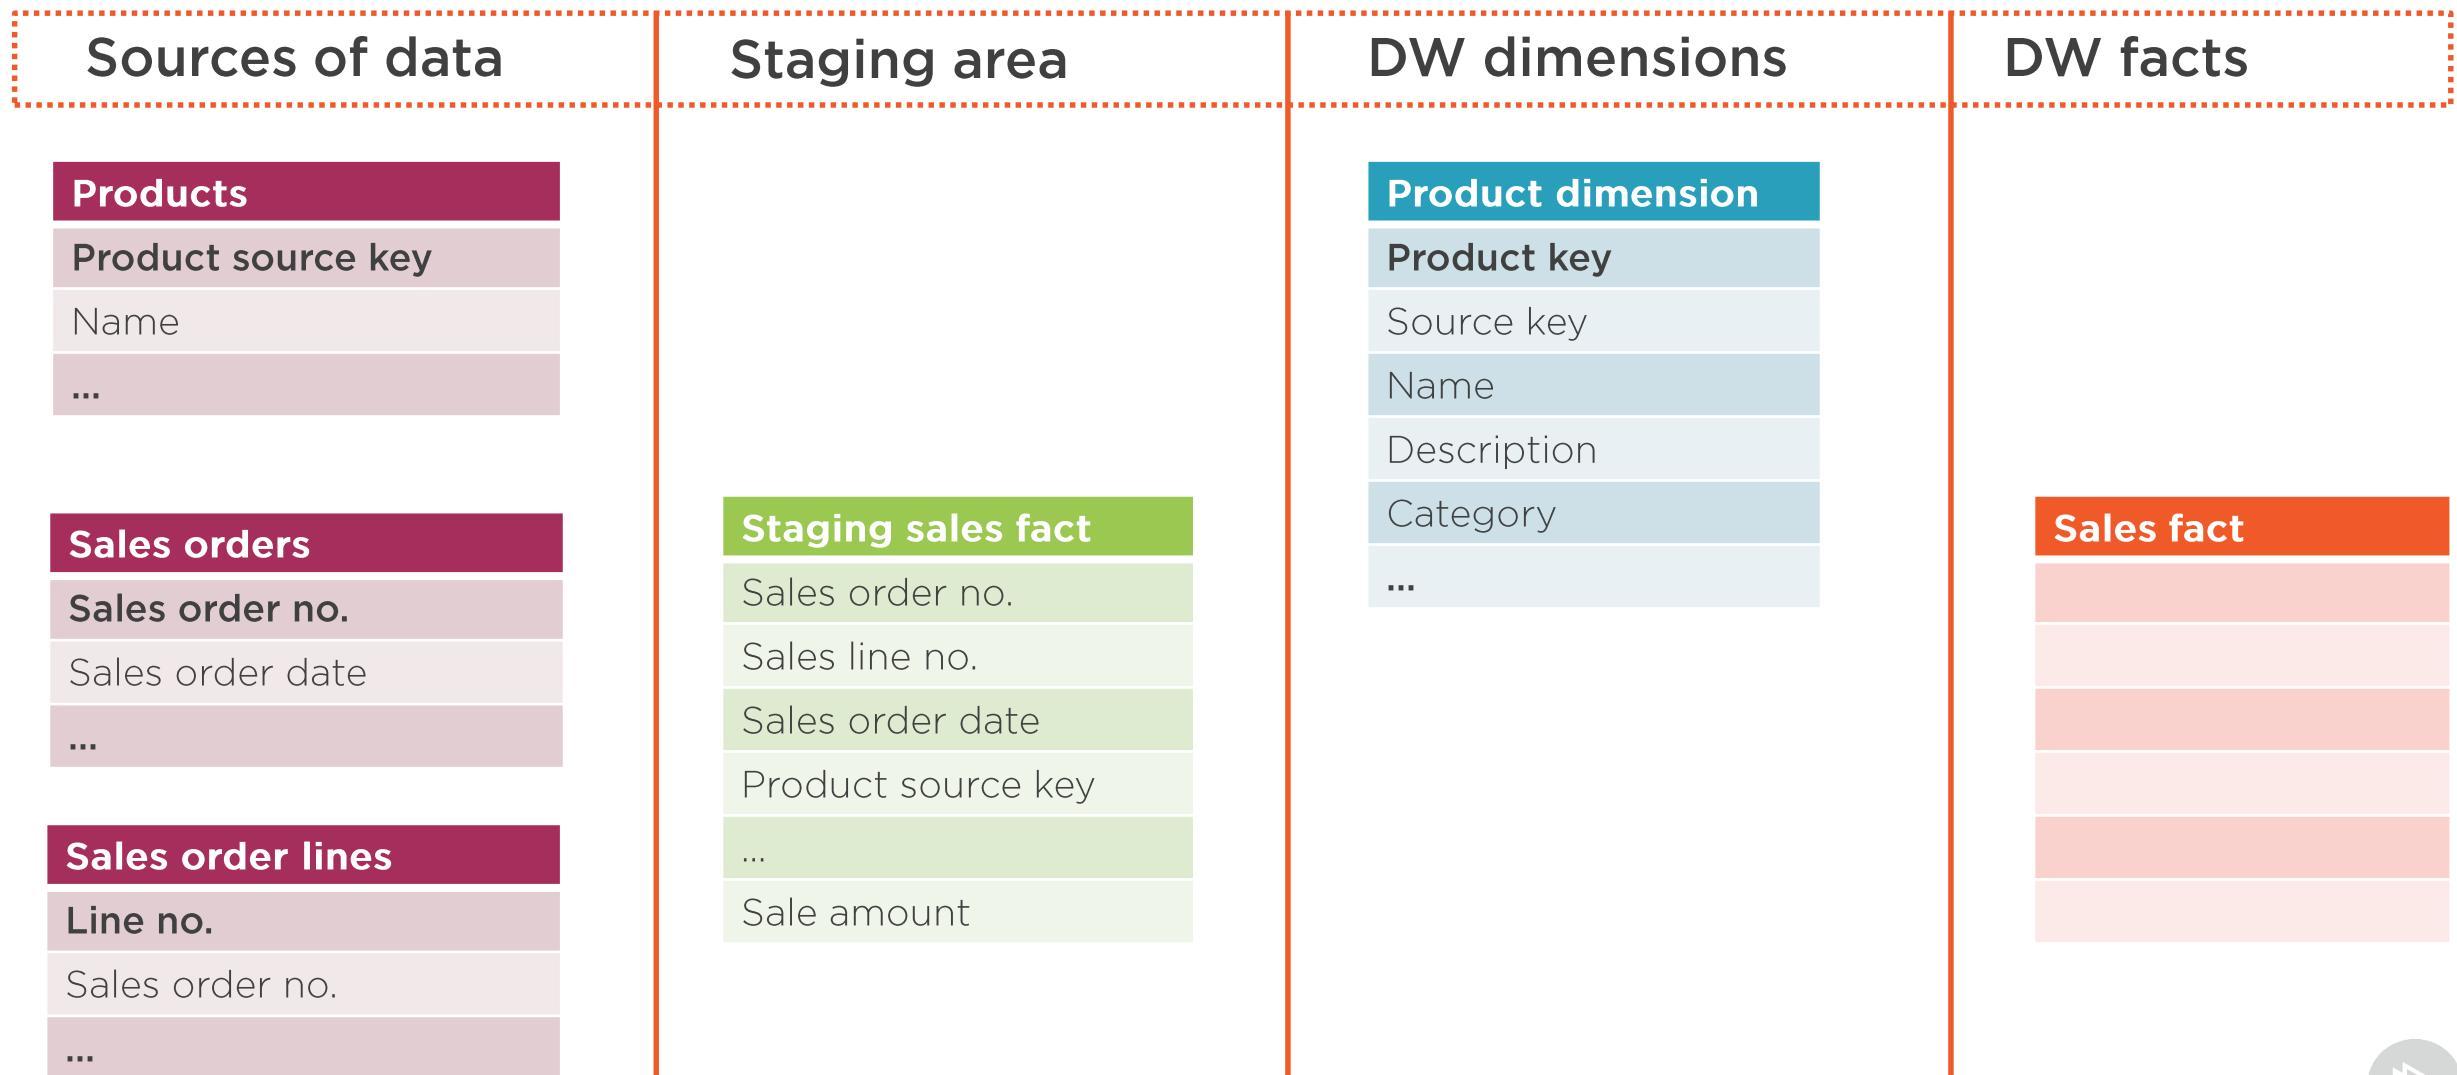
- A primary key
- Foreign keys to its dimensions, linked to the dimensions' surrogate key

Relationship between the fact and the dimension

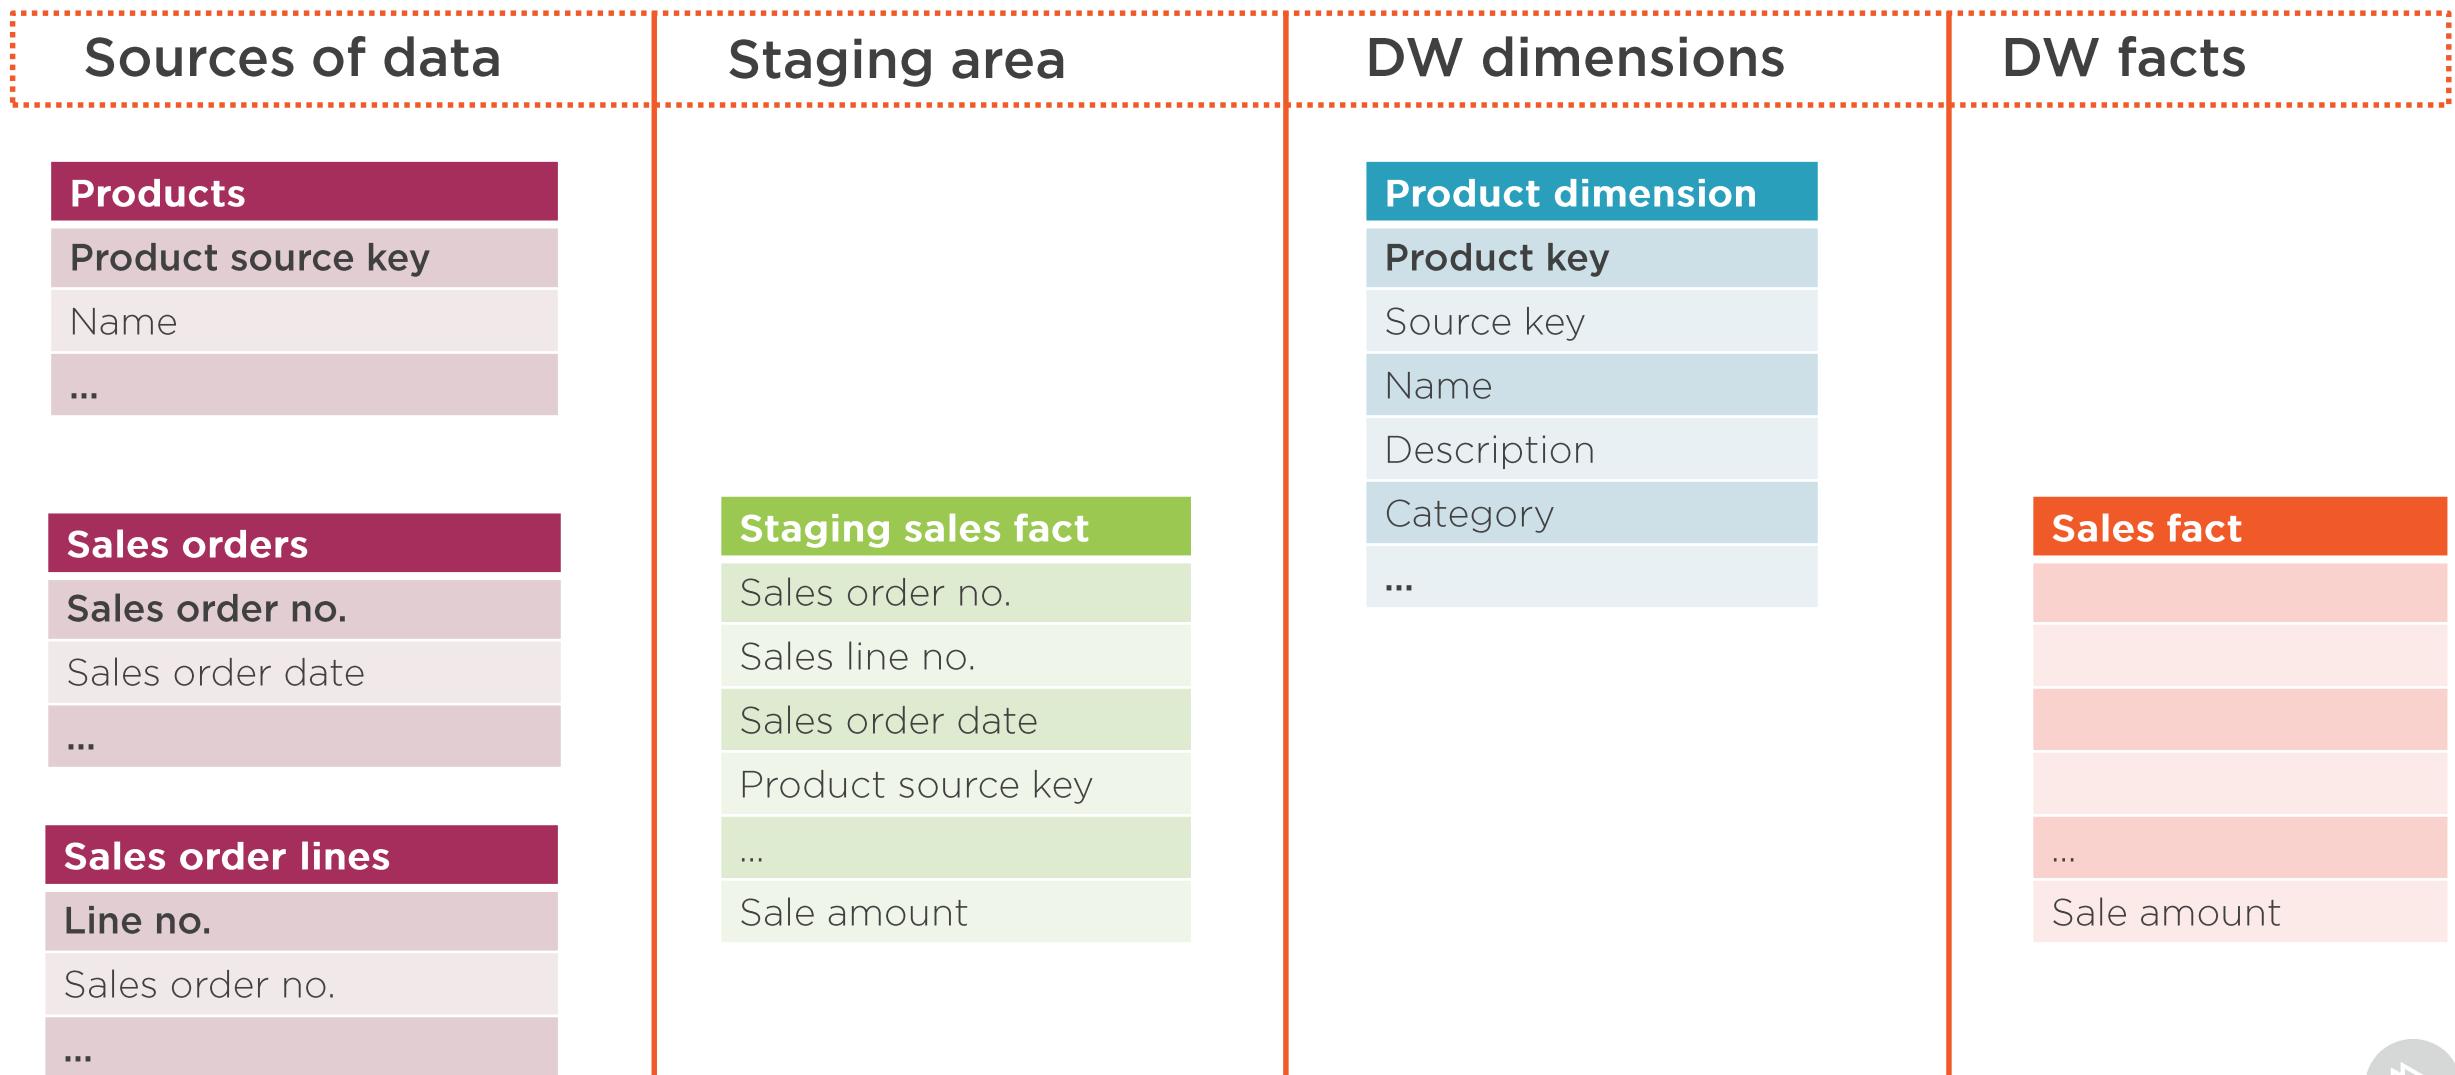
- The business key from the dimension
- An intermediary (staging) fact table
 - Fact table data is joined with dimension table data on the real join conditions
 - Staging table stores the PKs from the source tables for dimensions



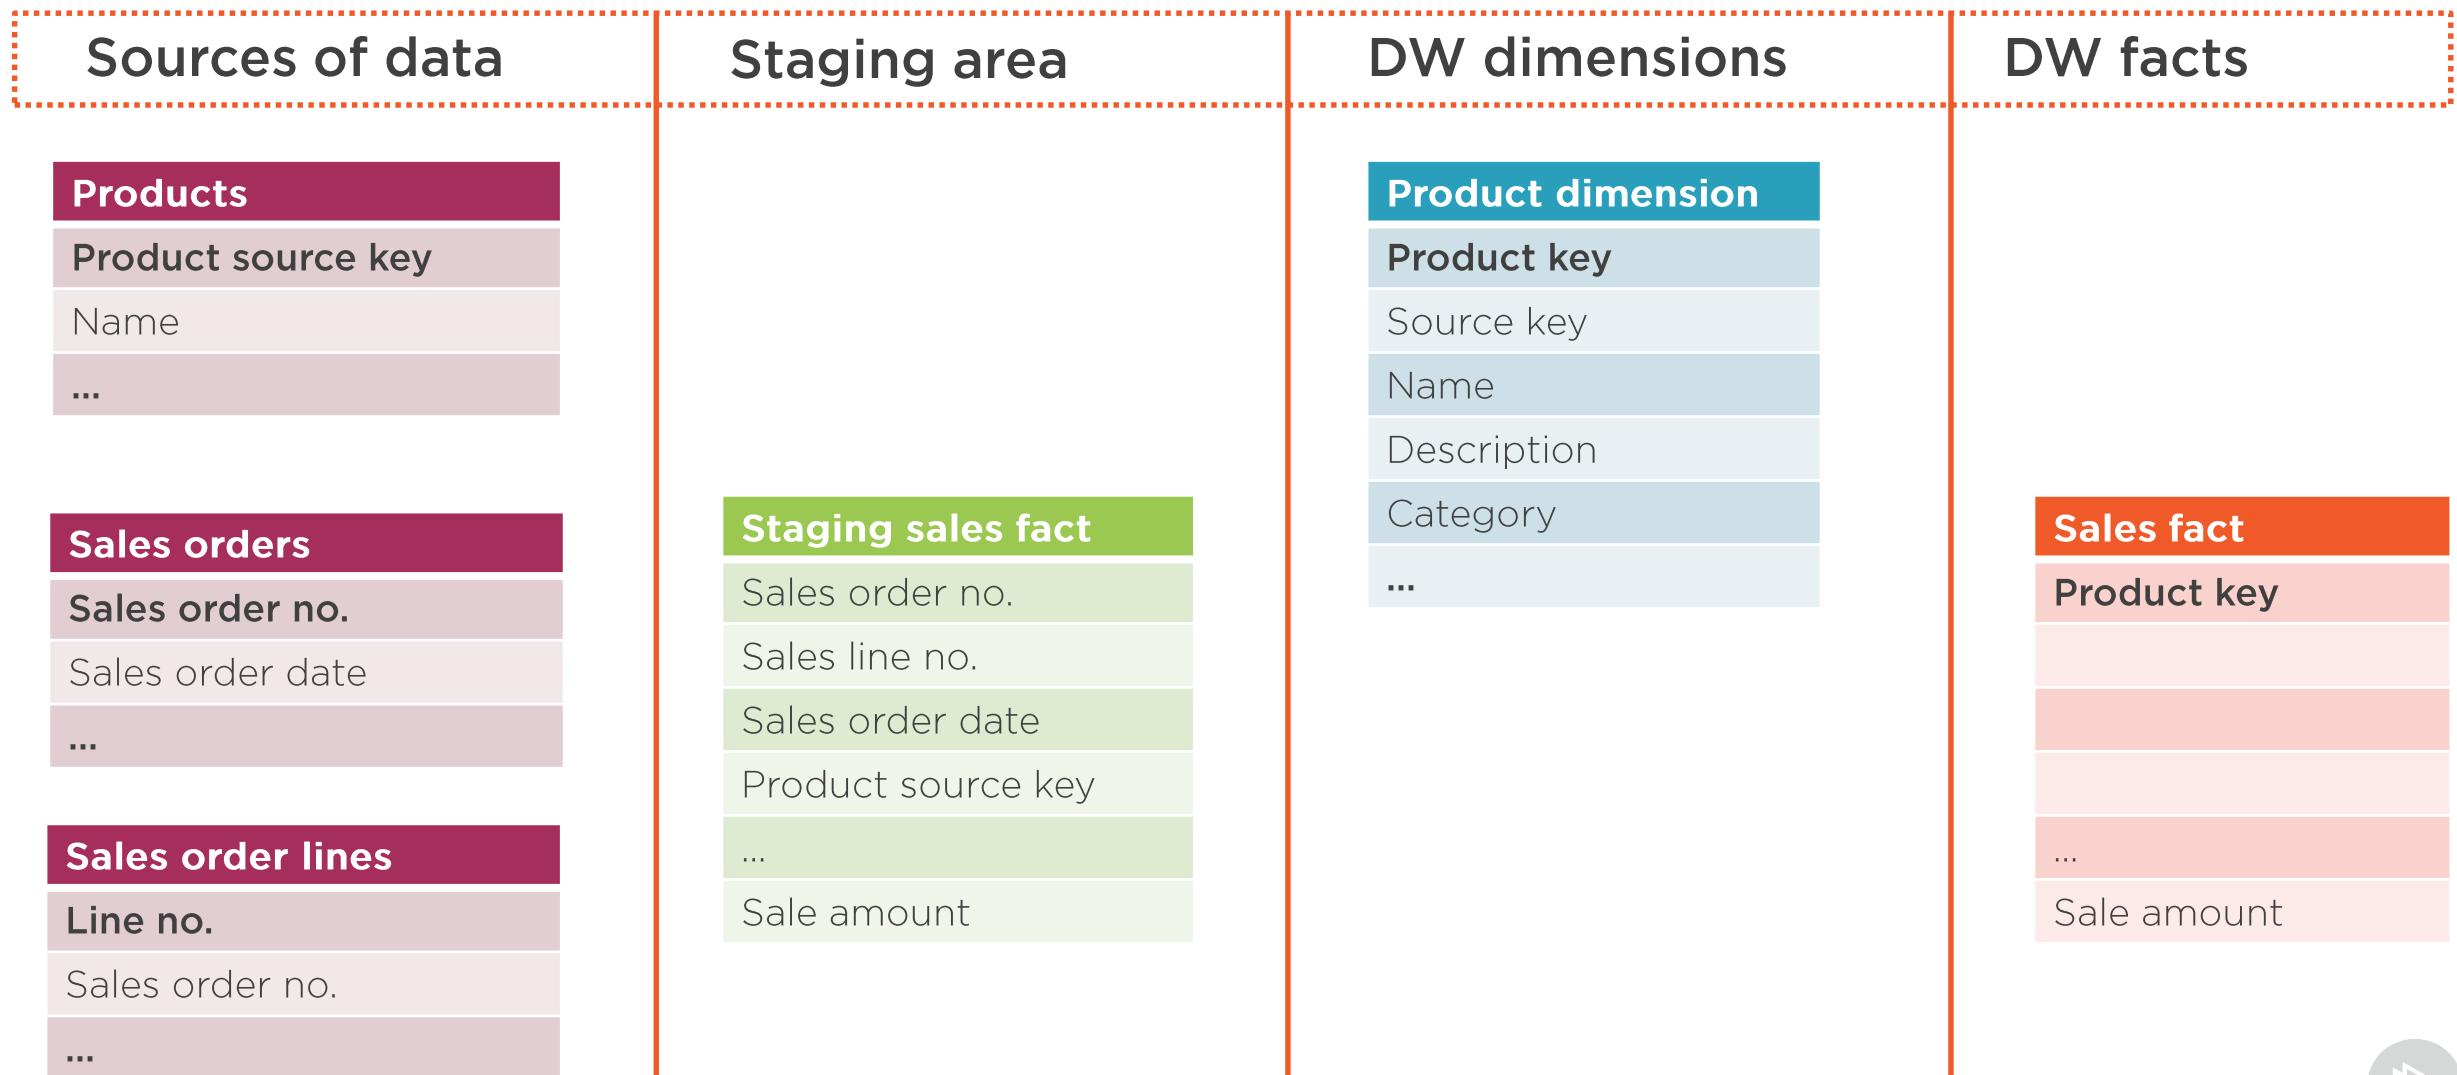
Creating the Foreign Keys



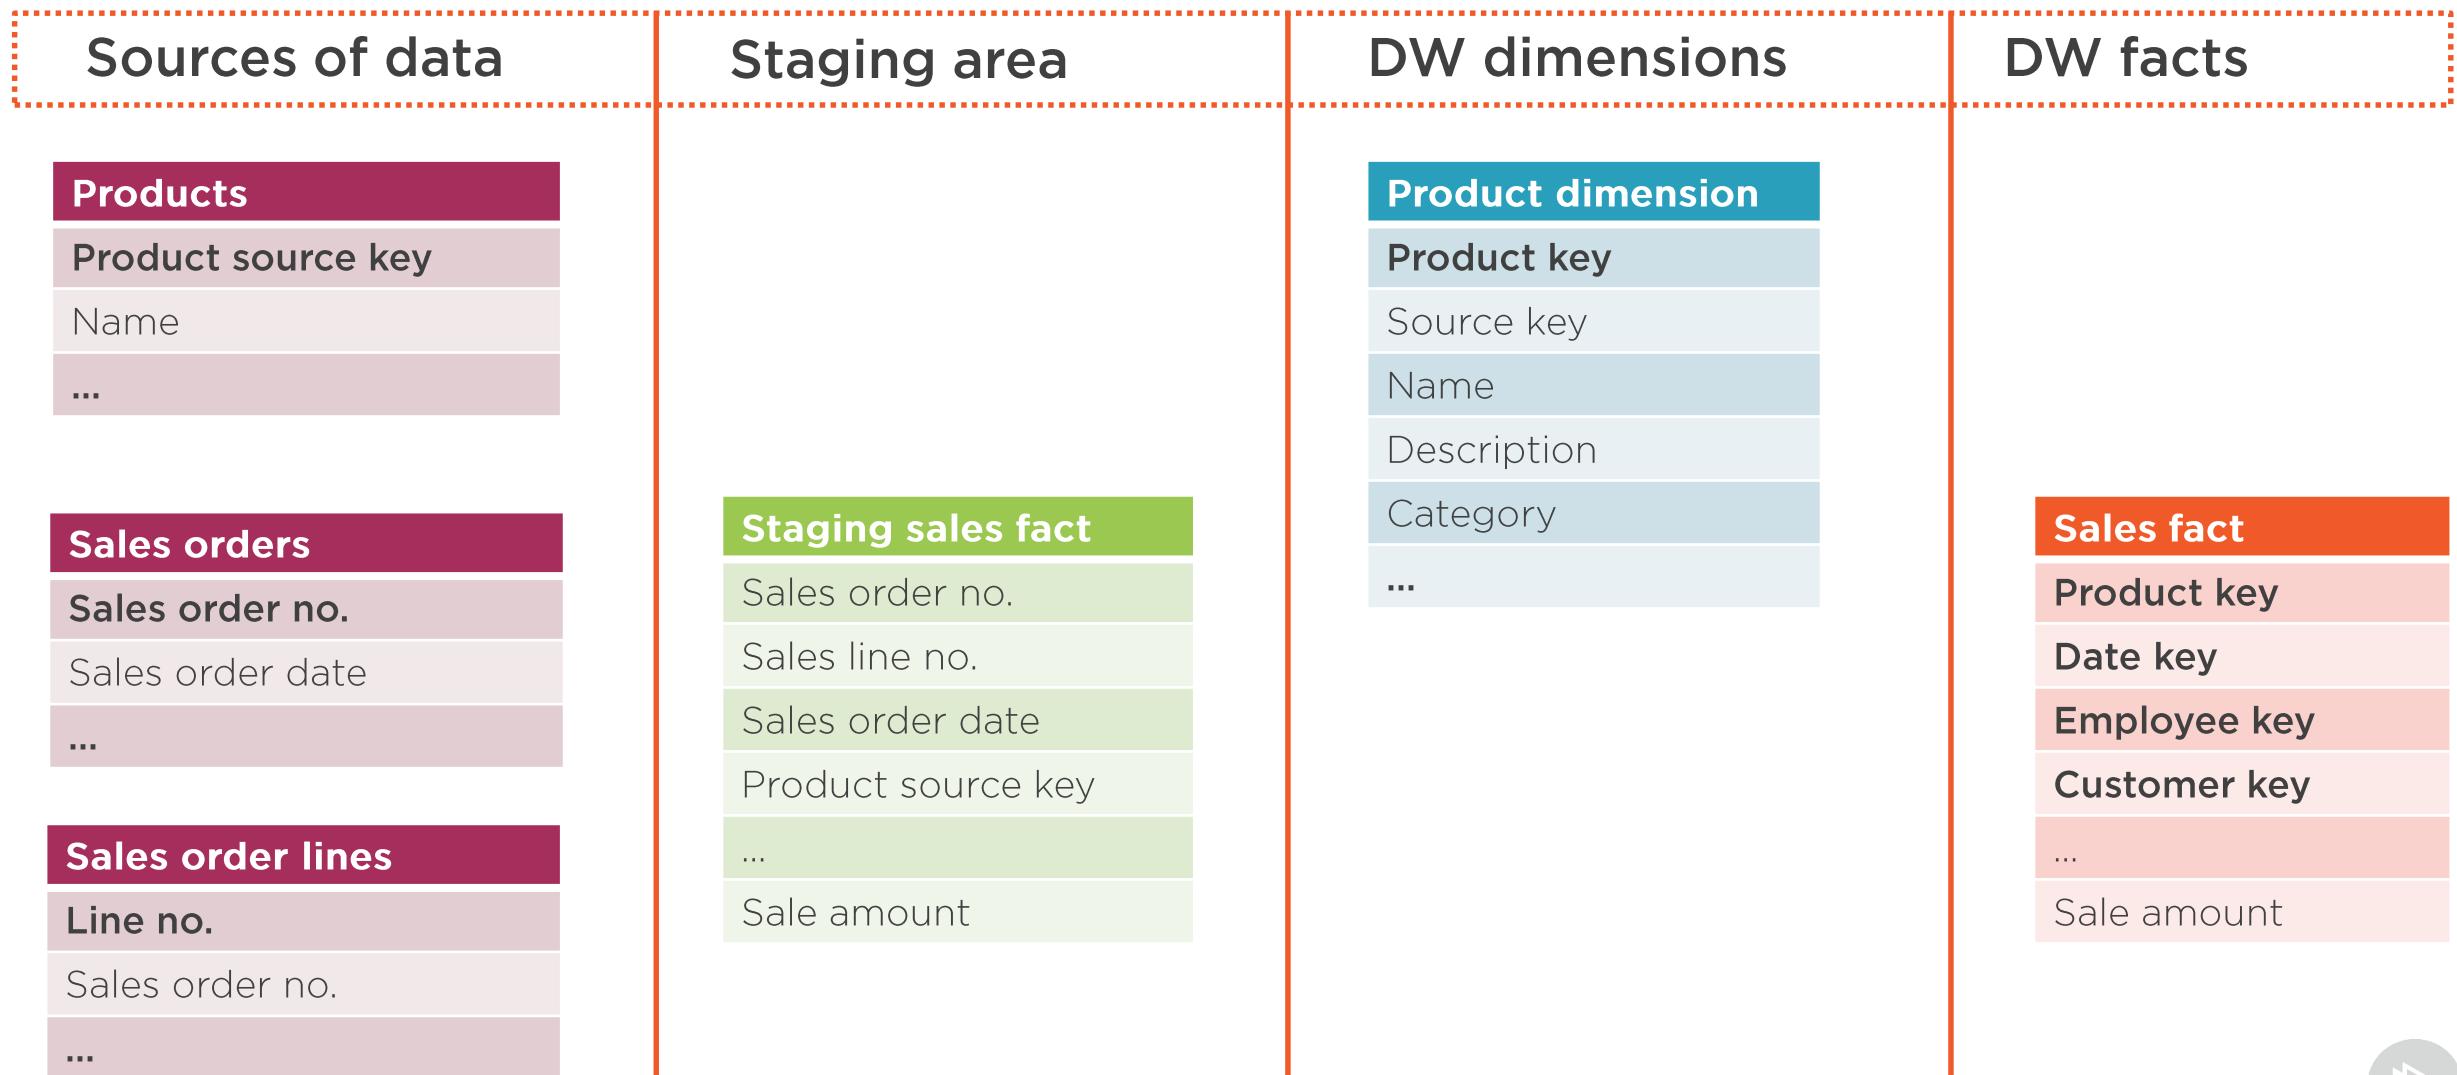
Creating the Foreign Keys



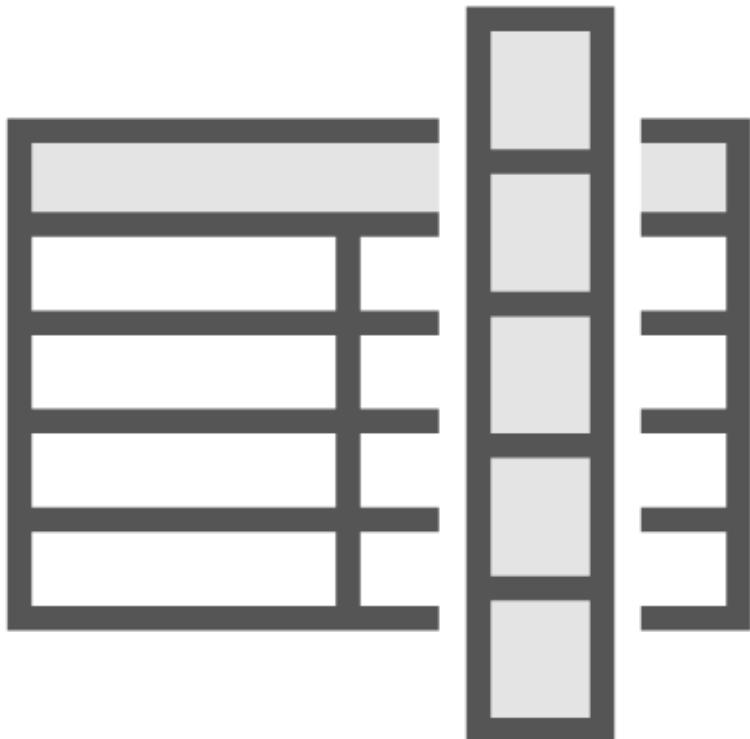
Creating the Foreign Keys



Creating the Foreign Keys



Degenerate Dimensions



Columns added to fact tables

- Different than facts
- Different than keys

Similar to foreign keys, but don't point to any dimension

Examples:

- Sales order number
- Invoice number
- Other transaction numbers

Advantages of degenerate dimensions

- Usually part of the fact's PK
- Group all rows that were part of the same transaction
- Track data back to the operational system

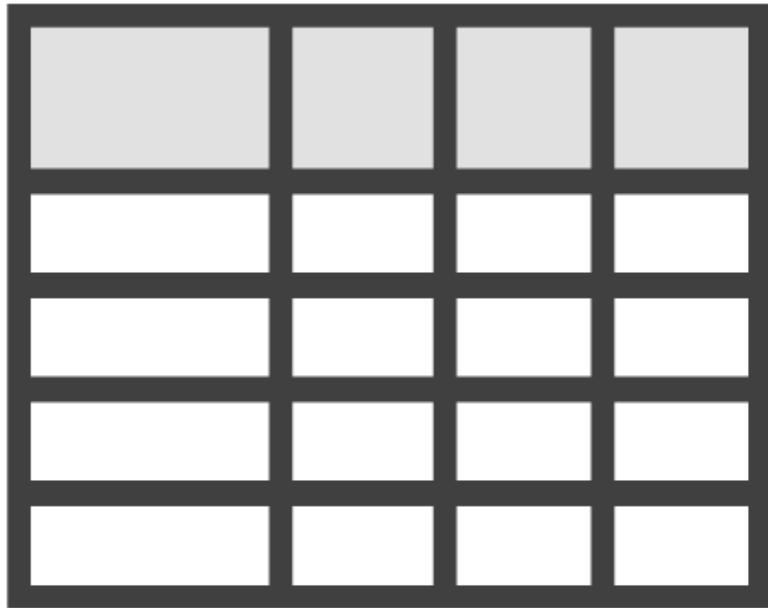
Storing too many may impact the performance of the fact table



Types of Fact Tables



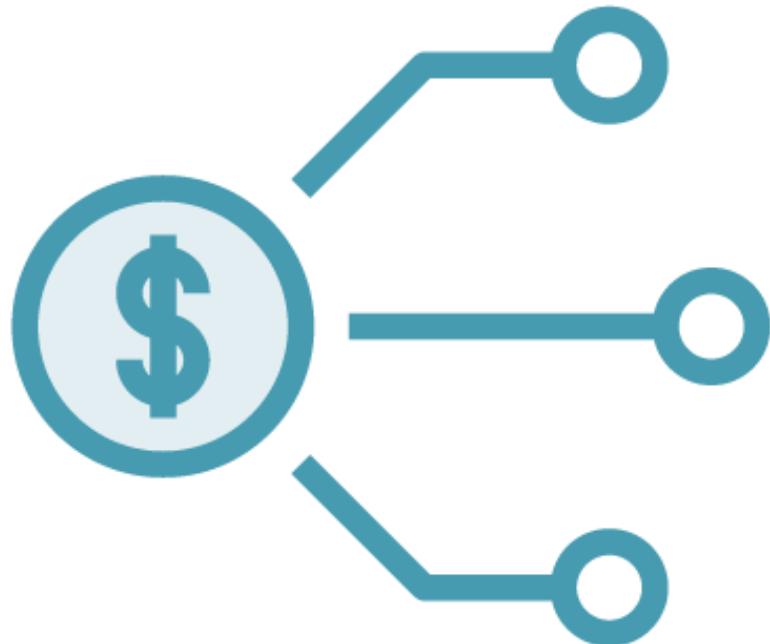
Most Common Types of Fact Tables



Transaction fact tables
Periodic snapshot
Accumulating snapshot



Transaction Fact Tables



The most common type of fact table in a data warehouse

Data inside it represents a process that happened at a certain point in time

A row exists for a dimension member only if the member was involved in a transaction



Properties of Transaction Fact Tables

The granularity is
the transaction or
transaction line

Tendency to
become very large

Relationships with
many dimensions

Sparingly
populated

The facts are
additive



Periodic Snapshot Fact Tables



Show information as it was at the end of a time interval

- End of each day
- End of each month, etc.

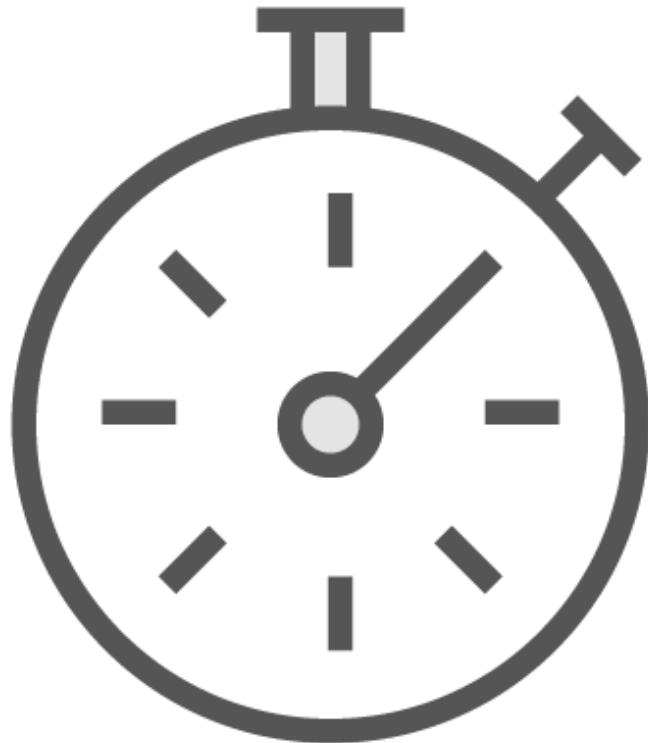
Are not periodically refreshed

- The new snapshots are added on top of the existing ones
- Are useful for analyzing trends in time

Transactional fact tables can be a starting point for the periodic snapshots



Accumulating Snapshot Fact Tables



Store information about a process that has

- A clear beginning
- A clear end
- A number of intermediary steps

Example: pipeline or workflow processes

- Fulfilling an order
- Progress of a mortgage application
- Status of a support ticket



Accumulating Snapshot Fact Tables

Handling support calls



Accumulating Snapshot Fact Tables

Handling support calls



Accumulating Snapshot Fact Tables

Handling support calls



Created date key	Assigned date key	In progress date key	In testing date key	Done date key
20190404	20190405	20190406	20190412	20190419



Accumulating Snapshot Fact Tables

Handling support calls

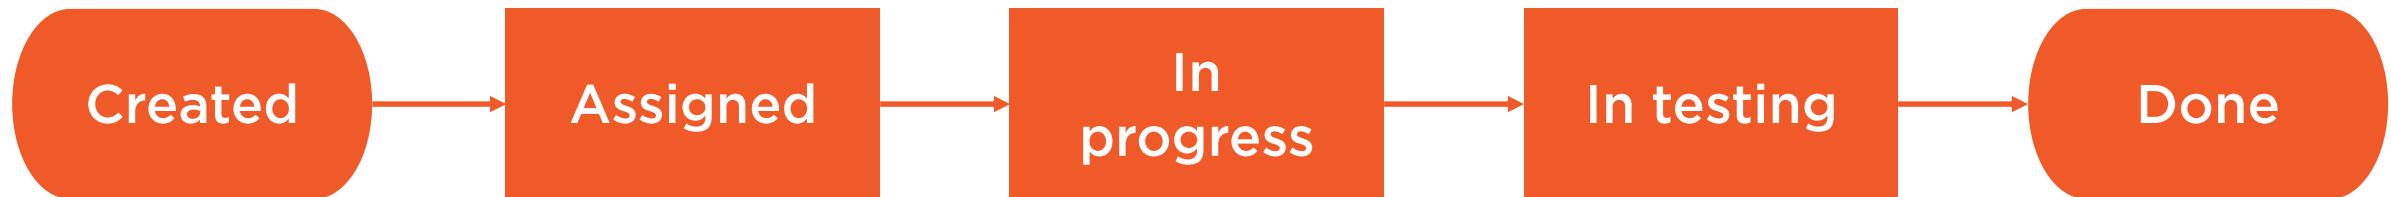


Created date key	Assigned date key	In progress date key	In testing date key	Done date key
20190404	20190405	20190406	20190412	20190419
20190407	17530101	17530101	17530101	17530101



Accumulating Snapshot Fact Tables

Handling support calls



Created date key	Assigned date key	In progress date key	In testing date key	Done date key
20190404	20190405	20190406	20190412	20190419
20190407	20190407	17530101	17530101	17530101



Accumulating Snapshot Fact Tables

Handling support calls



Created date key	Assigned date key	In progress date key	In testing date key	Done date key
20190404	20190405	20190406	20190412	20190419
20190407	20190407	20190407	17530101	17530101



Accumulating Snapshot Fact Tables

Handling support calls



Created date key	Assigned date key	In progress date key	In testing date key	Done date key	Duration development	Duration testing
20190404	20190405	20190406	20190412	20190419	6	3
20190407	20190407	20190407	17530101	17530101	NULL	NULL



Accumulating Snapshot Fact Tables

Handling support calls

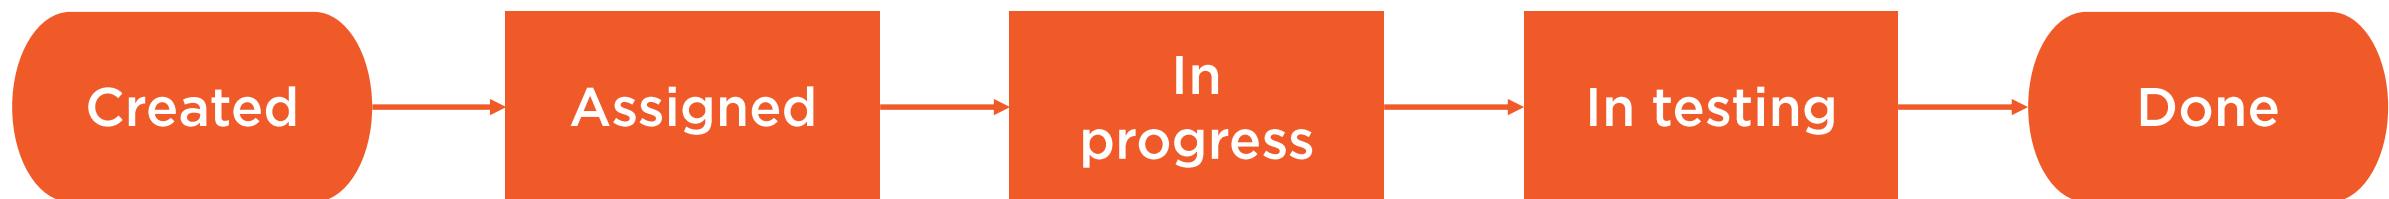


Created date key	Assigned date key	In progress date key	In testing date key	Done date key	Duration development	Duration testing	Is done
20190404	20190405	20190406	20190412	20190419	6	3	1
20190407	20190407	20190407	17530101	17530101	NULL	NULL	0



Accumulating Snapshot Fact Tables

Handling support calls



Created date key	Assigned date key	In progress date key	In testing date key	Done date key	Duration development	Duration testing	Is done	Current status key
20190404	20190405	20190406	20190412	20190419	6	3	1	6
20190407	20190407	20190407	17530101	17530101	NULL	NULL	0	4



Handling Null Values in Fact Tables

Nulls in the facts

It is alright to have nulls in the facts

The aggregate functions can handle null values

Replacing null with zeros would influence the calculations performed

Value
2
2
null

Average: 2

Value
2
2
0

Average: 1

Nulls in the foreign keys

Will produce a referential integrity violation error

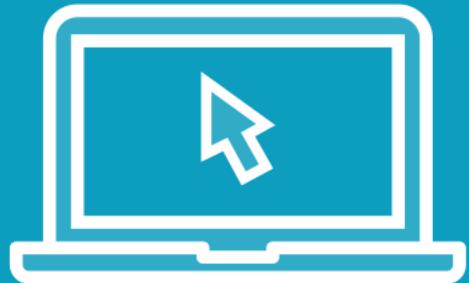
This situation must be avoided

Use the “empty row” technique:

- Every dimension should have an empty row, with its own primary key
- When there is no link between fact and dimension, the foreign key column will store the PK of the empty row



Demo



Creating and working with a transaction fact table

- Populate staging table with source data
- Populate fact table with data from staging
- Use the fact table in reports



Summary



Components of a fact table:

- Facts (or numeric measurements)
 - Additive
 - Semi-additive
 - Non-additive
- The primary key
 - Surrogate key
 - Composite key
- Foreign keys for dimensions
- Degenerate dimensions

The most common types of fact tables:

- Transaction fact table
- Periodic snapshot
- Accumulating snapshot

Be aware of the null values in fact tables

- Measures can have null values
- Foreign key columns must not have nulls



Loading Data in a Data Warehouse



Ana Voicu
@ana_voicu



Review of the Previous Steps



Discuss with the business people

- Processes
- Analysis needs

Investigate data sources

Implement the dimensional model

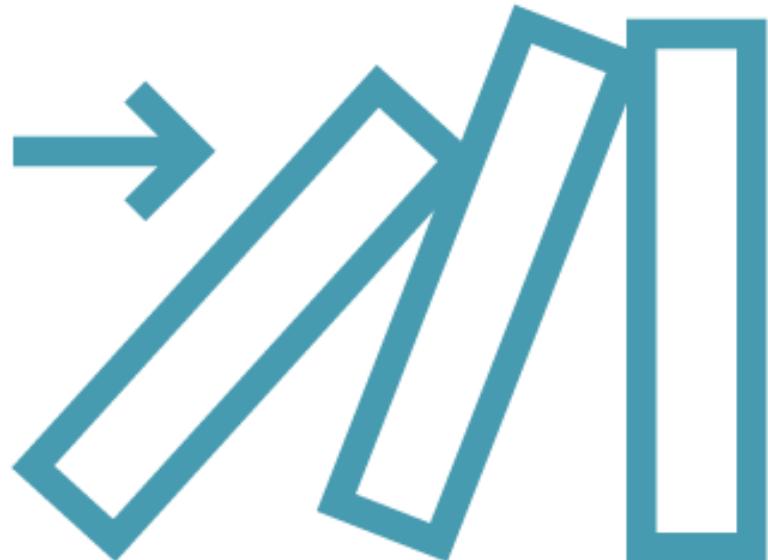
- End result will be the physical design
- Fact and dimension tables are defined

Populate the data warehouse tables with source data

- Also called “loading” the data warehouse
- Implement the ETL system



The ETL system



Similar to a “black box” for the data warehouse users

Data is:

- Integrated from multiple sources
- Transformed into meaningful reports
- Helping people make decisions

The ETL should be properly designed and implemented



Overview



What does ETL mean?

- Extraction
- Transformation
- Load

Types of loads

- Full (initial) load
- Incremental load

Data lineage

- What is data lineage?
- Why is it important?
- How can you implement it?



Demos



Perform the needed steps to load a dimension table

Create and work with the auxiliary objects involved in a load

- Staging tables
- Lineage table
- Incremental loads table



Overview of an ETL System



What Is ETL?



A system that incorporates all operations performed on data

- Since it's selected from data sources
- Until it reaches the presentation area

Consists of three phases

- Extraction (E)
- Transformation (T)
- Loading (L)



Overview of an ETL System



Overview of an ETL System

Sources of data

Products



Subcategories



Categories



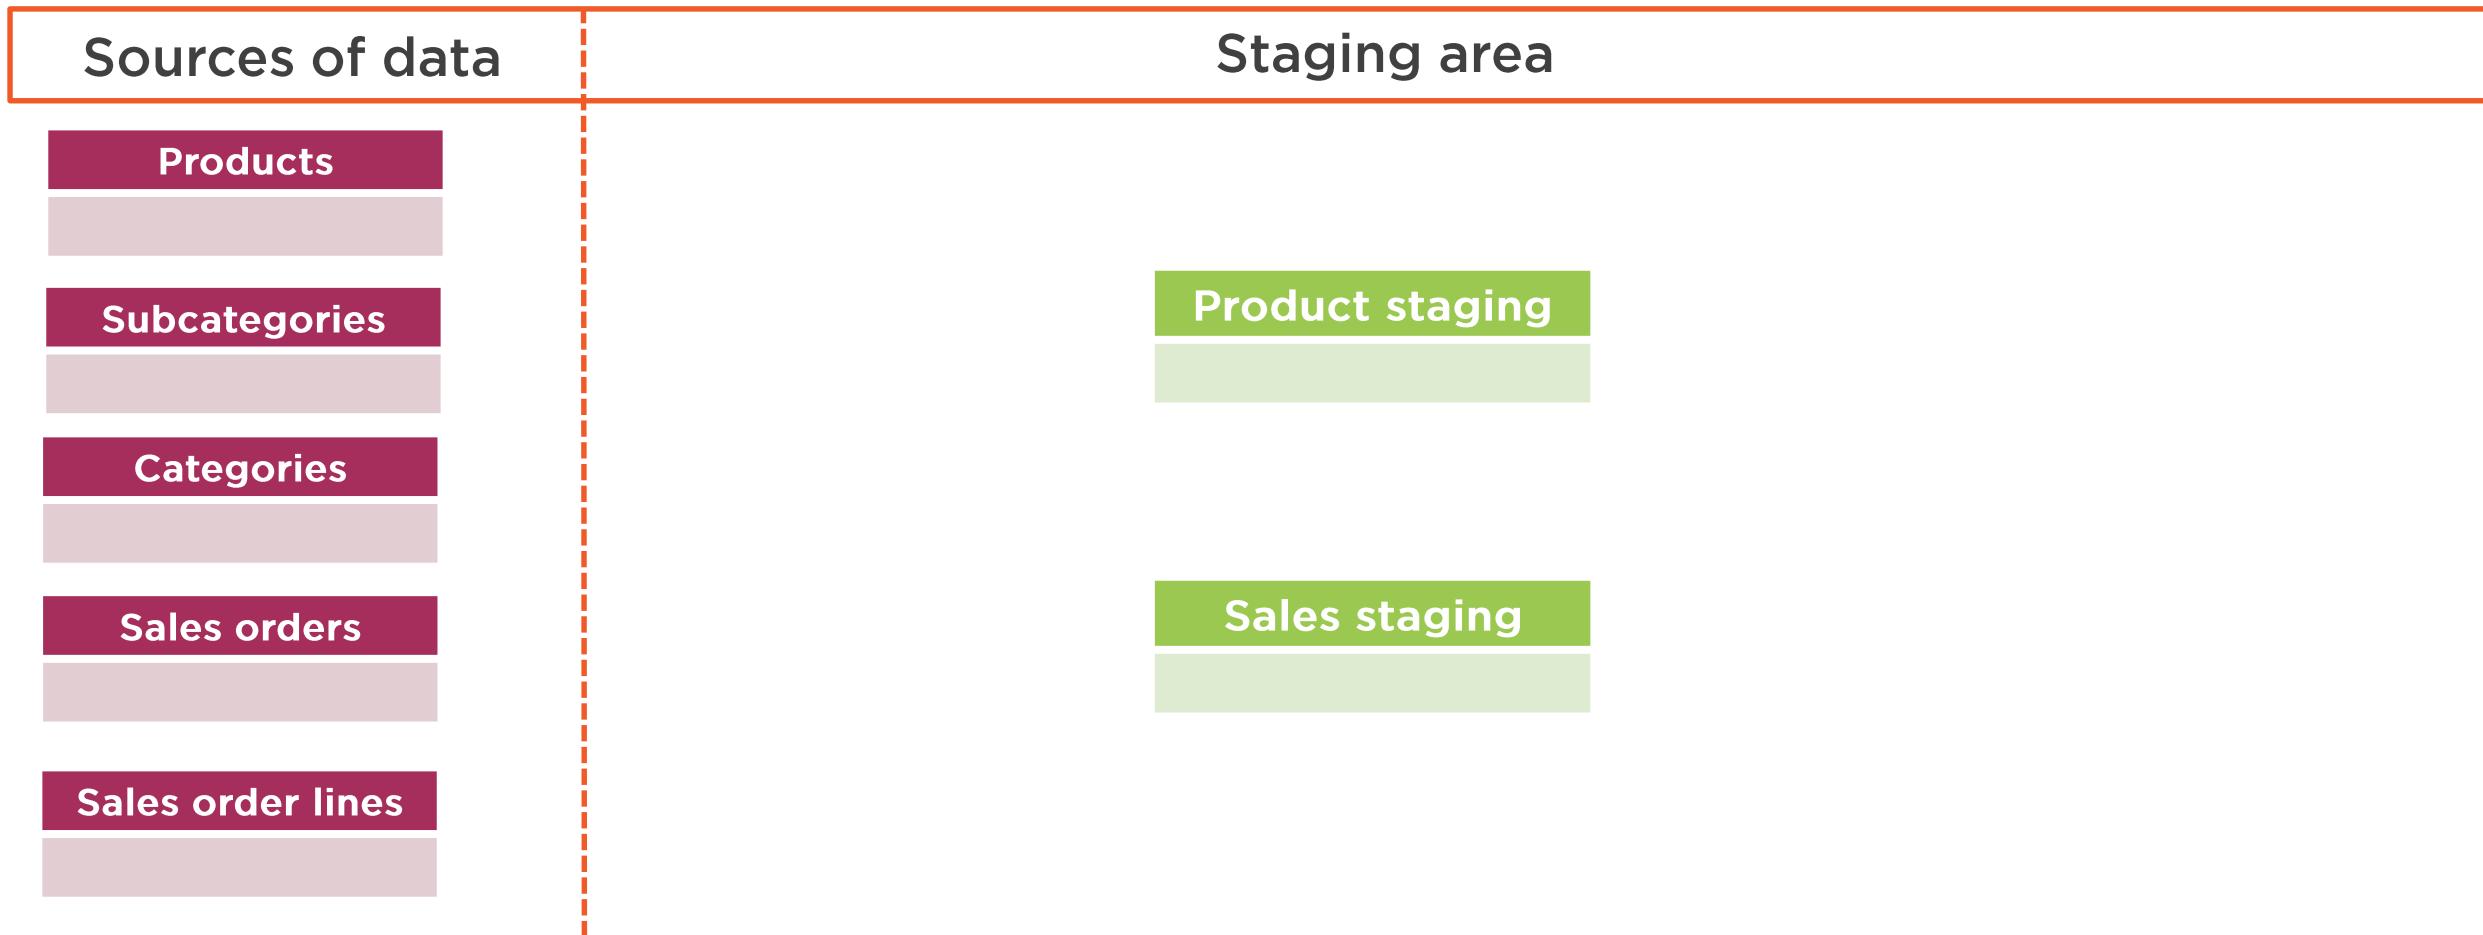
Sales orders



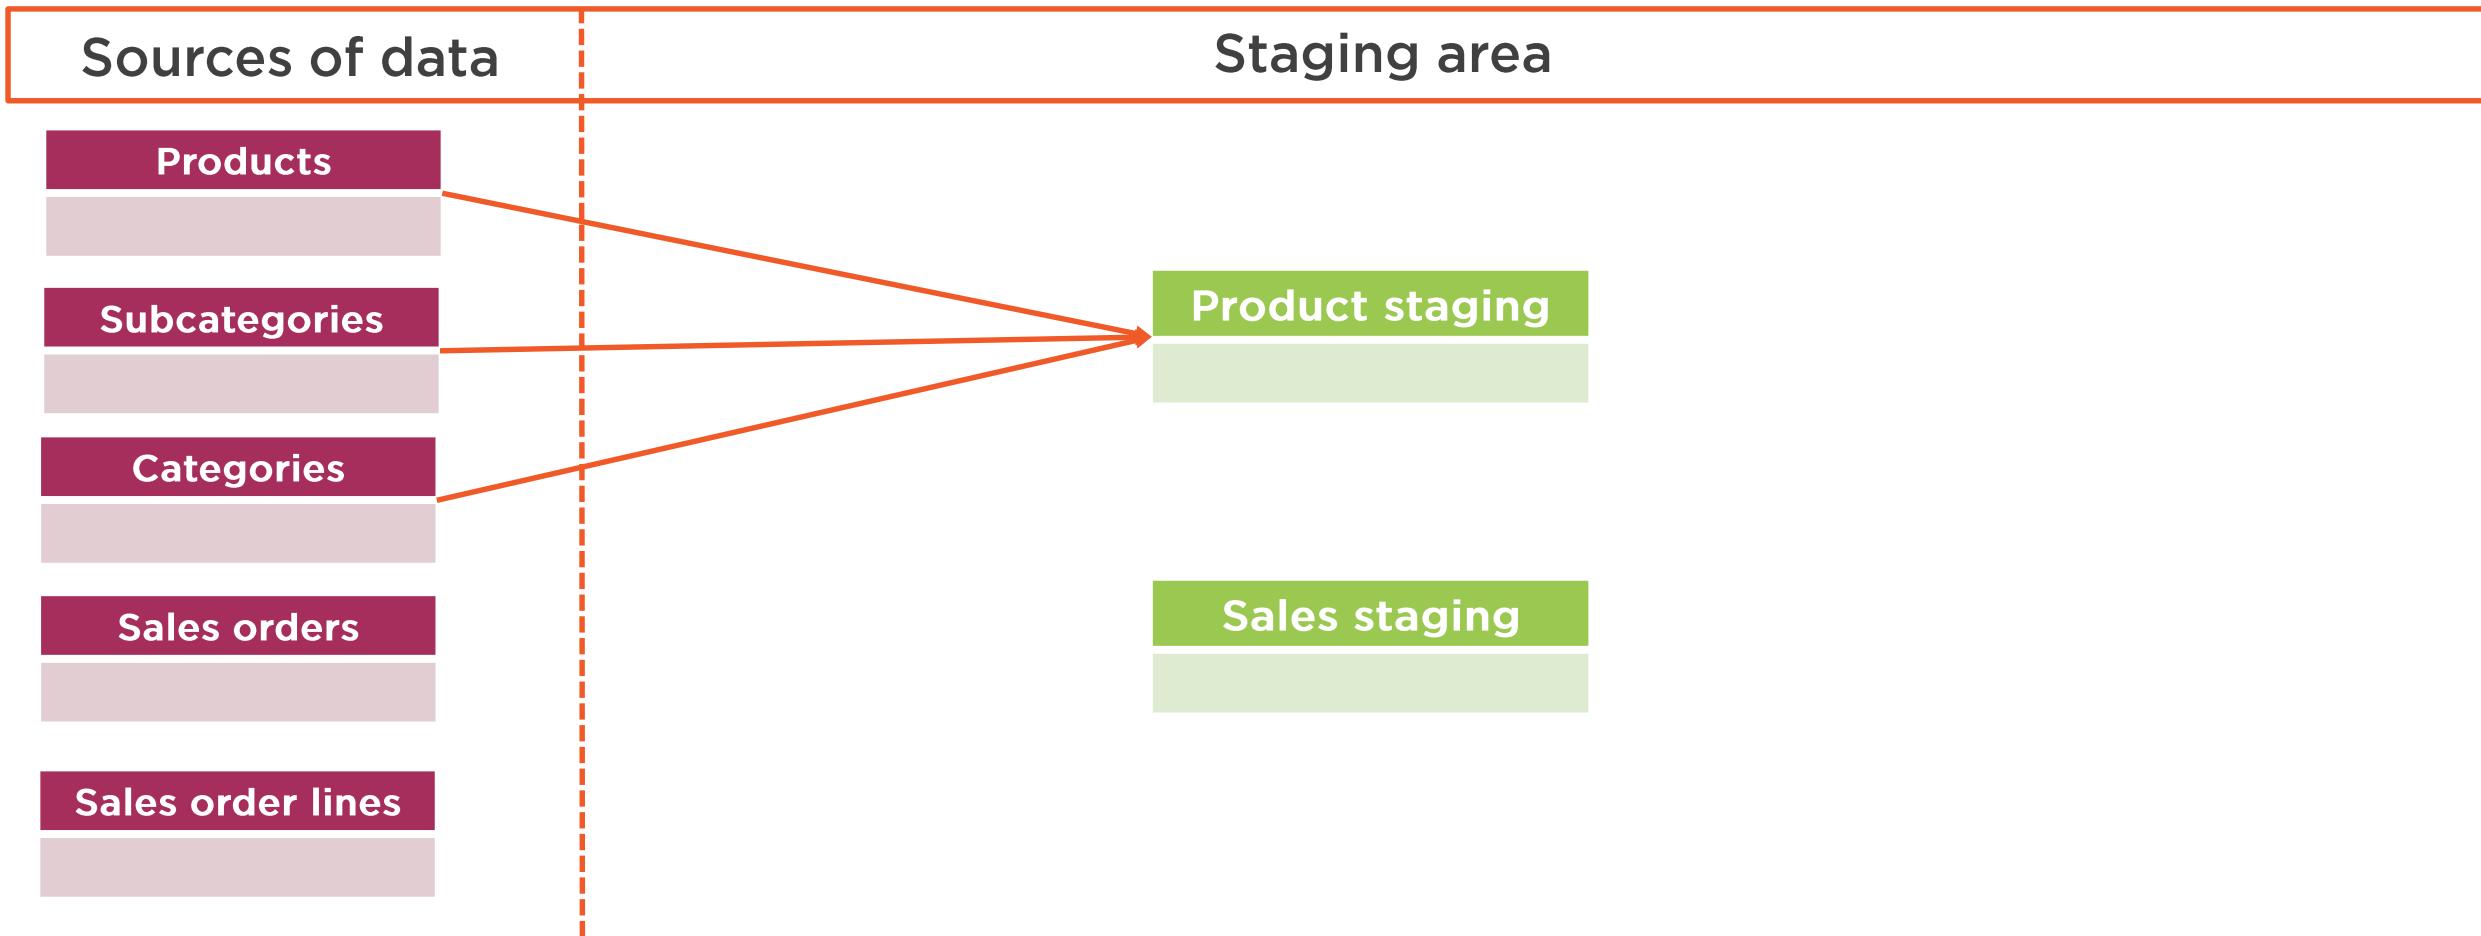
Sales order lines



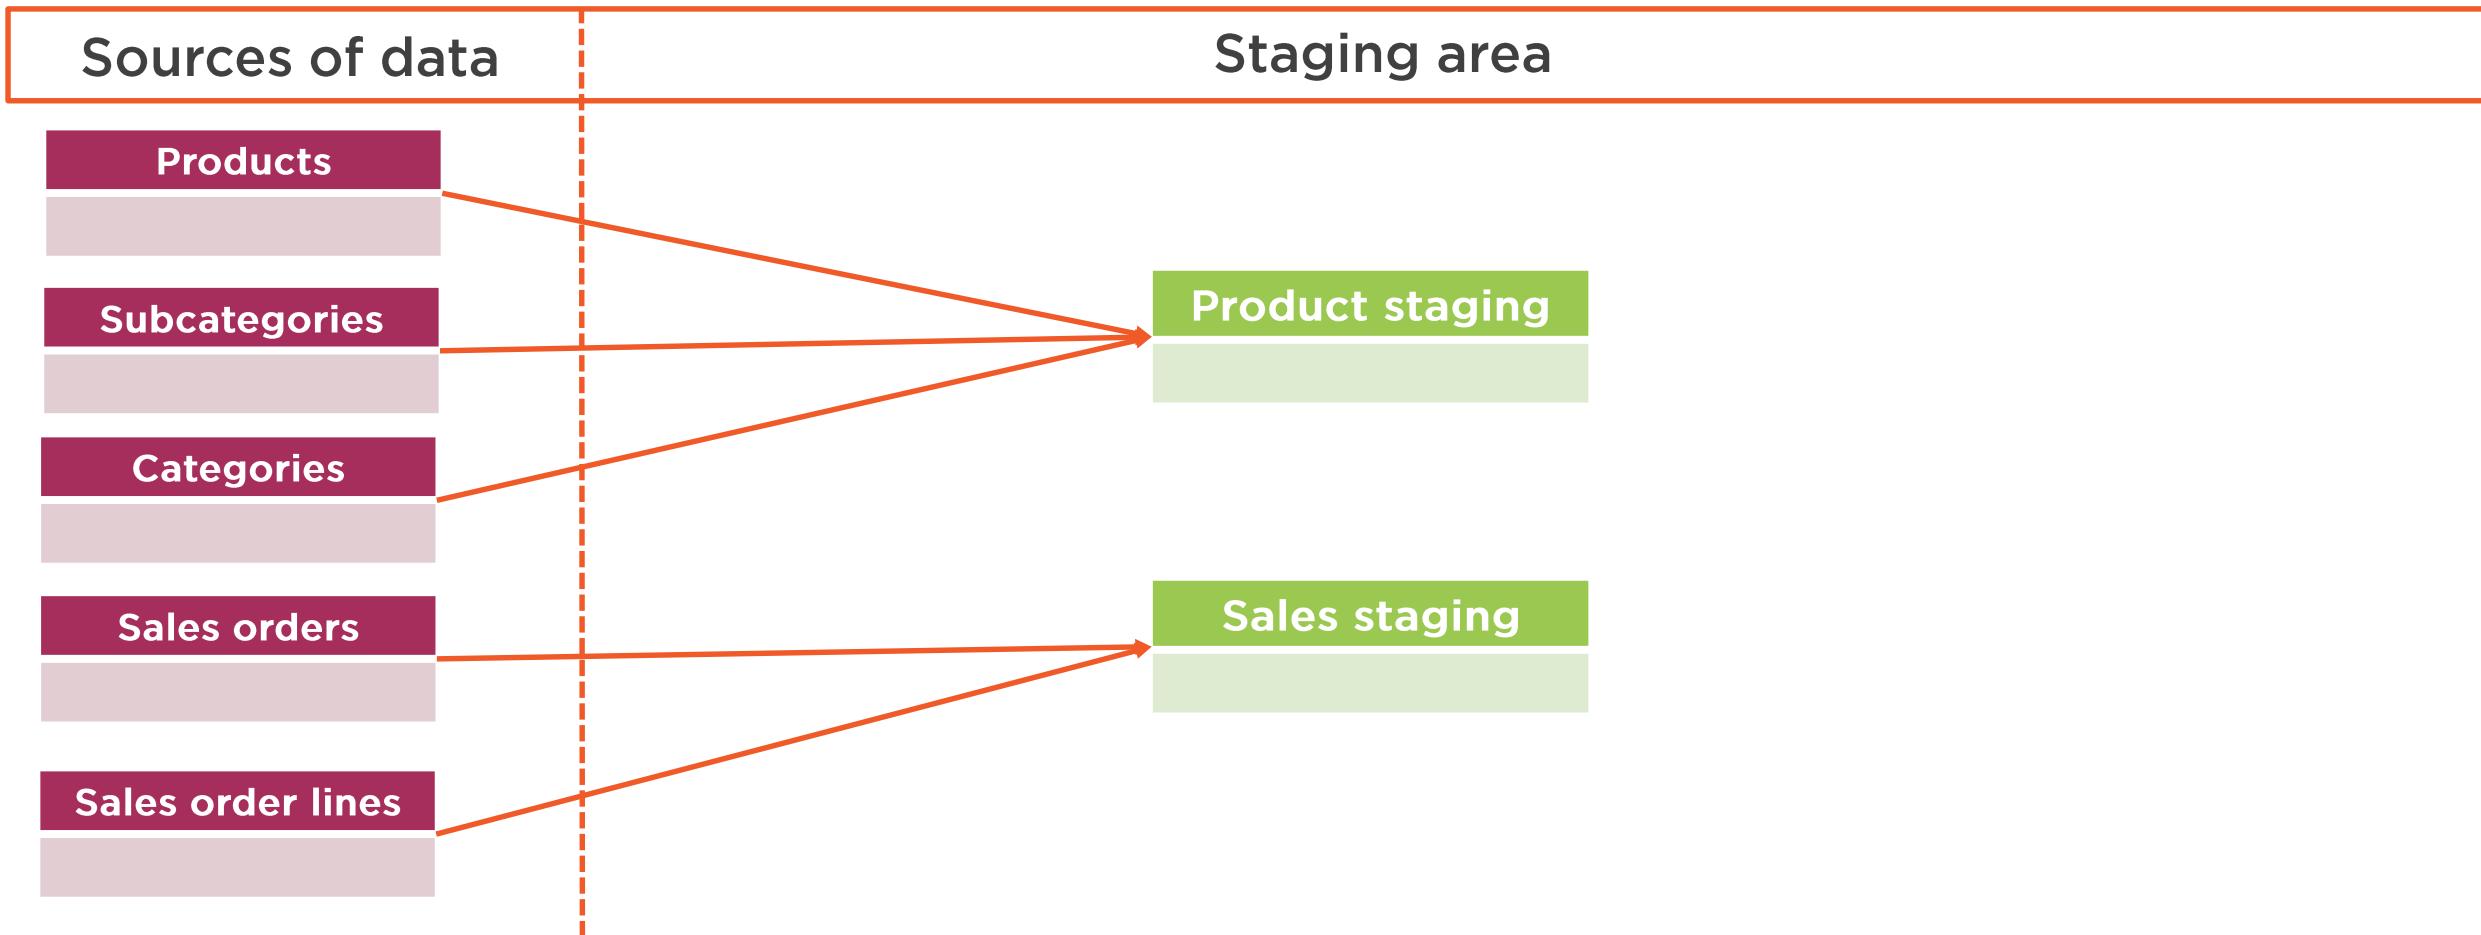
Overview of an ETL System



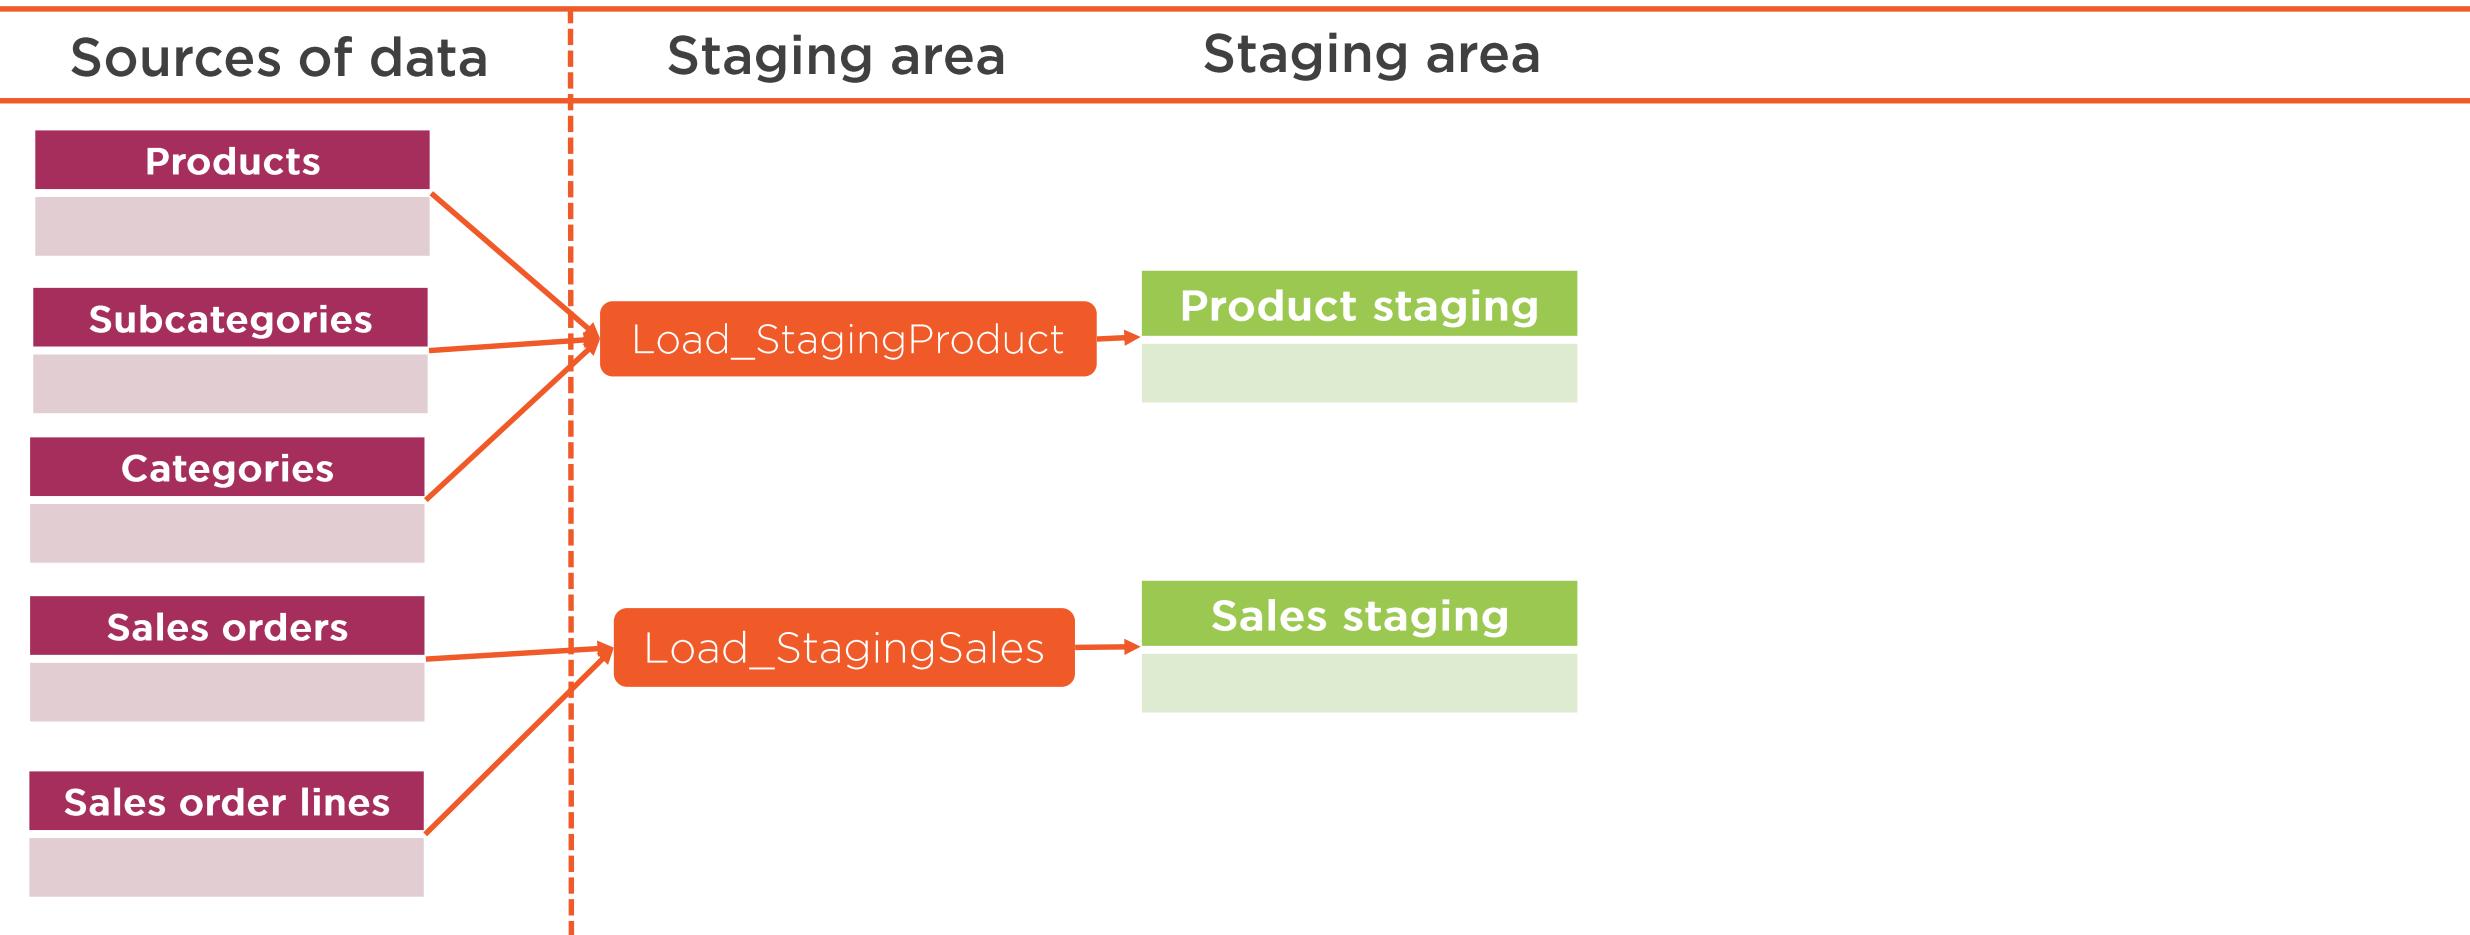
Overview of an ETL System



Overview of an ETL System

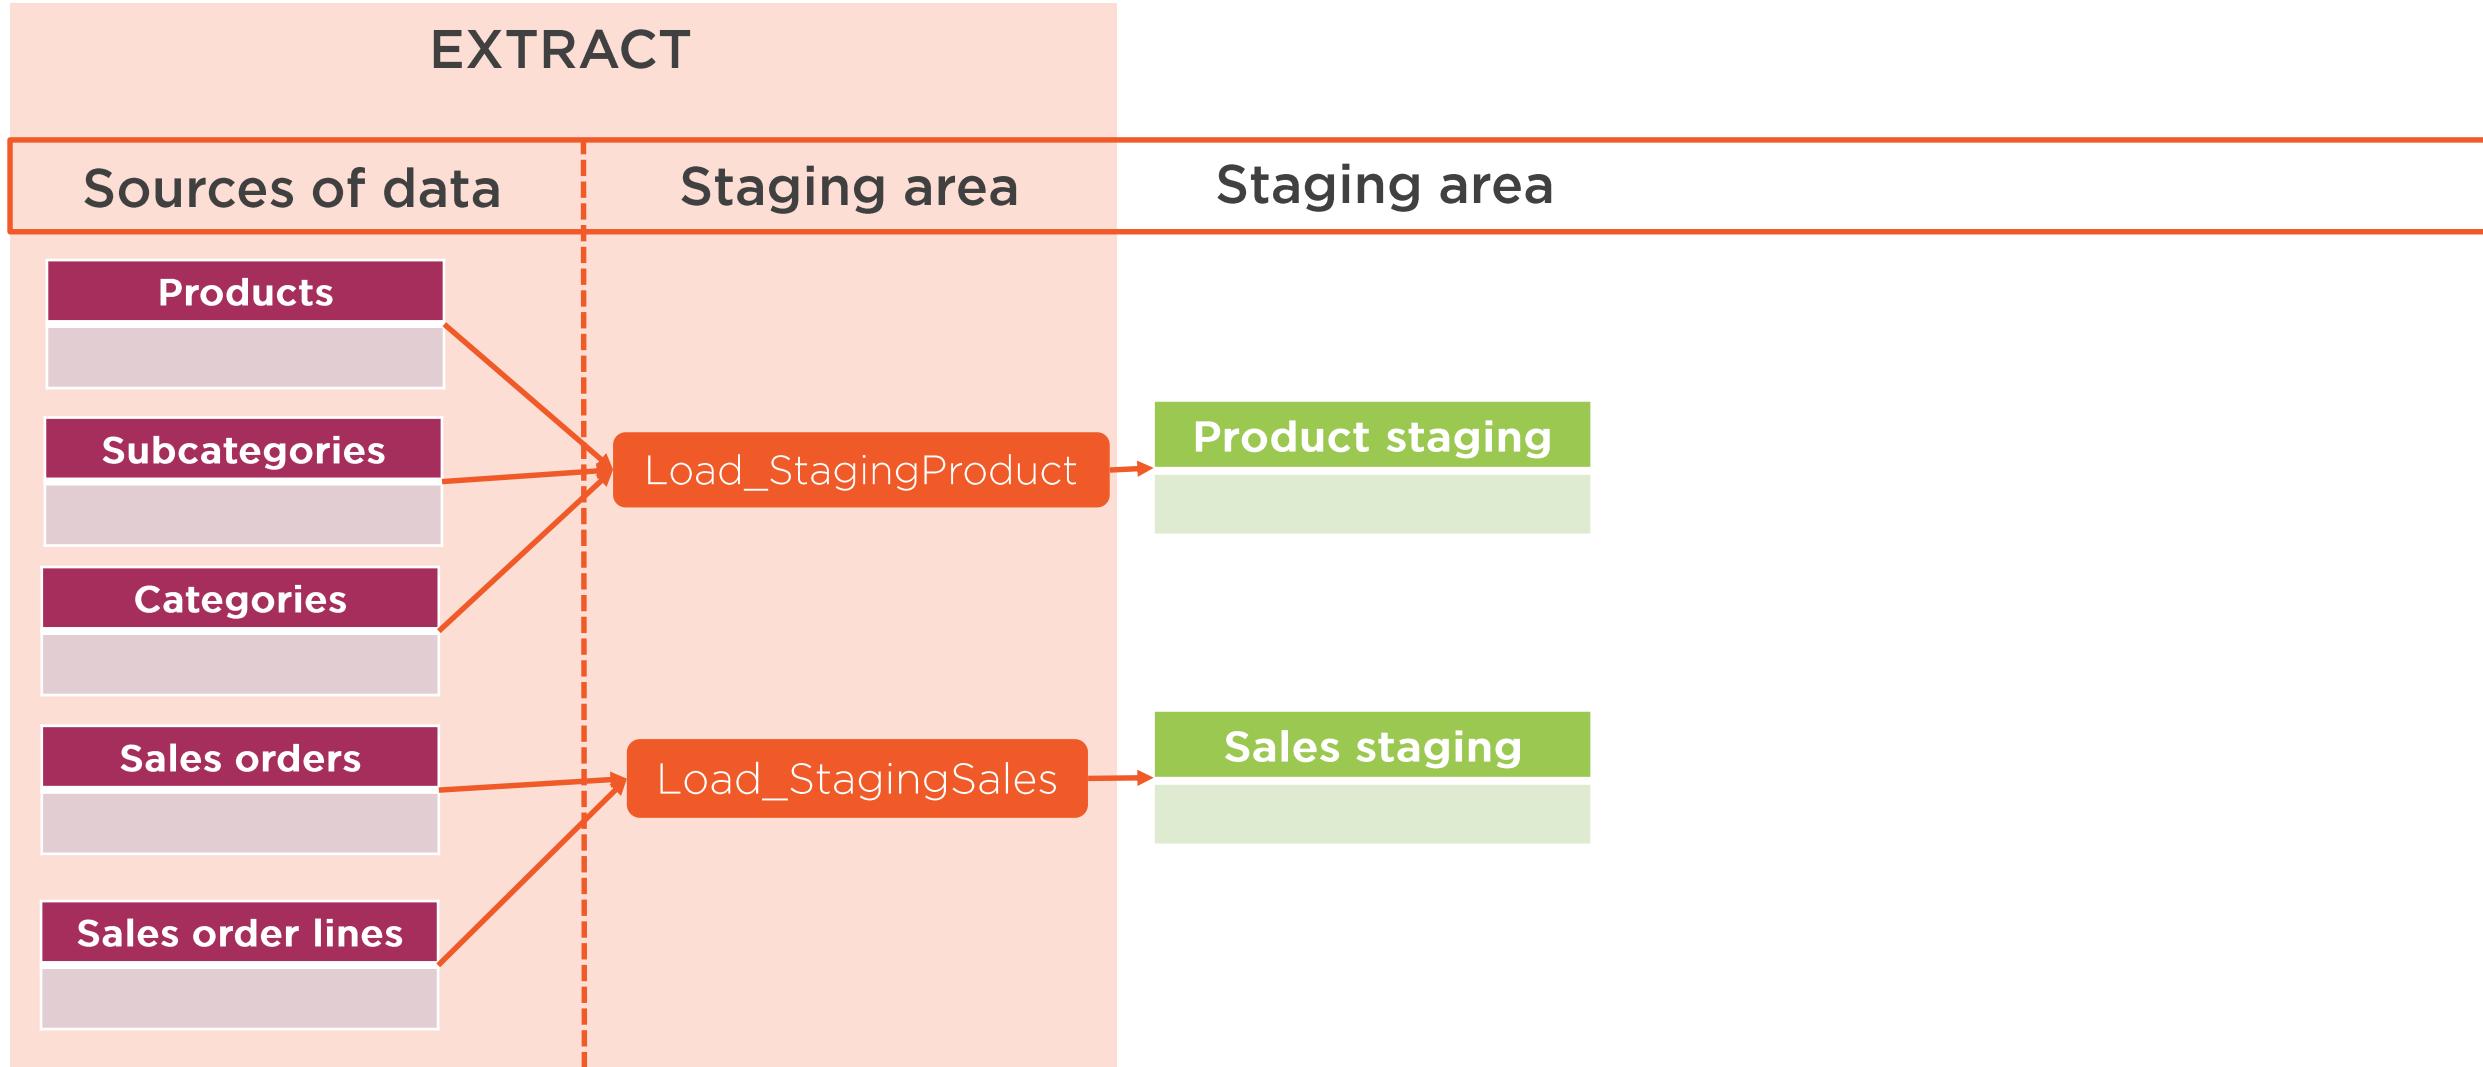


Overview of an ETL System

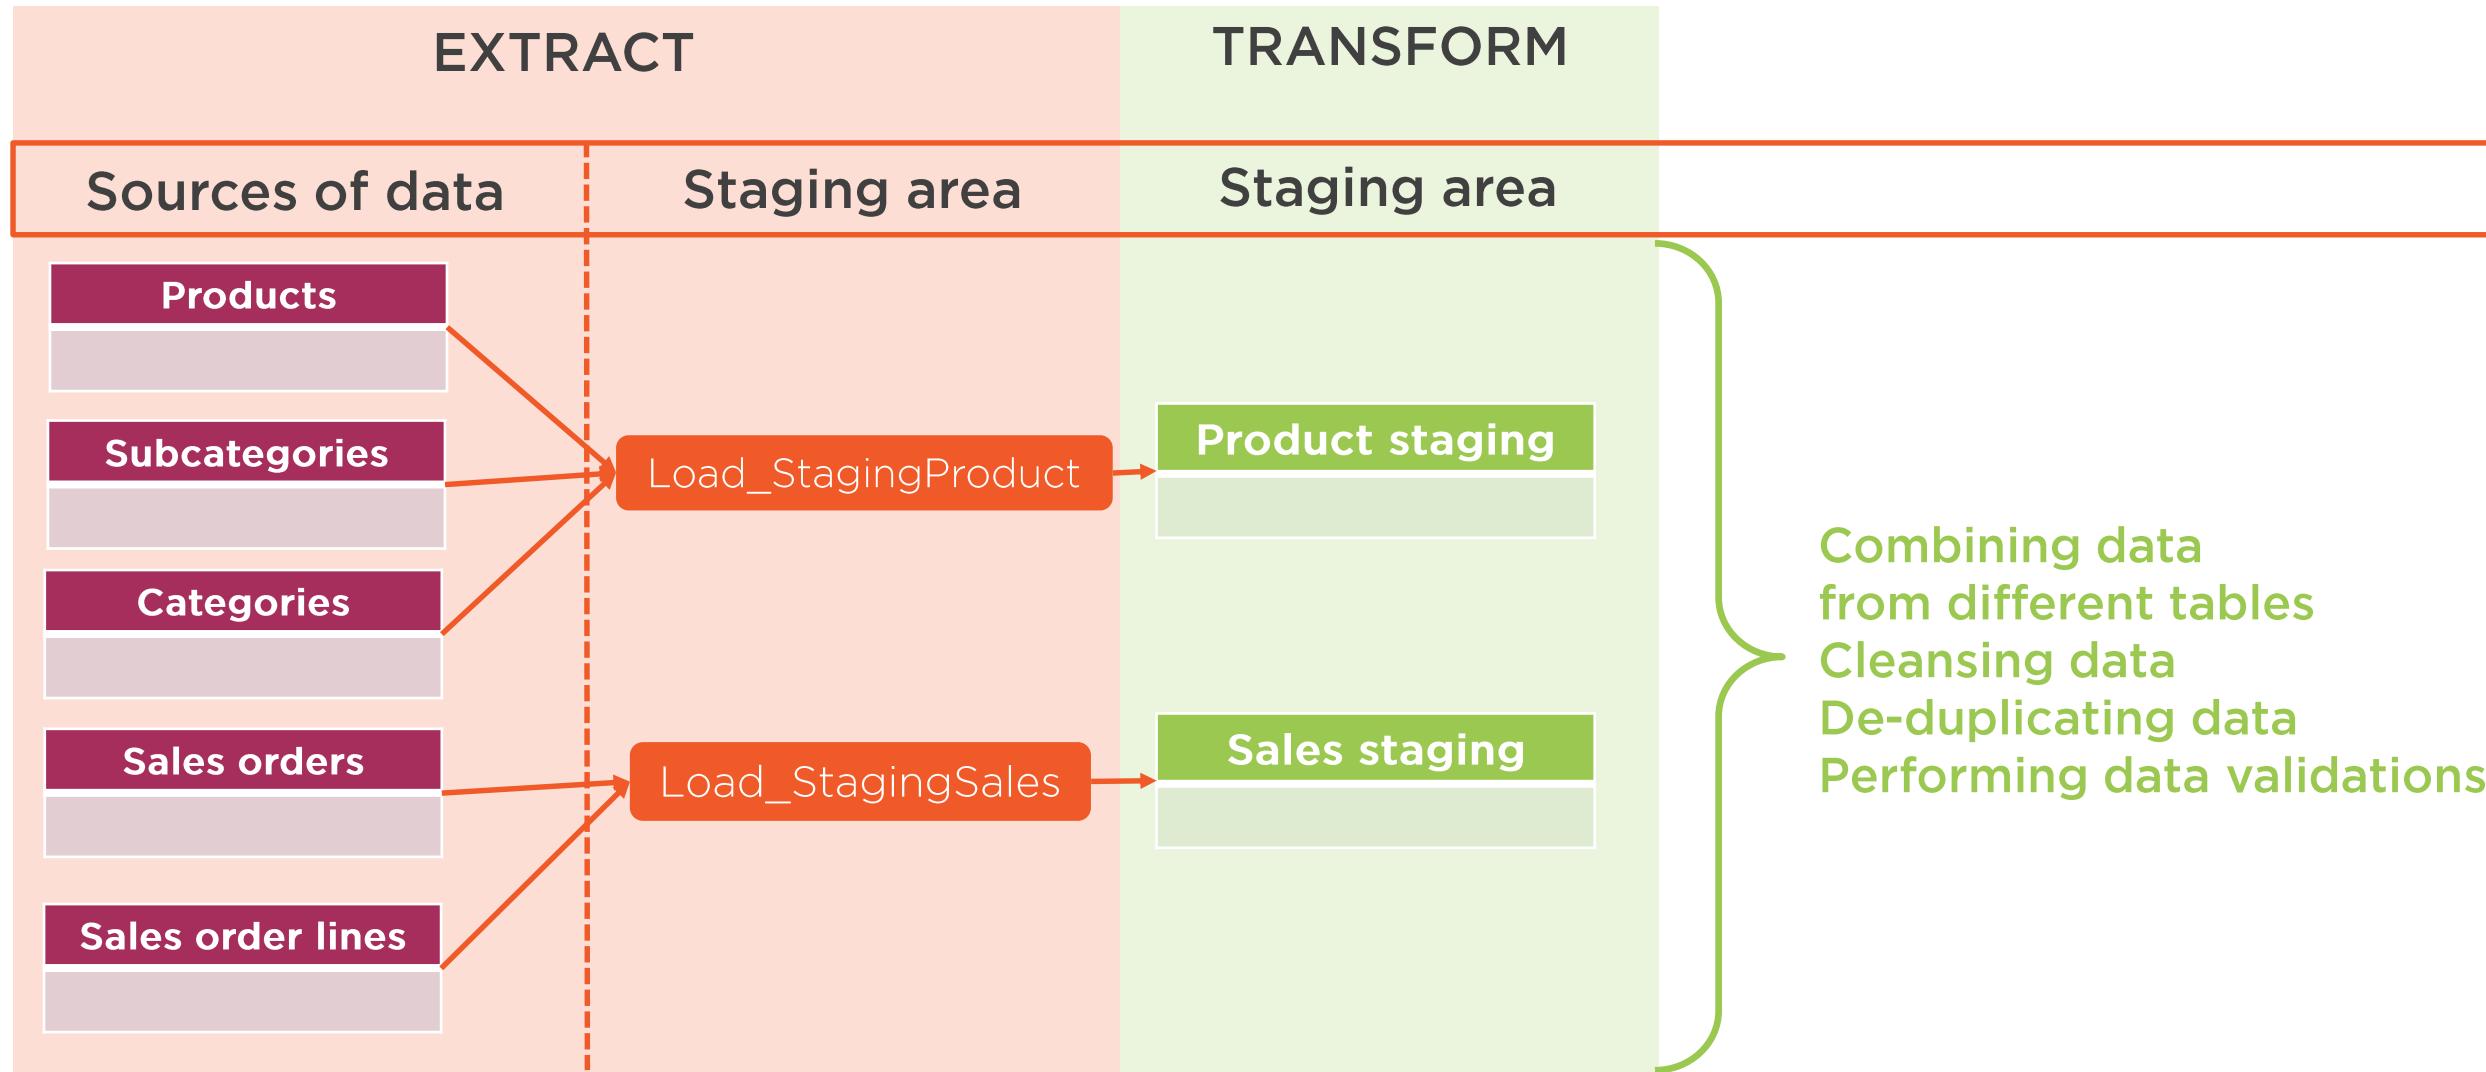


Overview of an ETL System

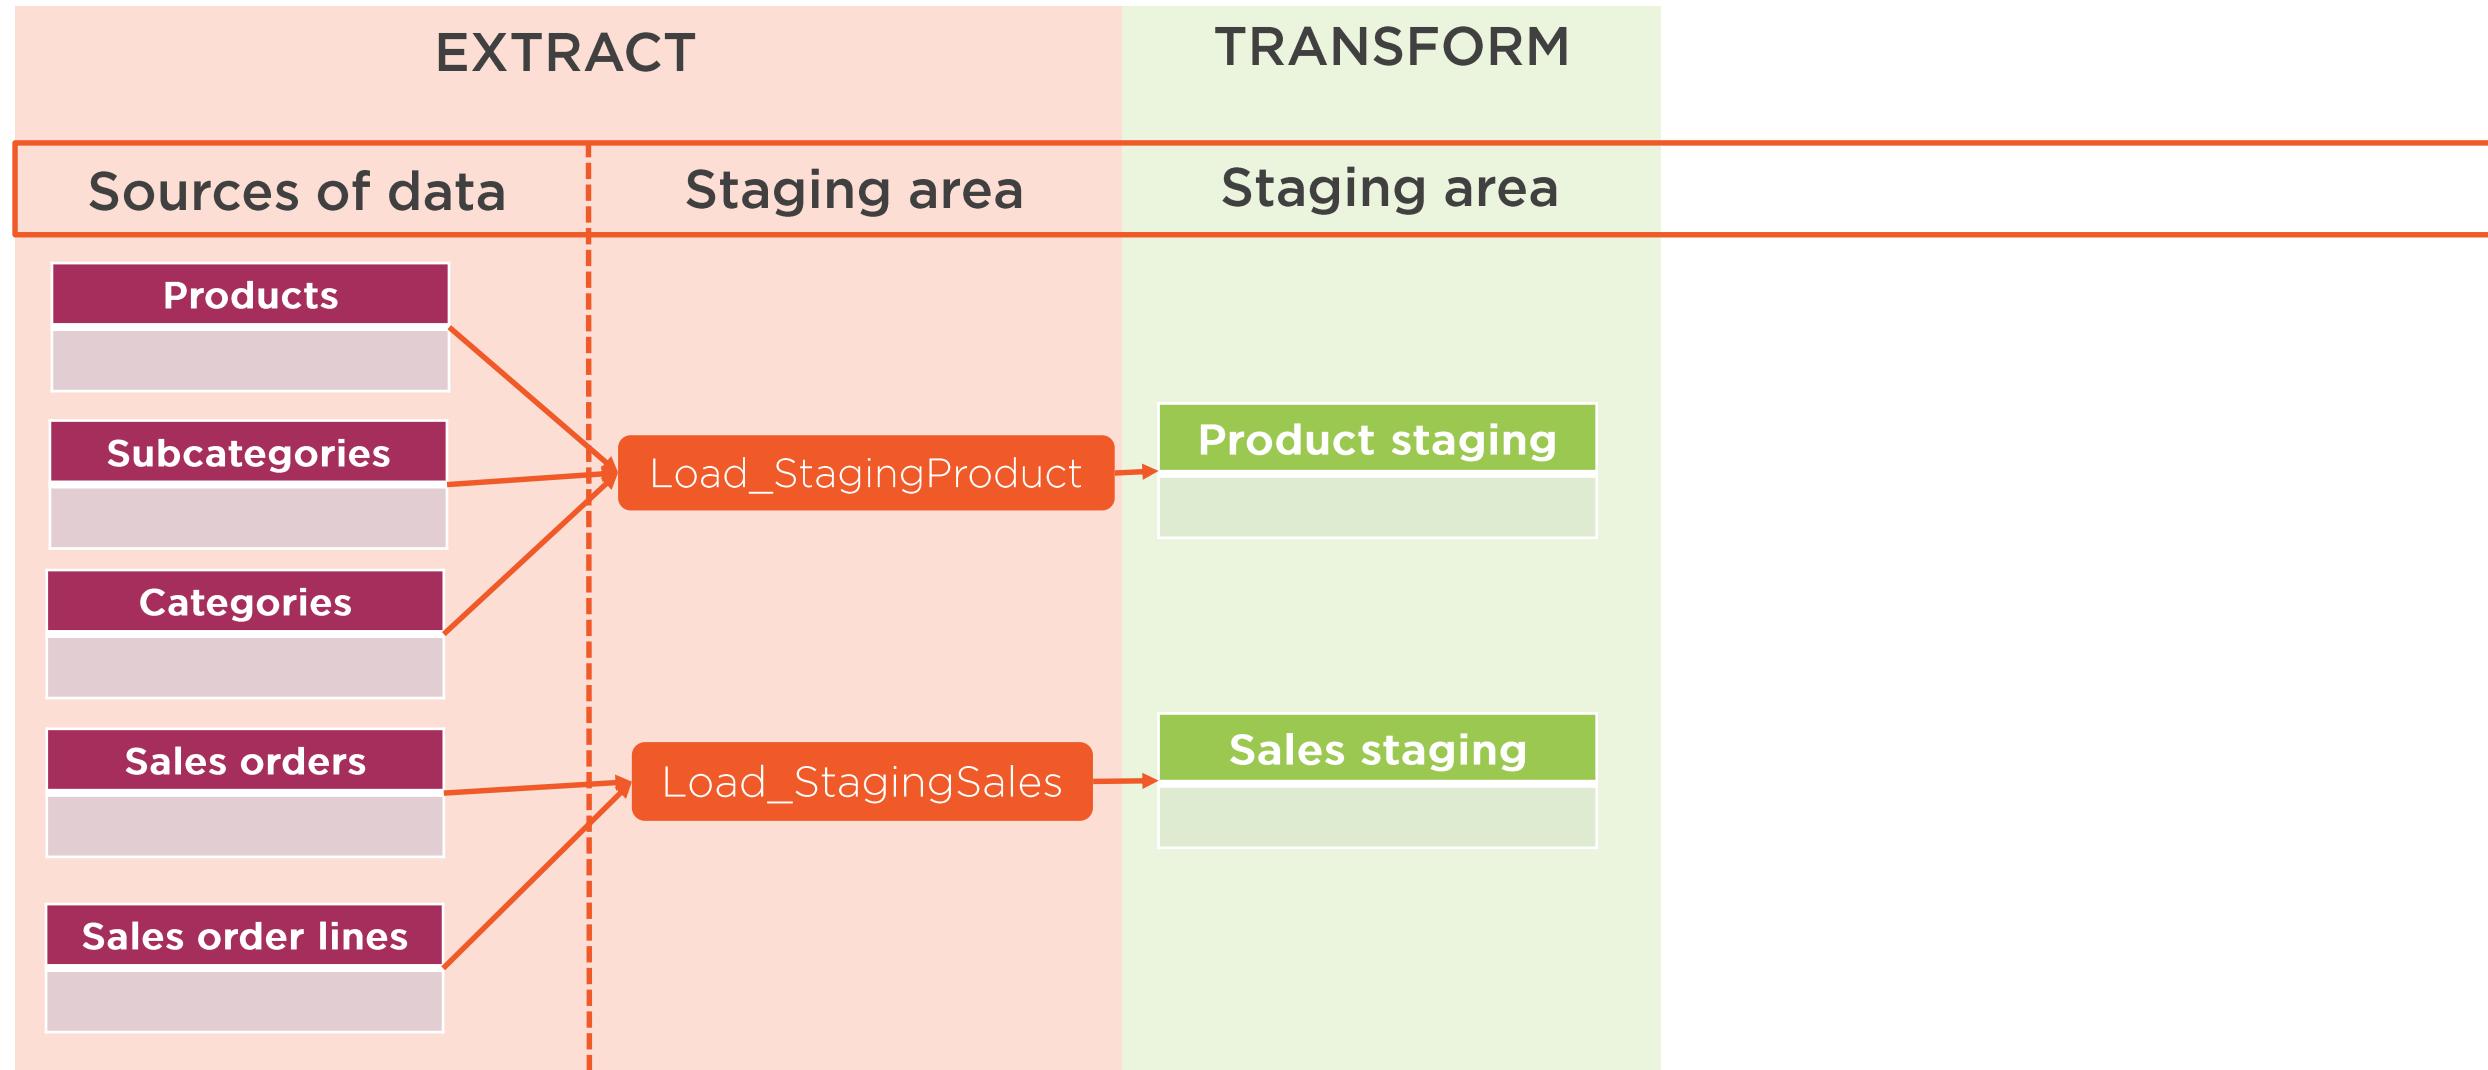
EXTRACT



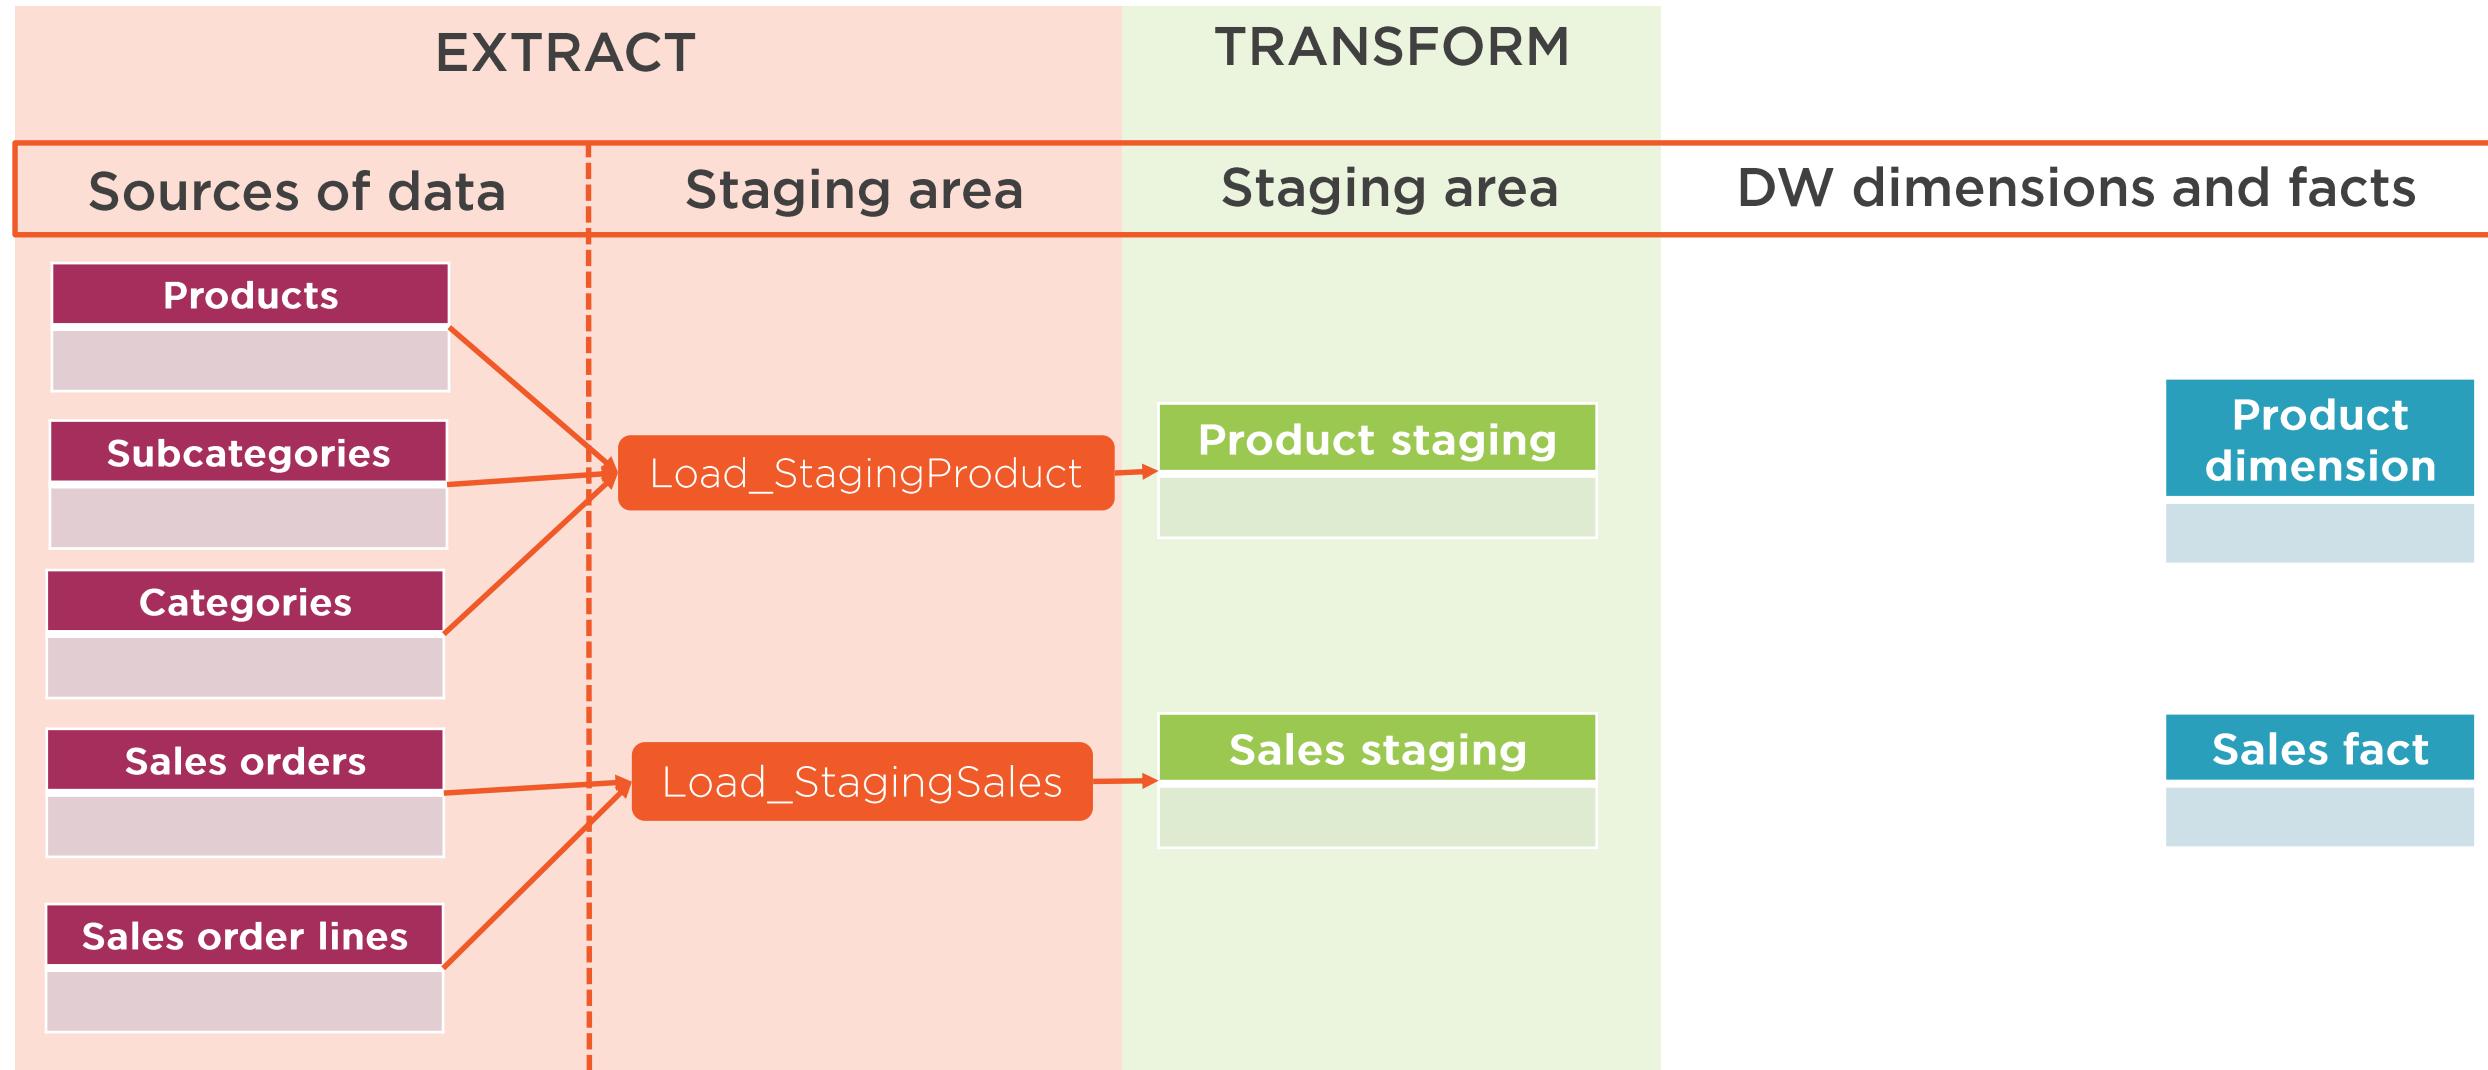
Overview of an ETL System



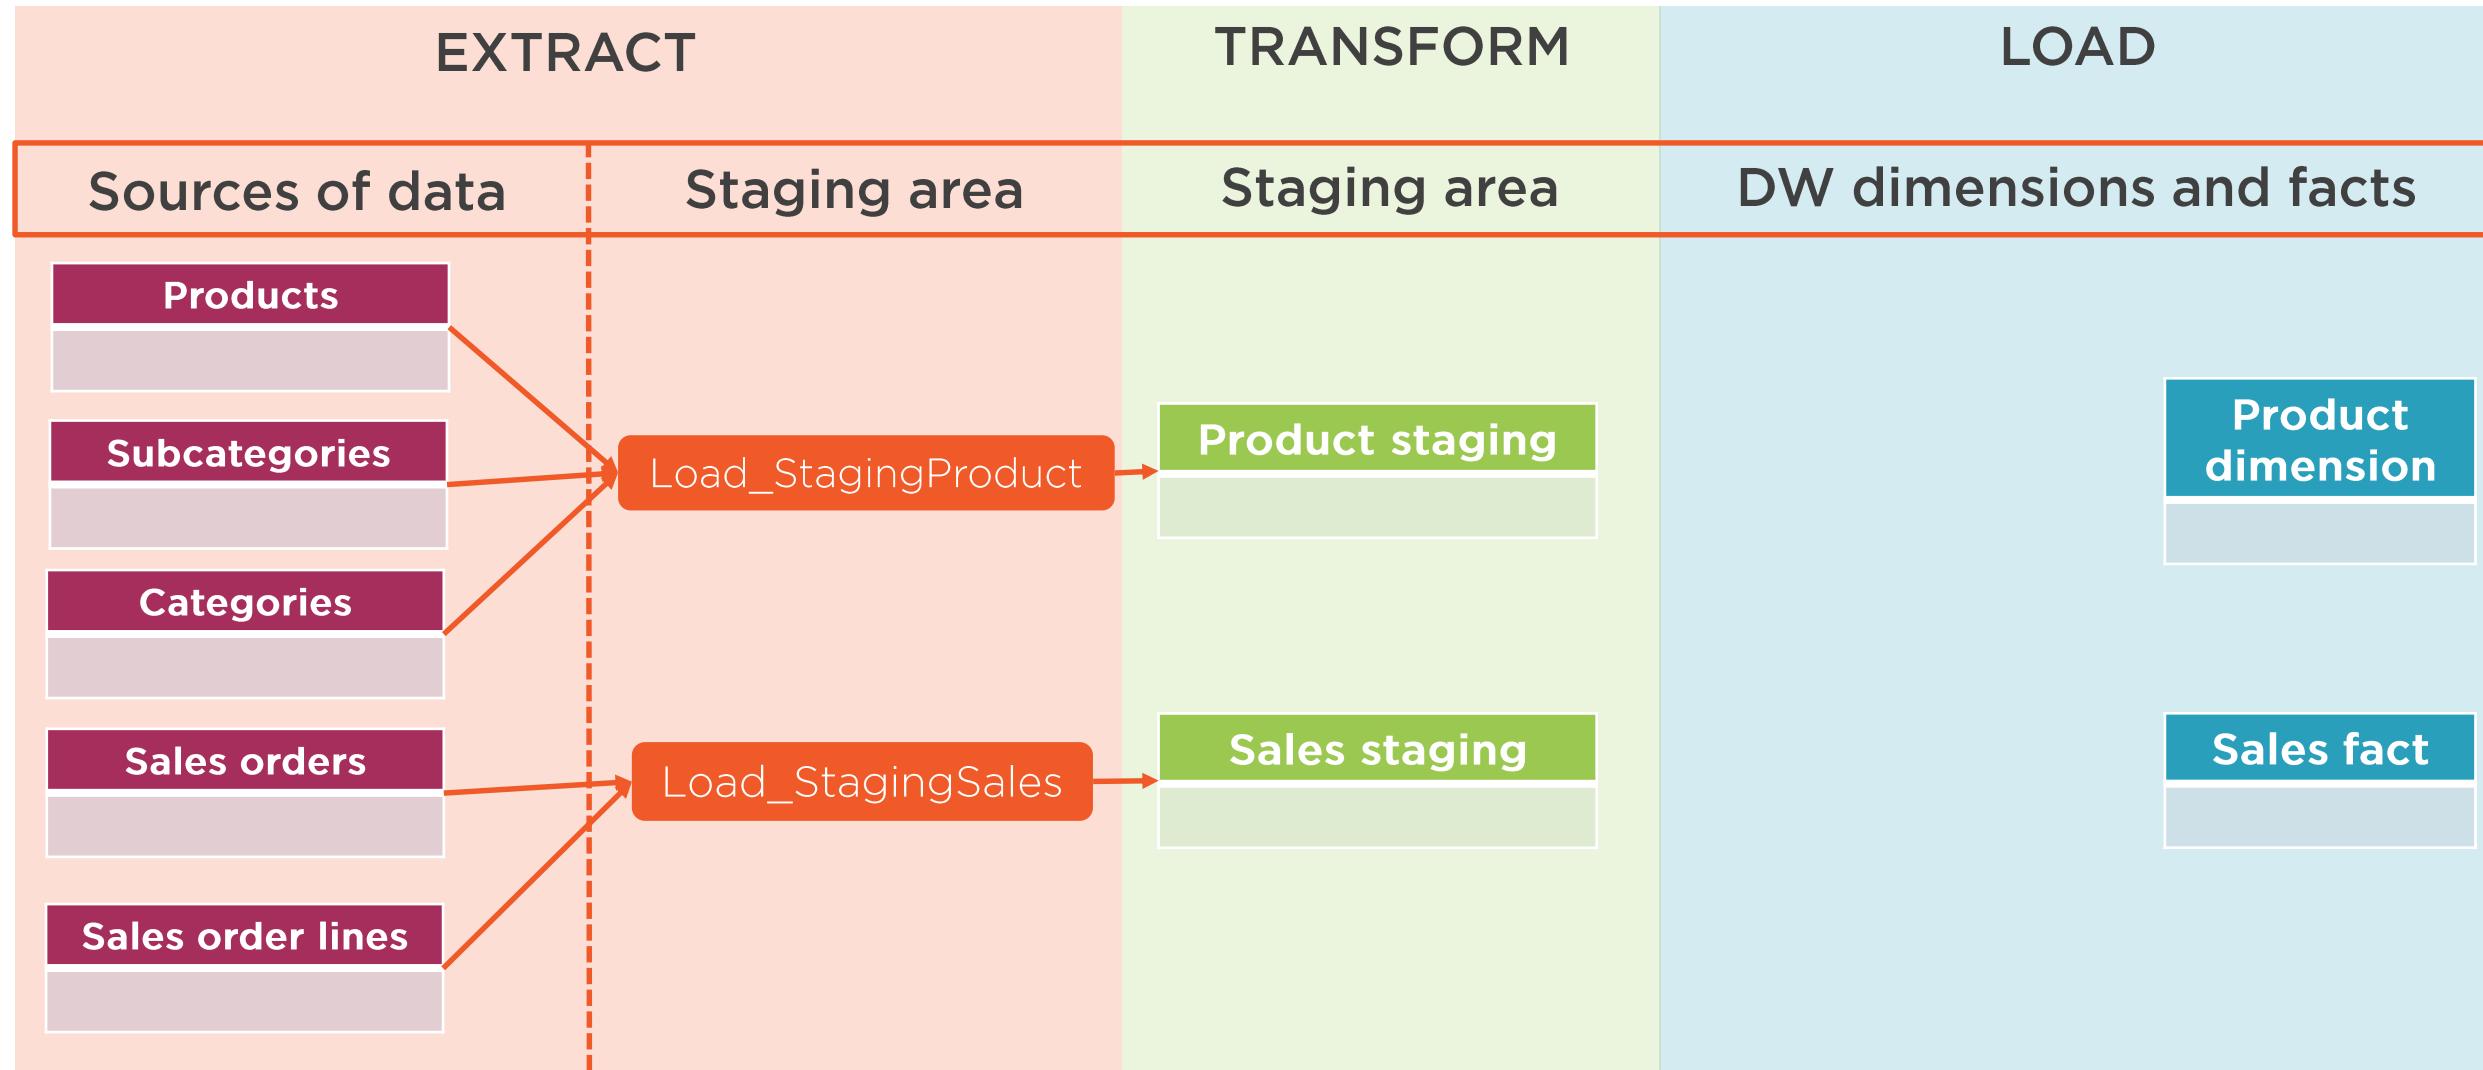
Overview of an ETL System



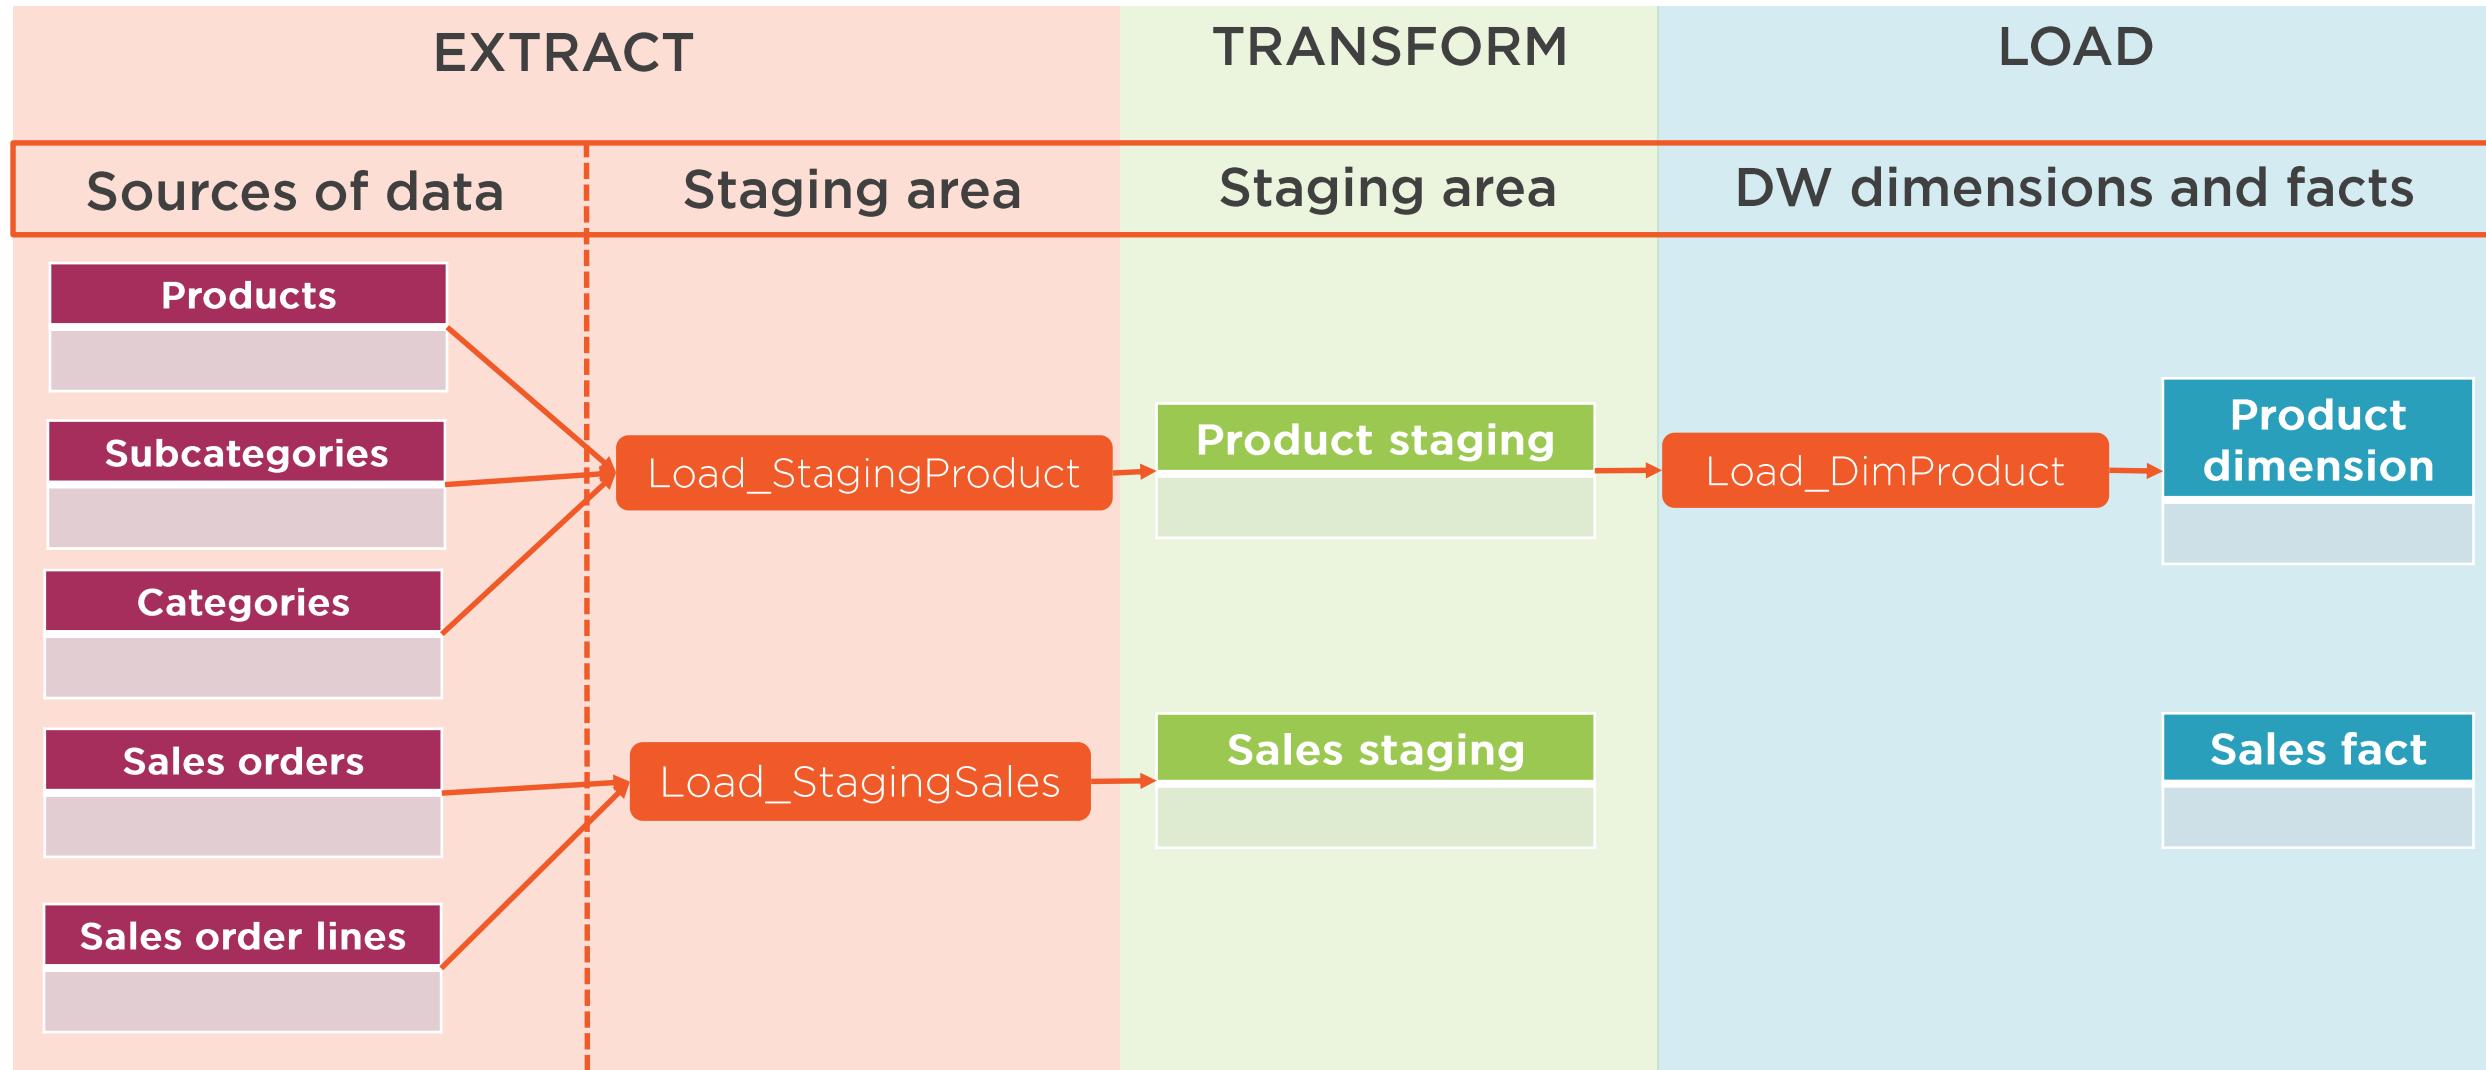
Overview of an ETL System



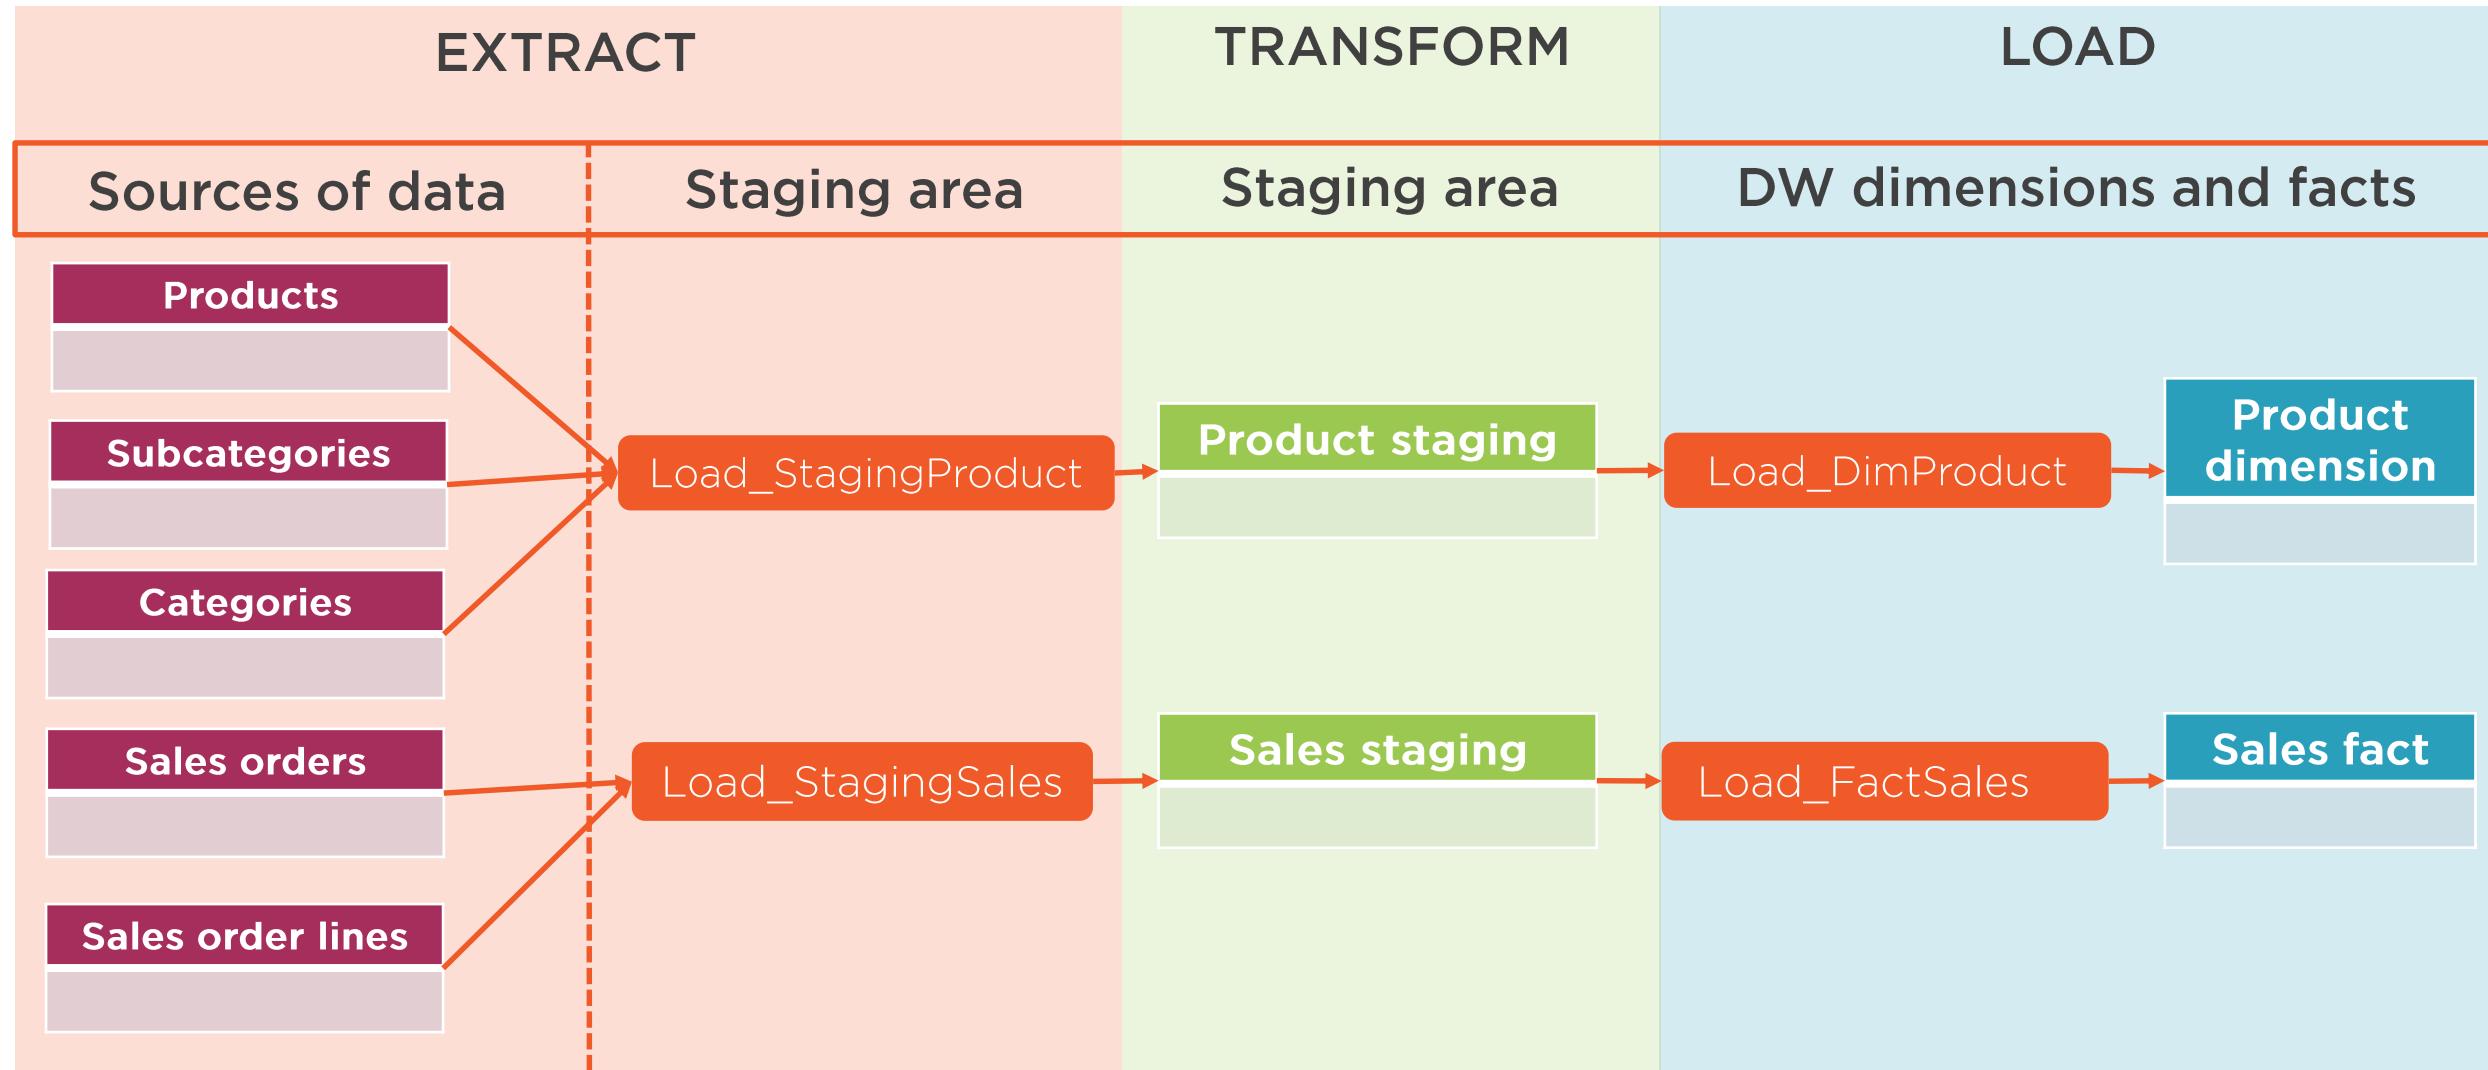
Overview of an ETL System



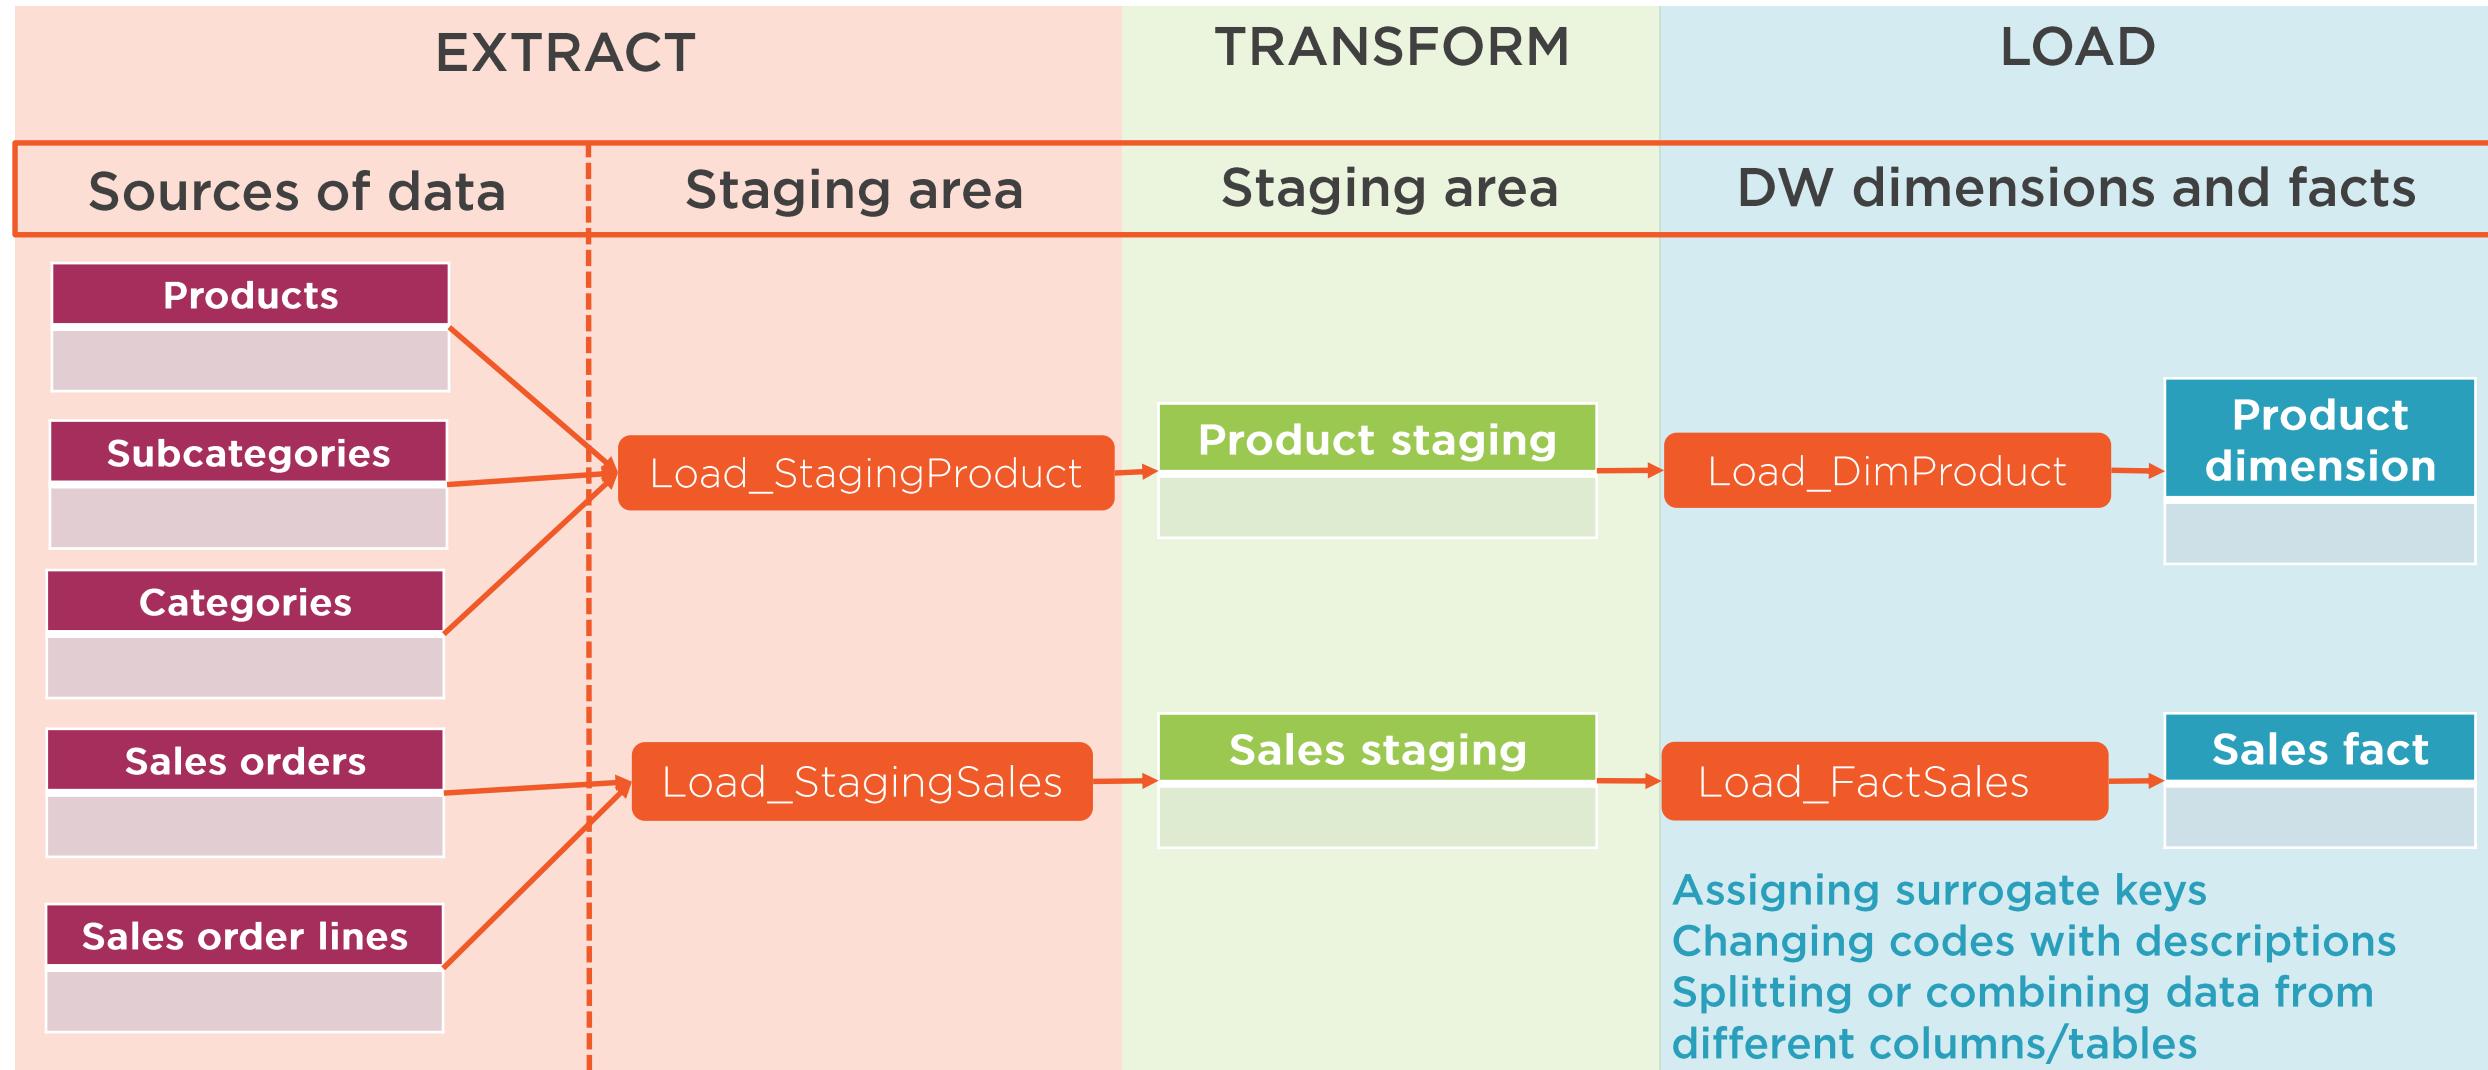
Overview of an ETL System



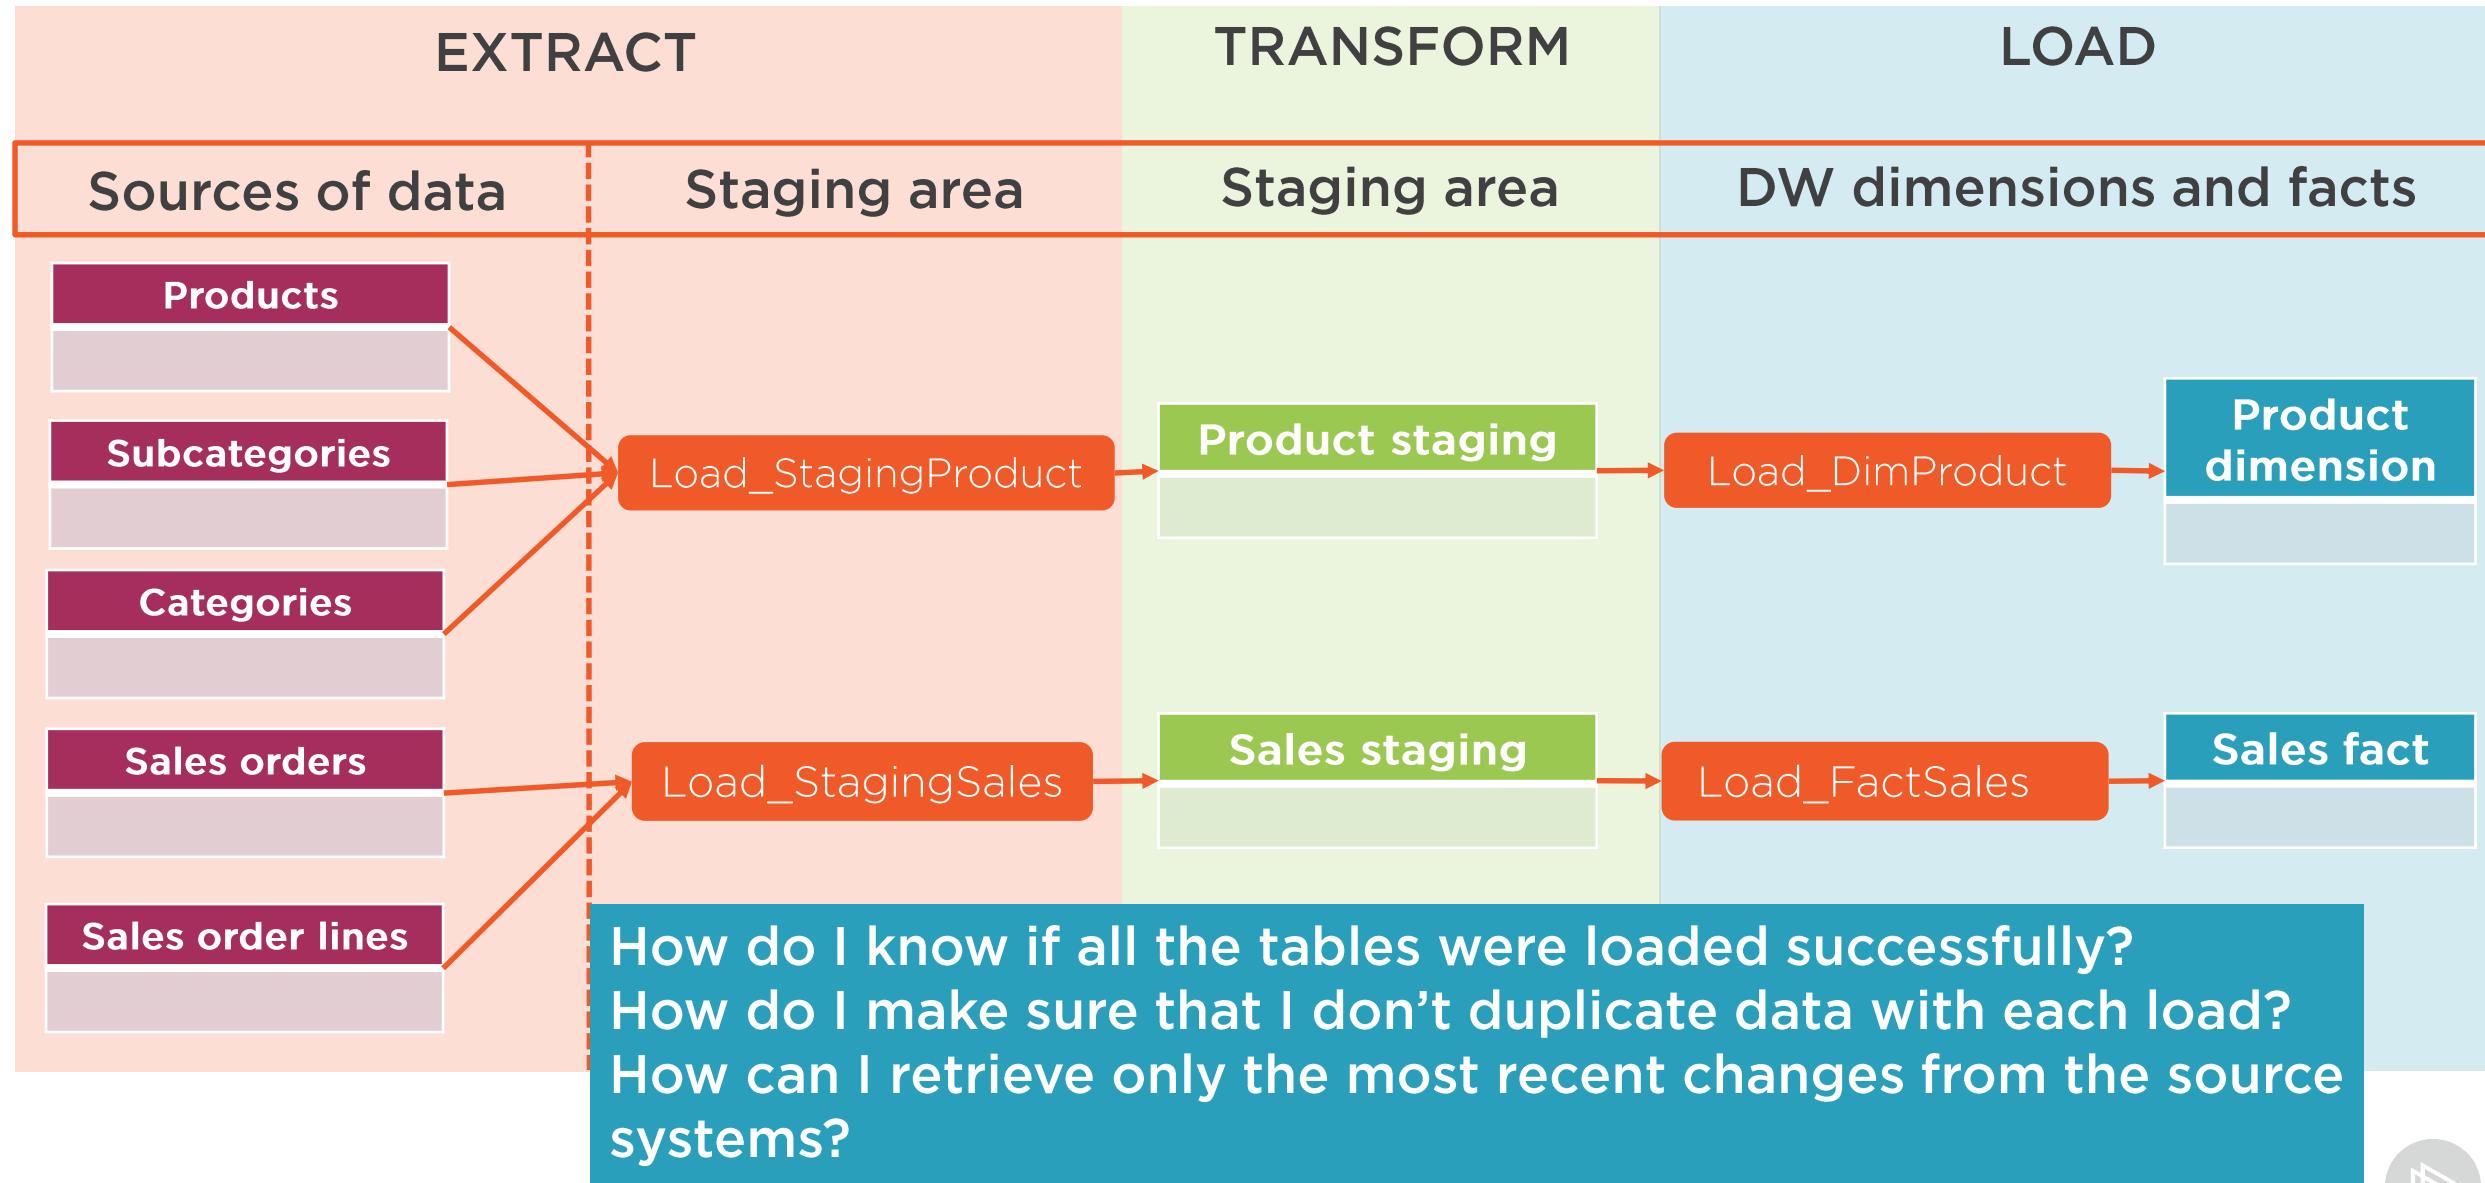
Overview of an ETL System



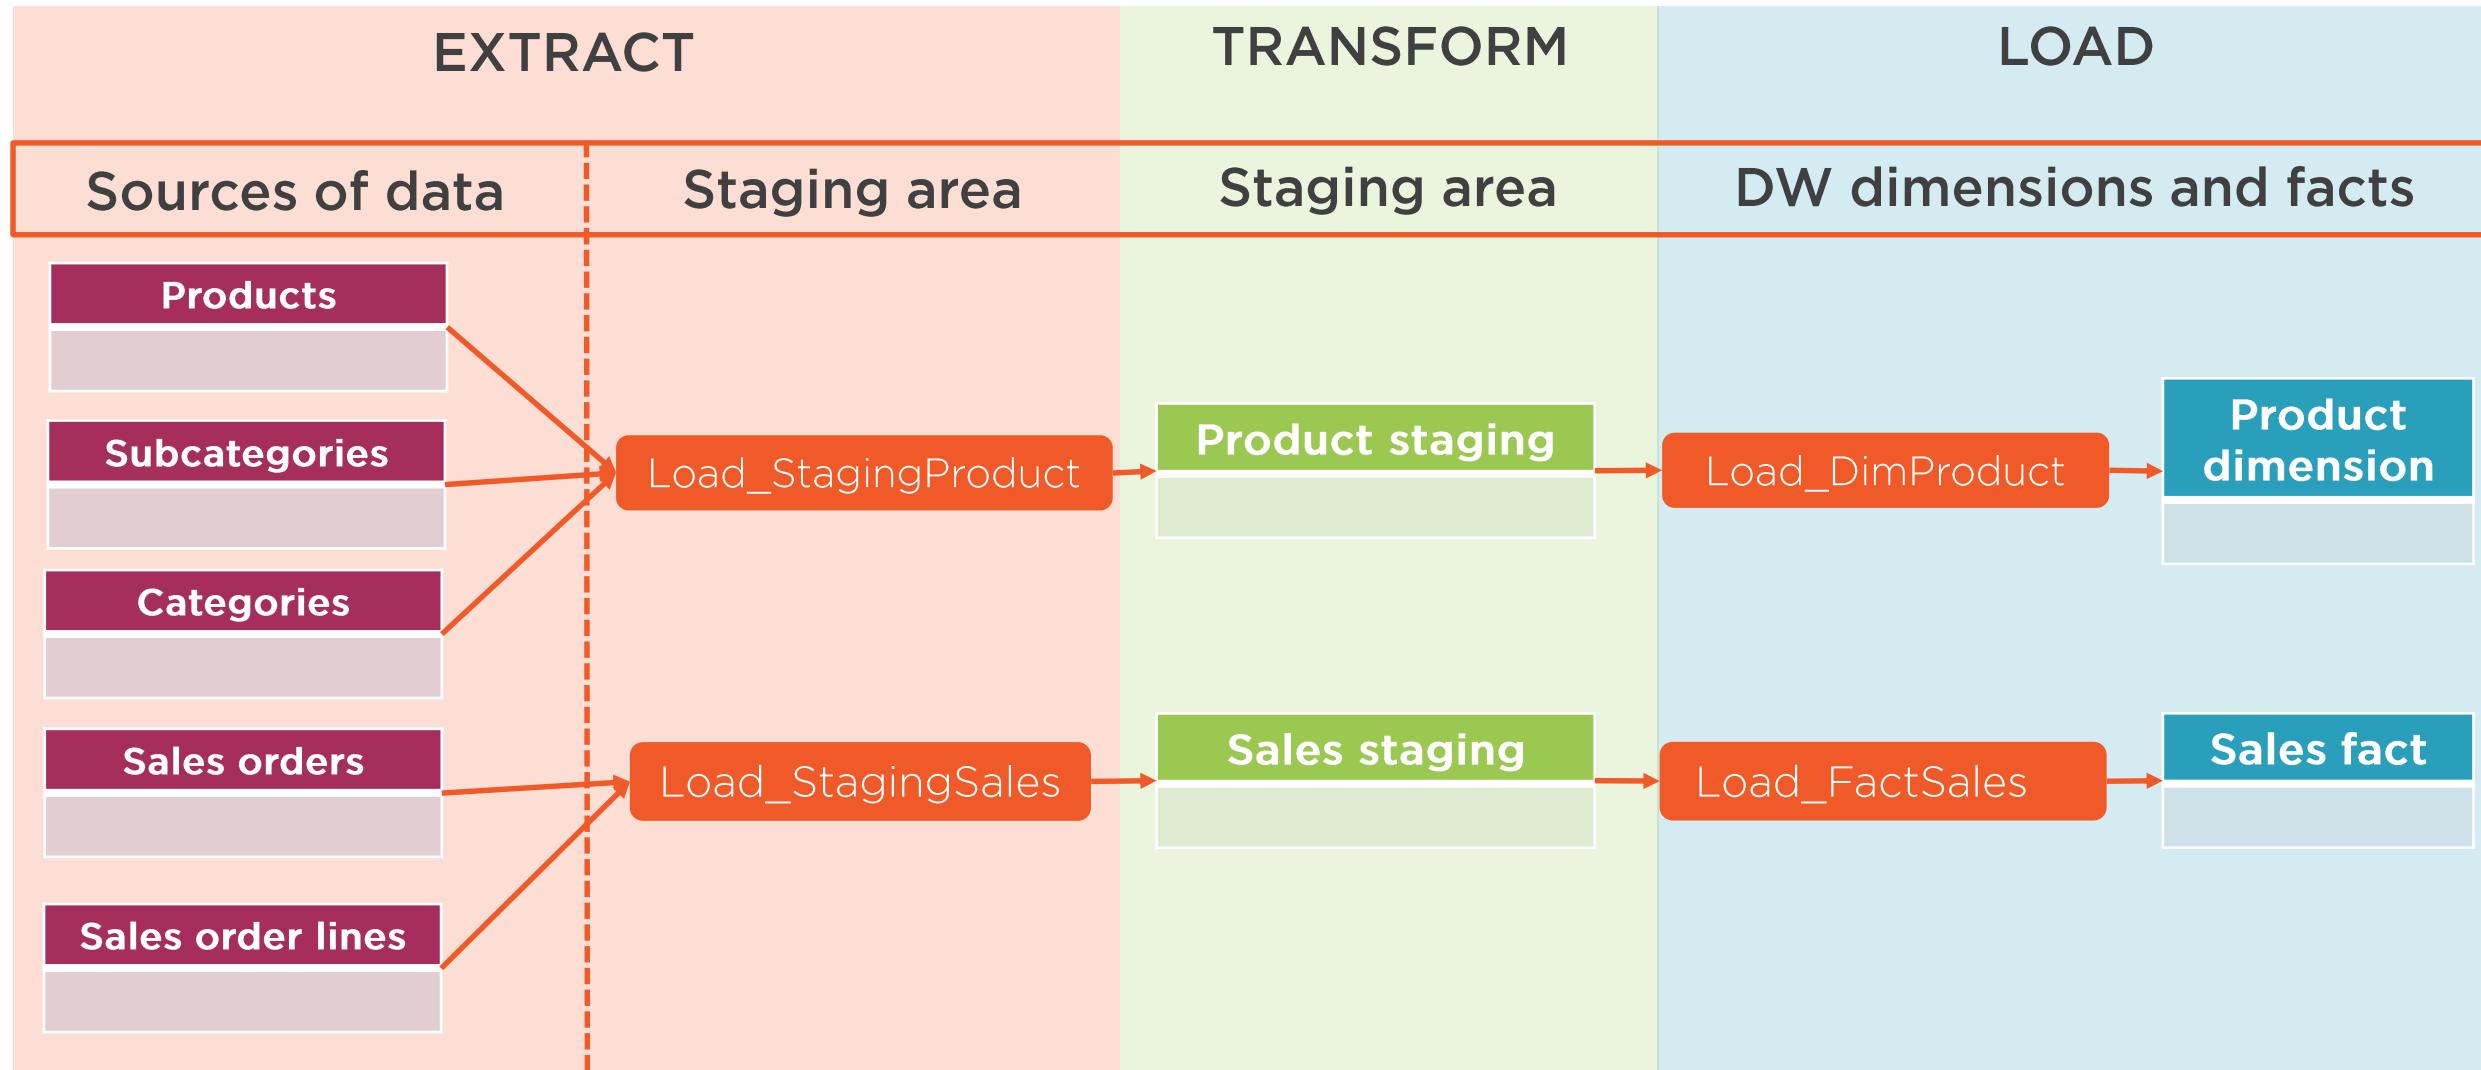
Overview of an ETL System



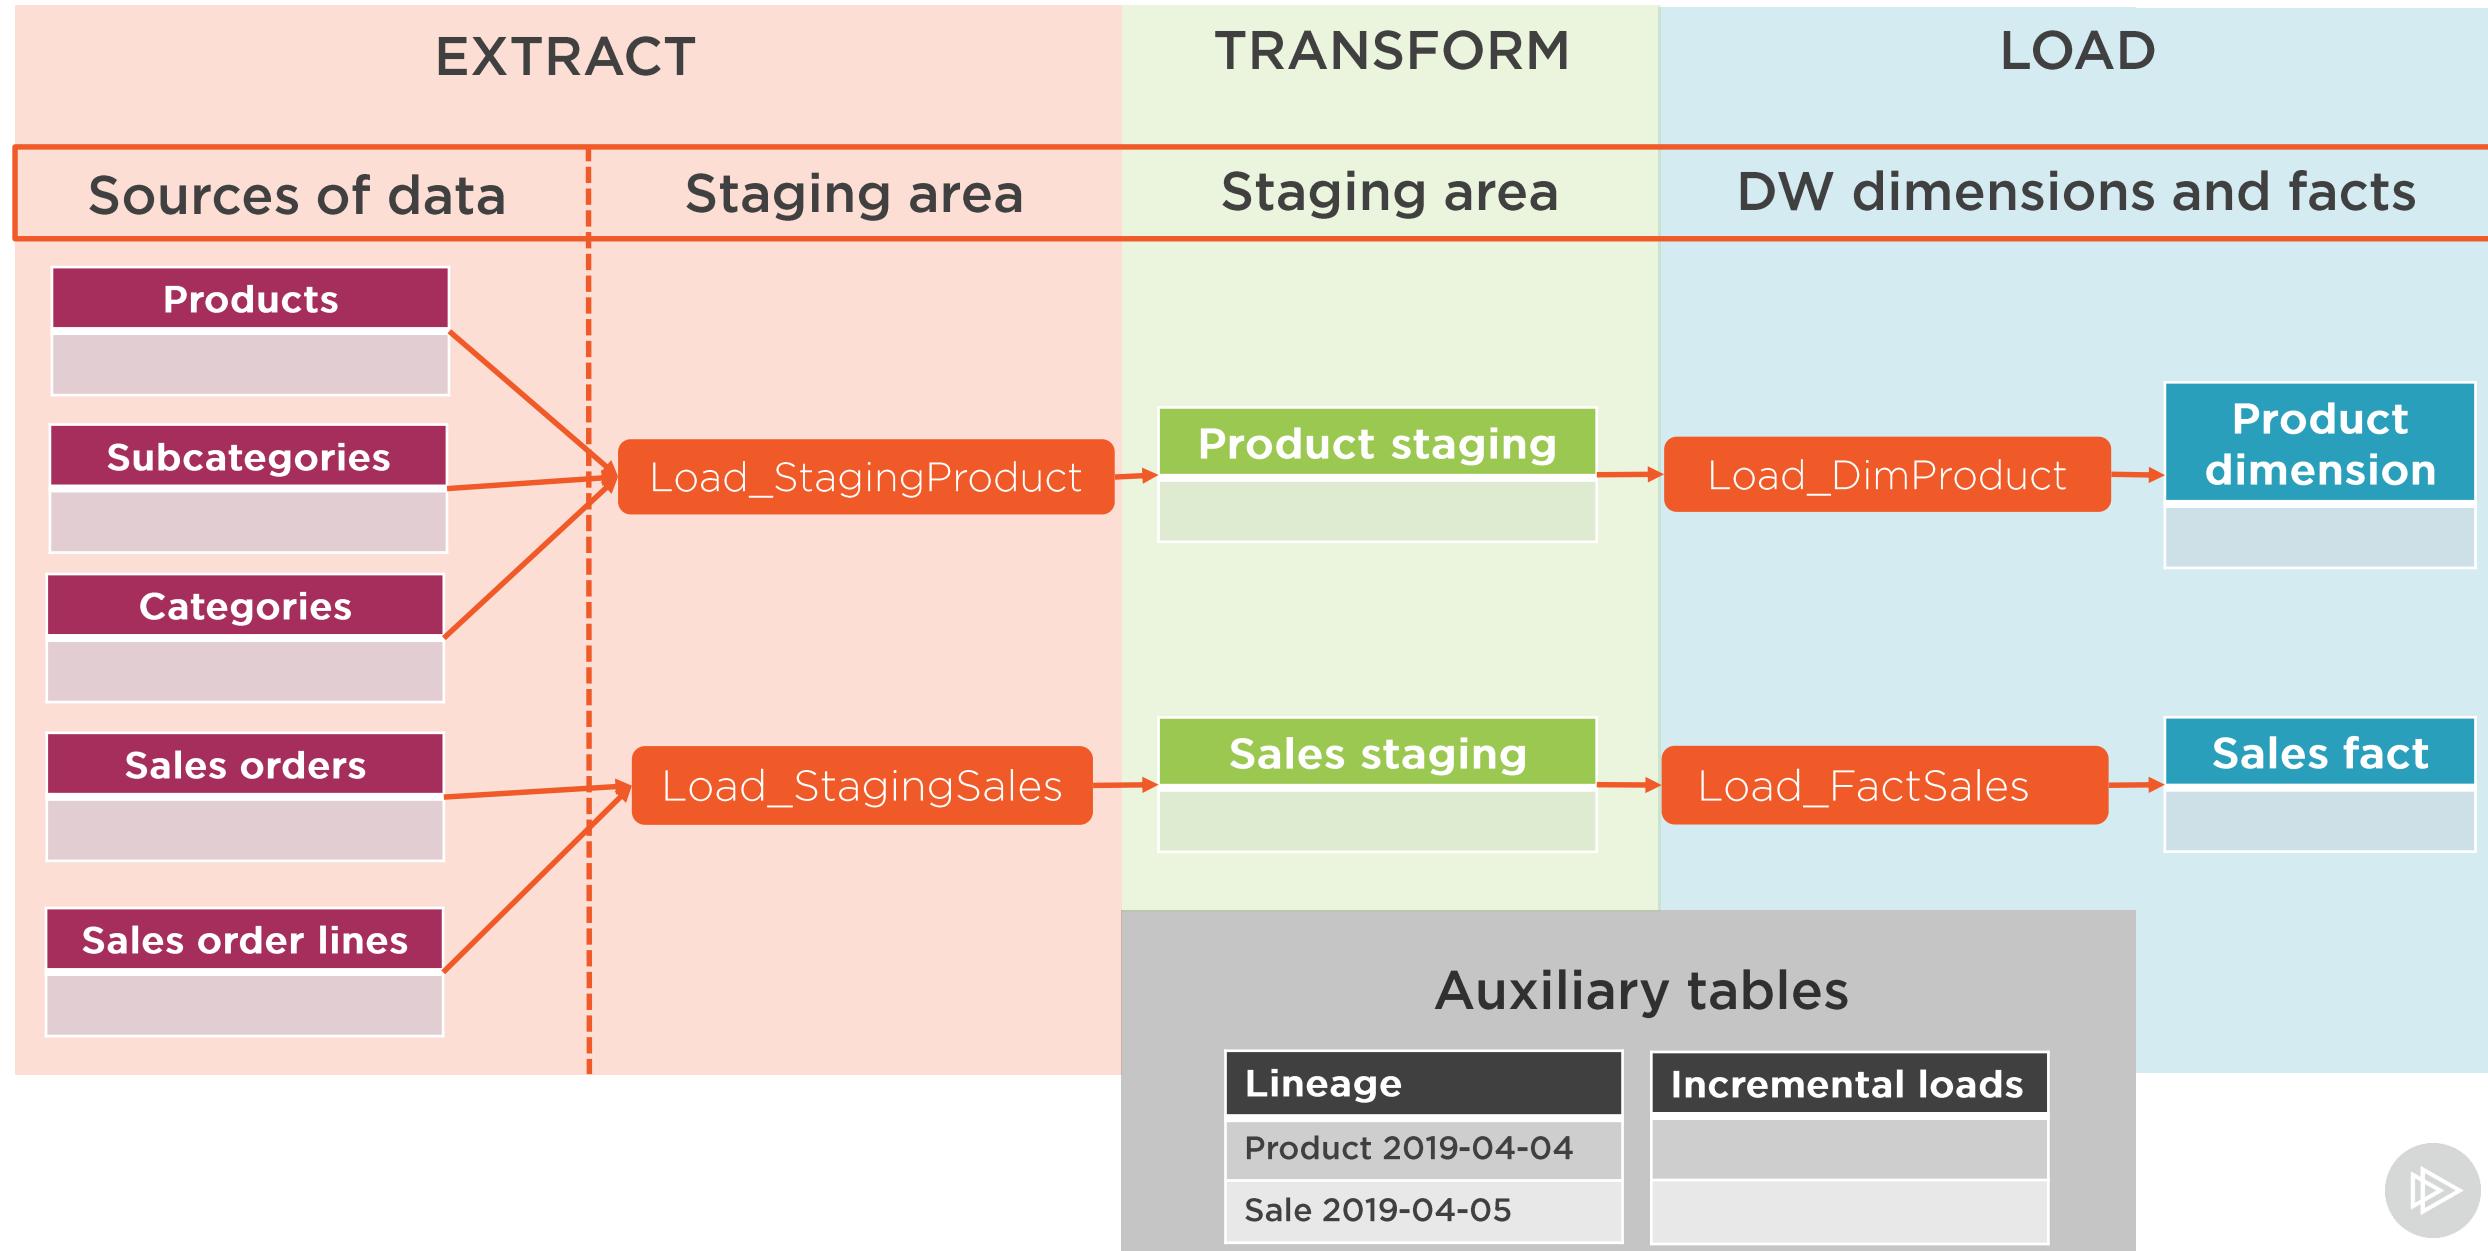
Overview of an ETL System



Overview of an ETL System



Overview of an ETL System



Types of Data Warehouse Loads



Types of Data Warehouse Loads

Full/initial load

The process of populating the data warehouse for the first time with data from the operational system

All tables are truncated and reloaded

Old data is lost

Takes a lot of time to finish

Easy to implement

Incremental load

The process of updating the data warehouse with the operational system changes

Tables are updated with new data

Old data is preserved

Takes less time than the initial load

Implementation is more complex

- Keep track of the previous load date
- Store multiple versions of the same row



Incremental Load Elements

Source key	Name	Valid from	Valid to
387	Cherry toffee	2019-01-01	2019-09-16
387	Super cherry toffee	2019-09-16	9999-12-31
104	Banana bread	2019-01-01	9999-12-31
105	Grape juice	2019-01-01	9999-12-31

The “Valid from” and “Valid to” columns

Load date key	Table name	Load date
1	Dim_Product	2019-04-13
2	Dim_Employee	2019-01-01
...	...	
16	Dim_Product	2019-04-15

The “Incremental loads” table



Incremental Load Elements

Source key	Name	Valid from	Valid to
387	Cherry toffee	2019-01-01	2019-09-16
387	Super cherry toffee	2019-09-16	9999-12-31
104	Banana bread	2019-01-01	9999-12-31
105	Grape juice	2019-01-01	9999-12-31

The “Valid from” and “Valid to” columns

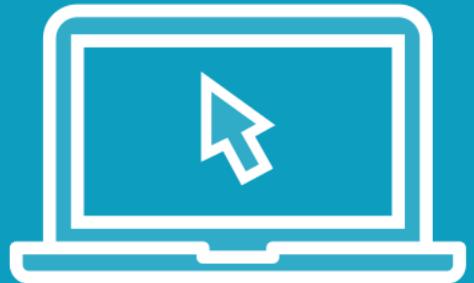
Load date key	Table name	Load date
1	Dim_Product	2019-04-13
2	Dim_Employee	2019-01-01
...	...	
16	Dim_Product	2019-04-15

The “Incremental loads” table

```
SELECT *
FROM Products
WHERE
ModifiedDate > '2019-04-13'
AND ModifiedDate <= '2019-04-15'
```



Demo



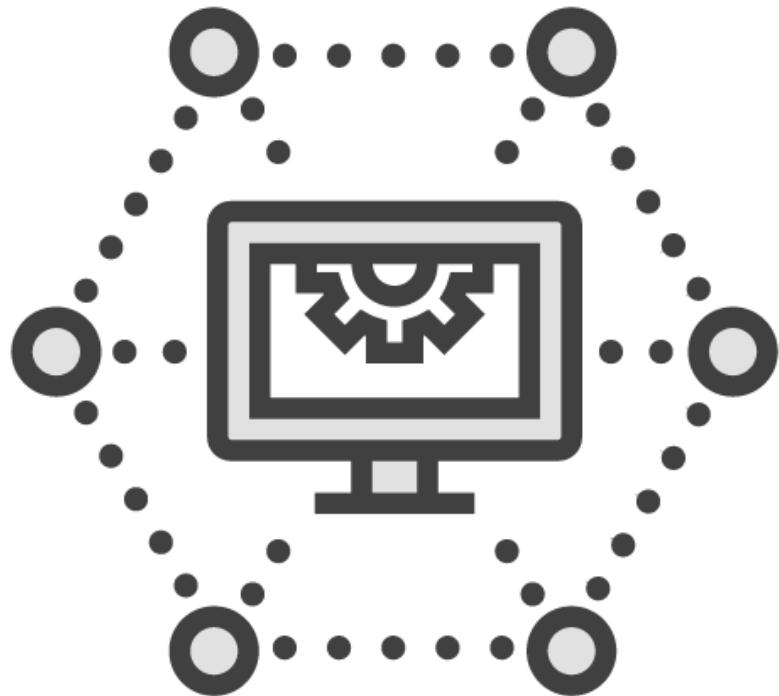
**Creating and working with the
“Incremental loads” table**



Setting up Data Lineage



What Is Data Lineage?



Data lineage tracks the movement of data

- What are the origins of data
- Where is the data going to
- When was the data loaded or updated
- What transformations are applied to the data

The complexity of data lineage implementations varies

- Depends on the necessities of the project

It is important to set up data lineage in data warehouse projects



Advantages of Data Lineage



Advantages of Data Lineage

Troubleshooting



Advantages of Data Lineage

Troubleshooting

Data warehouse

Sale number	Product
123	Cherry toffee

Source system

Sale number	Product
123	Super cherry toffee



Advantages of Data Lineage

Troubleshooting



Advantages of Data Lineage

Troubleshooting

Data trust

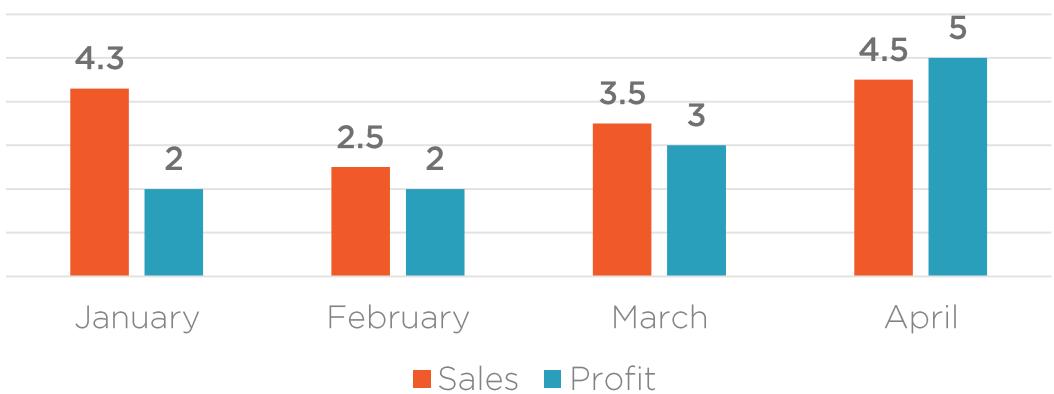


Advantages of Data Lineage

Troubleshooting

Data trust

Sales and profit



Advantages of Data Lineage

Troubleshooting

Data trust



Advantages of Data Lineage

Troubleshooting

Data trust

Transparent business rules



Advantages of Data Lineage

Troubleshooting

Data trust

Transparent business rules

Rate	% of VAT	Applies to
Standard	20%	Mixed ice cream, frozen yogurt
Reduced	5%	Some goods and services
Zero	0%	Herbal tea, pita bread, cold sandwiches, crocodile meat



Advantages of Data Lineage

Troubleshooting

Data trust

Transparent business rules



Advantages of Data Lineage

Troubleshooting

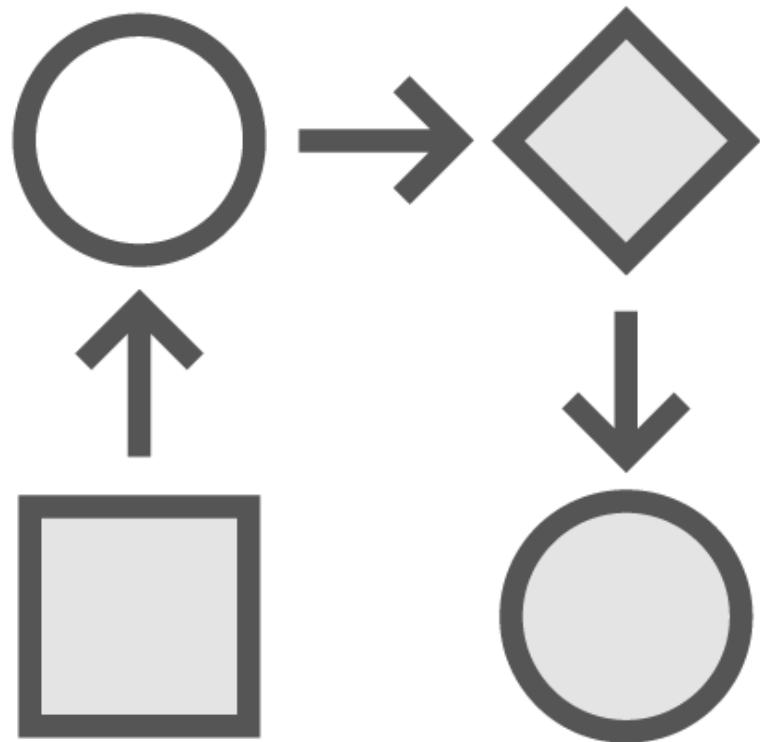
Data trust

Transparent business rules

Data audit



Implementing Data Lineage



Ensure row uniqueness

- Use surrogate keys
- Data passes multiple stages from source to destination
- Keeping the same key through all the stages helps tracing back the data

Keep track of the operation that loaded each row

- Create a “Lineage” column in all tables from the data warehouse



Implementing Data Lineage

Lineage key	Table name	Start	End	Type	Status



Implementing Data Lineage

Lineage key	Table name	Start	End	Type	Status

Product key	Product name	...	Lineage key

Employee key	Employee name	...	Lineage key



Implementing Data Lineage

Lineage key	Table name	Start	End	Type	Status
10	Dim_Product	2019-03-04 20:39:10.000	null	F	P

Product key	Product name	...	Lineage key

Employee key	Employee name	...	Lineage key



Implementing Data Lineage

Lineage key	Table name	Start	End	Type	Status
10	Dim_Product	2019-03-04 20:39:10.000	null	F	P

Product key	Product name	...	Lineage key
1	Cotton candy		10
2	Green tea		10

Employee key	Employee name	...	Lineage key



Implementing Data Lineage

Lineage key	Table name	Start	End	Type	Status
10	Dim_Product	2019-03-04 20:39:10.000	2019-03-04 20:54:35.000	F	S

Product key	Product name	...	Lineage key
1	Cotton candy		10
2	Green tea		10

Employee key	Employee name	...	Lineage key



Implementing Data Lineage

Lineage key	Table name	Start	End	Type	Status
10	Dim_Product	2019-03-04 20:39:10.000	2019-03-04 20:54:35.000	F	S
11	Dim_Employee	2019-03-04 20:45:21.000	null	F	P

Product key	Product name	...	Lineage key
1	Cotton candy		10
2	Green tea		10

Employee key	Employee name	...	Lineage key



Implementing Data Lineage

Lineage key	Table name	Start	End	Type	Status
10	Dim_Product	2019-03-04 20:39:10.000	2019-03-04 20:54:35.000	F	S
11	Dim_Employee	2019-03-04 20:45:21.000	null	F	P

Product key	Product name	...	Lineage key
1	Cotton candy		10
2	Green tea		10

Employee key	Employee name	...	Lineage key
1	Mary Baker		11
2	Jack Peanut		11
3	Gigi Knopper		11
...



Implementing Data Lineage

Lineage key	Table name	Start	End	Type	Status
10	Dim_Product	2019-03-04 20:39:10.000	2019-03-04 20:54:35.000	F	S
11	Dim_Employee	2019-03-04 20:45:21.000	2019-03-04 20:59:12.000	F	S

Product key	Product name	...	Lineage key
1	Cotton candy		10
2	Green tea		10

Employee key	Employee name	...	Lineage key
1	Mary Baker		11
2	Jack Peanut		11
3	Gigi Knopper		11
...



Implementing Data Lineage

Lineage key	Table name	Start	End	Type	Status
10	Dim_Product	2019-03-04 20:39:10.000	2019-03-04 20:54:35.000	F	S
11	Dim_Employee	2019-03-04 20:45:21.000	2019-03-04 20:59:12.000	F	S
...

Product key	Product name	...	Lineage key
1	Cotton candy		10
2	Green tea		10

Employee key	Employee name	...	Lineage key
1	Mary Baker		11
2	Jack Peanut		11
3	Gigi Knopper		11
...



Implementing Data Lineage

Lineage key	Table name	Start	End	Type	Status
10	Dim_Product	2019-03-04 20:39:10.000	2019-03-04 20:54:35.000	F	S
11	Dim_Employee	2019-03-04 20:45:21.000	2019-03-04 20:59:12.000	F	S
...
28	Dim_Product	2019-03-15 20:07:10.000	null	I	P

Product key	Product name	...	Lineage key
1	Cotton candy		10
2	Green tea		10

Employee key	Employee name	...	Lineage key
1	Mary Baker		11
2	Jack Peanut		11
3	Gigi Knopper		11
...



Implementing Data Lineage

Lineage key	Table name	Start	End	Type	Status
10	Dim_Product	2019-03-04 20:39:10.000	2019-03-04 20:54:35.000	F	S
11	Dim_Employee	2019-03-04 20:45:21.000	2019-03-04 20:59:12.000	F	S
...
28	Dim_Product	2019-03-15 20:07:10.000	null	I	P

Product key	Product name	...	Lineage key
1	Cotton candy		10
2	Green tea		10
3	Grape juice		28
4	Banana bread		28

Employee key	Employee name	...	Lineage key
1	Mary Baker		11
2	Jack Peanut		11
3	Gigi Knopper		11
...



Implementing Data Lineage

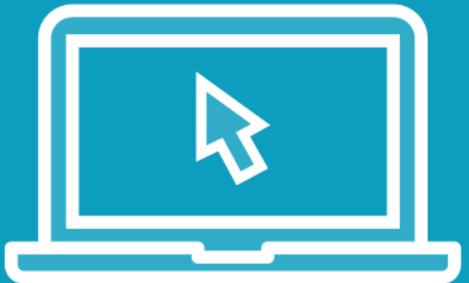
Lineage key	Table name	Start	End	Type	Status
10	Dim_Product	2019-03-04 20:39:10.000	2019-03-04 20:54:35.000	F	S
11	Dim_Employee	2019-03-04 20:45:21.000	2019-03-04 20:59:12.000	F	S
...
28	Dim_Product	2019-03-15 20:07:10.000	2019-03-15 20:12:33.000	I	S

Product key	Product name	...	Lineage key
1	Cotton candy		10
2	Green tea		10
3	Grape juice		28
4	Banana bread		28

Employee key	Employee name	...	Lineage key
1	Mary Baker		11
2	Jack Peanut		11
3	Gigi Knopper		11
...



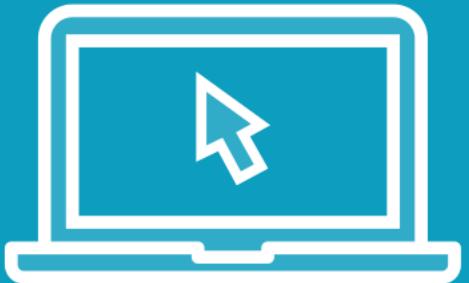
Demo



Creating and working with the Lineage table



Demo



Populating a dimension table with the help of stored procedures

- Load_StagingProduct
- Load_DimProduct



Objects Participating in a Load Process

Object (table/stored procedure)	Was discussed
Fact/dimension tables	✓
Staging tables	✓
SP for loading the staging table	✓
SP for loading the fact/dim table	✓
Lineage table	✓
Incremental loads table	✓
SPs for updating the log tables	✓



Summary



Overview of an ETL system

Types of data warehouse loads

- Full (initial) load
- Incremental load

Setting up data lineage

Demos

- Creating and using the auxiliary tables
 - Lineage
 - Incremental loads
- Populating with data the staging and dimension tables



Setting up an ETL Project in SSIS



Ana Voicu
@ana_voicu



Overview



What is an ETL tool

Advantages of using an ETL tool

Overview of SSIS

Demos

- Creating a new SSIS project with Visual Studio
- Adding connection managers
- Creating control flow components
- Linking tasks with precedence constraints
- Adding project variables
- Executing SQL tasks with parameters



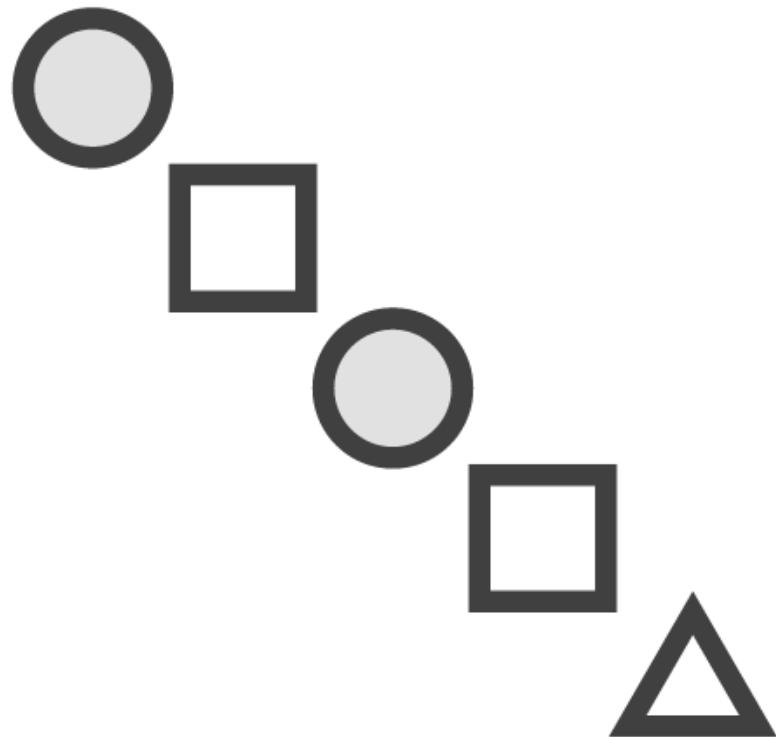
Disclaimer: This is not an intensive course on SSIS



Advantages of Using an ETL Tool



What Are ETL Tools?

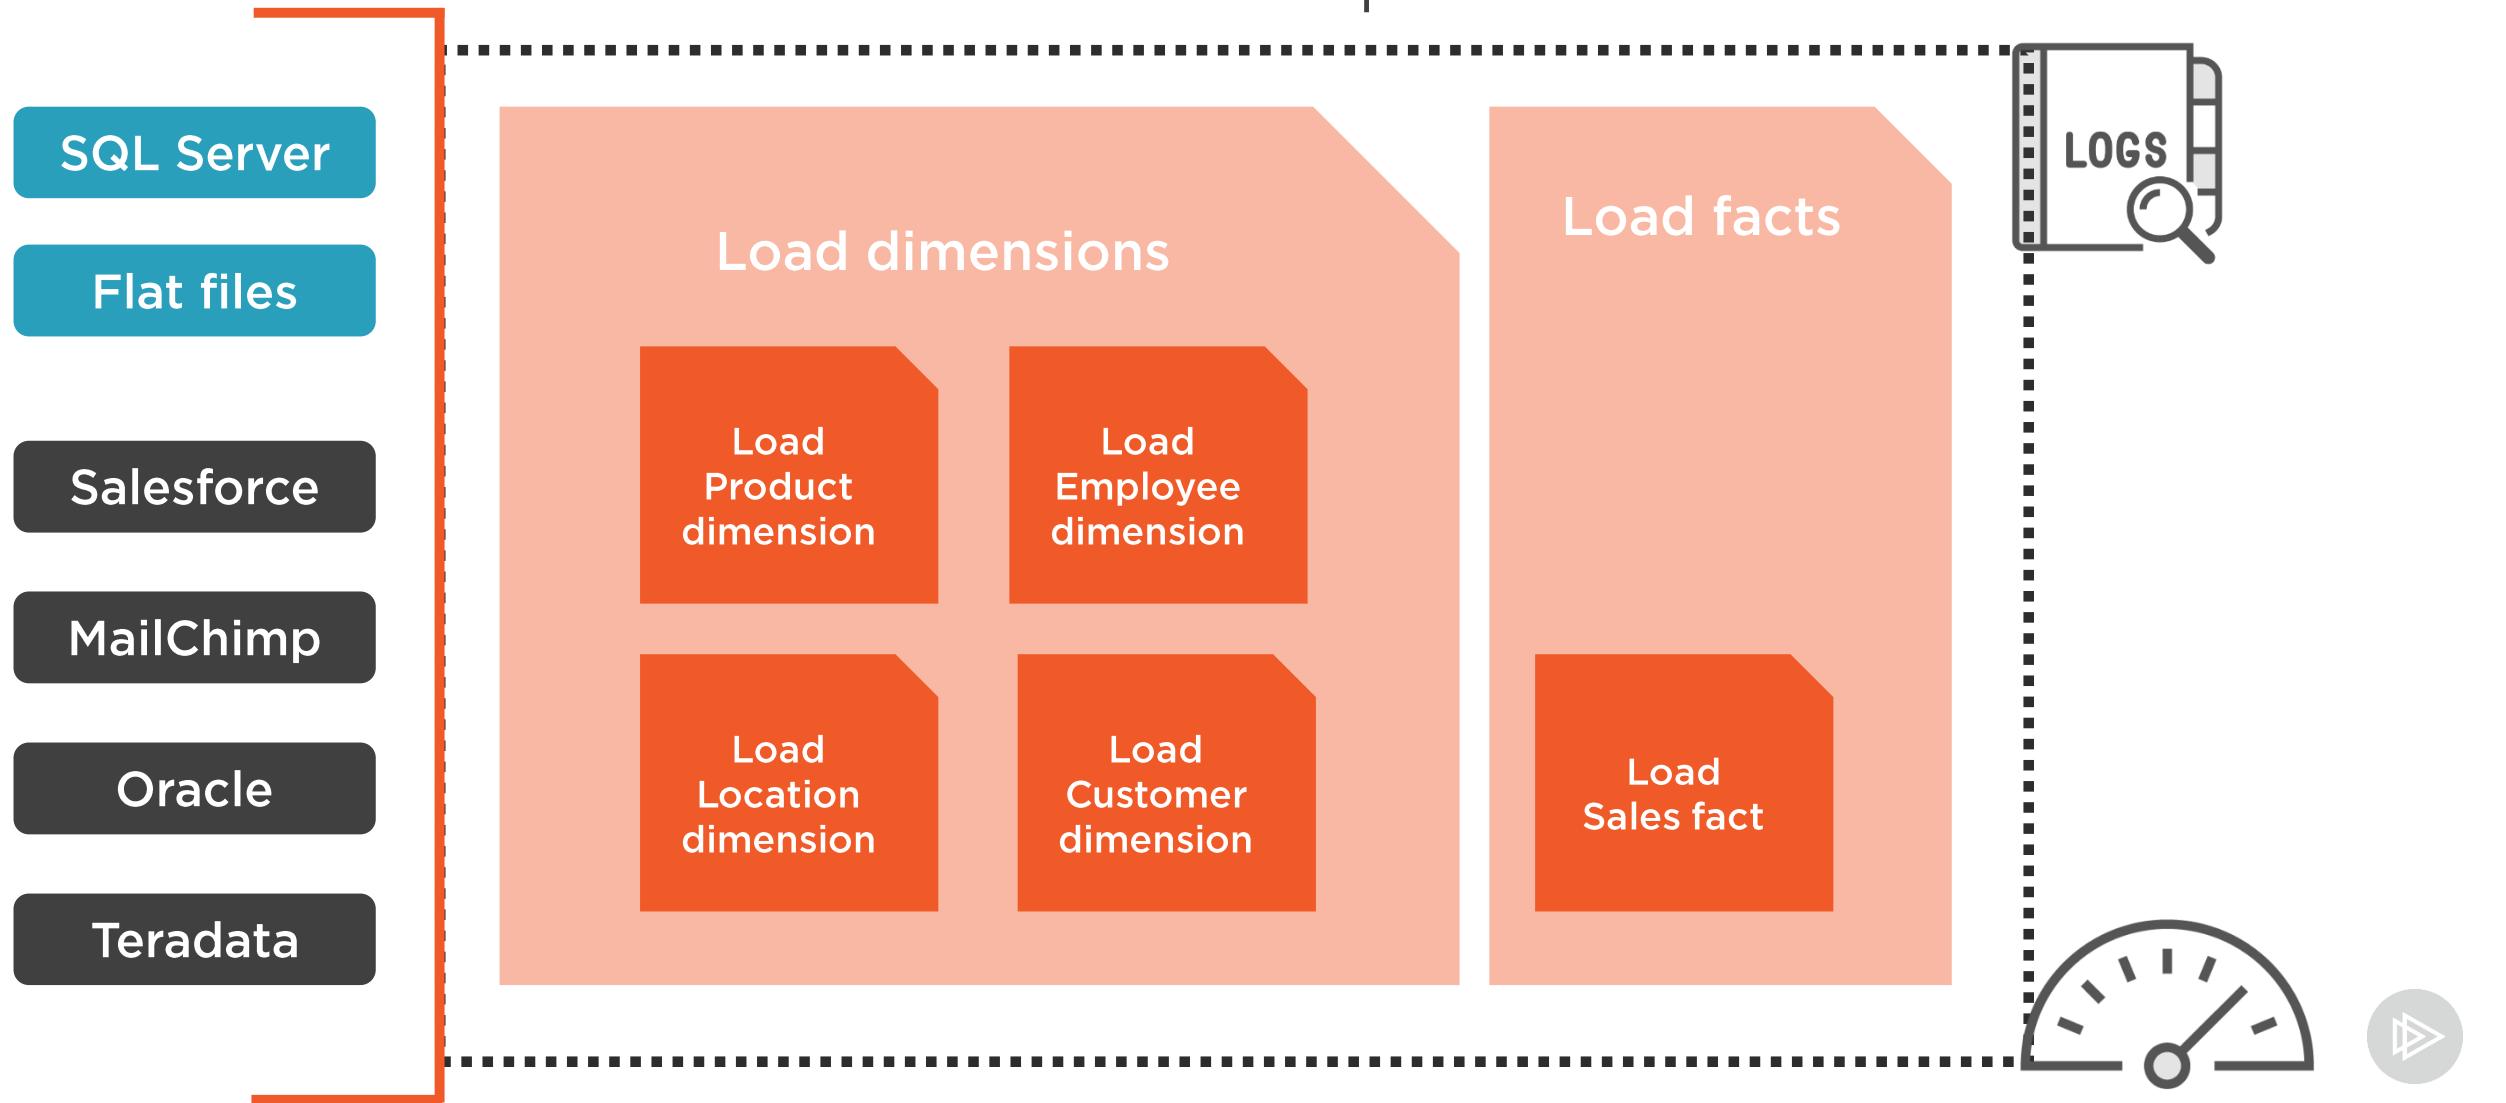


Software systems for building and automating ETL processes

Offer an alternative to creating all the code for a DW solution by hand



Loading a Data Warehouse Only Using SQL Scripts



What Is an ETL Tool?



A system comprised of all the components needed to build ETL processes

Components:

- Data flow management
- Connectors to other systems
- Data transformation components
- Logging and auditing
- Debugging
- Deployment



Advantages of Using an ETL Tool

Flexibility

Maintainability

- Easier than working with hundreds of stored procedures
- ETL tools are very graphical and easy to interpret
- Easy to learn and use



Advantages of Using an ETL Tool

Flexibility

Maintainability



Advantages of Using an ETL Tool

Flexibility

Maintainability

Performance

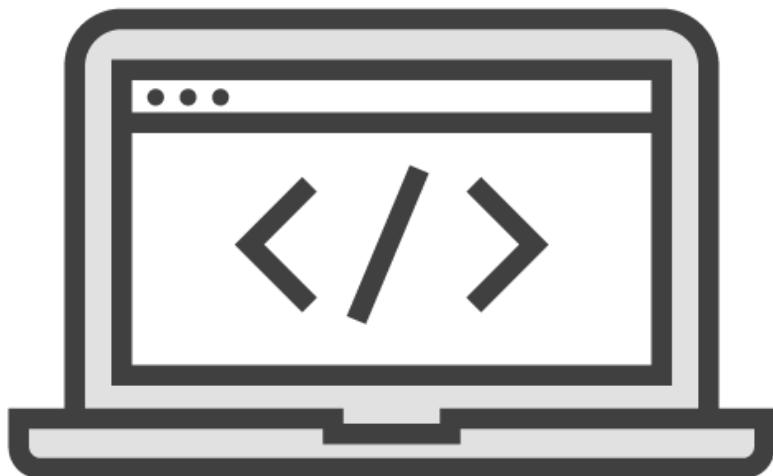
Logging and audit



Overview of SSIS



What Is SSIS?



Component of SQL Server

Used for data integration and automating processes

Provides all the components necessary for creating an ETL process

Graphical tool

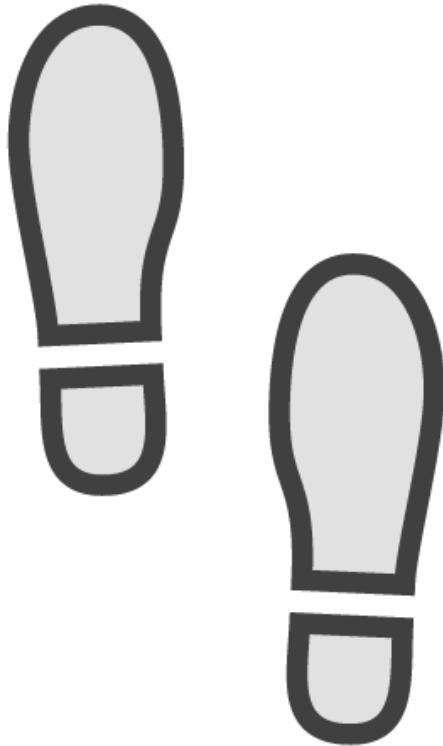
- Doesn't require writing complex code
- Based on drag and drop actions

The output is a workflow, or pipeline

The unit of work in SSIS is called a package



Steps for Loading the Product Dimension



1. Save the current load date
2. Insert new row in the Lineage table
3. Truncate the staging table
4. Populate the staging table
5. Transfer data from staging to dimension

These steps are part of a flow or process

- Executed together
- In a specified order
- Executed multiple times



Overview of a SSIS Package



Get date of current load



Overview of a SSIS Package



Get date of current load



`@[User::LoadDate] = GETDATE()`



Overview of a SSIS Package



Get date of current load

..... @User::LoadDate = GETDATE()



Update Lineage table



Overview of a SSIS Package



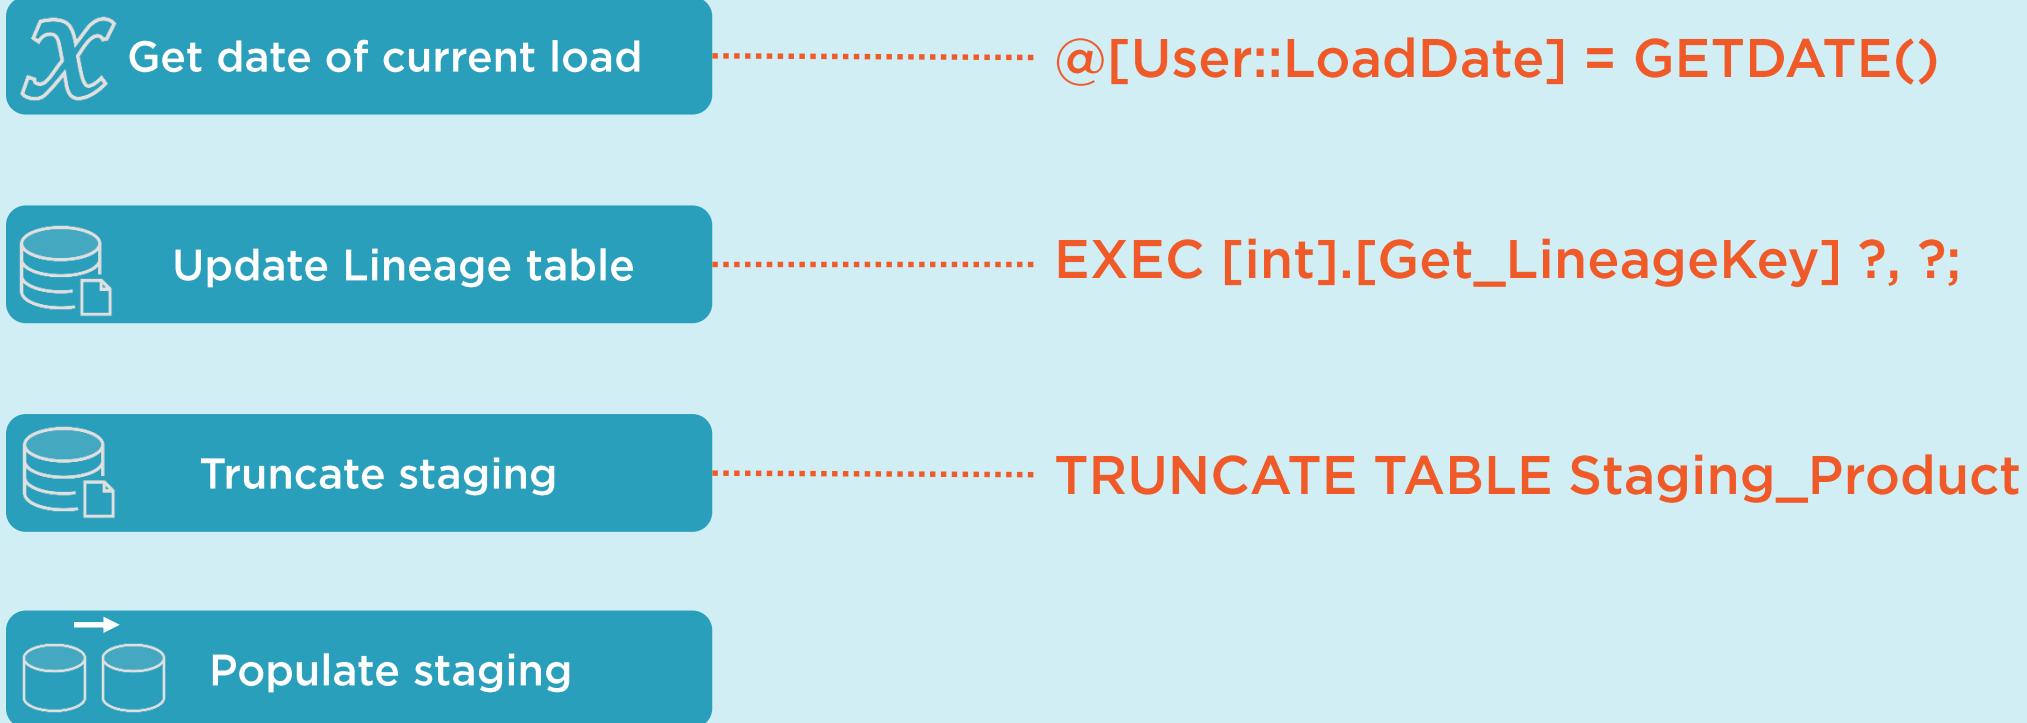
Overview of a SSIS Package



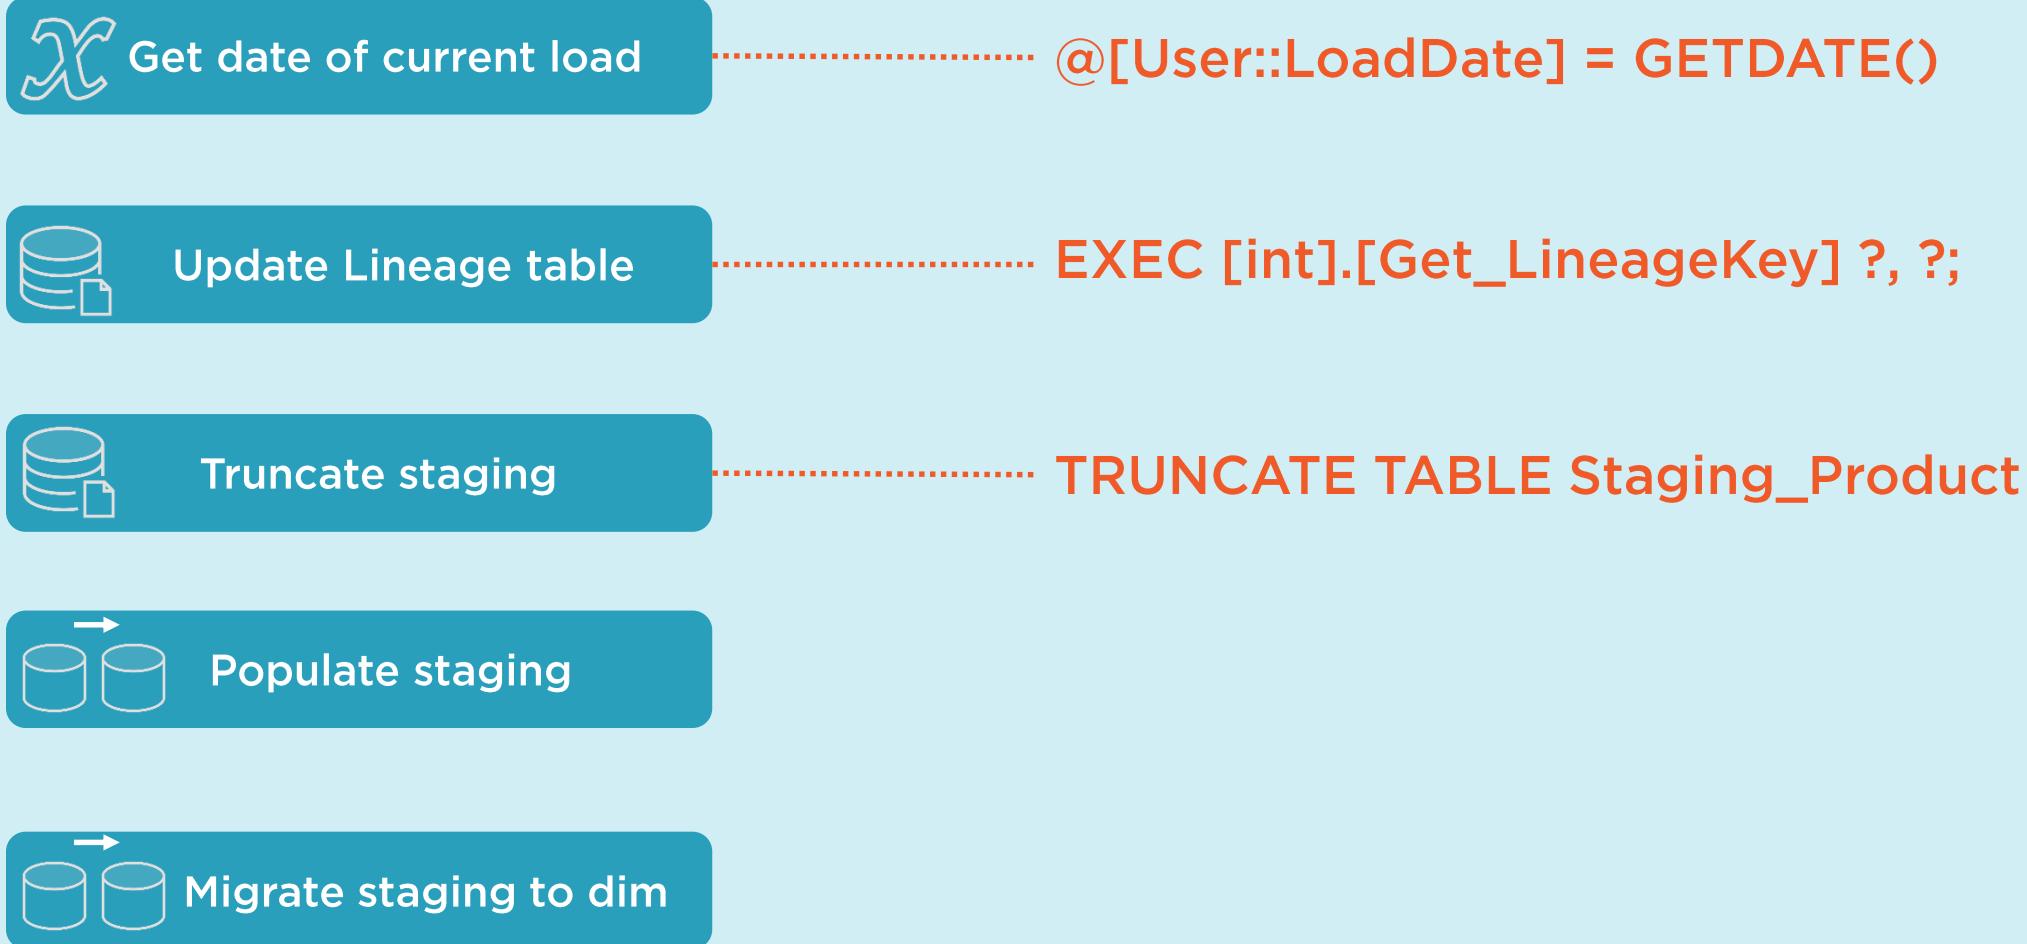
Overview of a SSIS Package



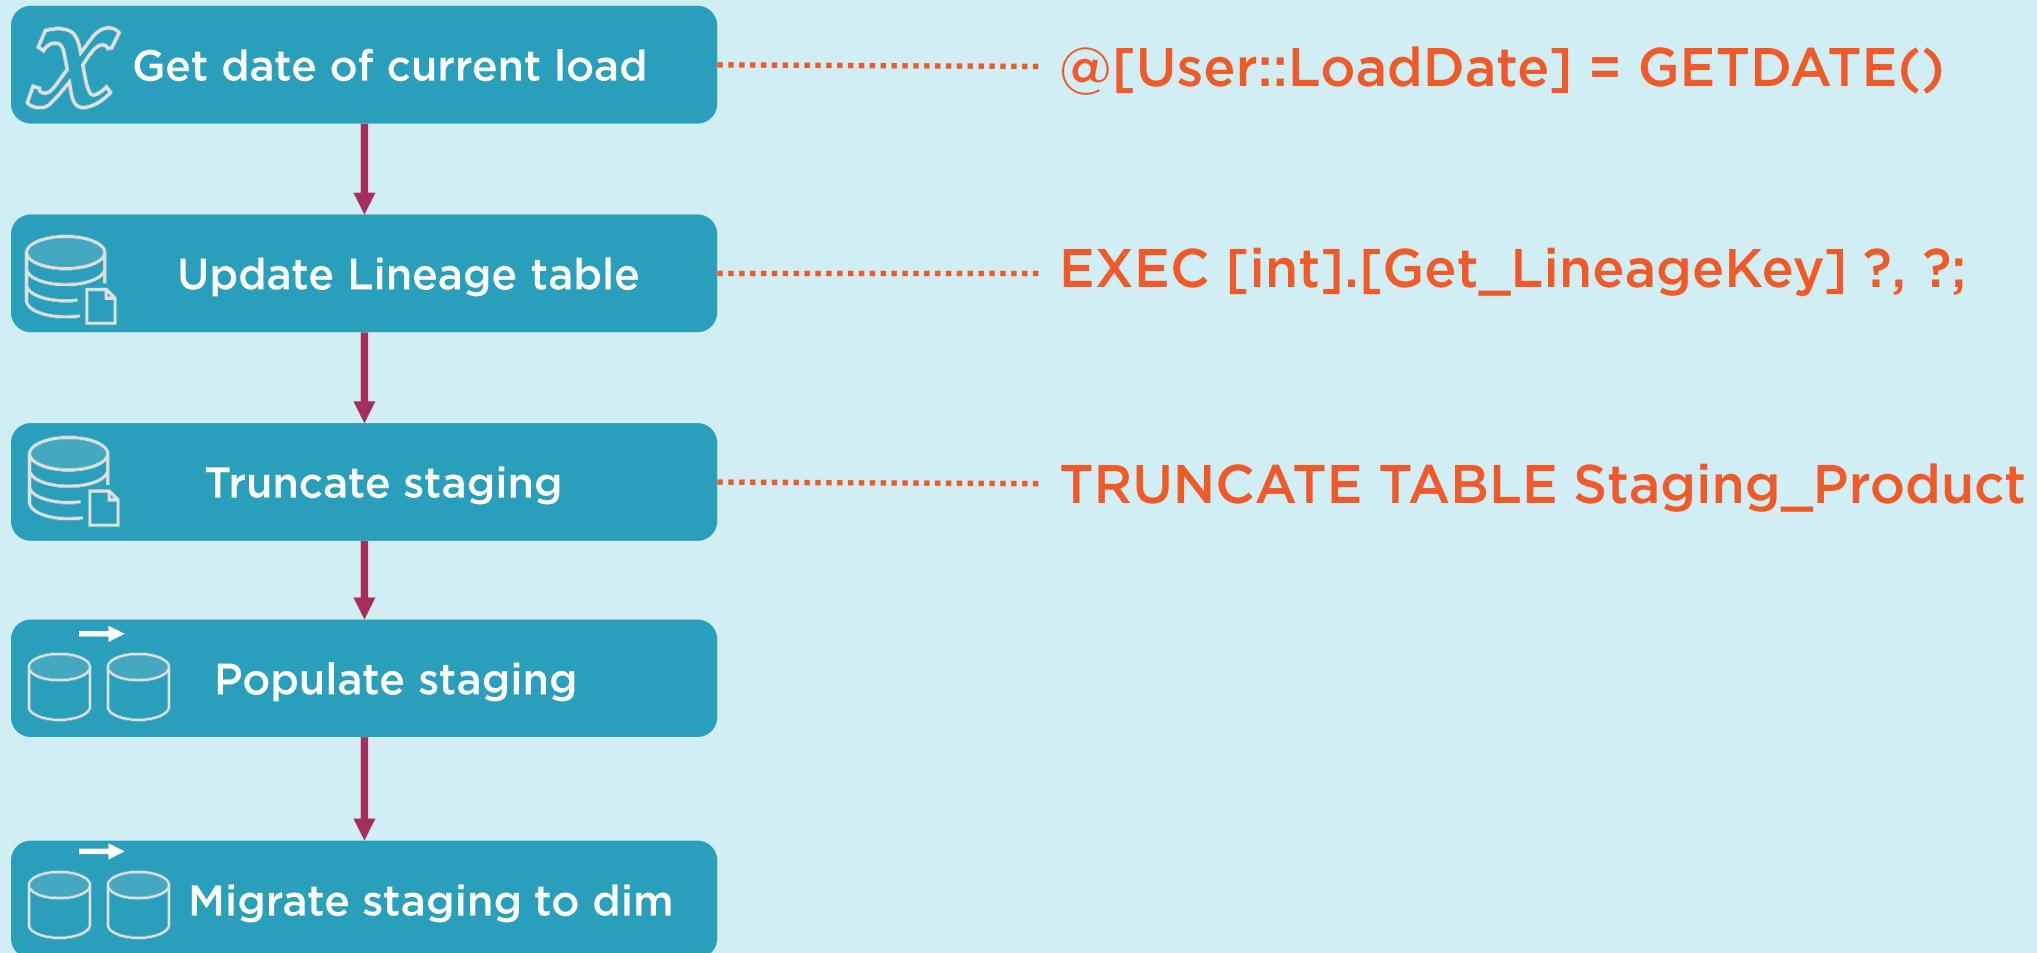
Overview of a SSIS Package



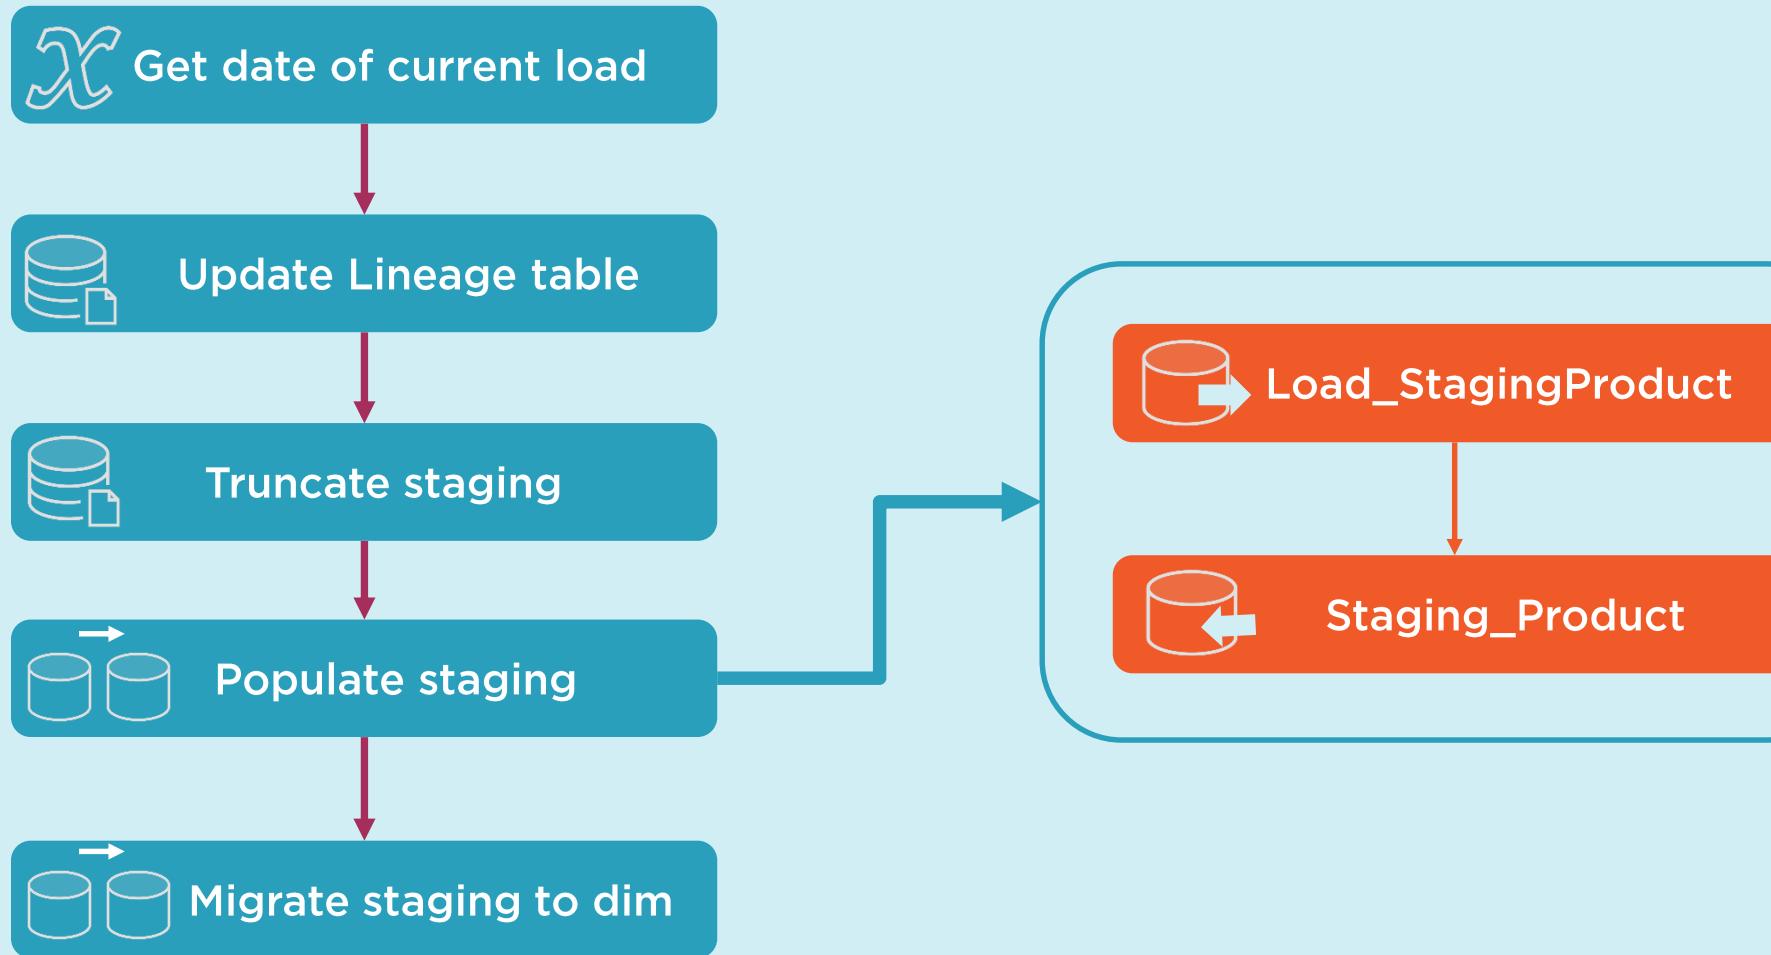
Overview of a SSIS Package



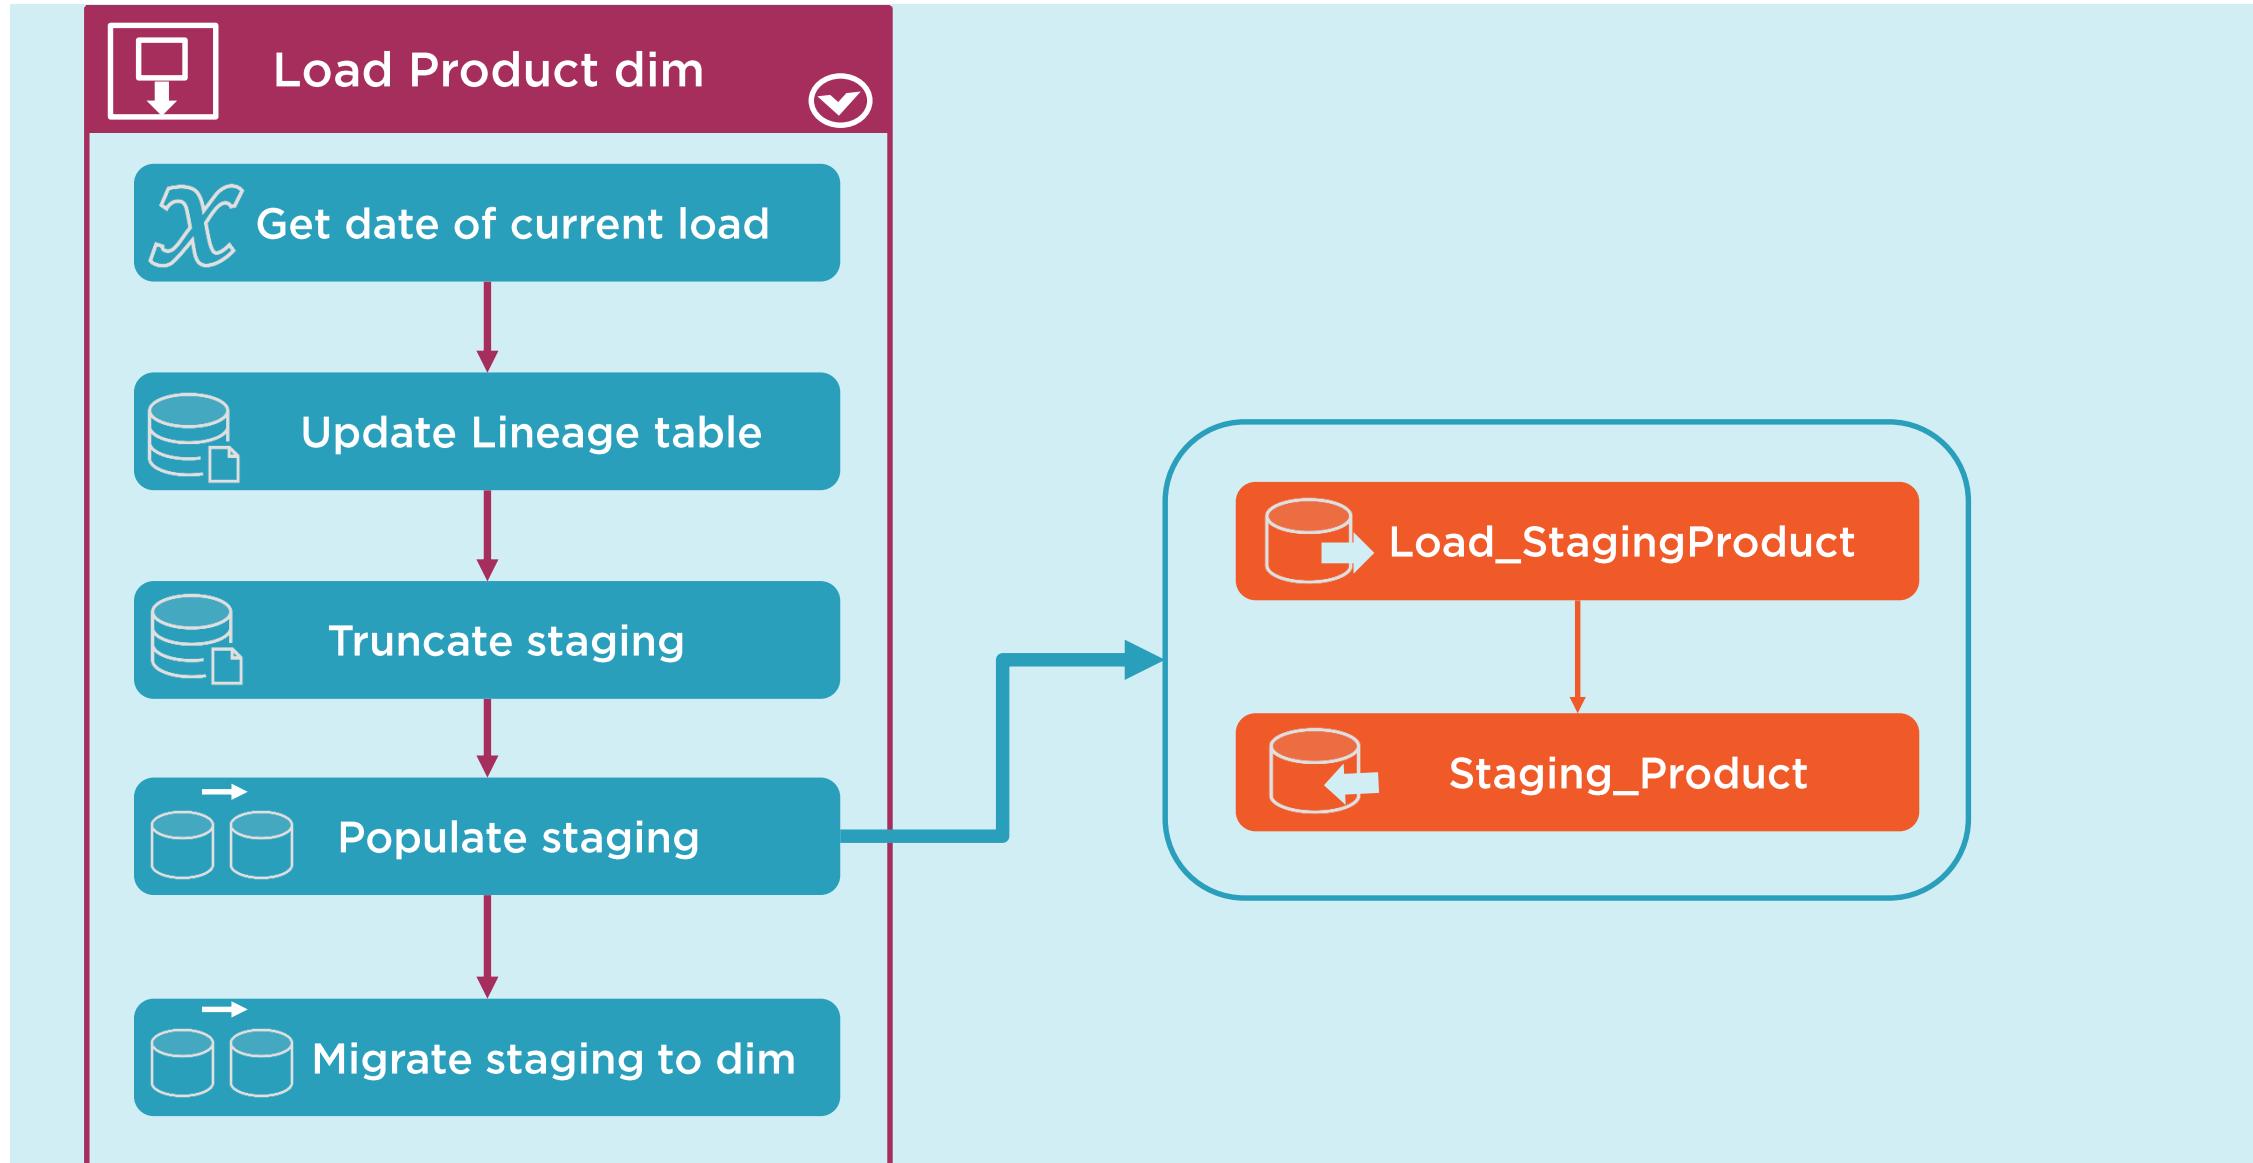
Overview of a SSIS Package



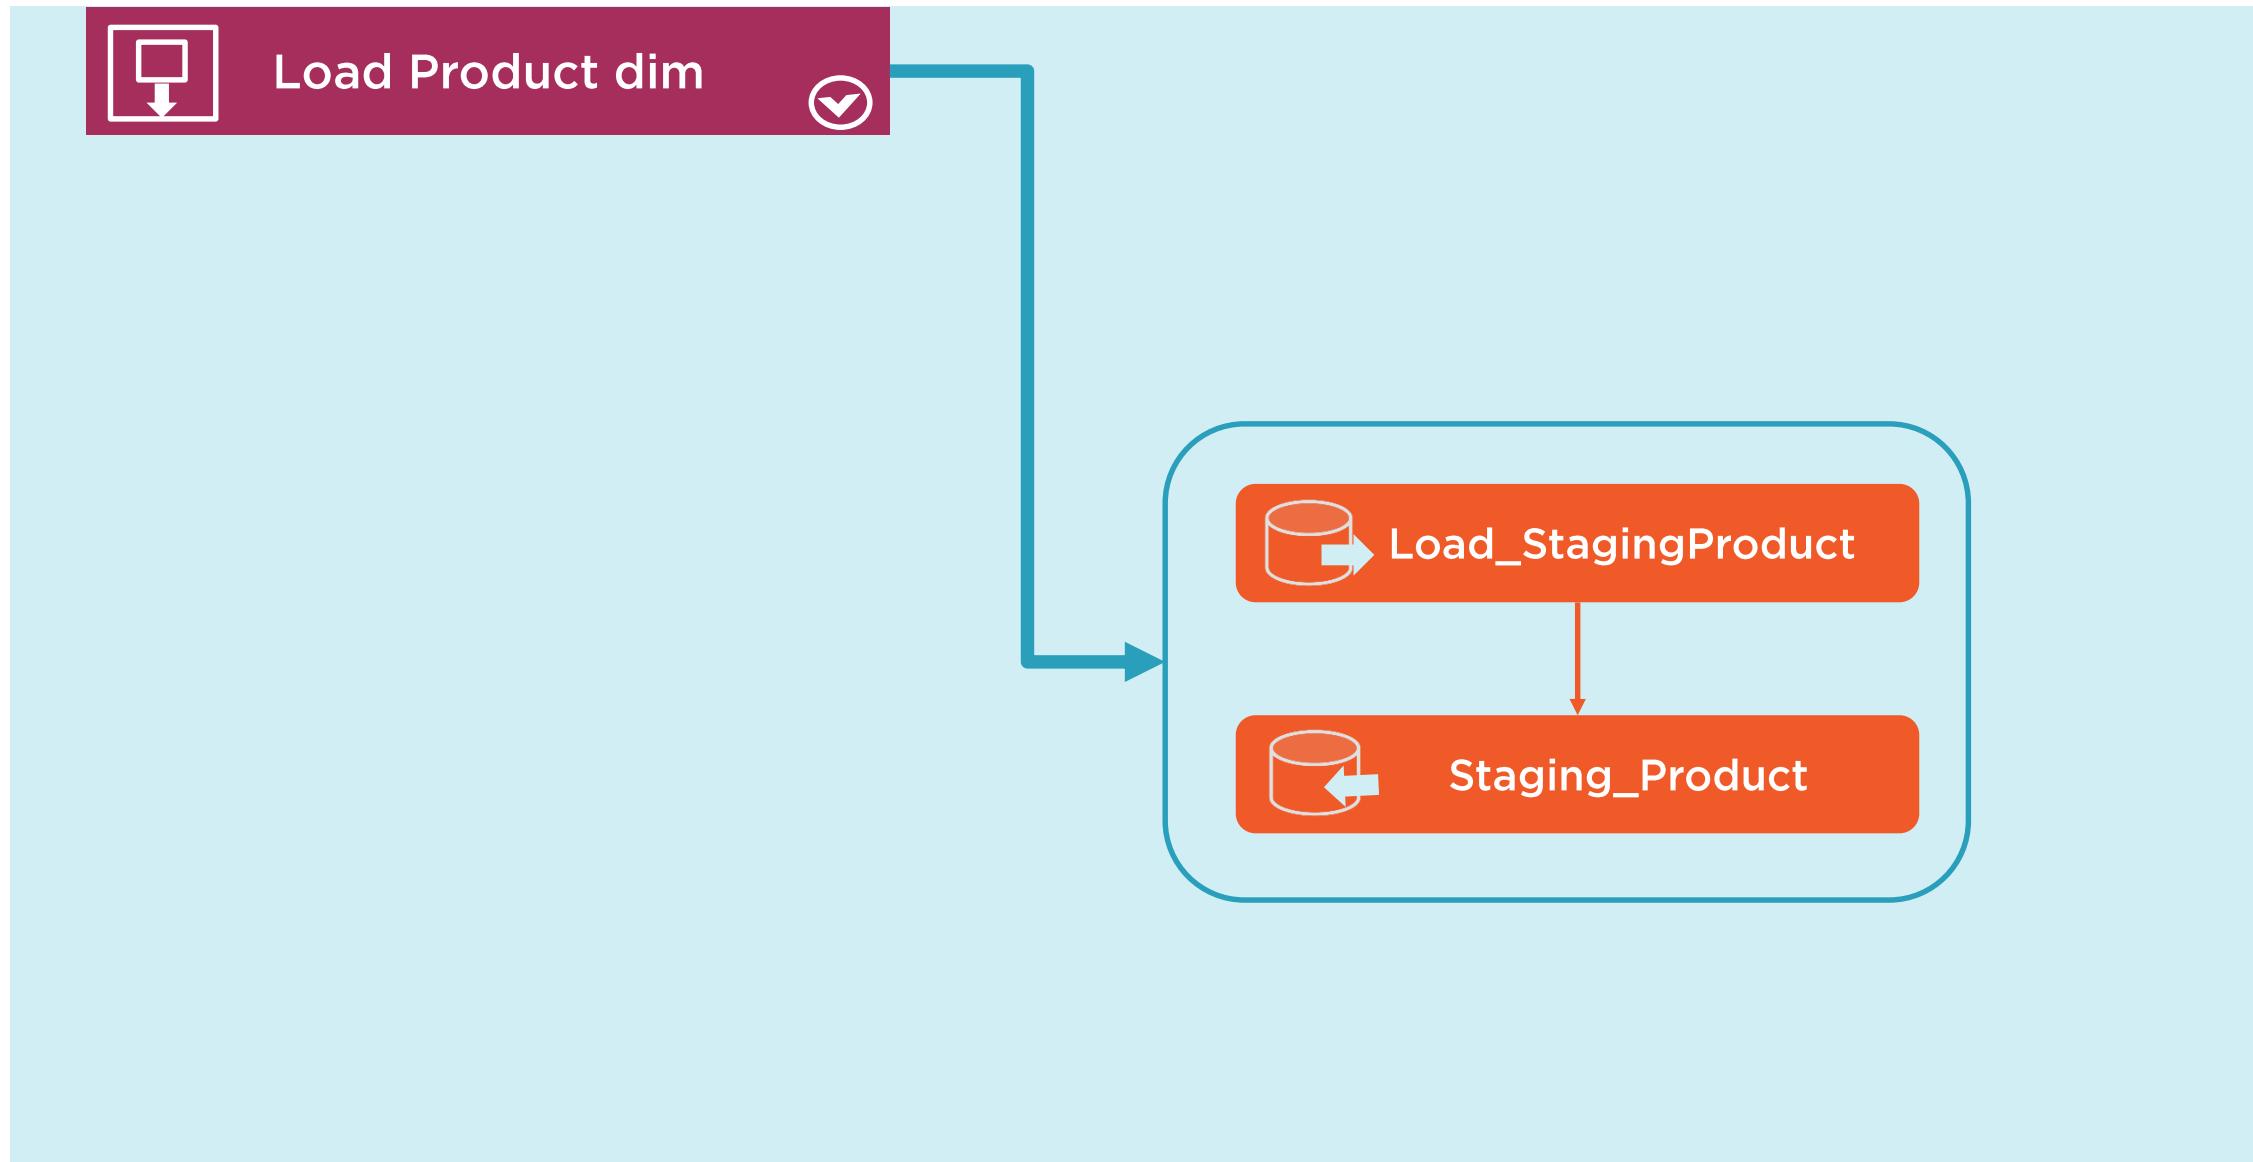
Overview of a SSIS Package



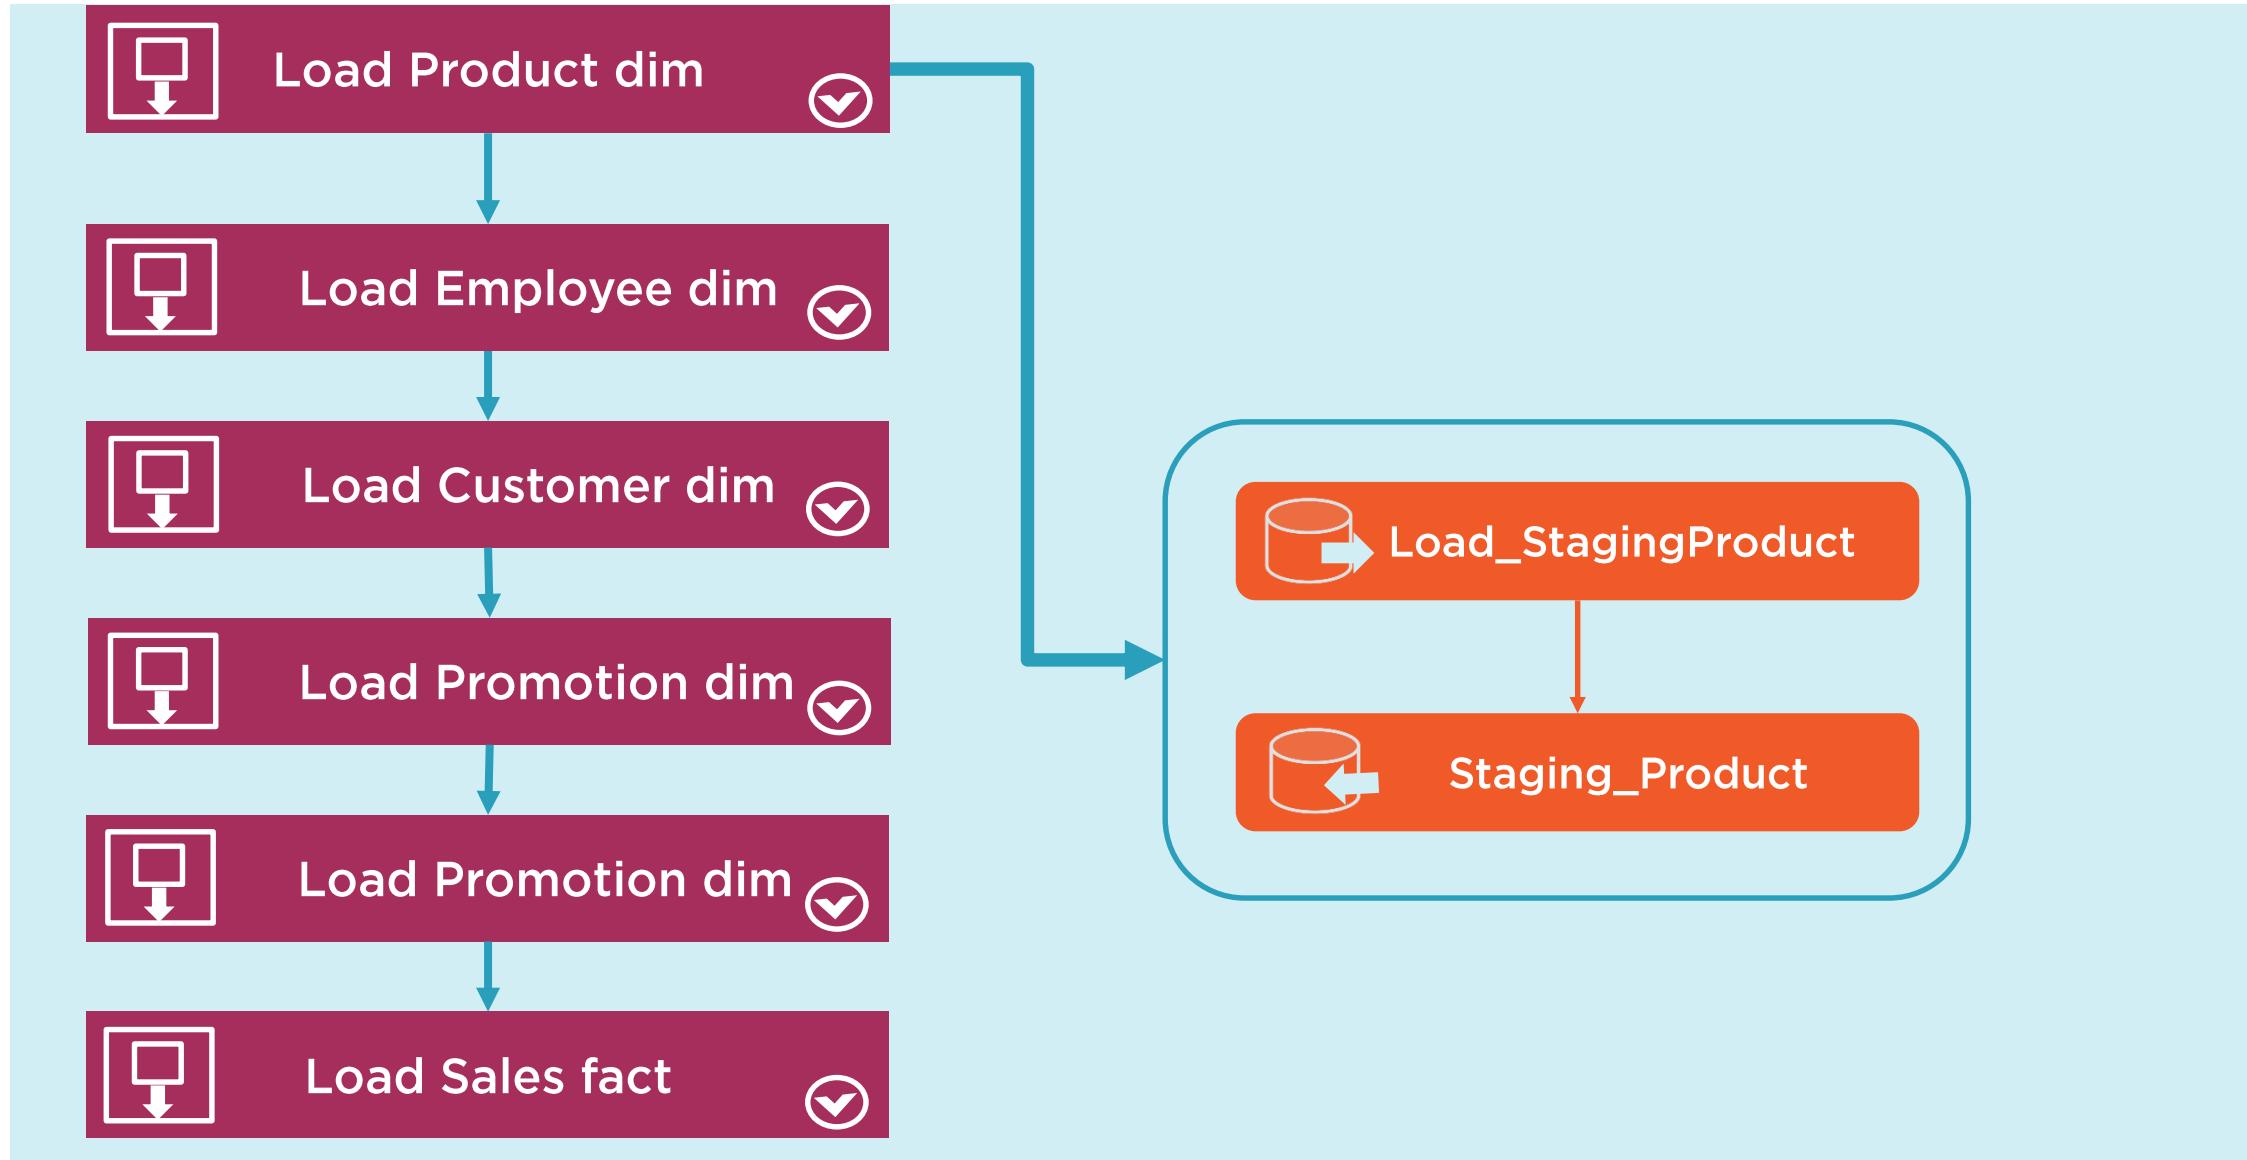
Overview of a SSIS Package



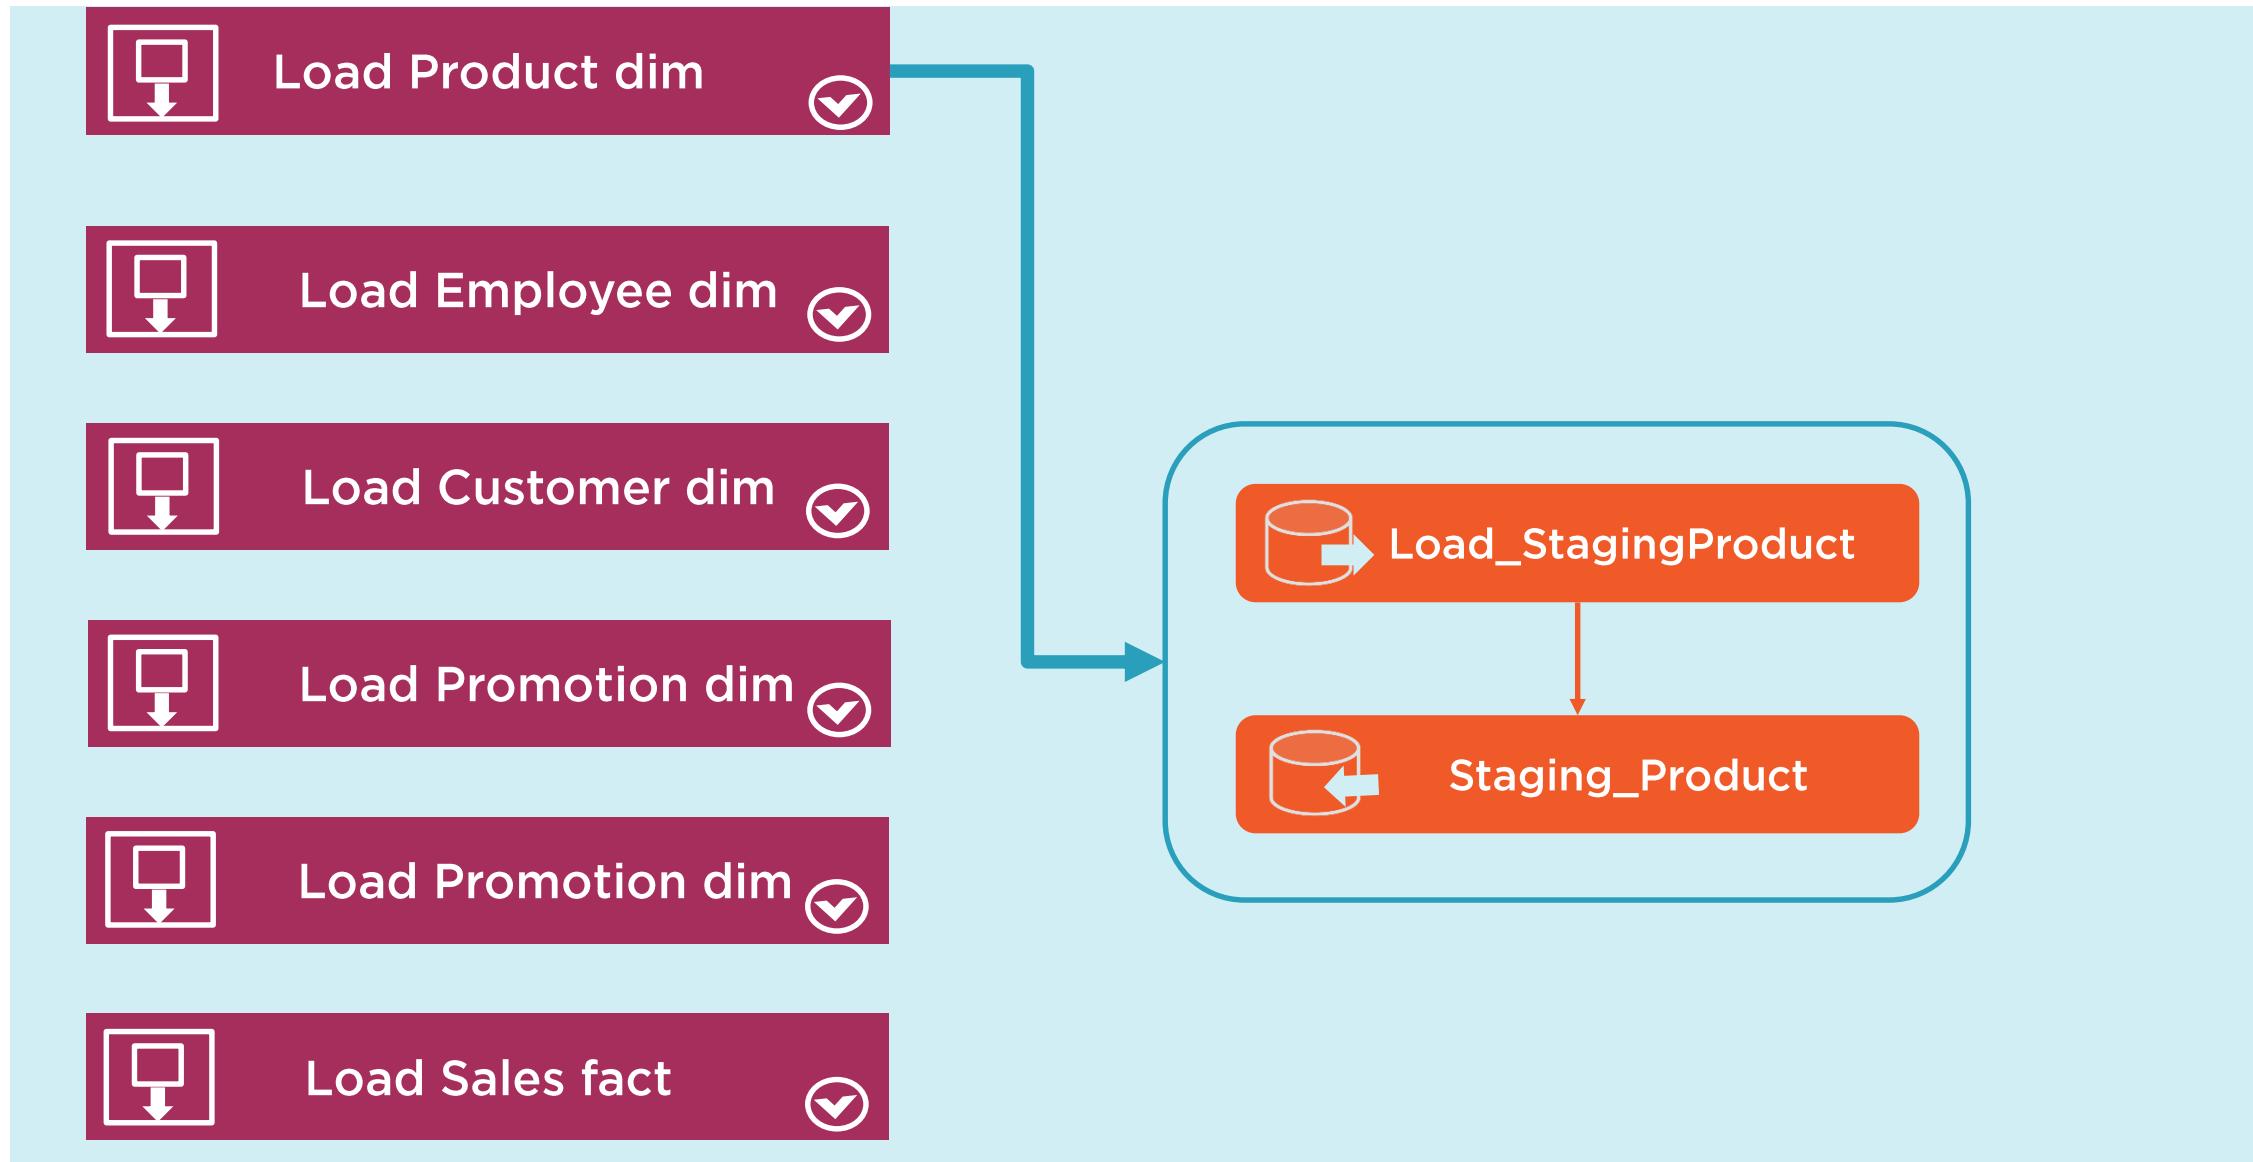
Overview of a SSIS Package



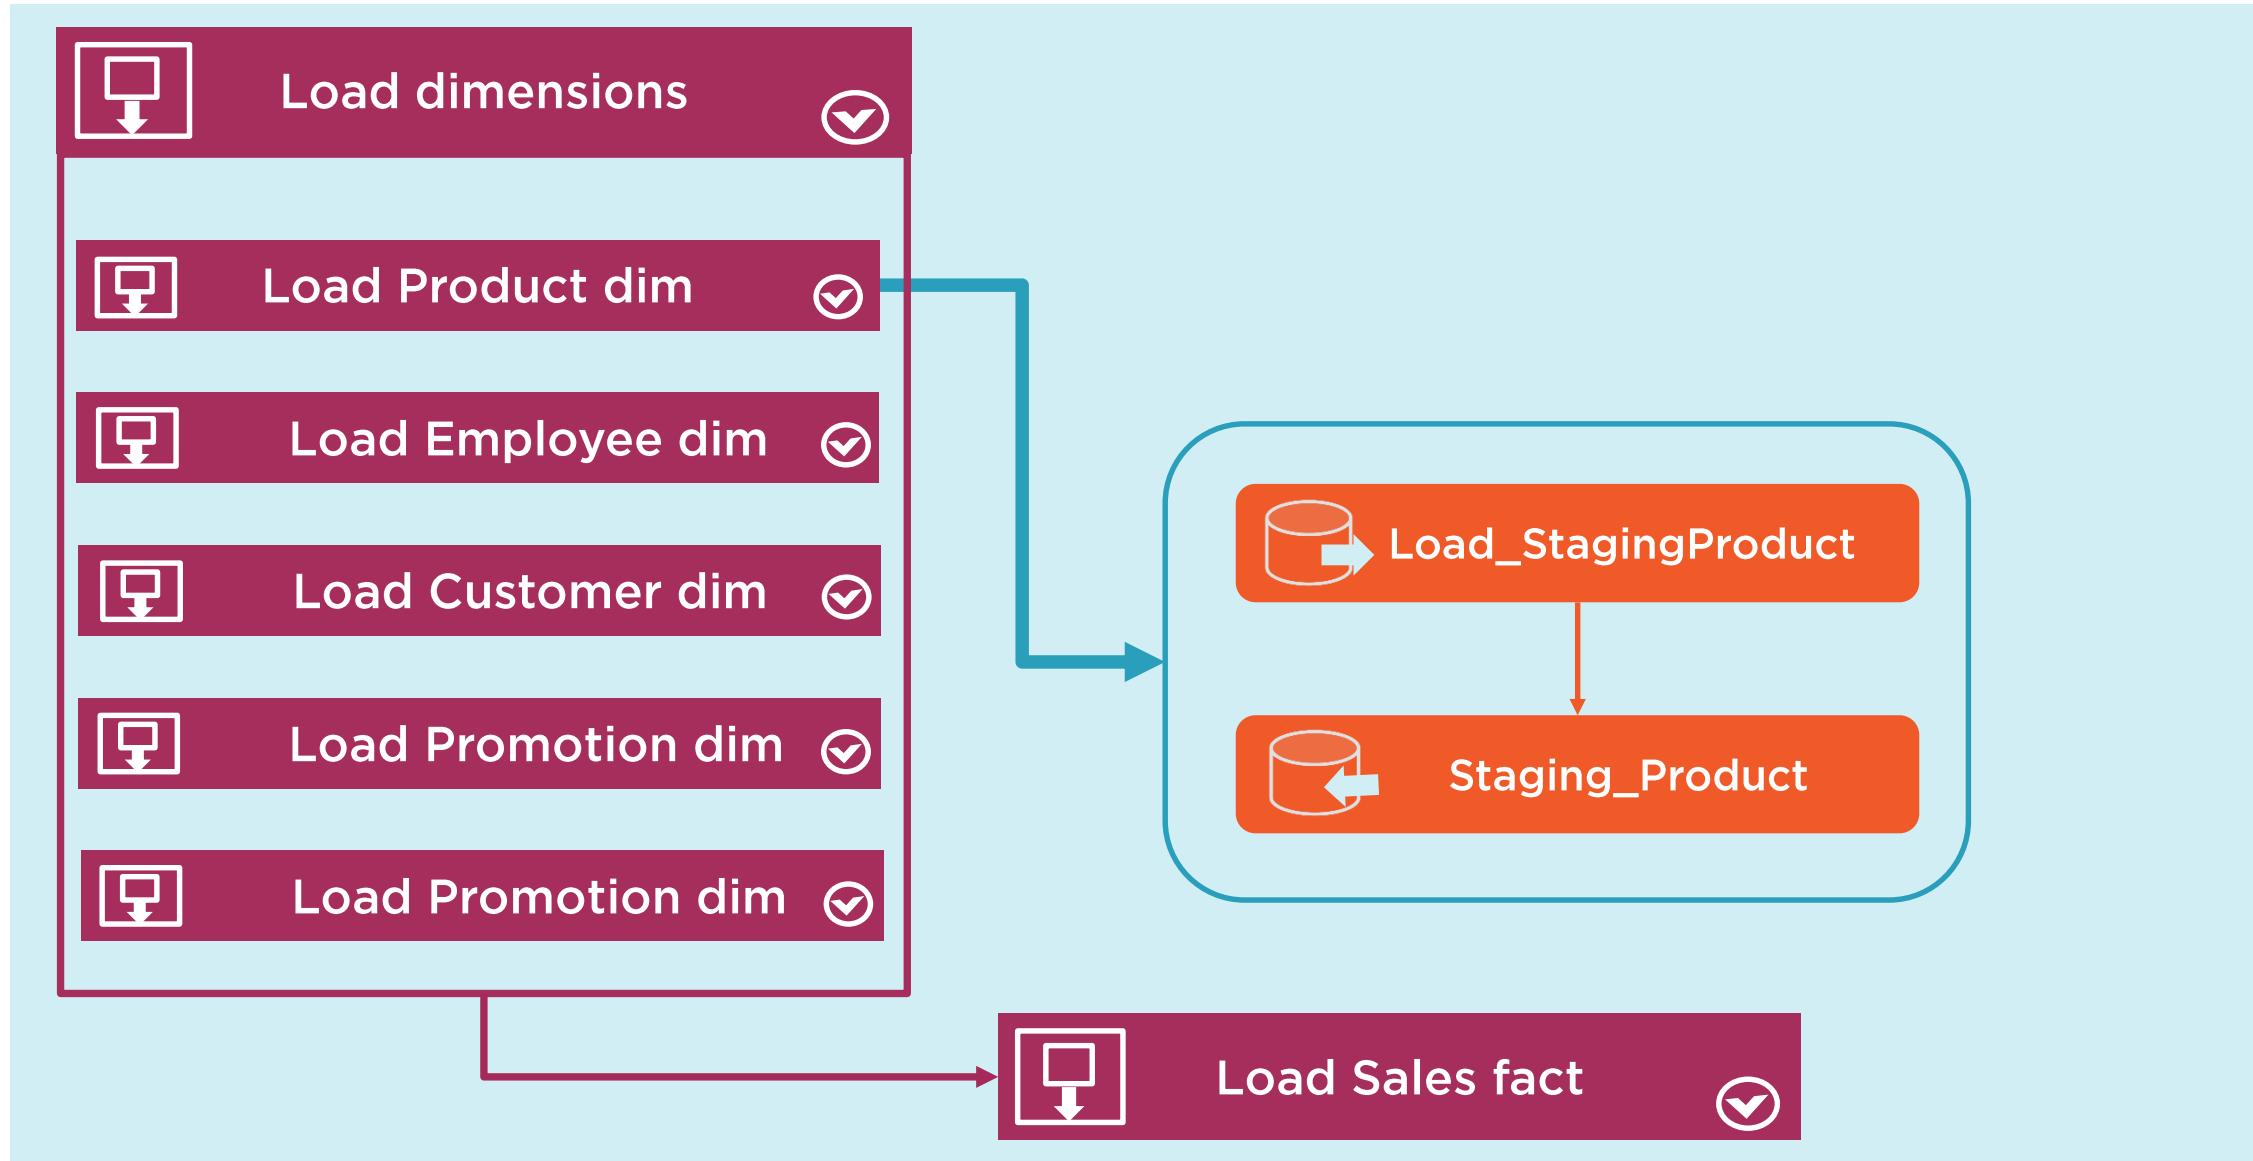
Overview of a SSIS Package



Overview of a SSIS Package



Overview of a SSIS Package



SSIS Terminology



SSIS Terminology

SSIS package	



SSIS Terminology

SSIS package	
Task	



SSIS Terminology

SSIS package	
Task Control flow task	



SSIS Terminology

SSIS package	
Task Control flow task Data flow task	



SSIS Terminology

SSIS package	
Task Control flow task Data flow task	
Control flow	



SSIS Terminology

SSIS package	
Task Control flow task Data flow task	
Control flow	
Data flow	



SSIS Terminology

SSIS package	
Task Control flow task Data flow task	
Control flow	
Data flow	
Connection manager	



SSIS Terminology

SSIS package	Container
Task <ul style="list-style-type: none">Control flow taskData flow task	
Control flow	
Data flow	
Connection manager	



SSIS Terminology

SSIS package	Container
Task Control flow task Data flow task	Precedence constraint
Control flow	
Data flow	
Connection manager	



SSIS Terminology

SSIS package	Container
Task Control flow task Data flow task	Precedence constraint
Control flow	Variable
Data flow	
Connection manager	

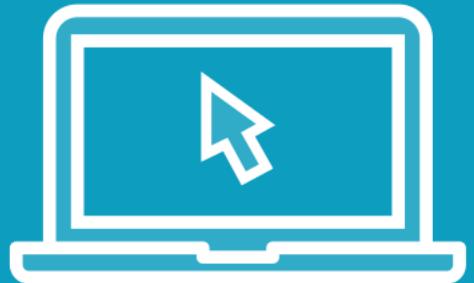


SSIS Terminology

SSIS package	Container
Task Control flow task Data flow task	Precedence constraint
Control flow	Variable
Data flow	Parameter
Connection manager	



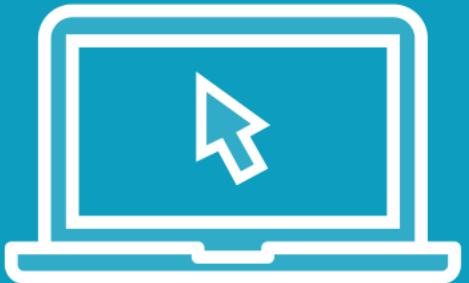
Demo



Creating a new SSIS project



Demo



Creating connection managers

- This is one of the first tasks to do in a new project
- Almost all components in SSIS need a connection



Demo



Creating SSIS components for loading Product dimension

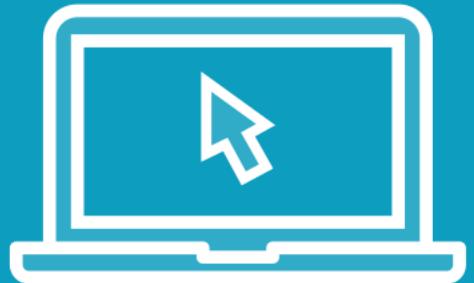
- Connection managers
- Tasks
- Precedence constraints
- Containers

The package will do the following:

- Update Lineage and Incremental loads tables
- Transfer data into the staging table
- Update the dimension table



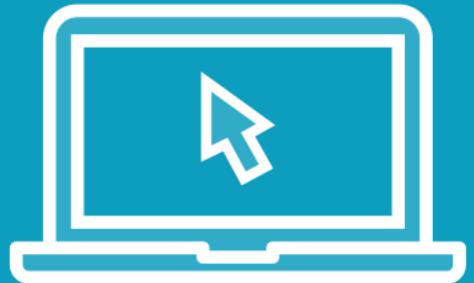
Demo



Creating and working with variables



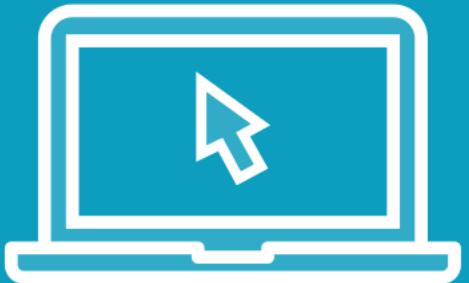
Demo



Creating precedence constraints



Demo



Executing a SQL task with parameters



Practical Exercise

Retrieve the date of the previous load

- Create a SQL task that executes the stored procedure **[int].[Get_LastLoadedDate]**
- Save the result in the variable: **[User::PreviousLoadDate]**
- Execute the task after the staging table is truncated
- Rename the procedure “SQL Get date of the previous load”



Summary



Advantages of using ETL tools

High-level overview of SSIS

Creating an SSIS project

Configuring connections to the data sources

Creating control flow tasks

Setting up the order of execution with precedence constraints

Creating and working with package variables



Loading a Data Warehouse with SSIS



Ana Voicu
@ana_voicu



Overview



Design data flow elements

Organize control flow tasks in containers

Add package parameters

Change the properties of a task using expressions

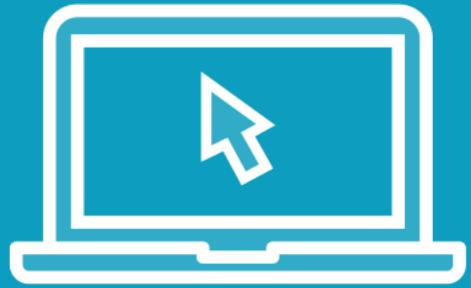
Deploy and execute the SSIS project in SQL Server

The result is a project that:

- Loads the data warehouse (full or incremental load)
- Can be customized for different customers with parameters
- Can be executed automatically, from a SQL Agent job



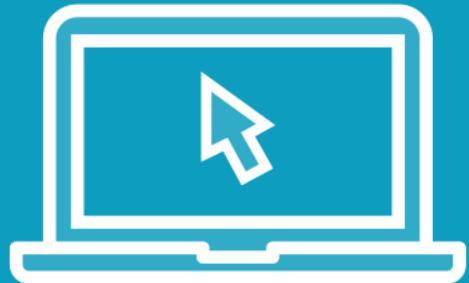
Demo



Designing data flow elements



Demo



Organizing data with containers



What Are Containers?



SSIS components that organize tasks into units

Three types of containers:

- For loop
- Foreach loop
- Sequence container

Description of Each Container Type



For loop

- The tasks are executed repeatedly, based on an expression

Foreach loop

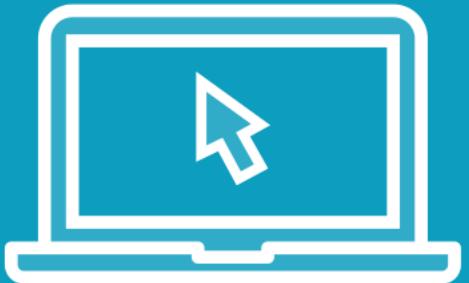
- Iterate through a collection of objects

Sequence container

- Treats the elements as a whole
- If a task fails, the entire container fails
- You cannot add a precedence constraint between a task from the container and a task outside of it
- “All for one and one for all” behavior



Demo



Working with package parameters

- What are parameters?
- Types of parameters
- How to create parameters?
- The difference between parameters and variables



SSIS Parameters



Assign values to properties at package execution

Can be set:

- At package level
- At project level

Choosing what type to use depends on the deployment strategy for the project

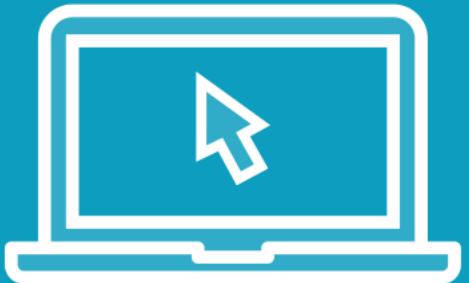
Create the link between the outside world and the content within SSIS

Examples:

- Database server name
- Type of load (full or incremental)



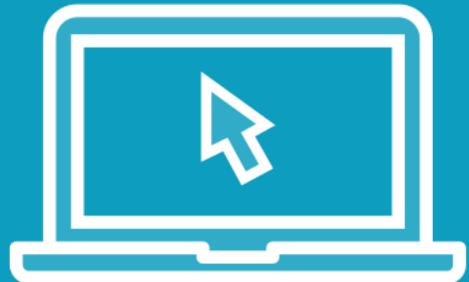
Demo



Adding expressions to tasks



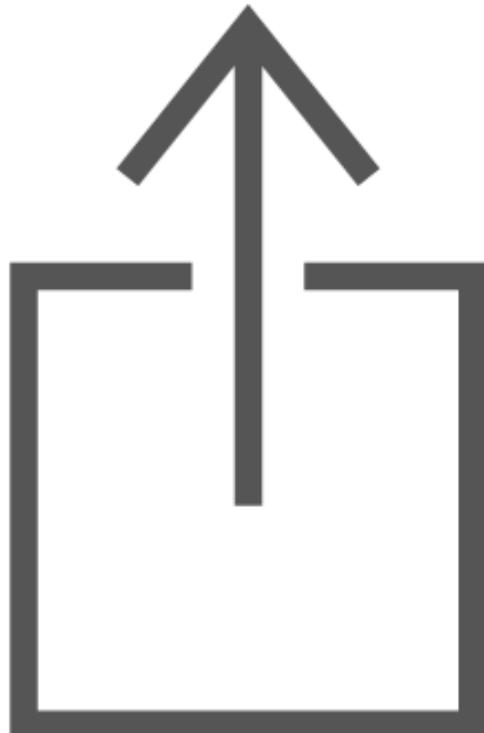
Demo



Deploying and configuring a SSIS project



SSIS Deployment Model



The current deployment model is available starting with SQL Server 2012

Everything SSIS-related is centralized in one place

- Project and package parameters
- Execution runtime values
- Execution status and other troubleshooting information
- General information about SSIS

The centralized location is called SSISDB catalog



Summary



Understand the difference between control flow and data flow

Created data flow components

Created tasks to load all the dimensions and fact tables

Organized tasks in containers

Added package parameters

Used expressions to change the properties of a task at package execution

Deployed an SSIS project in SQL Server



Visualizing Data from the Data Warehouse



Ana Voicu
@ana_voicu



Project Status Update

Task	Status
------	--------



Project Status Update

Task	Status
Create data warehouse	Done



Project Status Update

Task	Status
Create data warehouse	Done
Create load process in SSIS	Done



Project Status Update

Task	Status
Create data warehouse	Done
Create load process in SSIS	Done
Expose data and reports to end-users	Not done



Overview



Adding a semantic layer on top of the data warehouse

- A semantic layer is a structure that:
 - Reduces the complexity of the data warehouse
 - Shows only what is relevant
 - Creates pre-calculated measures, to improve performance
- SQL Server Analysis Services is used
 - Multidimensional model
 - Tabular model

Microsoft reporting tools:

- SQL Server Reporting Services
- Power BI
- Power BI Report Server



Creating a Semantic Layer



Stages of Creating a BI Solution



Summary of the previous steps

- Gathered business requirements
- Created the dimensional model
- Implemented the data warehouse in SQL Server
- Extended the DW with auxiliary tables
- Created the load process

Options to move forward

- a) Create the visualization layer directly from the data warehouse
- b) Create a semantic layer



Semantic layer

“It is a business representation of corporate data that helps end-users access data autonomously using common business terms.

A semantic layer maps complex data into familiar business terms such as product, customer, or revenue to offer a unified, consolidated view of data across the organization.” (Wikipedia)



Clarifying the Confusion



Isn't that what the dimensional design is doing?

- Yes, it is

The dimensional design

- Is an abstraction layer on top of different data sources
- Can be extended and simplified even more with a semantic model on top of it
- There are some advantages for creating the semantic model



Advantages of Using a Semantic Layer



Advantages of Using a Semantic Layer



Advantages of Using a Semantic Layer

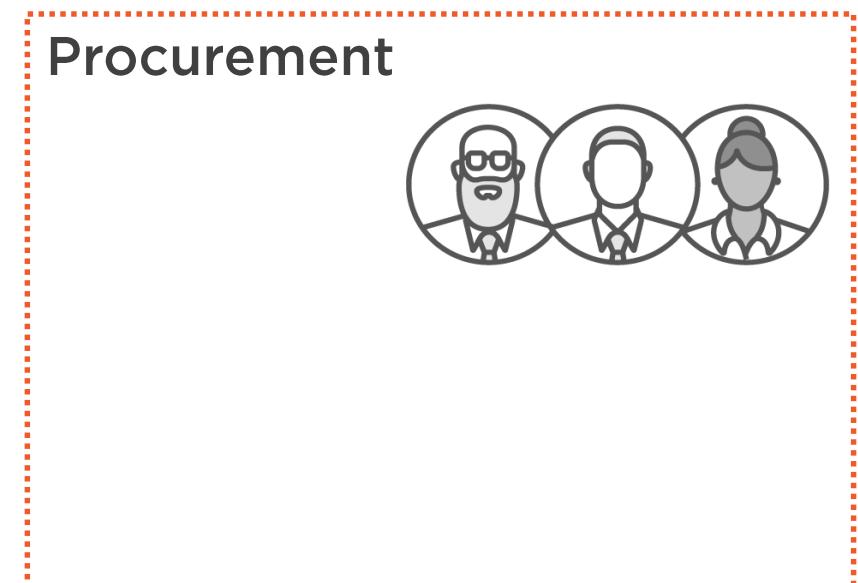


Sales

Procurement



Advantages of Using a Semantic Layer



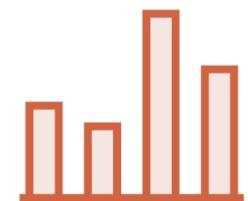
Advantages of Using a Semantic Layer



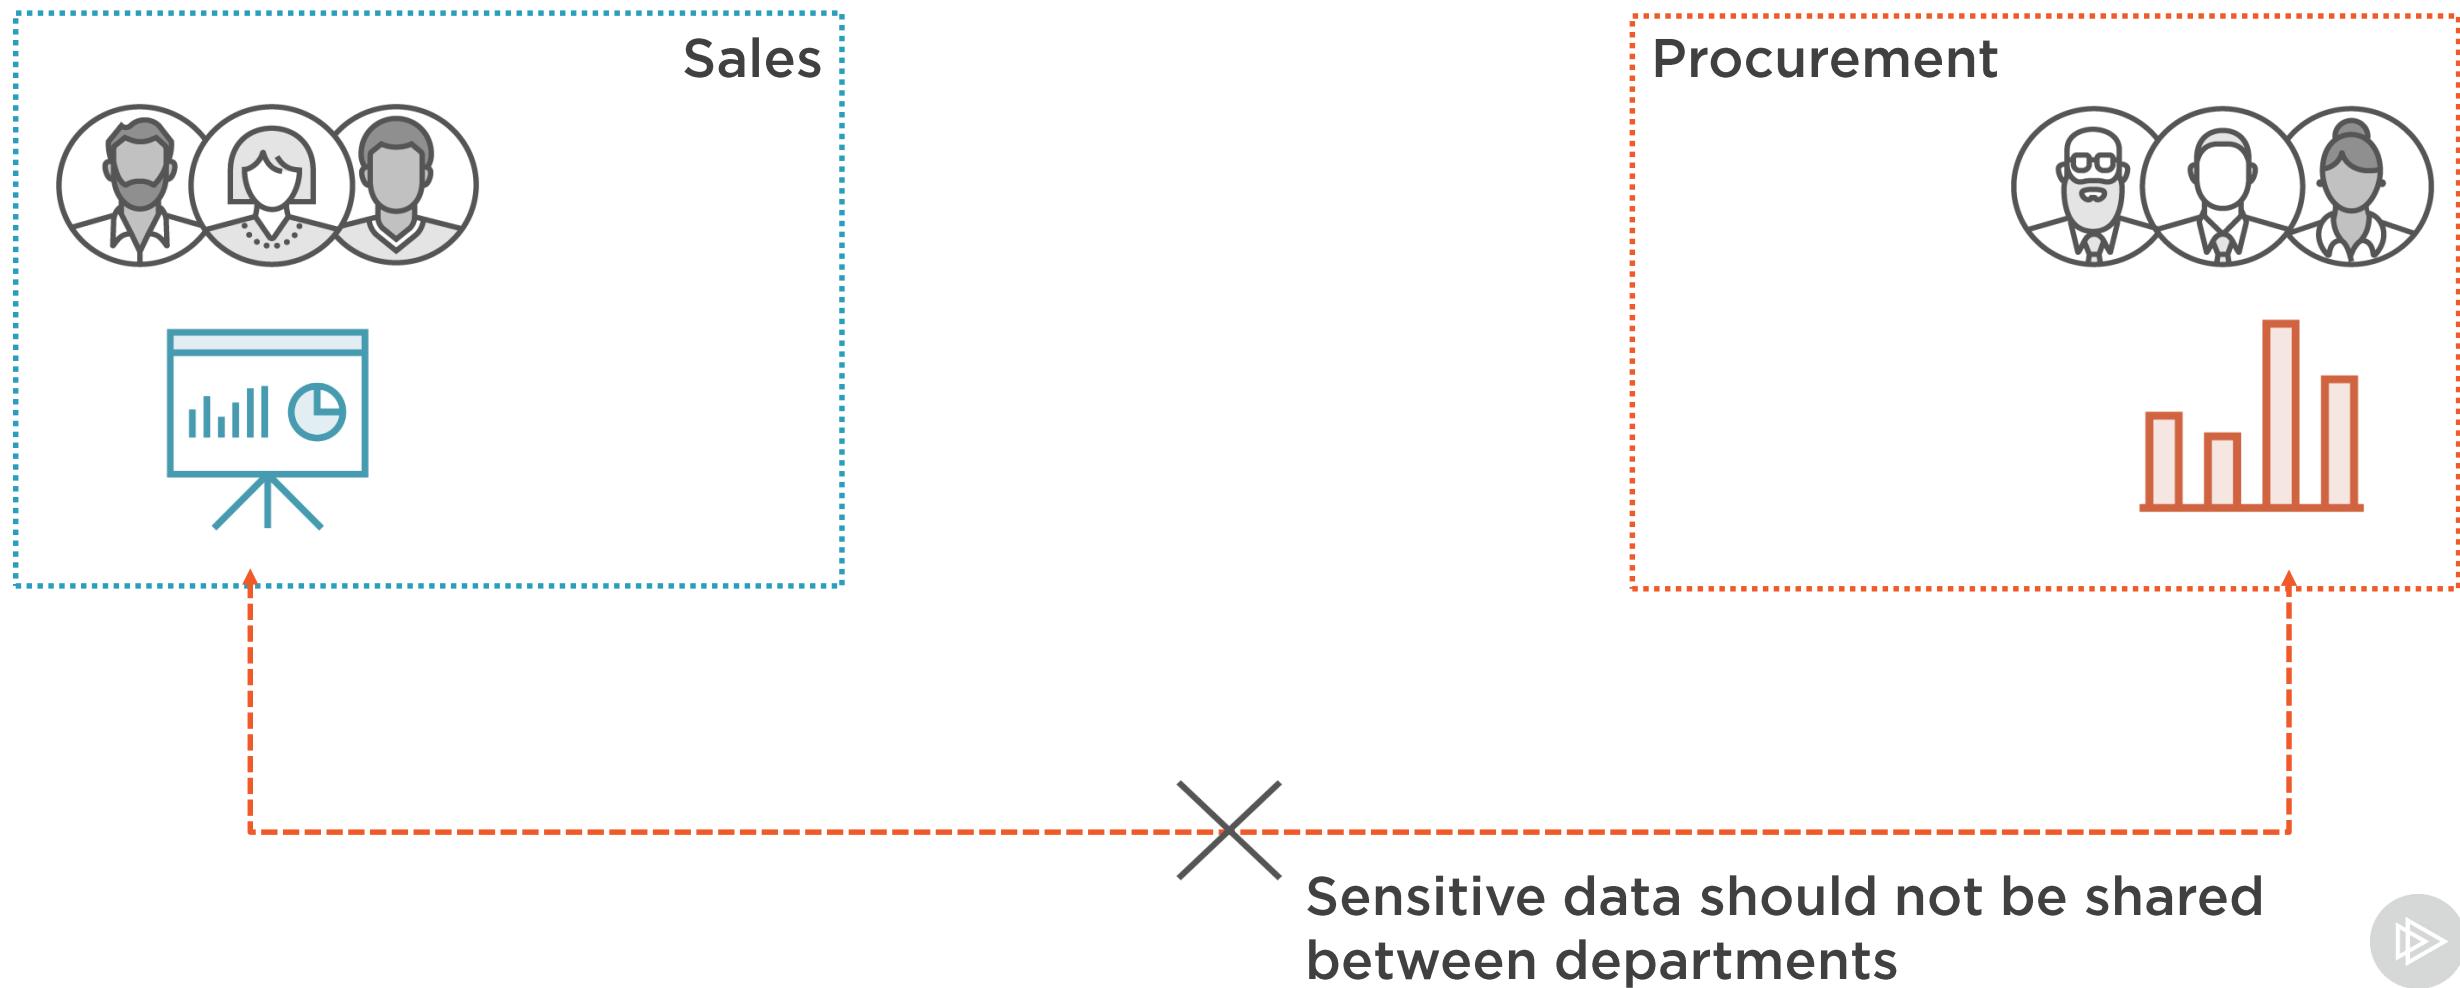
Sales



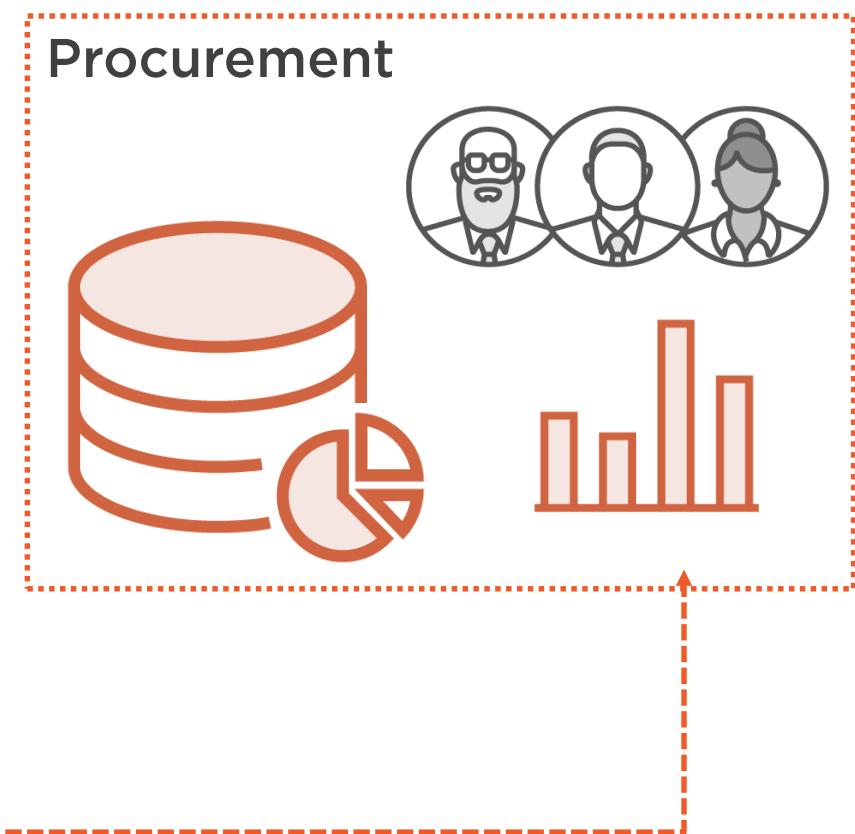
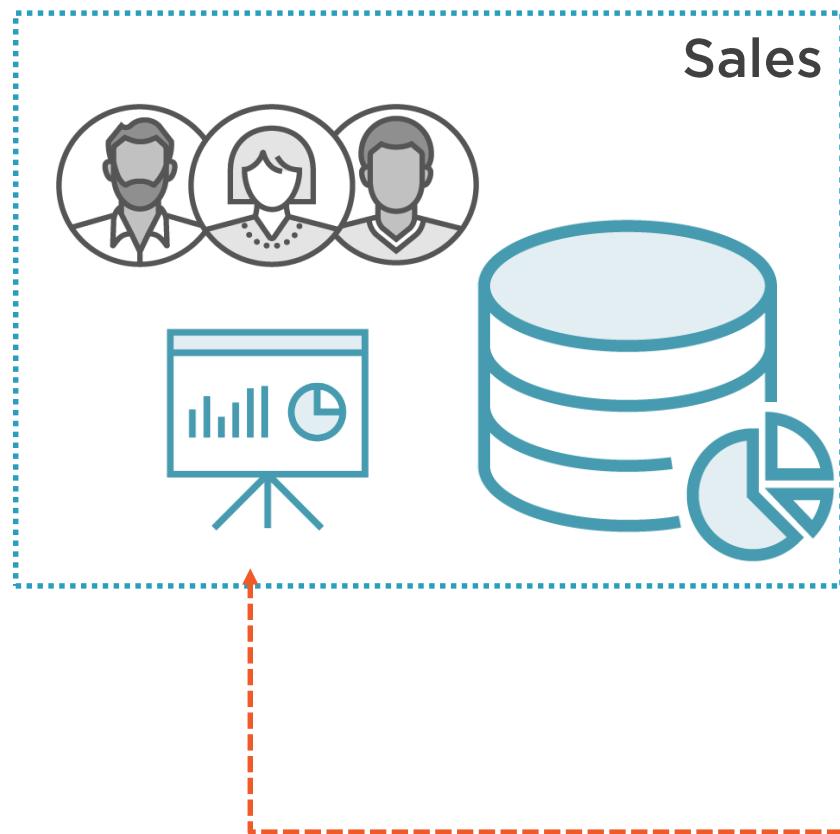
Procurement



Advantages of Using a Semantic Layer



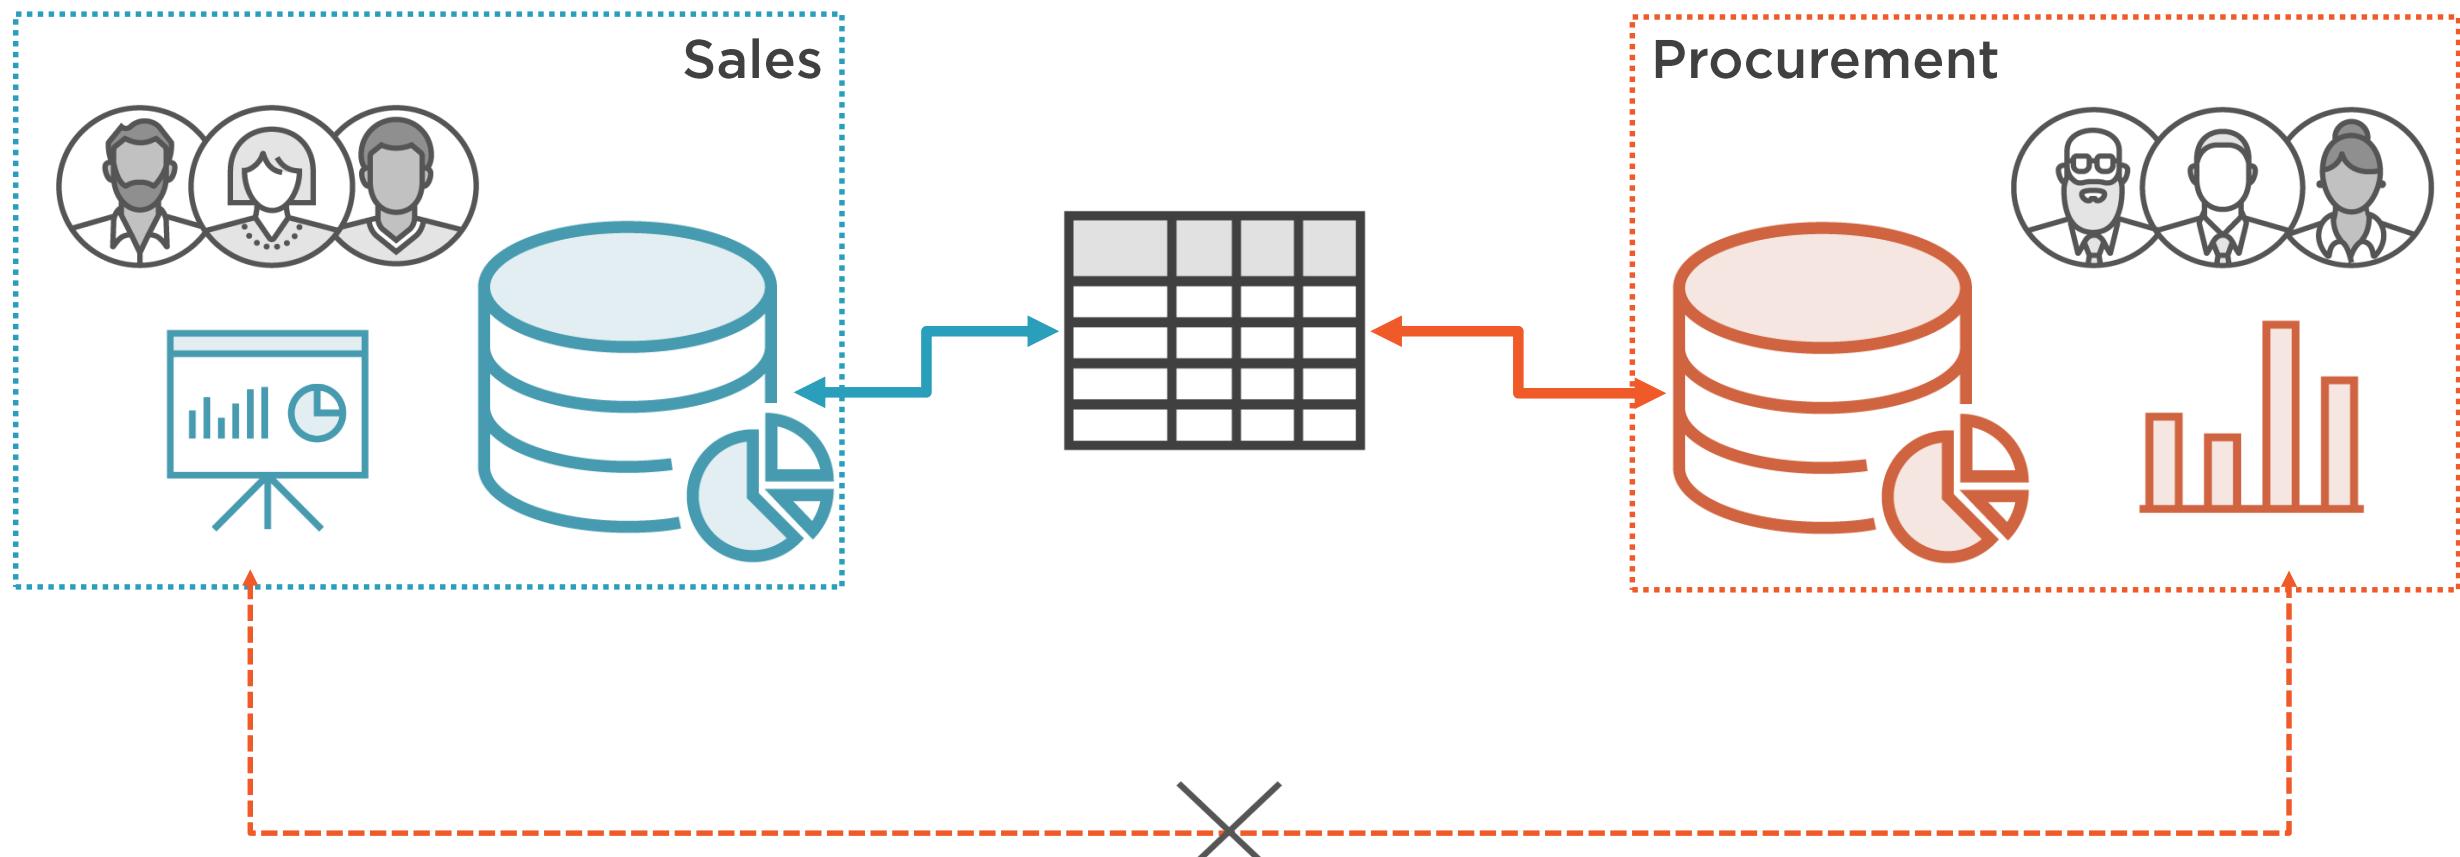
Advantages of Using a Semantic Layer



Sensitive data should not be shared
between departments



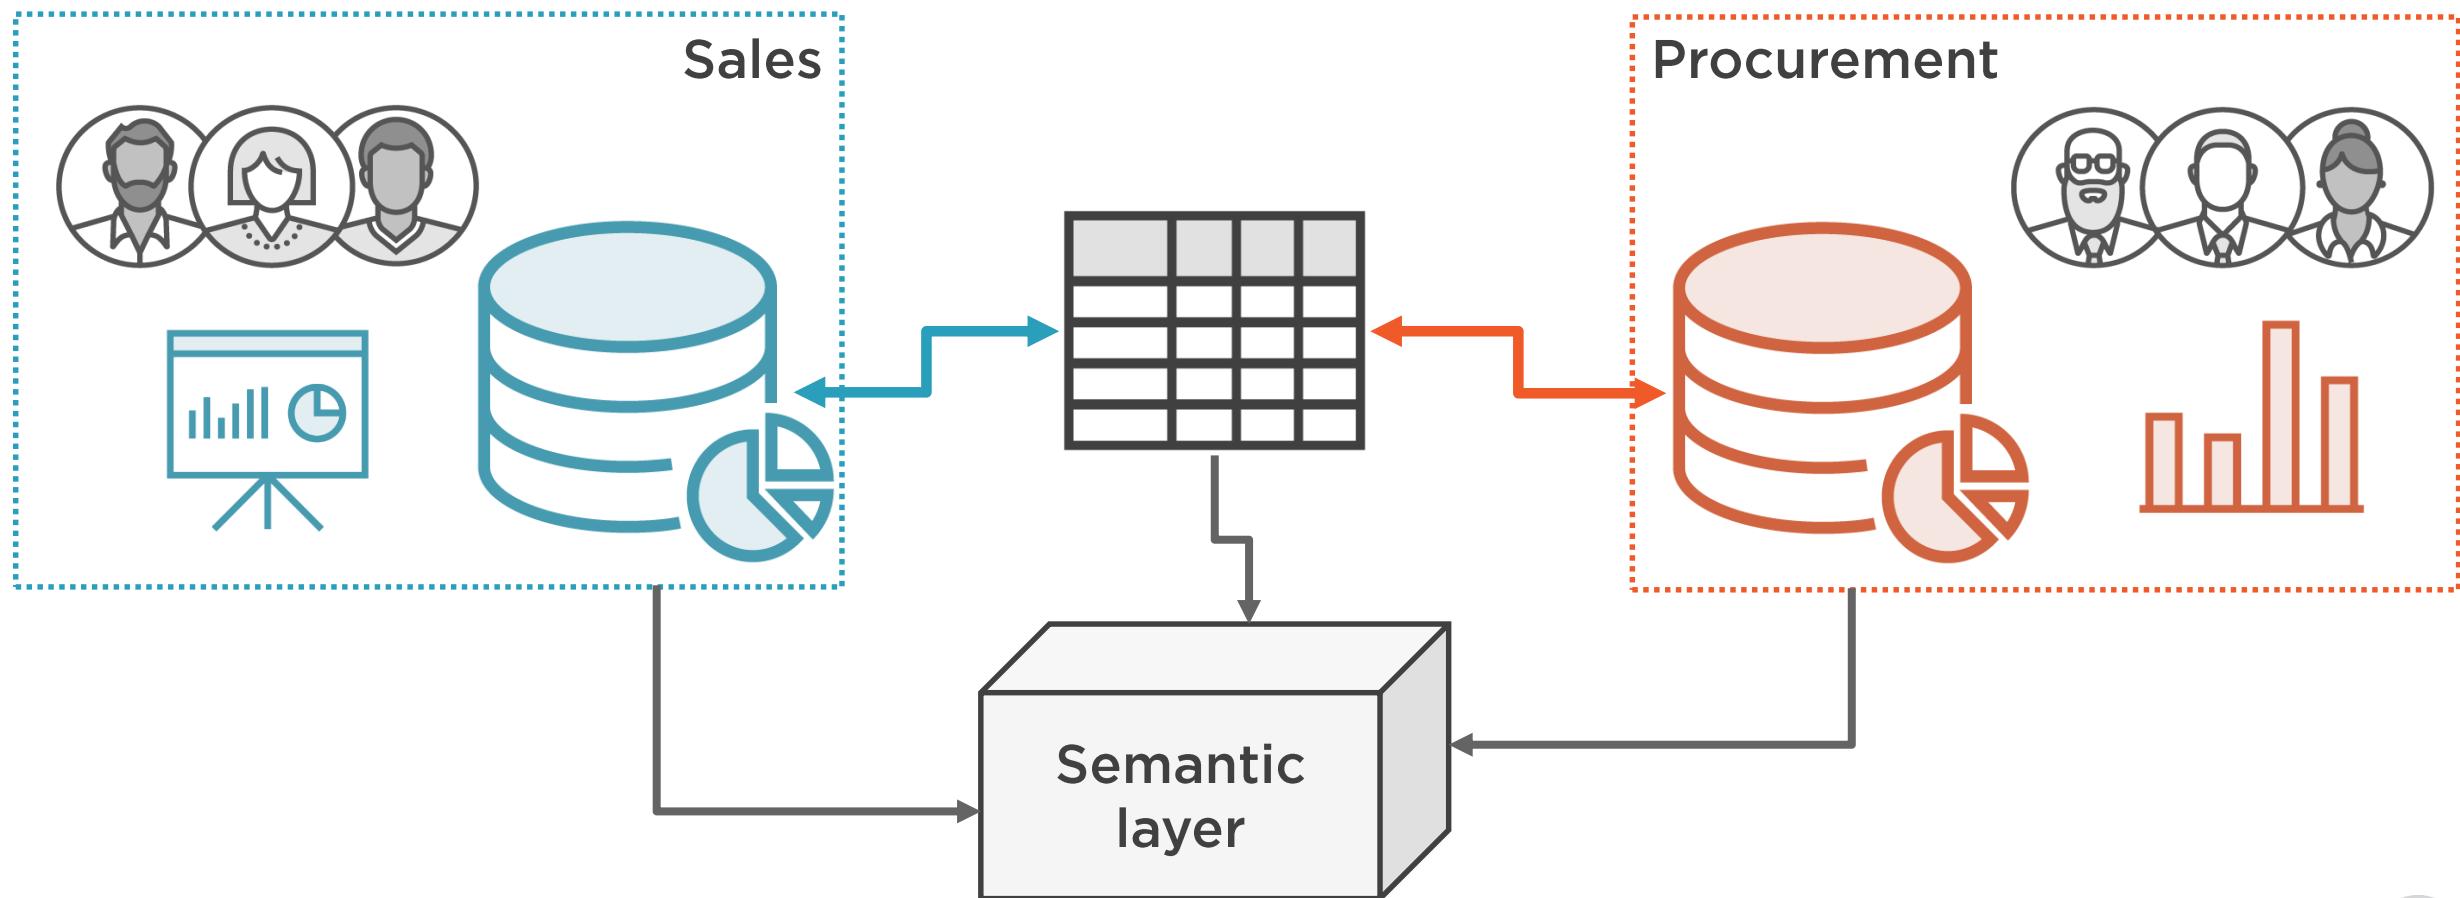
Advantages of Using a Semantic Layer



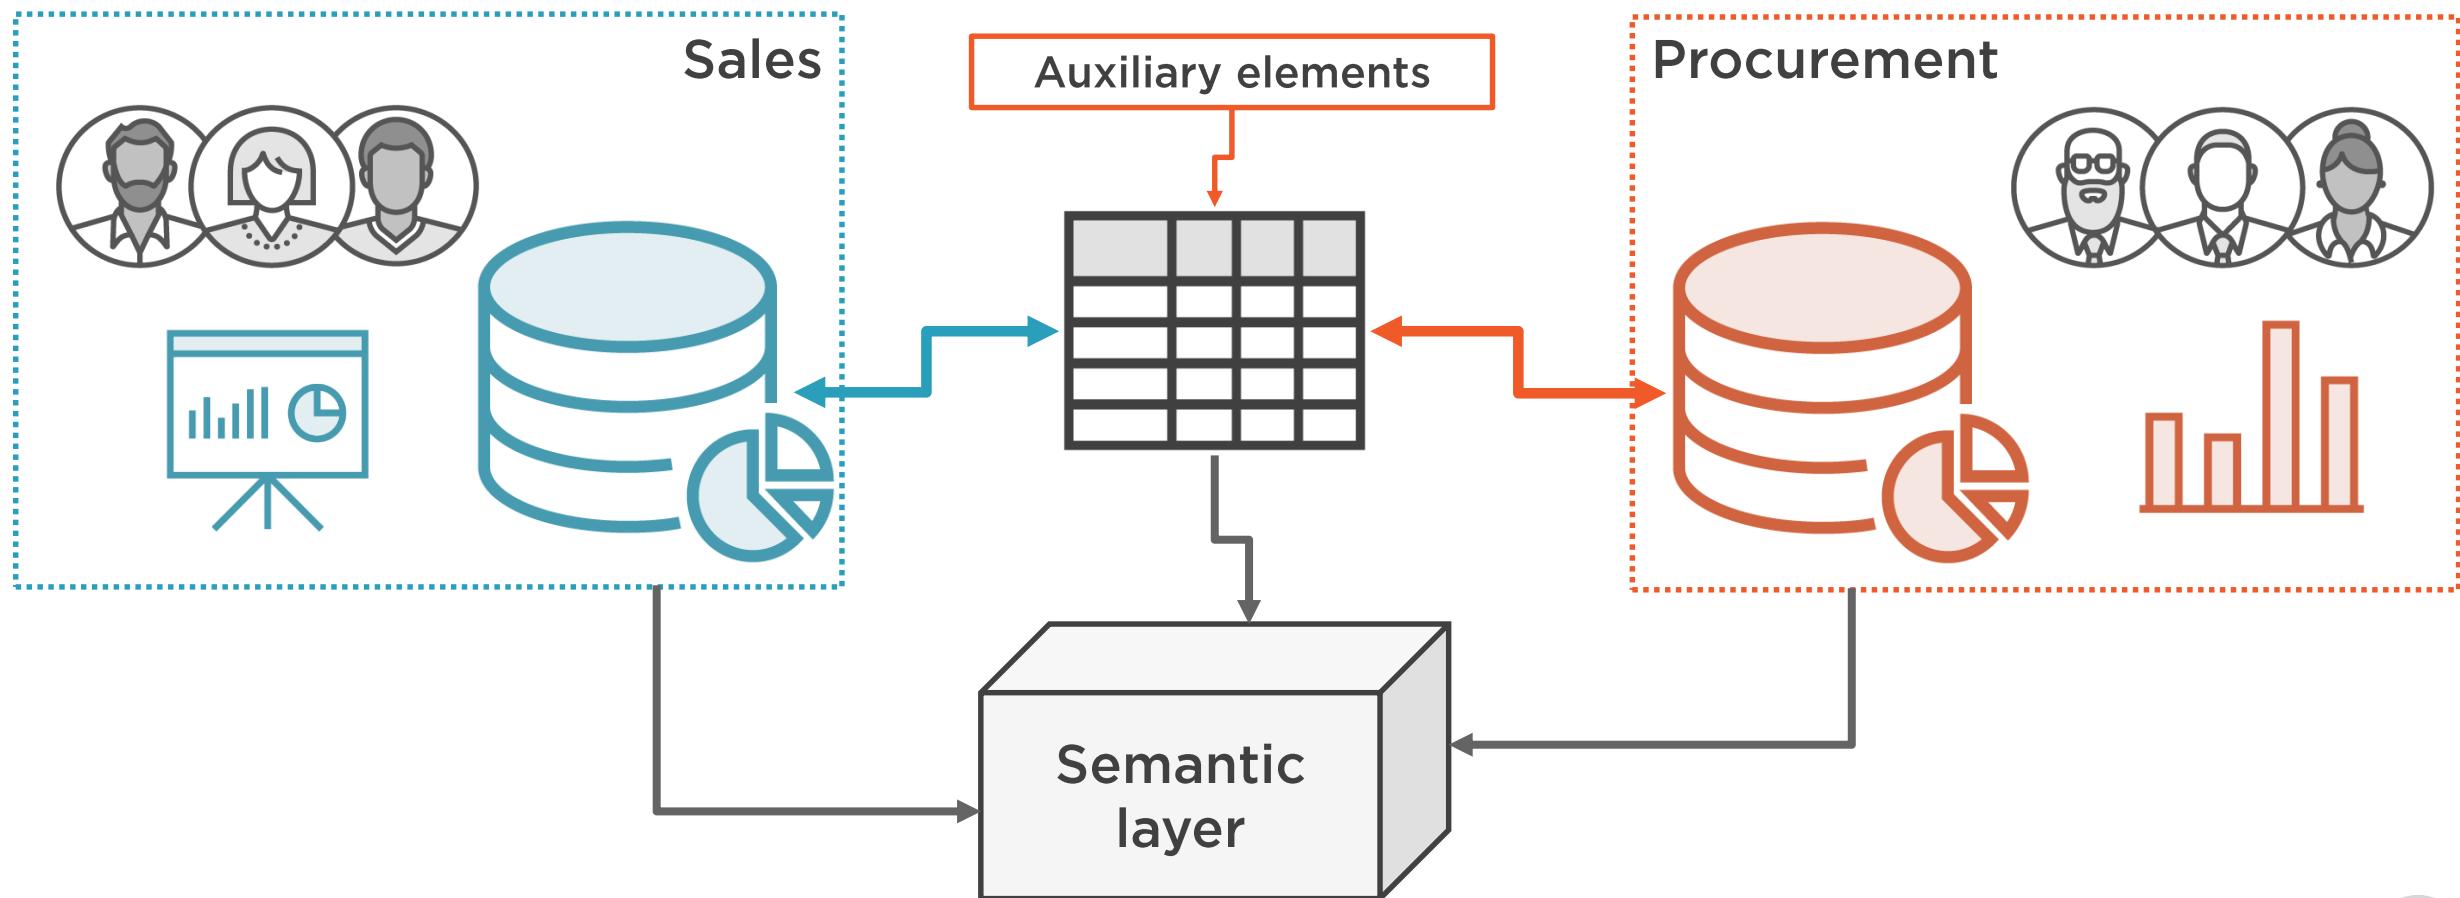
Sensitive data should not be shared
between departments



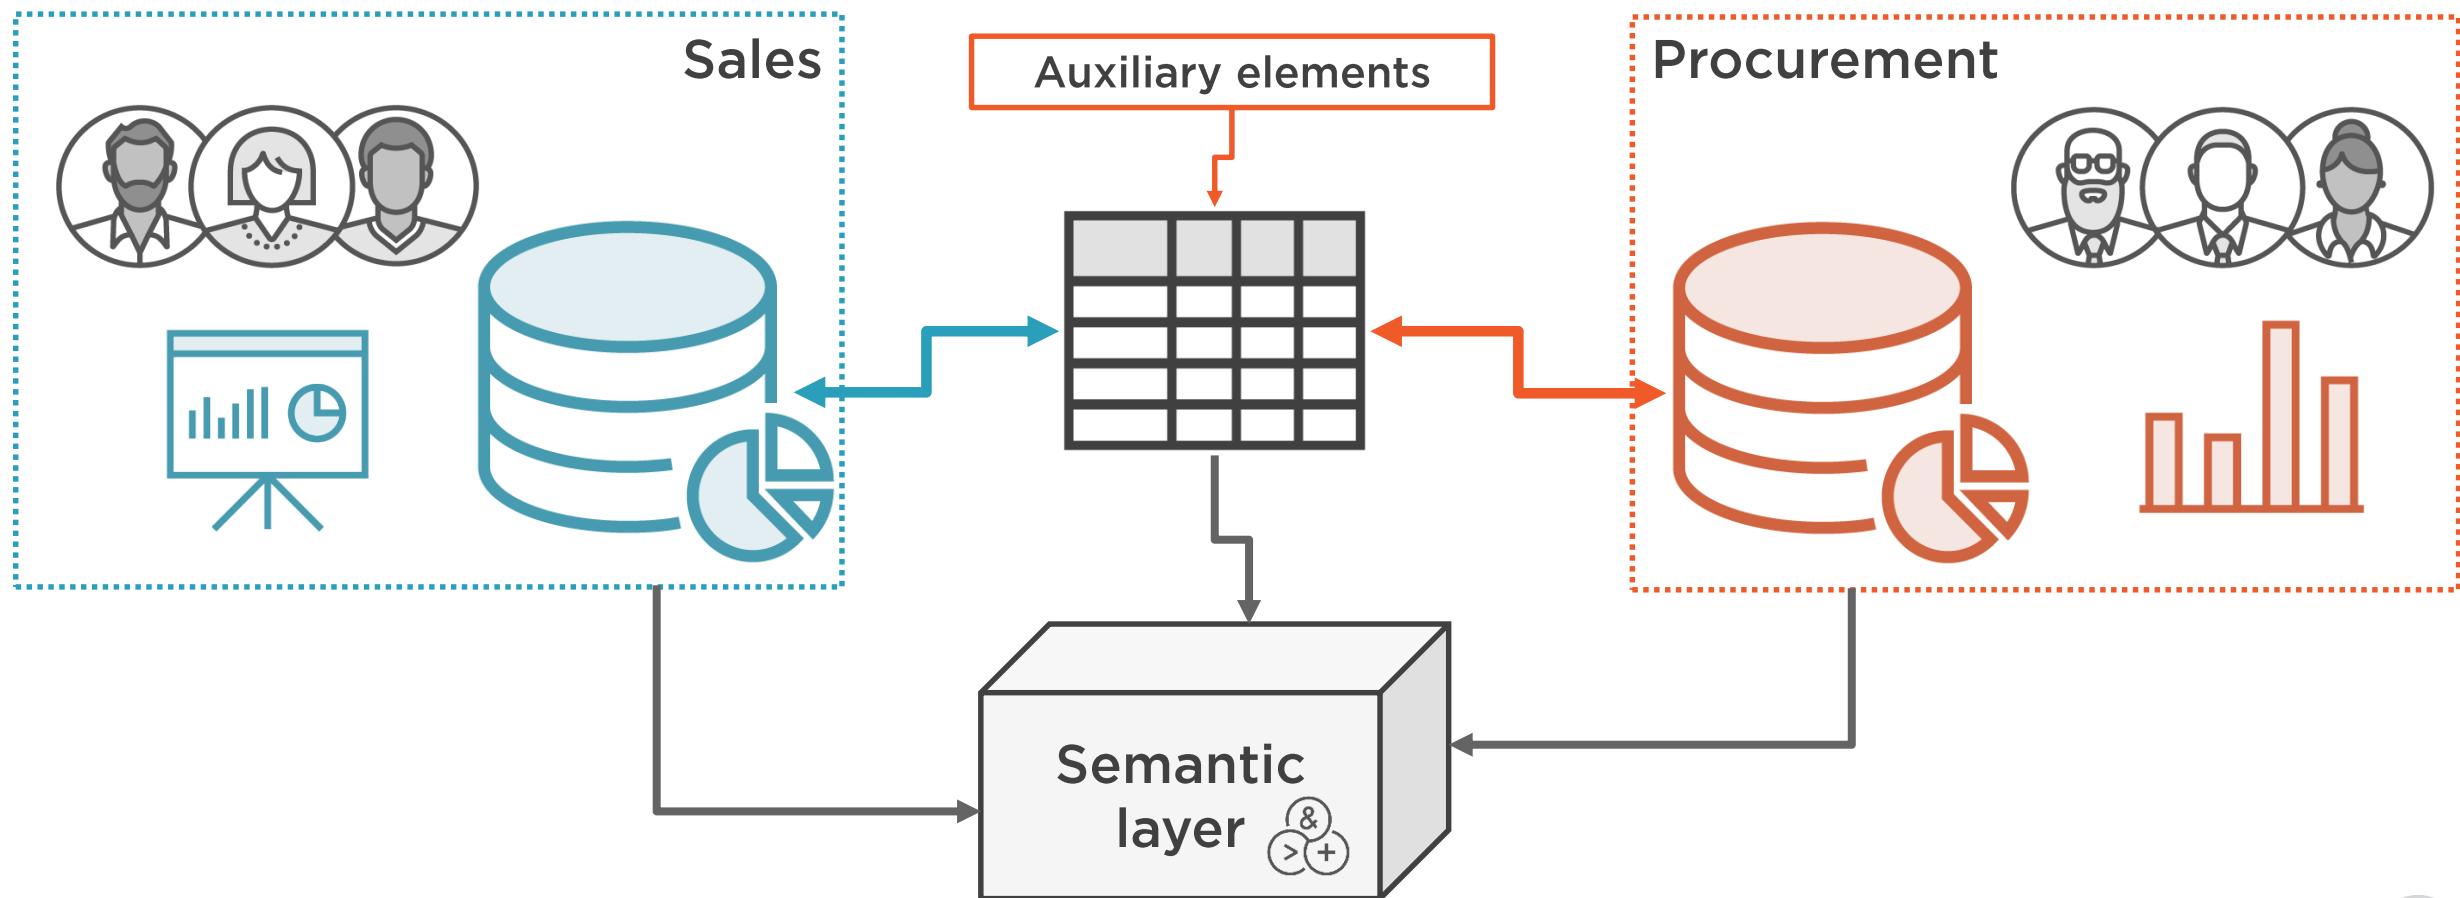
Advantages of Using a Semantic Layer



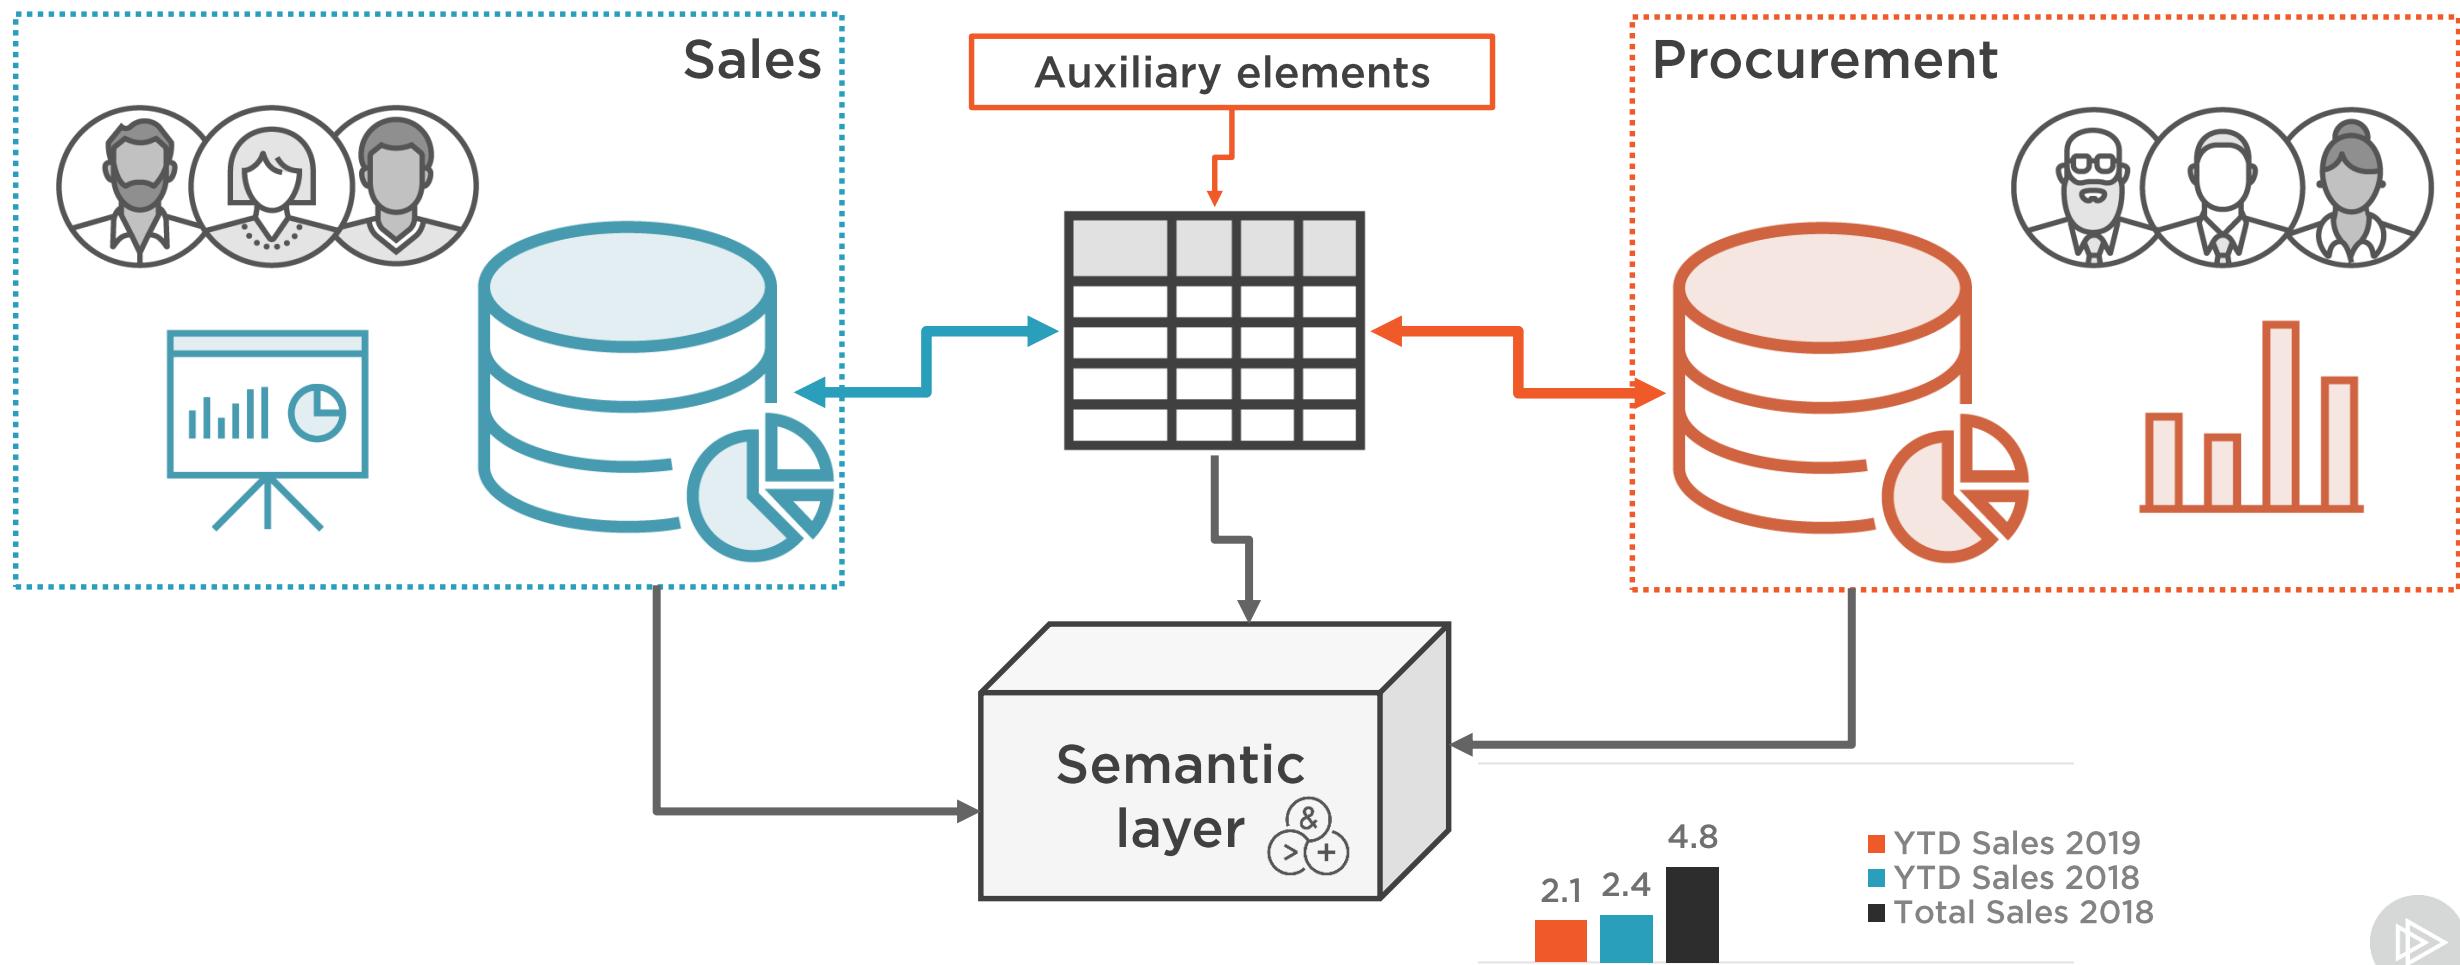
Advantages of Using a Semantic Layer



Advantages of Using a Semantic Layer



Advantages of Using a Semantic Layer



Advantages of Using a Semantic Layer

Show only
relevant data

Faster query
performance

Complex
aggregations

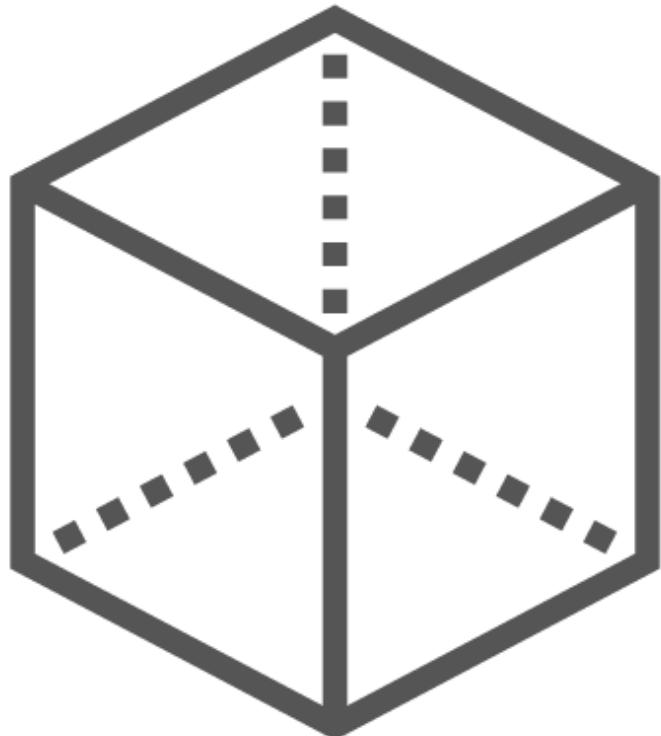
Readable names
for objects

Consolidate data
from multiple
source

Integration with
many BI tools



Types of Models in SSAS



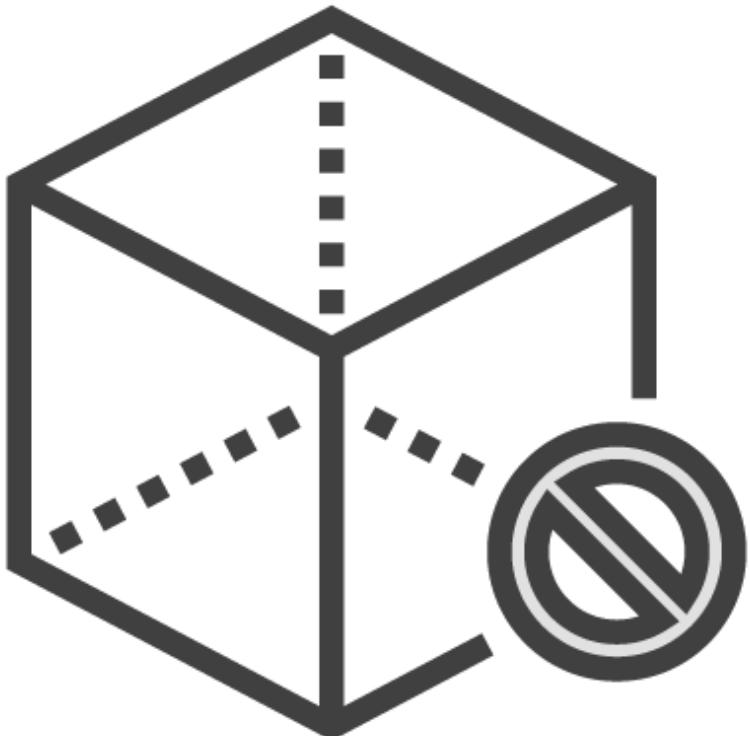
Multidimensional model
Tabular model



The Multidimensional and Tabular Models



The Multidimensional Model



Similar to the star schema design

The main structure used for data analysis is the cube

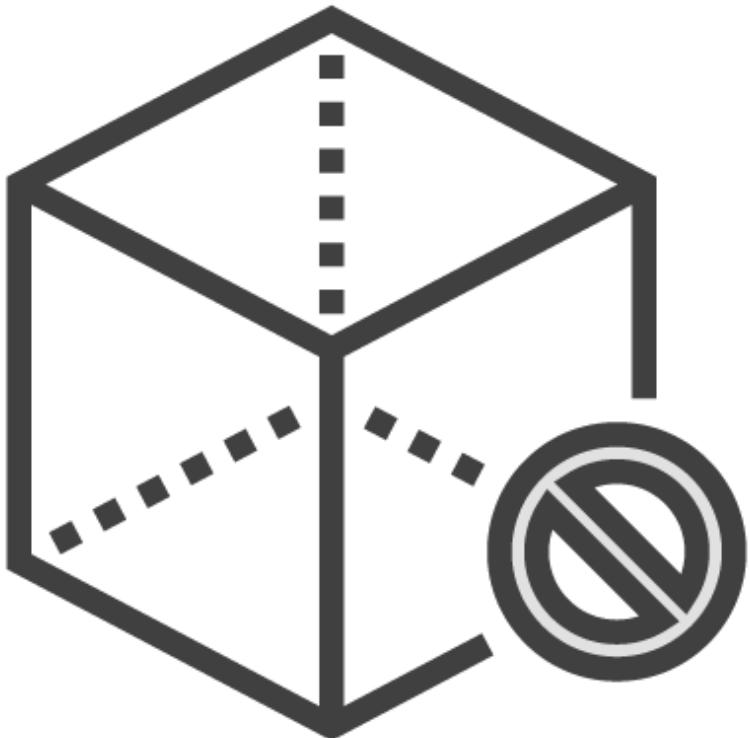
A cube consists of

- Dimensions
- Measure groups

Data sources must be defined before creating other elements



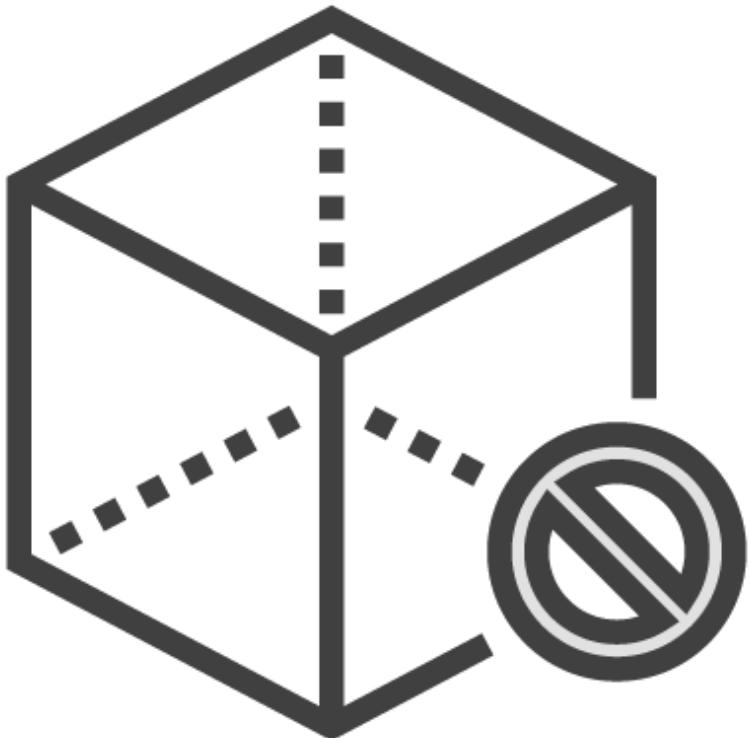
The Multidimensional Model



Multidimensional object	Data warehouse object
Dimension	Dimension table



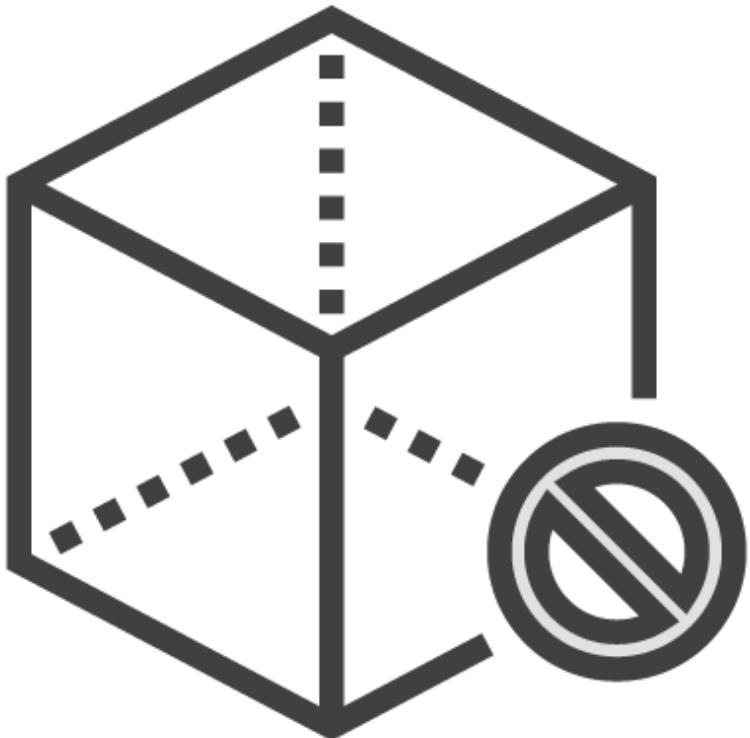
The Multidimensional Model



Multidimensional object	Data warehouse object
Dimension	Dimension table
Measure group	Fact table



The Multidimensional Model

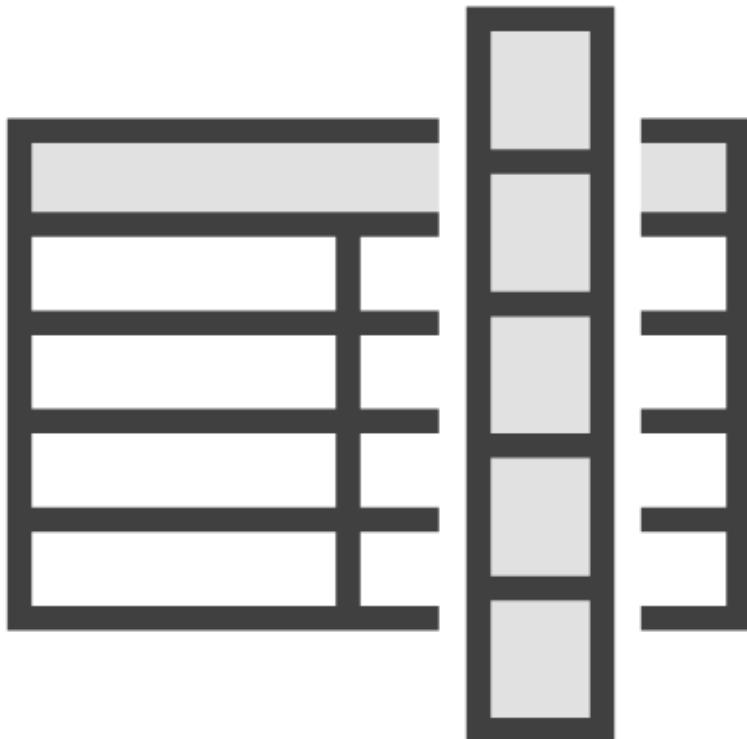


Multidimensional object	Data warehouse object
Dimension	Dimension table
Measure group	Fact table
Measure	Column in the fact table

New measures can be calculated in this database



The Tabular Model



In-memory database

Connects directly to a relational source of data

Similar to both SQL Server relational databases and a SSAS cube

- Consists of tables and relationships between tables
- You can also create measures and KPIs

The database can be used in Excel, Reporting Services, Power BI, etc.



Multidimensional vs. Tabular Model

Multidimensional

Better option when working with large amounts of data

Better performance in terms of scalability

Easily create and work with parent-child hierarchies

Writeback capabilities

Tabular

Easier for developers to understand and implement the model

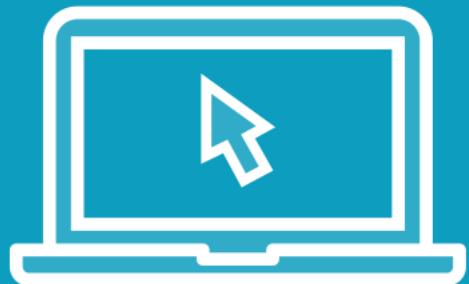
DAX language

Doesn't require strong hardware (disk) capabilities. It is only memory-dependent

The database is much smaller than the original source of data



Demo



Creating a multidimensional database in SSAS

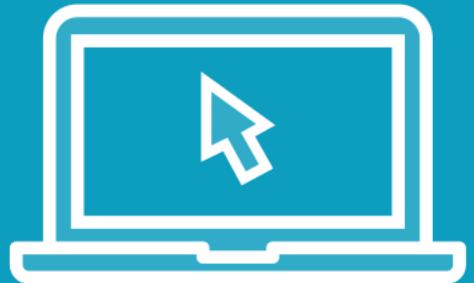
- This is a separate type of database

Prerequisite for this demo

- Install SQL Server Analysis Services in multidimensional mode



Demo



Analyzing data from a SSAS database



Creating Interactive Reports



SQL Server Reporting Services

Paginated reports

- The “traditional” reports
- Data is organized in tables, spread on multiple pages
- Optimized for printing or saving as Word or PDF

Mobile reports

- Compatible with several mobile devices
- Connect to different data sources (including SQL Server and SSAS)
- The visualizations are not very diverse, but they still have a modern look and feel

Web portal

- The portal can be accessed from any browser
- Types of reports: Power BI, mobile, paginated, KPIs, Excel files
- The data sets used in reports can be accessed from the portal



Power BI



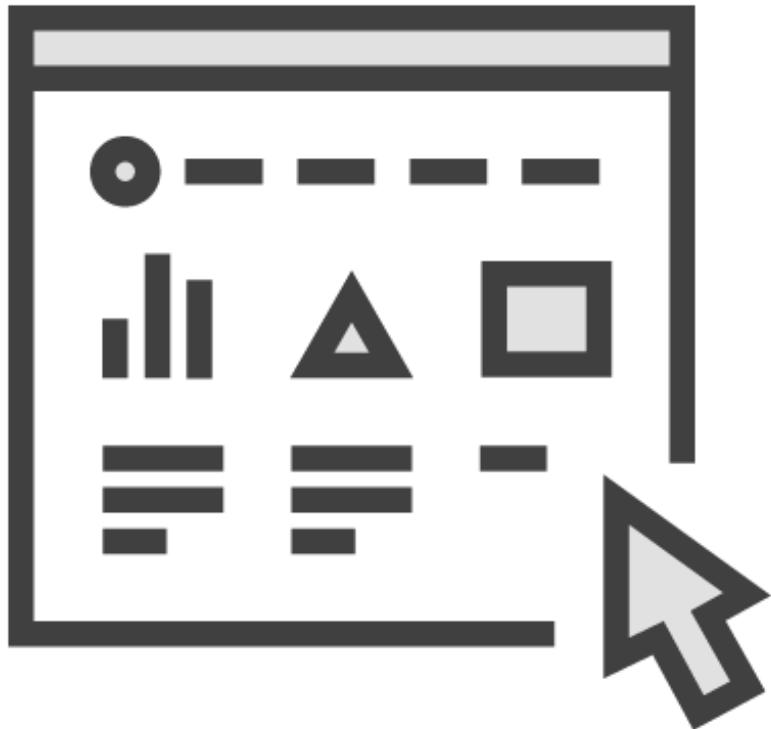
Tool used to create rich and interactive reports

Components:

- **Power BI Desktop**
 - Download and use for free, to create reports
 - Collaborative editing is not supported
- **Power BI Service**
 - Cloud service
 - Share reports with people within/external to the organization
- **Power BI mobile app**



Power BI Report Server



Reporting capabilities of Power BI

Data remains on-premises

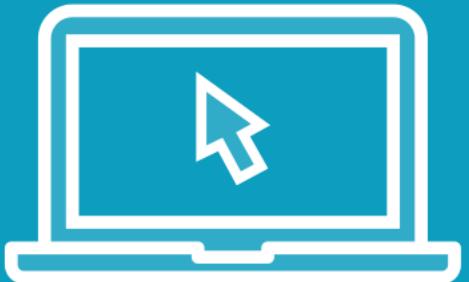
Users can visualize the reports

- In the web portal
- On mobile devices
- Shared by email

**Available in SQL Server Enterprise Edition,
or with Power BI Premium license**



Demo



Creating and publishing a report with Power BI Report Server

You need to install two components

- Power BI Report Server
- Power BI Desktop for Report Server
- Download them from the Microsoft Download Center

Power BI Report Server is available in

- SQL Server Enterprise Edition
- SQL Server Developer Edition



Course Summary



Fundamental concepts of a data warehouse

- Purpose of a data warehouse
- Facts and fact tables
- Dimensions and dimension tables
- Steps of the dimensional design process

Dimension table techniques

Hierarchies

Fact tables

- Difference between fact and dimension tables
- Common types of fact tables



Course Summary



Creating the SQL scripts for loading a dimension table

- Create the Lineage and Incremental loads tables
- Create the stored procedures for loading data

Implementing a load process in SSIS

- Create a new SSIS project in Visual Studio
- Work with control flow and data flow tasks
- Add parameters
- Deploy the project in SQL Server
- Create a SQL Agent job to run the project automatically

Multidimensional and tabular models

Creating reports on top of the data warehouse



Thank you!

